

Fundamentals

"Anyone can write code which a machine can understand - the trick is to write code which another human can understand." -- Martin Fowler

Introduction

- What is computer programming?
- What is a computer program?
- What is a programming language?
- Differences between languages ([Currently popular languages](#) Don't take it too seriously!)
- Compilers vs. Interpreters
- Creating a program
 1. Editing source files
 2. Compiling source files into object files
 3. Linking object files into executable files
- Executing the program
- Central Processing Unit (CPU)
 - Arithmetic and Logic Units
 - Registers
 - Cache memory
- Main memory

First C Programs (The Chicken and The Egg)

- A C program is essentially a collection of one or more functions.
- There must be a function named `main`; it must be in all lowercase.
- `main` is the program's starting point; it can call other functions by name
- There must be *exactly* one function named `main` in every program.

The general form of a C program looks like this: (The parts in **bold** are required)

```
include files

function declarations (prototypes)

data declarations (global)

main function header
{
    data declarations (local)
    statements
}

other functions
```

Therefore, the simplest C program you can write:

```
int main(void)
{
}
```

Technically, you should have a `return` statement:

```
int main(void)
{
    return 0;
}
```

It looks simple because it is. It does nothing of interest. But, nevertheless, it produces a "functional" program. This simple program demonstrates many characteristics of a C program.

Students that think that `int main(void)` can be changed to `int main()` are required [to read this](#), as I'm not going to spend any more time on it. (Yes, I know that [our textbook](#) interchanges them, but you shouldn't.)

A second program in C that actually does something:

```
#include <stdio.h>

/* Say hi to the world */
int main(void)
{
    printf("Hello, World!\n");
    return 0;
}
```

Output:

Hello, World!

Adding line numbers for clarity. They are not (and cannot be) present in the actual code.

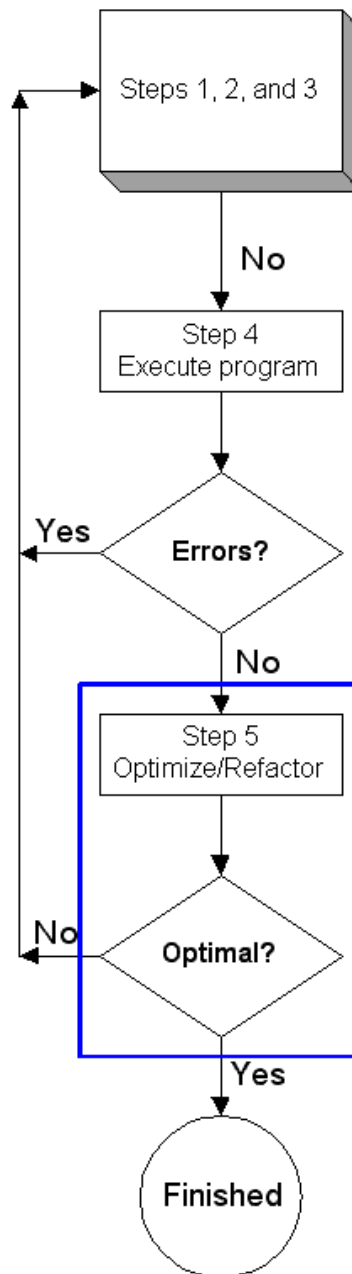

```
1. #include <stdio.h>
2.
3. /* Say hi to the world */
4. int main(void)
5. {
6.     printf("Hello, World!\n");
7.     return 0;
8. }
```

Each non-blank line above has significant meaning to the C compiler. Look at [stdio.h](#) to see what the pre-processor adds.

Editing, Compiling, Linking, and Executing

Edit/Compile/Execute Loop

Edit/Compile/Execute Loop Extended



Contrasting Languages at 3 Levels

Let's say we have 4 variables (just like in algebra), where $a = 1$, $b = 2$, $c = 3$, $d = 4$. We want to set a fifth one, e , to $ab(c + d)$. We can rewrite the multiplication explicitly as

$$a \times b \times (c + d)$$

This will result in e having the value 14.

C code (partial program)

```

int a = 1;
int b = 2;
int c = 3;
int d = 4;
int e = a * b * (c + d);

```

Assembly language (Intel x86)

Assembly language (compiler-generated)	Assembly language (with comments)
<pre> _main proc far mov word ptr [bp-2],1 mov word ptr [bp-4],2 mov word ptr [bp-6],3 mov word ptr [bp-8],4 mov ax,word ptr [bp-2] imul word ptr [bp-4] mov dx,word ptr [bp-6] add dx,word ptr [bp-8] imul dx mov word ptr [bp-10],ax _main endp </pre>	<pre> _main proc far mov word ptr [bp-2],1 ;the address of a is bp-2 ;the value is 1 mov word ptr [bp-4],2 ;the address of b is bp-4 ;the value is 2 mov word ptr [bp-6],3 ;the address of c is bp-6 ;the value is 3 mov word ptr [bp-8],4 ;the address of d is bp-8 ;the value is 4 mov ax,word ptr [bp-2] ;put a's value in ax reg imul word ptr [bp-4] ;multiply ax reg by b, ;put result back in ax mov dx,word ptr [bp-6] ;put c's value in dx reg add dx,word ptr [bp-8] ;add d to the dx reg ;put result back in dx imul dx ;multiply ax reg by dx ;put result back in ax mov word ptr [bp-10],ax ;the address of e is bp-10 ;the value of e is 14 _main endp </pre>

Hand-coded assembler (80386 using GNU's assembler, comments start with # and are in **bold**)

```

.section .data
a: .long 1
b: .long 2
c: .long 3
d: .long 4
e: .long

.section .text
.globl _start

_start:
    movl a, %eax    #    a -> %eax
    movl b, %ebx    #    b -> %ebx
    imull %eax, %ebx # a * b -> %ebx

    movl c, %eax    #    c -> %eax
    movl d, %ecx    #    d -> %ecx
    addl %eax, %ecx  # c + d -> %ecx

    imull %ebx, %ecx # (a * b) * (c + d) -> %ecx
    movl %ecx, e     # put result in e

```

Here are [two other assembler samples](#) output from Microsoft's compiler and GNU's compiler.

Machine language (Intel 80386)

Hexadecimal dump

```
014c 0003 0000 0000 0110 0000 000c 0000
```

Octal dump

```
000514 000003 000000 000000 000420 000000 000014 000000
```

000000	000404	072056	074145	000164	000000	000000	000000
000000	000000	000160	000000	000214	000000	000374	000000
000000	000000	000002	000000	000040	060000	062056	072141
000141	000000	000000	000000	000000	000000	000000	000000
000000	000000	000000	000000	000000	000000	000000	000000
000100	140000	061056	071563	000000	000000	000000	000000
000000	000000	000000	000000	000000	000000	000000	000000
000000	000000	000000	000000	000200	140000	104525	101745
014354	162203	134360	000000	000000	140203	101417	007700
164301	140404	002340	042611	105750	164105	000350	000000
164000	000000	000000	042707	000774	000000	143400	174105
000002	000000	042707	001764	000000	143400	170105	000004
000000	042613	104774	007702	052657	105770	170105	042403
007764	141257	042611	134354	000000	000000	141711	110220
110220	110220	110220	110220	110220	110220	000041	000000
000013	000000	000024	000046	000000	000011	000000	000024
063056	066151	000145	000000	000000	000000	177776	000000
000547	064563	070155	062554	061456	000000	000000	000000
000000	000000	066537	064541	000156	000000	000000	000000
000001	000040	000002	072056	074145	000164	000000	000000
000000	000001	000000	000403	000142	000000	000002	000000
000000	000000	000000	000000	000000	062056	072141	000141
000000	000000	000000	000002	000000	000403	000000	000000
000000	000000	000000	000000	000000	000000	000000	061056
071563	000000	000000	000000	000000	000003	000000	000403
000000	000000	000000	000000	000000	000000	000000	000000
000000	057537	066537	064541	000156	000000	000000	000000
000040	000402	000000	000000	000000	000000	000000	000000
000000	000000	000000	057537	066141	067554	060543	000000
000000	000000	000000	000002	000004	000000		

[illegible]

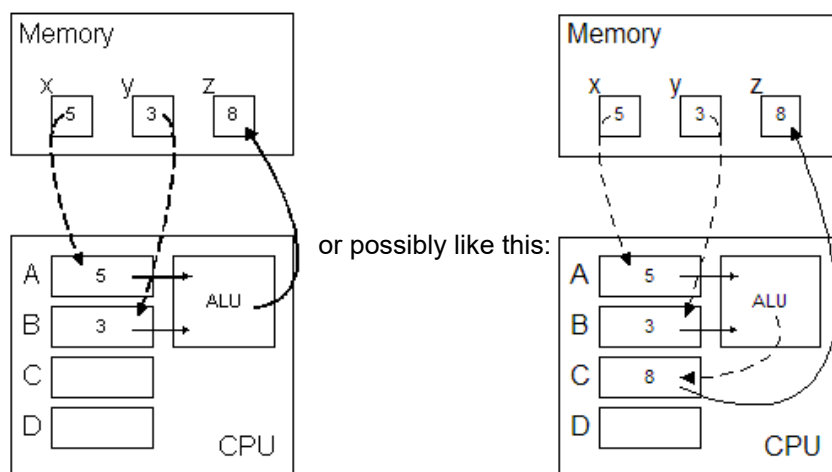
[illegible]

Now, which language would you rather work with?

Simple calculation:

```
int x = 5;
int y = 3;
int z = x + y;
```

Simplified view at runtime showing 3 variables in memory, a CPU with 4 registers and an Arithmetic-Logic Unit:



Putting It All Together

Step 1: Edit

Create a text file for this C code named `simple.c`. The size of this file is 119 bytes. (You can use any text editor like Notepad++ or, preferably, the [Crimson Editor](#).)

```
int main(void)
{
    int a = 1;
    int b = 2;
    int c = 3;
    int d = 4;
    int e = a * b * (c + d);
    return 0;
}
```

Step 2: Compile

The *source* code (text) is compiled into *object* code (binary) and saved in a file named `simple.o`. The size of this file is about 450 bytes (depends on the compiler and version).

```
gcc -c simple.c -o simple.o
```

Step 3: Link

The object file is linked (combined) with other object code and saved in a file named `simple.exe`. The size of this file is about 10,000 bytes. (Again, depends on the compiler.)

```
gcc simple.o -o simple.exe
```

Step 3: Execute

Run the executable file by simply typing the name of the executable file (the `.exe` extension is optional under Windows):

```
simple
```

Of course, nothing *appears* to happen. The program *did* run and it *did* perform the calculations. There just aren't any instructions in the program that display anything for you to see. (Also, the sizes of the object files and executable files will vary with the version of the compiler you are using.)

We can modify it to use the `printf` function and display the value of `e` after the calculations:

```
#include <stdio.h> /* printf */

/* Calculate some values */
int main(void)
{
    int a = 1;
    int b = 2;
    int c = 3;
    int d = 4;

    int e = a * b * (c + d);
    printf("%i\n", e);
    return 0;
}
```

The C program above is a complete program that, when executed, will print the value 14 on the screen.

Additional Compile Switches

To enable the compiler to perform a more thorough check of your source code, use the `-Wall` *command line switch* like so:

```
gcc -Wall -c simple.c -o simple.o
```

Now, if the compiler detects any potential problems or misuse of the C language, it will alert you with a warning message. For example, if I add the variable `f` like this:

```
int a = 1;
int b = 2;
int c = 3;
int d = 4;
int f = 5; /* Add this variable, but don't use it anywhere */

int e = a * b * (c + d);
printf("%i\n", e);
return 0;
```

I get this warning from the compiler:

```
simple.c: In function 'main':
simple.c:10: warning: unused variable 'f'
```

or, if I add something like this:

```
a + b - c; /* Perform some calculation, but discard the value */
```

I get this warning from the compiler:

```
simple.c: In function 'main':
simple.c:11: warning: statement with no effect
```

Compile and Link in One Step

If you want to perform both compile and link steps with one command, don't provide the `-c` switch. This will compile and then link the program:

```
gcc simple.c -o simple.exe
```

However, if the compile step fails for any reason, the link step is skipped.

Other Useful Switches

If you want to generate the assembly output, use `-S` switch:

```
gcc -S simple.c
```

This will produce a text file called `simple.s` which you can view with any text editor.

If you want to generate the preprocessor output, use `-E` switch:

```
gcc -E simple.c
```

This will produce a ton of information to the screen. To capture the output so you can view it more easily, redirect the output to a file:

```
gcc -E simple.c > simple.out
```

The

```
> simple.out
```

causes the output to be written to a file (named `simple.out`) instead of to the screen. This file is a text file that you can view with any text editor.

Another Example

This example will show constructs such as identifiers, literal constants, defines, expressions, and several others.

The C code: `marathon.c`: (with line numbers for clarity)

```
1. #include <stdio.h> /* printf */
2.
3. #define YARDS_PER_MILE 1760
4. #define KILOS_PER_MILE 1.609
5.
6. /* A marathon is 26 miles, 385 yards */
7. /* Prints the distance of a marathon in kilometers */
8. int main(void)
9. {
10.     int miles, yards;
11.     double kilometers;
12.
13.     miles = 26;
14.     yards = 385;
15.     kilometers = (miles + (double)yards / YARDS_PER_MILE) * KILOS_PER_MILE;
16.     printf("A marathon is %f kilometers\n", kilometers);
17.
18.     return 0;
19. }
20.
```

The program will output: A marathon is 42.185969 kilometers

The preprocessed file named `marathon.pre` (generated by: `gcc -E marathon.c > marathon.s`)

The assembly file named `marathon.s` (generated by: `gcc -S marathon.c`)

The code above uses a few arithmetic operators. Many operators correspond with the ones you've seen in algebra. Here are a few *binary* operators:

Operator	Meaning
+	Add
-	Subtract
*	Multiply
/	Divide
%	Modulo (Remainder)

A full [list of operators](#) including their *precedence* and *associativity*.

Computer Data Storage (Refresher)

This section is just a short refresher on the binary, decimal, and hexadecimal number systems.

- Computers represent information as patterns of bits (0's and 1's), known as base 2; e.g. 1011101000010101000111100011
- People use base 10; doesn't translate well into base 2
- Hexadecimal (base 16) works well with binary; computer programmers often work in base 16
- Group the bits into sets of 4 and translate into hexadecimal using the chart below

Relationship between hex and binary numbers:

0	0000	4	0100	8	1000	C	1100
1	0001	5	0101	9	1001	D	1101
2	0010	6	0110	A	1010	E	1110
3	0011	7	0111	B	1011	F	1111

Character representations

The number above, 1011101000010101000111100011, translates into hex (BA151E3) and decimal (195,121,635) as:

Binary	1011	1010	0001	0101	0001	1110	0011
Hexadecimal	B	A	1	5	1	E	3
Decimal	195121635						

More information on [Binary numbers](#).

[Binary/Decimal converter](#) (BinConverter.exe)

Lexical Conventions

C programs are typically stored in one or more files on the disk. These files are given a rather fancy name: *translation units*. After the pre-processor has removed the directives and performed the appropriate action, the compiler starts its job. The first thing the compiler needs to do is to *parse* through all of the *tokens* in the files.

There are different classes of tokens (lexical elements). In no particular order they are:

1. keywords
2. identifiers
3. constants
 - integers
 - characters
 - floating point
 - enumerations
 - string literals
4. operators (Lots!)
5. punctuators and separators

The standard actually names these 6:

1. keywords
2. identifiers
3. constants
4. string literals
5. operators
6. punctuators

White space includes things like blank spaces, tab, newlines, etc. Comments are a form of whitespace since they are stripped out and replaced by a single space.

We will spend the entire course studying these aspects of the C language.

Identifiers

- Variables, constants, function names, and other elements of a C program are called *identifiers*; they are used to name things.
- Identifiers must contain only letters (upper- or lower-case), digits (0-9), and underscores (_).
- They must begin with a letter or underscore. (Not a digit)
- There are certain words in C that have special meanings (keywords) and can't be used as identifiers. (e.g. **int**, **float**, **const**)
- C is a case sensitive language; upper- and lower-case letters are considered to be different. (e.g. SUM, Sum, sum are all different identifiers)
- Most compilers allow identifiers to be 31 characters or more in length. (You probably won't need more than that.)
- Use meaningful names, you will be glad you did (and your grade in this course will reflect it!)

Some examples:

Valid	Invalid	Invalid Reason
foo1	1foo	Doesn't start with a letter or underscore
_foo	\$foo	1. Doesn't start with a letter or underscore 2. \$ is illegal character
foo	foo\$	\$ is an illegal character
void_identifiier	invalid-identifier	- is an illegal character
a_long_and_valid_name	foo bar	Can't have spaces in identifier names
Int	int	int is a keyword

Examples:

Good Identifier Names	Bad Identifier Names
<pre>int rate; int time; int distance; rate = 60;</pre>	<pre>int x; int y; int z; x = 60;</pre>

time = 20; distance = rate * time;	y = 20; z = x * y;
#define PI 3.1416F float radius = 5.25F; float volume; volume = 4.0 / 3.0 * PI * radius * radius * radius;	#define A 3.1416F float id1 = 5.25F; float id2; id2 = 4.0 / 3.0 * A * id1 * id1 * id1;
double base = 2.75, height = 4.8; double area_of_triangle = 0.5 * base * height;	double table = 2.75, chair = 4.8; double couch = 0.5 * table * chair;

Keywords

Keywords are identifiers that are reserved for the compiler. You can't use any of these as identifiers:

auto	const	double	float	int	short	struct	unsigned
break	continue	else	for	long	signed	switch	void
case	default	enum	goto	register	sizeof	typedef	volatile
char	do	extern	if	return	static	union	while

Remember that C is *case-sensitive* so these keywords must be typed exactly as shown. `int` is not the same as `Int` or `INT`.

Constants

A literal value is a constant just as you type it in the code:

```
int a = 1;
float f = 3.14F;
double d = 23.245;
```

Examples:

Constant	Type
5	int
3.14	double
3.14F	float
'A'	int
"hello"	string

The `marathon.c` program has 6 literal values. There are 4 integers, 1 double, and 1 string.

Algorithms

A large part of your programming career will deal with *algorithms*.

- One of the most fundamental concepts in computer programming
- A set a steps that defines how a task is accomplished; a recipe
- A computer program represents an algorithm
- Programs = software
- Algorithms have been known and studied for centuries; they were not invented by computer scientists
- Algorithms *do not* require computers, although computers require algorithms
- There are many famous algorithms; Euclid's Greatest Common Divisor (GCD) is a simple one (e.g. the GCD of 27 and 12 is 3, the GCD of 68 and 12 is 4)

Euclid's algorithm in English (with some algebra thrown in):

Step	Actions to be Performed
1	Assign the larger number to M, and the smaller number to N.
2	Divide M by N (M/N) and assign the remainder to R.
3	If R is not 0, then assign the value of N to M, assign the value of R to N, and return to step 2. If R = 0, then the GCD is N and the algorithm terminates.

Notice that in Step 3 there are two different possibilities. This is typically how algorithms work. There usually needs to be some *terminating condition*, otherwise the algorithm (program) runs forever.

Here is how we might write the algorithm in *pseudo-code*:

1. Assign larger value to M
2. Assign smaller value to N
3. Divide M by N and assign remainder to R
4. While remainder, R, is not 0
 1. Assign N to M
 2. Assign R to N
 3. Divide M by N and assign remainder to R
5. End While
6. The algorithm has terminated and the GCD is N

Coding this algorithm in an assembler language might look like the code below. This version is using memory locations that are named `M` and `N` for easier understanding.

```
.section .data
M: .long 45      # put 45 in location named M
N: .long 12      # put 12 in location named N

.section .text
.globl _start

_start:
    movl M, %eax    # put M in %eax
    movl N, %ecx    # put N in %ecx
    movl $0, %edx   # zero out for idivl
    idivl %ecx      # divide M/N, (%edx:%eax/%ecx)
                    # result -> %eax, remainder -> %edx

start_loop:
    cmpl $0, %edx   # is remainder 0?
    je loop_exit    # if 0, we're done

    movl %ecx, %eax  # put N in M
    movl %edx, %ecx  # put R in N
    movl $0, %edx   # zero out for idivl
    idivl %ecx      # divide M/N, (%edx:%eax/%ecx)
                    # result -> %eax, remainder -> %edx
    jmp start_loop  # check again
loop_exit:
```

This second version doesn't use any memory locations (like `M` and `N` above). All values are stored directly in the registers on the CPU.

```
.section .data
.section .text
.globl _start

_start:
    movl $45, %eax  # put 45 in eax (M)
    movl $0, %edx   # set to 0 (high word of divisor)
    movl $12, %ecx  # put 12 in ecx (N)
    idivl %ecx      # divide M/N, (%edx:%eax/%ecx)
```

```
# result -> %eax, remainder -> %edx
```

```
start_loop:
    cml $0, %edx    # Is remainder 0?
    je loop_exit    # if 0, we're done

    movl %ecx, %eax  # put N into M
    movl %edx, %ecx  # put R into N
    movl $0, %edx    # set high word
    idivl %ecx       # divide M/N, (%edx:%eax/%ecx)
                    # result -> %eax, remainder -> %edx
    jmp start_loop   # continue algorithm
loop_exit:
```

Euclid's GCD algorithm as high-level computer programs (assuming that M and N already have values and that $M > N$)

C	Pascal	BASIC
<pre>r = m % n; while (r != 0) { m = n; n = r; r = m % n; }</pre>	<pre>r := m Mod n; While r <> 0 Do Begin m := n; n := r; r := m Mod n; End;</pre>	<pre>r = m MOD n WHILE r <> 0 m = n n = r r = m MOD n WEND</pre>

The algorithm reprinted:

Step	Actions to be Performed
1	Assign the larger number to M, and the smaller number to N.
2	Divide M by N (M/N) and assign the remainder to R.
3	If R is not 0, then assign the value of N to M, assign the value of R to N, and return to step 2. If R = 0, then the GCD is N and the algorithm terminates.