

# Pointers

## Pointer Basics

- Memory is a vast collection of *bits*, either 0 or 1.
- Bits are grouped into 8-bit blocks called a *bytes*. (8 is arbitrary but most modern computers do this.)
- Bytes are grouped into *words* depending on the *wordsize* of the computer:
  - A 16-bit machine has 2 bytes per word
  - A 32-bit machine has 4 bytes per word
  - A 64-bit machine has 8 bytes per word
- Every byte in a computer's memory is uniquely identified by it's location in memory.
- Each location is a *cell* and has two attributes: an *address* (the cell's location in memory) and a *value* (contents of cell)
- Addresses and contents are known as *l-values* and *r-values*, respectively.
- Note that even though a word is made up of multiple bytes, it only has a single address.
- All data in memory has an address from 0 up to the size of memory. 2 GB of memory would look like this (the contents are arbitrary)

□

- We can find out the address of a variable by simply using the *address operator* (which we've already seen with `scanf`)

```
int i = 10;

printf("The address of i is %p\n", &i); /* Use & to get the address */
printf("The contents of i is %i\n", i);
```

Output:

```
The address of i is 0x22cca4
The contents of i is 10
```

On computers where integers and addresses have the same range, it's possible to interchange them. However, in general, you can't do this.

- Addresses *look* like integers, but they are different. (e.g. just like 3.0 is not the same as 3.0F or 3)

[illegible]

- The compiler will complain if you treat an address like an integer and vice-versa.
- The type of the address depends on what is stored there. (e.g. the address of an integer is not the same type as an address of a float)
- Up until now, all of the addresses the compiler has been using for data have contained integer, float, and double values.
- An address that holds the address of another variable is called a *pointer*.

We can declare pointer variables easily:

```
void foo(void)
{
    int i; /* i can only store integer values */
           /* The value of i is undefined at this point */

    int *p; /* p can only store the address of an integer */
            /* The value of p is undefined at this point */

    p = &i; /* The value of p is now the address of i */
           /* The value of p is now defined */

    i = 10; /* The value of i is now 10 */
           /* The value of i is now defined */
}
```

This is the notation that will be used when talking about identifiers in memory:

□

- **identifier** - The name of the identifier
- **address** - The arbitrary address (the actual values are meaningless, but are useful for discussion purposes)
- **contents** - The value stored at this location. ??? means it is undefined.
- When looking at arrays, since each cell is anonymous (unnamed), I'll only show the address above the cell.

Visualizing the code above:

After declarations for i and p	After assignment to p	After assignment to i

□

□

□

**Note:** Be sure not to confuse the asterisk ( \* ) used for multiplication with the asterisk used to declare pointers. They are the

same token, but have very different meanings.

Other examples:

```
int *pi;    /* pi can only "point" to integers */
float *pf;  /* pf can only "point" to floats  */
double *pd; /* pd can only "point" to doubles */
short *ps;  /* ps can only "point" to shorts  */
long *pL;   /* pL can only "point" to longs   */

unsigned short *pus; /* pus can only "point" to unsigned shorts */
unsigned int *pui;   /* pui can only "point" to unsigned ints  */

/* etc.... you get the idea */
```

Unfortunately, we're not done abusing the asterisk ( \* ). It has a third meaning as the *indirection operator* also known as the *dereference operator*. [Operator precedence chart](#)

```
void foo(void)
{
    int i = 10;
    int *p = &i;

    printf("The value of i is %i\n", i);
    printf("The address of i is %p\n", &i);

    printf("The value of p is %p\n", p);
    printf("The address of p is %p\n", &p);
    printf("The value of what p is pointing at is %i\n", *p);
}
```

**Output:**

```
The value of i is 10
The address of i is 0x22cca4
The value of p is 0x22cca4
The address of p is 0x22cca0
The value of what p is pointing at is 10
```

Once we have a pointer to some data, we can read/write the data *through* the pointer:

```
void f5(void)
{
    int i = 10; /* i has the value 10 */
    int *p = &i; /* p has the address of i */
    int j;      /* j has undefined value */

    printf("The value of i is %i\n", i);
    printf("The value of *p is %i\n", *p);

    j = *p; /* j has the value of i */
    *p = 20; /* i has the value 20 */

    printf("The value of j is %i\n", j);
    printf("The value of i is %i\n", i);
}
```

**Output:**

```
The value of i is 10
The value of *p is 10
The value of j is 10
The value of i is 20
```

Visually:

After declarations for i, p, j    After j = \*p    After \*p = 20

□

□

□

Watch out for these illegal expressions:

```
void foo(void)
{
    int i = 10; /* i has the value 10 */
    int *p = &i; /* p has the address of i */
    int j;      /* j has undefined value */

    j = p; /* ILLEGAL: different types */
    p = 20; /* ILLEGAL: different type */
}
```

Also, watch out for this as well. The compiler may not even warn you about it.

```
void f7(void)
{
    int *p; /* p has undefined value (some random memory location) */
    *p = 10; /* BAD: writing the value 10 to some random memory location */
}
```

You can have multiple pointers to the same address. The code below gives us three different ways to modify the contents of `i`:

```
void foo(void)
{
    int i = 10; /* i gets value 10 */
    int *p;     /* undefined */
    int *q;     /* undefined */

    p = &i; /* p gets address of i */
    q = p;  /* q gets value of p (address of i) */

    *p = 20; /* i gets value 20 (*p and *q evaluate to 20) */
    *q = 30; /* i gets value 30 (*p and *q evaluate to 30) */
}
```

Visually:

After declarations for `p, q, i`    After `p = &i`    After `q = p`

□

□

□

After `*p = 20`    After `*q = 30`

□

□

Given the code below, which of the assignment statements are illegal and why? What does the compiler say about the illegal statements?

```
int i;
int *pi;

pi = i; /* 1. ??? */
pi = 5; /* 2. ??? */
*pi = i; /* 3. ??? */
*pi = 5; /* 4. ??? */
*pi = &i; /* 5. ??? */
pi = &i; /* 6. ??? */
```

Given these declarations:

```
/* Assume addresses of variables are */
/* a:100, b:104, c:108, d:112, e:116 */
int a = 10;
int b = 108;
float c = 3.14F;
int *d = &a;
float *e = &c;
```

Diagram: (looks like an array, but it's not)

□

Showing the pointer arrows:

□

We can evaluate the expressions as such:

Expression	Type	Value
a	int	10
b	<b>int</b>	<b>108</b>
c	float	3.14
d	int *	100
e	<b>float *</b>	<b>108</b>
*d	int	10
*e	float	3.14
&a	int *	100
&d	int * *	112
*a	illegal	----
*b	illegal	----
*c	illegal	----

## Passing Pointers to Functions

A major benefit of using pointers is realized when we need another function to modify a value. The simple example is `scanf`:

```
int a, b, c;
scanf("%d%d%d", &a, &b, &c);
```

We want `scanf` to modify our local variables, but `scanf` is unable to access them (local scope). So, we need to pass them to `scanf`. We need to pass the addresses of the local variables. If we just passed the values to `scanf`, any changes made to them would be lost after returning from the function.

Suppose we want to write a function that exchanges the values of two integers. Here's our first attempt:

```
int main(void)
{
    int x = 10;
    int y = 20;

    printf("Before: x = %i, y = %i\n", x, y);
    swap(x, y);
    printf(" After: x = %i, y = %i\n", x, y);

    return 0;
}

/* Exchanges the values of the parameters */
void swap(int a, int b)
{
    int temp = a; /* Save a for later */
    a = b;        /* a gets value of b */
    b = temp;     /* b gets old value of a */
}

Output:
Before swap: x = 10, y = 20
After swap: x = 10, y = 20
```

What went wrong?

Visually:

**Initialize values, pass to swap    Make copy of a    Exchange a,b**



Try it again, the correct way:

```
/* Exchanges the values of the parameters */
void swap(int *a, int *b)
{
    int temp = *a; /* Save a for later */
    *a = *b;       /* a gets value of b */
    *b = temp;     /* b gets old value of a */
}

int main(void)
{
    int x = 10;
    int y = 20;

    printf("Before swap: x = %i, y = %i\n", x, y);
    swap(&x, &y);
    printf(" After swap: x = %i, y = %i\n", x, y);

    return 0;
}

Output:
Before swap: x = 10, y = 20
After swap: x = 20, y = 10
```

Visually:

**In main before swap    Call swap, pass addresses    Make copy of a for later**



**Copy \*b to \*a (y to x in main)    Copy temp to \*b (temp to y in main)    After swap**



Notes to remember:

- The correct swap function takes *pointers to integers* as parameters:

```
void swap(int *a, int *b); /* a and b are pointers to integers */
```

- You must pass the address of the integers you wish to swap:

```

int x = 10;
int y = 20;
swap(&x, &y); /* Must pass the addresses (pointers) of a and b */

```

---

Protecting data passed to a function:

This works as expected:

```

int main(void)
{
    int a[] = {4, 5, 3, 9, 5, 2, 7, 6};
    int largest;

    printf("Array before:\n");
    print_array(a, 8);

    largest = find_largest(a, 8);
    printf("Largest value: %i\n", largest);

    printf("Array after:\n");
    print_array(a, 8);

    return 0;
}

/* Assumes there is at least */
/* one element in the array */
int find_largest(int a[], int size)
{
    int i;
    int max = a[0]; /* assume 1st is largest */

    for (i = 1; i < size; i++)
        if (a[i] > max)
            max = a[i]; /* found a larger one */

    return max;
}

```

Output:

```

Array before:
4 5 3 9 5 2 7 6
Largest value: 9
Array after:
4 5 3 9 5 2 7 6

```

Let's modify the function to do something unexpected:

```

/* Modifies the array the was passed in!! */
int find_largest_BAD(int a[], int size)
{
    int i;
    int max = a[0]; /* assume 1st is largest */
    a[0] = 0; /* change first element! */
    for (i = 1; i < size; i++)
    {
        if (a[i] > max)
            max = a[i]; /* found a larger one */
        a[i] = 0; /* set element to 0!! */
    }
    return max;
}

```

Output:

```

Array before:
4 5 3 9 5 2 7 6
Largest value: 9
Array after:
0 0 0 0 0 0 0 0

```

We need to make the array a *constant*, so we modify the declaration in the function:

```

/* Unable to modify the array since it's const */
int find_largest_BAD(const int a[], int size)
{
    int i;
    int max = a[0]; /* assume 1st is largest */
    a[0] = 0; /* ILLEGAL: elements are const */
    for (i = 1; i < size; i++)
    {
        if (a[i] > max)
            max = a[i]; /* found a larger one */
        a[i] = 0; /* ILLEGAL: elements are const */
    }
    return max;
}

```

Compiler errors:

```

main.c: In function `find_largest_BAD':
main.c:161: error: assignment of read-only location
main.c:166: error: assignment of read-only location

```

Going back to the original non-const version of this function:

```

int find_largest(int a[], int size);

```

If we really don't want our array to change, we declare it with the **const** keyword:

```

int main(void)
{

```

```

const int a[] = {4, 5, 3, 9, 5, 2, 7, 6}; /* Elements are constant (can't be changed) */
int largest = find_largest(a, 8);        /* ILLEGAL: Function expects non-const array */

return 0;
}

```

Compiler error:

```
main.c:185: warning: passing arg 1 of `find_largest' discards qualifiers from pointer target type
```

**Important:** When you create a function that will accept arrays as parameters, be sure to mark them as `const` if you do not intend to modify them. If you don't make them `const`, a lot of code will not be able to use your function.

The `const` keyword can be used to protect the pointer as well as what's being pointed at (the pointee?).

---

Never return a pointer to a local variable. Ever. In other words, don't do this:

```

/* Function returns a pointer to an int */
int *foo(void)
{
    int i = 10; /* i is local (on the stack) */
    return &i; /* This will be undefined!! */
}

```

Fortunately, most compilers today will point out the sheer insanity of your code:

```
main.c: In function `foo':
main.c:174: warning: function returns address of local variable
```

More on `const`

Sometimes you want to make sure that you don't accidentally change values in your programs. The safest way to do this is to mark the identifiers with the `const` keyword.

With pointers, you have more flexibility with `const`. You can make the pointer itself constant, which means once it points at something, it can never be changed to point at something else. Or, you can make the data pointed at constant, which means that although you can change the pointer to point at something else, you can't change what's being pointed at.

Here are the four cases:

1. Neither the pointer nor the data being pointed at (the pointee) is `const`. Both can be changed:

```

int *pi; /* pi is a (non-const) pointer to a (non-const) int */

```

2. The pointer is not `const`, but the data pointed at (the pointee) is `const`. The data is protected. Only the pointer can change:

```

const int *pci; /* pci is a pointer to a const int */

```

3. The pointer is `const` but the data being pointed at is non-const. The pointer is protected. The data can be changed through the pointer.

```

int * const cpi = &i; /* cpi is a const pointer to an int */

```

4. Both the pointer and the data being pointed at are `const`. Both are protected. Neither can be changed:

```

const int * const cpci = &ci; /* cpci is a const pointer to a const int */

```

Points to remember:

- If you mark something as `const`, you are indicating that it should not change.
- If you DO NOT mark something as `const`, you are indicating that it should change.
- Therefore, if you are not going to change some data in your program, make sure to mark it `const`.
- Note that with function parameters, these rules really only apply to pointers (addresses) and arrays. Why?

Here is an example that shows the `const` keyword in action in various ways:

```

void foo(void)
{
    int i = 5;          /* i is a non-constant int */
    int j = 6;          /* j is a non-constant int */
    const int ci = 10;   /* ci is a constant int */
    const int cj = 11;   /* cj is a constant int */

    int *pi;            /* pi is a pointer to an int */
}

```

```

const int *pci;           /* pci is a pointer to a const int */
int * const cpi = &i;     /* cpi is a const pointer to an int */
const int * const cpci = &ci; /* cpci is a const pointer to a const int */

i = 6;    /* Ok, i is not const */
j = 7;    /* Ok, j is not const */
ci = 8;   /* ERROR: ci is const */
cj = 9;   /* ERROR: cj is const */

pi = &i;   /* Ok, pi is not const */
*pi = 8;   /* Ok, *pi is not const */
pi = &j;   /* Ok, pi is not const */
*pi = 9;   /* Ok, *pi is not const */

pci = &ci; /* Ok, pci is not const */
*pci = 8;  /* ERROR: *pci is const */
pci = &cj; /* Ok, pci is not const */
*pci = 9;  /* ERROR: *pci is const */

cpi = &j;  /* ERROR: cpi is const */
*cpi = 10; /* Ok, *cpi is not const */
*cpi = 11; /* Ok, *cpi is not const */

cpci = &j; /* ERROR: cpci is const */
*cpci = 10; /* ERROR: *cpci is const */

pi = &ci; /* DANGER: constant ci can be changed through pi */
cpi = &ci; /* DANGER: constant ci can be changed through cpi */
}

```

## Short Array Review

Suppose we had an array of 8 chars (8-bits per char) and assigned a value to each char:

```

unsigned char bytes[8];

bytes[0] = 'J'; /* 0x4A */
bytes[1] = 'o'; /* 0x6F */
bytes[2] = 'h'; /* 0x68 */
bytes[3] = 'n'; /* 0x6E */
bytes[4] = 205; /* 0xCD */
bytes[5] = 204; /* 0xCC */
bytes[6] = 76; /* 0x4C */
bytes[7] = 62; /* 0x3E */

```

Remember that we can initialize the array as such:

```

unsigned char bytes[] = {'J', 'o', 'h', 'n', 205, 204, 76, 62};

```

If we assume that the address of the **bytes** variable is 1000, the raw bits of the array would look like this in memory:

□

Each of the 8 chars is uniquely identified by it's address:

Address (l-value) (shown in decimal)	Contents (r-value) (shown in binary)
&bytes[0] ==> 1000	bytes[0] ==> 01001010
&bytes[1] ==> 1001	bytes[1] ==> 01101111
&bytes[2] ==> 1002	bytes[2] ==> 01101000
&bytes[3] ==> 1003	bytes[3] ==> 01101110
&bytes[4] ==> 1004	bytes[4] ==> 11001101
&bytes[5] ==> 1005	bytes[5] ==> 11001100
&bytes[6] ==> 1006	bytes[6] ==> 01001100
&bytes[7] ==> 1007	bytes[7] ==> 00111110

We can draw the diagram any way that we like, the computer doesn't care.

As single printable characters:

□

As integer equivalents of the characters (ASCII values):

□

## Pointer Arithmetic

You can perform limited arithmetic with pointers:

1. pointer + integer (result is a pointer)
2. pointer - integer (result is a pointer)
3. pointer - pointer (result is a signed integer)

## When adding/subtracting integers:

The integers that are added to pointers are **scaled** by the size of the type of pointer. The sizes of `*p` assume a 32-bit machine:

Expression	Type of p	Size of *p	Value added
p + 1	char *	sizeof(char) == 1	1
p + 1	short *	sizeof(short) == 2	2
p + 1	int *	sizeof(int) == 4	4
p + 1	float *	sizeof(float) == 4	4
p + 1	double *	sizeof(double) == 8	8
p + 1	Foo *	sizeof(Foo) == X	X
p + 2	char *	sizeof(char) == 1	2
p + 2	short *	sizeof(short) == 2	4
p + 2	int *	sizeof(int) == 4	8
p + 2	float *	sizeof(float) == 4	8
p + 2	double *	sizeof(double) == 8	16
p + 2	Foo *	sizeof(Foo) == X	2 * X
p + 3	char *	sizeof(char) == 1	3
p + 3	short *	sizeof(short) == 2	6
p + 3	int *	sizeof(int) == 4	12
p + 3	float *	sizeof(float) == 4	12
p + 3	double *	sizeof(double) == 8	24
p + 3	Foo *	sizeof(Foo) == X	3 * X

In the table above, **Foo** is some user-defined object that requires **X** bytes in memory. The compiler will know about **Foo** and will scale the value by the appropriate amount. For example, if **sizeof(Foo)** is 20, then  $2 * X$  would be 40 and  $3 * X$  would be 60.

## Using Pointers with Arrays

We can use pointers with arrays instead of subscripts. Given this array and integer pointer:

```
int a[5];
int *pi;
```

We have this layout:

□

Initializing each element of an integer array to 0:

```
for (pi = &a[0]; pi < &a[5]; pi++)
    *pi = 0;
    □
    After the first iteration
```

We can combine the increment and dereference operators. All of these loops do the same thing. You may want to refer to [the precedence chart](#).

```
for (pi = &a[0]; pi < &a[5];)
    *pi++ = 0;
```

or:

```
pi = &a[0];
while (pi < &a[5])
    *pi++ = 0;
```

or

```
for (pi = &a[5]; pi > a;)
    *--pi = 0;
```

## Subtracting pointers

- Subtracting two pointers (finding the difference) is only valid when both pointers point into the same array.
- The difference is in elements, not bytes, so:

```
int a[5];

int *pi1 = &a[1];
int *pi2 = &a[4];
int diff = pi2 - pi1; /* diff is 3 */
diff = pi1 - pi2;    /* diff is -3 */
```

Visually:

□

## Comparing Pointers



- You can use the operators: <, <=, >, >= with pointers.
- Again, both pointers must point into the same array if the result is to have any meaning.
- You can use the equality operators, ==, != with arbitrary pointers.

Notes on using pointers with arrays:

- The results of pointer arithmetic are only defined if the pointer is pointing at an element of an array.
- Pointing at an element *before* the start of the array is undefined (as per the Standard).
- Pointing at an element one past the end of the array is legal, but dereferencing it is not.

Pointers vs. Arrays

The short Q and A:

**Q:** Are pointers the same as arrays?

**A:** No. Never. Ever. Any questions? (Read all about the [historical accident or mistake](#) that confuses beginning programmers.)

Moving on...

□

Simple points:

- The address of the array, `a`, is 100.
- The address of the first element, `a[0]`, in the array is also 100.
- The type of `a` is *array of 5 ints*.
- The type of `a[0]` is *int*.
- The type of `&a` is *pointer to array of 5 ints*.
- The type of `&a[0]` is *pointer to int*.

Suppose we pass this array to some function:

#### Call the function

```
foo(a, 5);
```

#### The function definition

```
void foo(int x[], int elements)
{
    /* do whatever... */
}
```

We know we can't pass an array by value, we can only pass an address. (The address 100 in this example.) But which "100"? The *array* or the *int*?

**Important note:** When passing an array to a function, the compiler passes a *pointer to the first element* (i.e. the address of the first element) to the function.

Understanding the above statement is critical. *Read it. Learn it. Live it.*

This means that these two function prototypes have the exact same meaning:

#### Using array notation

```
void foo(int x[], int elements);
```

#### Using pointer notation

```
void foo(int *x, int elements);
```

Two functions that print an array of integers: (all functions do the same exact thing)

#### Using array notation

```
void print_array1(int x[], int elements)
{
    int i;
    for (i = 0; i < elements; i++)
        printf("%i ", x[i]);
}

void print_array1(int *x, int elements)
{
    int i;
    for (i = 0; i < elements; i++)
        printf("%i ", x[i]);
}
```

#### Using pointer notation

```
void print_array2(int x[], int elements)
{
    int i;
    for (i = 0; i < elements; i++)
        printf("%i ", *x++);
}

void print_array2(int *x, int elements)
{
    int i;
    for (i = 0; i < elements; i++)
        printf("%i ", *x++);
}
```

Of course, we'd really want to make sure not to modify the array elements:

### Using array notation

```
void print_array1(const int x[], int elements)
{
    int i;
    for (i = 0; i < elements; i++)
        printf("%i ", x[i]);
}
```

### Using pointer notation

```
void print_array2(const int *x, int elements)
{
    int i;
    for (i = 0; i < elements; i++)
        printf("%i ", *x++);
}
```

This means that, when dealing with arrays, we can use subscript notation or pointer notation.

The Basic Rule:

`array[index]`

is the same as

`*(array + index)`

where:

- *array* is an array of any type
- *index* is any integer expression

This means:

- Referencing an element of an array is done using a pointer and an offset.
- The *pointer* is usually the address of the start of the array (which is the address of the first element).
- The *offset* is scaled by the size of the element type. (Pointer arithmetic)
- The compiler converts all array references to a pointer/offset, so:

```
void f(void)
{
    int array[] = {5, 10, 15, 20, 25};
    int index = 3;

    printf("%i\n", array[index]); /* 20 */
}
```

According to the Basic Rule:

`array[index] ==> *(array + index)`

Don't forget that due to C's built-in pointer arithmetic, the addition (e.g. `array + 3`) is scaled:

`array[3] ==> *(array + 3 * sizeof(int)) ==> *(array + 12 bytes)`

The above means: *the fourth element is 12 bytes from the address of array*. You can see this from looking at a similar diagram:

□

## Relationship Between Subscripts and Pointers

Using the rule to convert a subscript to a pointer/offset, we get:

```
a[i] ==> *(a + i)
&a[i] ==> &*(a + i)
&a[i] ==> &*(a + i)
&a[i] ==> a + i
```

This shows that the address of any element is just the base address of the array plus the index (scaled).

```
char a[] = "abcdef";
char *p = a;

printf("%p, %p, %p, %p, %p\n", a, a + 2, &*(a + 2), p + 2, &*(p + 2));
```

Output:

`0012FED4, 0012FED6, 0012FED6, 0012FED6, 0012FED6`

Other equivalences:

```
a[i] ==> *(a + i)
a[0] ==> *(a + 0)
a[0] ==> *a
&a[0] ==> &*a
&a[0] ==> a
```

These calls are equivalent:

```
f(*a);    /* pass first element of array */  
f(a[0]); /* pass first element of array */
```

and so are these:

```
f(&a[0]); /* pass address of first element */  
f(a);    /* pass address of first element */
```

---

## Pointer Expressions and Arrays

Given this code:

```
int a[10] = {5, 8, 3, 2, 1, 9, 0, 4, 7, 6};  
int *p = a + 2;
```

Abstract diagram:

□

or shown with concrete values (addresses are arbitrary):

□

Give the equivalent expression using **a**. (Hint: Determine the type of each expression first)

1. p
  2. p[0]
  3. \*p
  4. p + 3
  5. \*p + 5
  6. \*(p + 6)
  7. p[6]
  8. &p
  9. p[-1]
  10. p[9]
-