

# HARDWARE SUPPORT FOR OS

Instructor: William Zheng

Email:

[william.zheng@digipen.edu](mailto:william.zheng@digipen.edu)

PHONE EXT: 1745


# Goals for this lecture

- HW and SW
- CPU - I/O Communication
  - Interrupt handling
  - DMA
  - More details in chapter 13.1-13.2 of textbook
- CPU modes
  - Backward-compatibility modes
  - Privilege modes

# HW & SW


- Difference?
- Important to know the distinction
  - OS (SW)
  - But what OS can do
    - Limited by the hardware it's running on

# Computer System Components

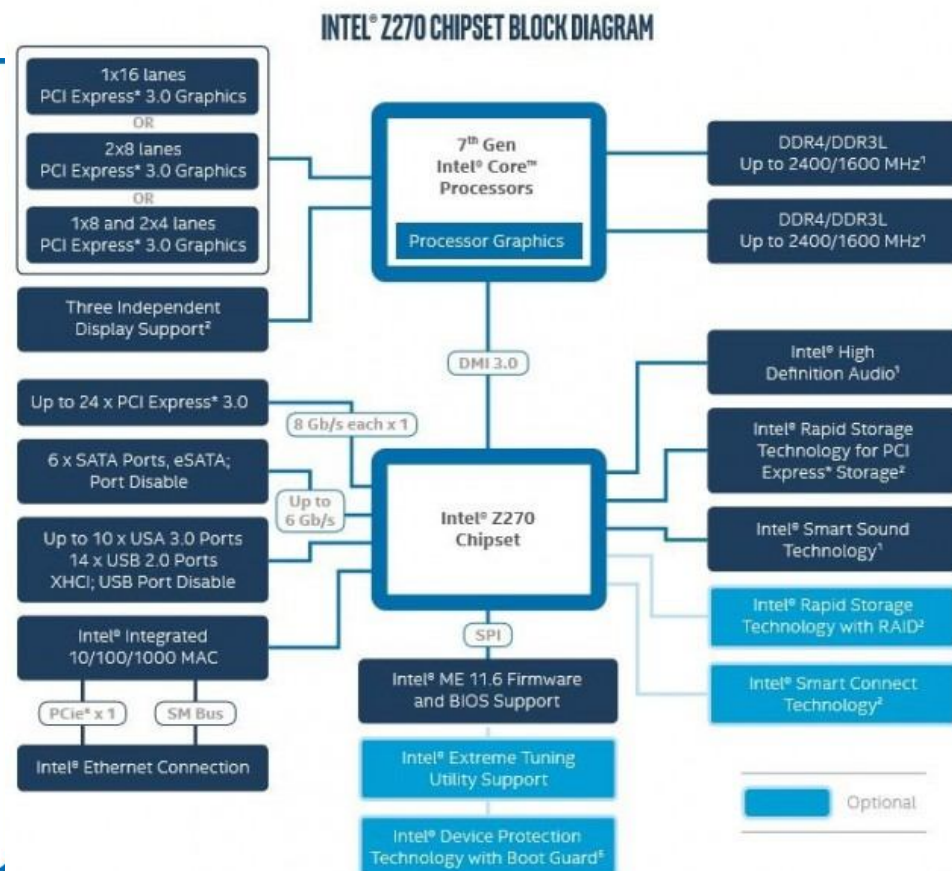
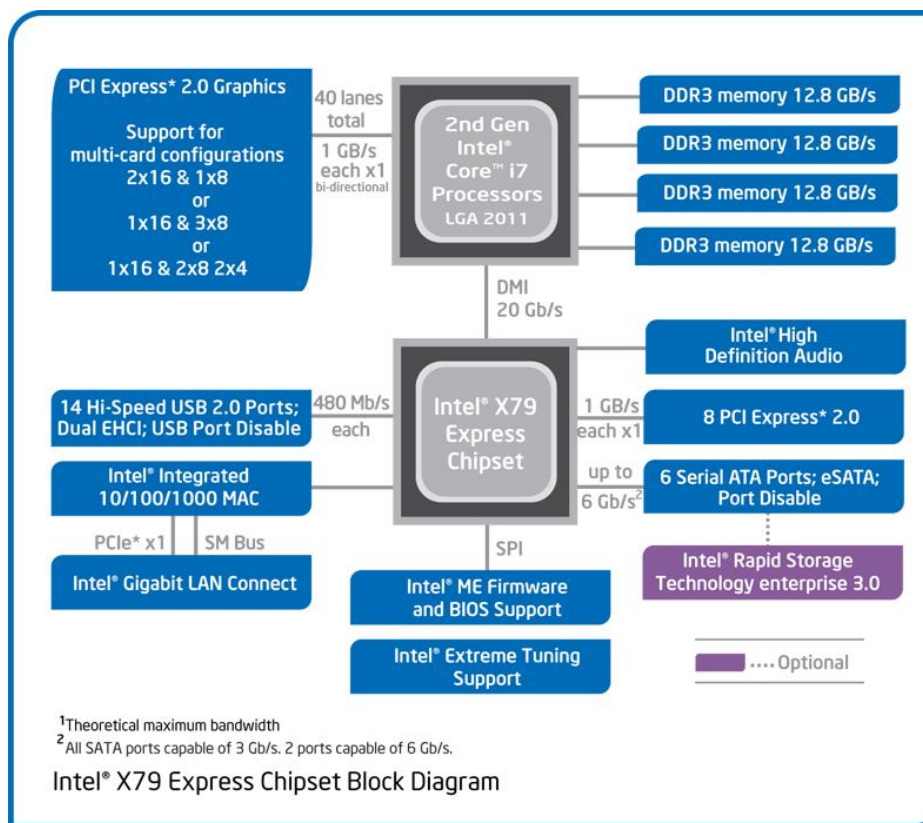


Main Components of a Computer System

# Motherboard

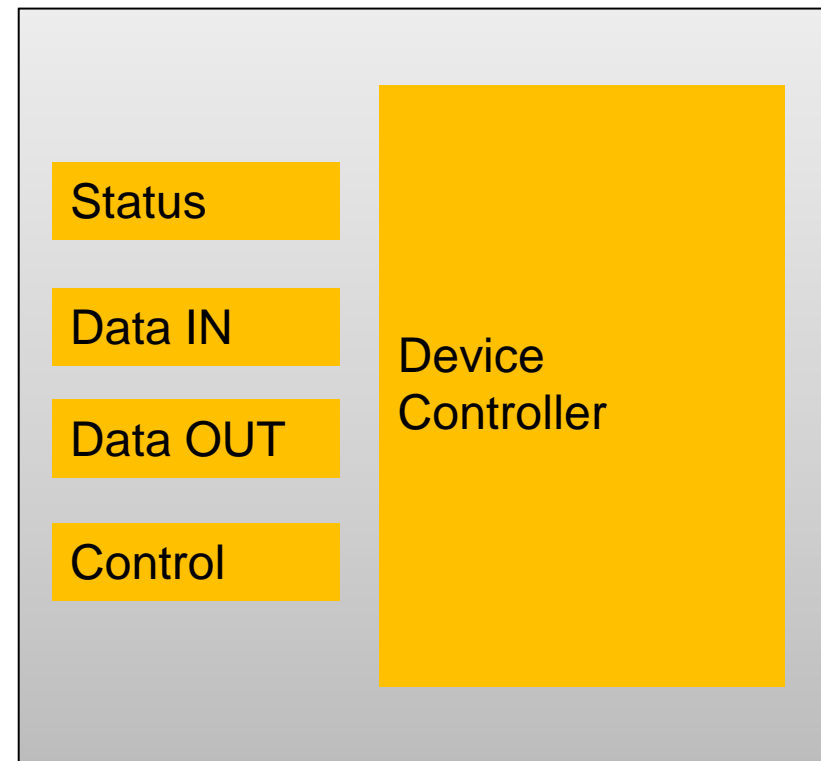


# Chipset Block Diagram



# Device Control Registers

- Status register
  - Done bit
  - Error bit
- Data register
  - Data to be printed
- Control registers
  - Commands



# CPU – I/O Communication

- From CPU to I/O
  - Special instructions
  - Memory-mapped I/O
    - Device data and command registers mapped to processor address space
    - Especially for large address spaces (graphics)
- From I/O to CPU
  - Polling
  - Interrupt

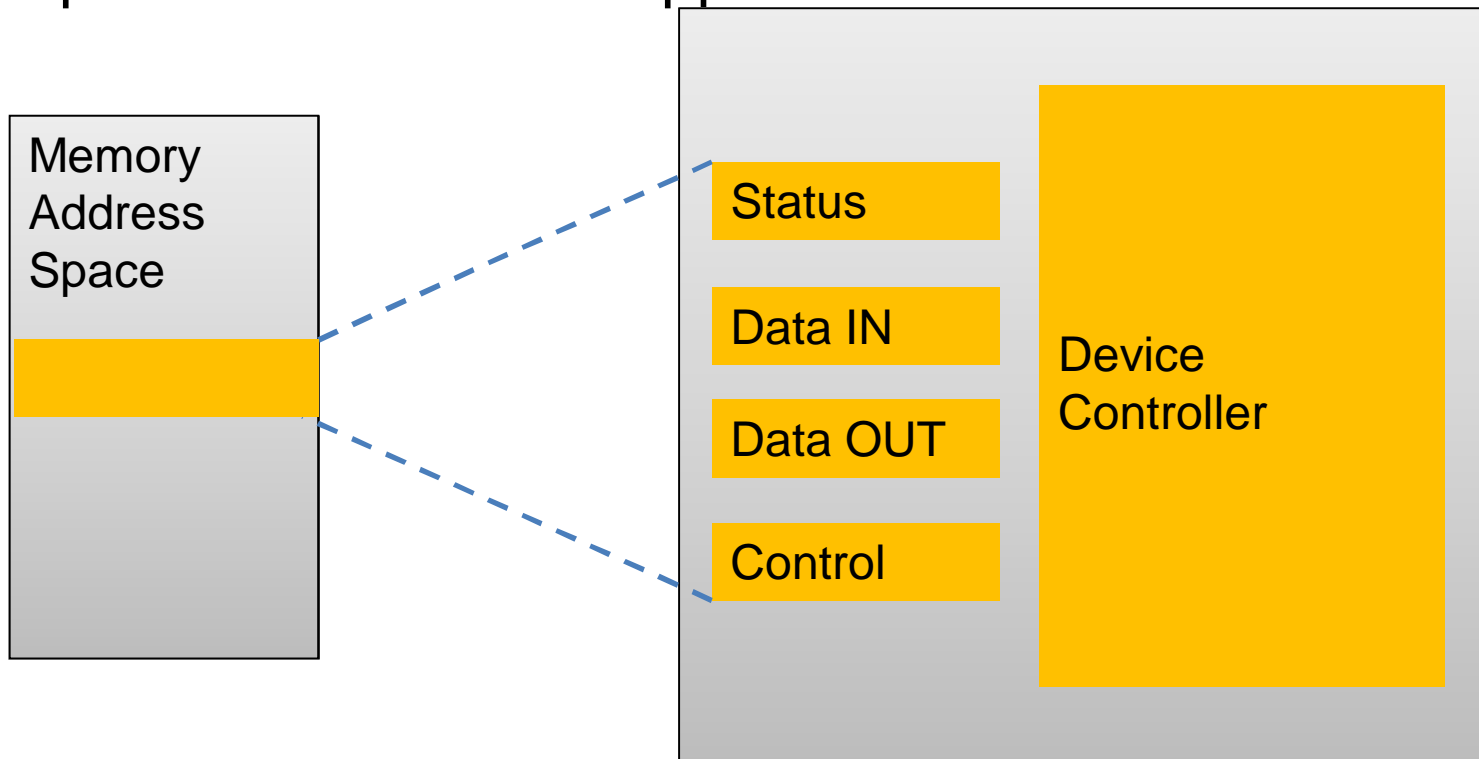


# Special I/O Instructions

- Instruction specially for I/O
  - I/O ports
    - Different address spaces for memory and I/O devices
  - E.g. Intel IN, OUT instructions
    - Register I/O instructions: IN/OUT move data between I/O parts and the EAX (32-bit I/O), AX (16-bit I/O) or AL (8-bit I/O) general registers.
    - Block I/O instructions: INS/OUTS move blocks of data between I/O ports and memory space

# Memory-mapped I/O - I

- Use memory read/write
  - E.g. intel MOV instructions
  - Specific addresses mapped to I/O devices

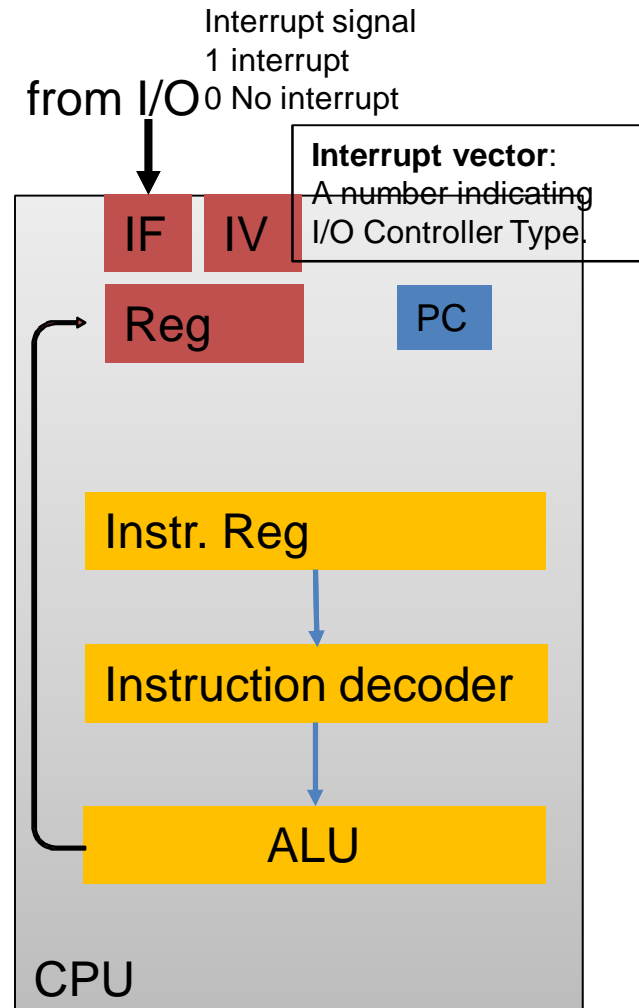


## Memory-mapped I/O - II

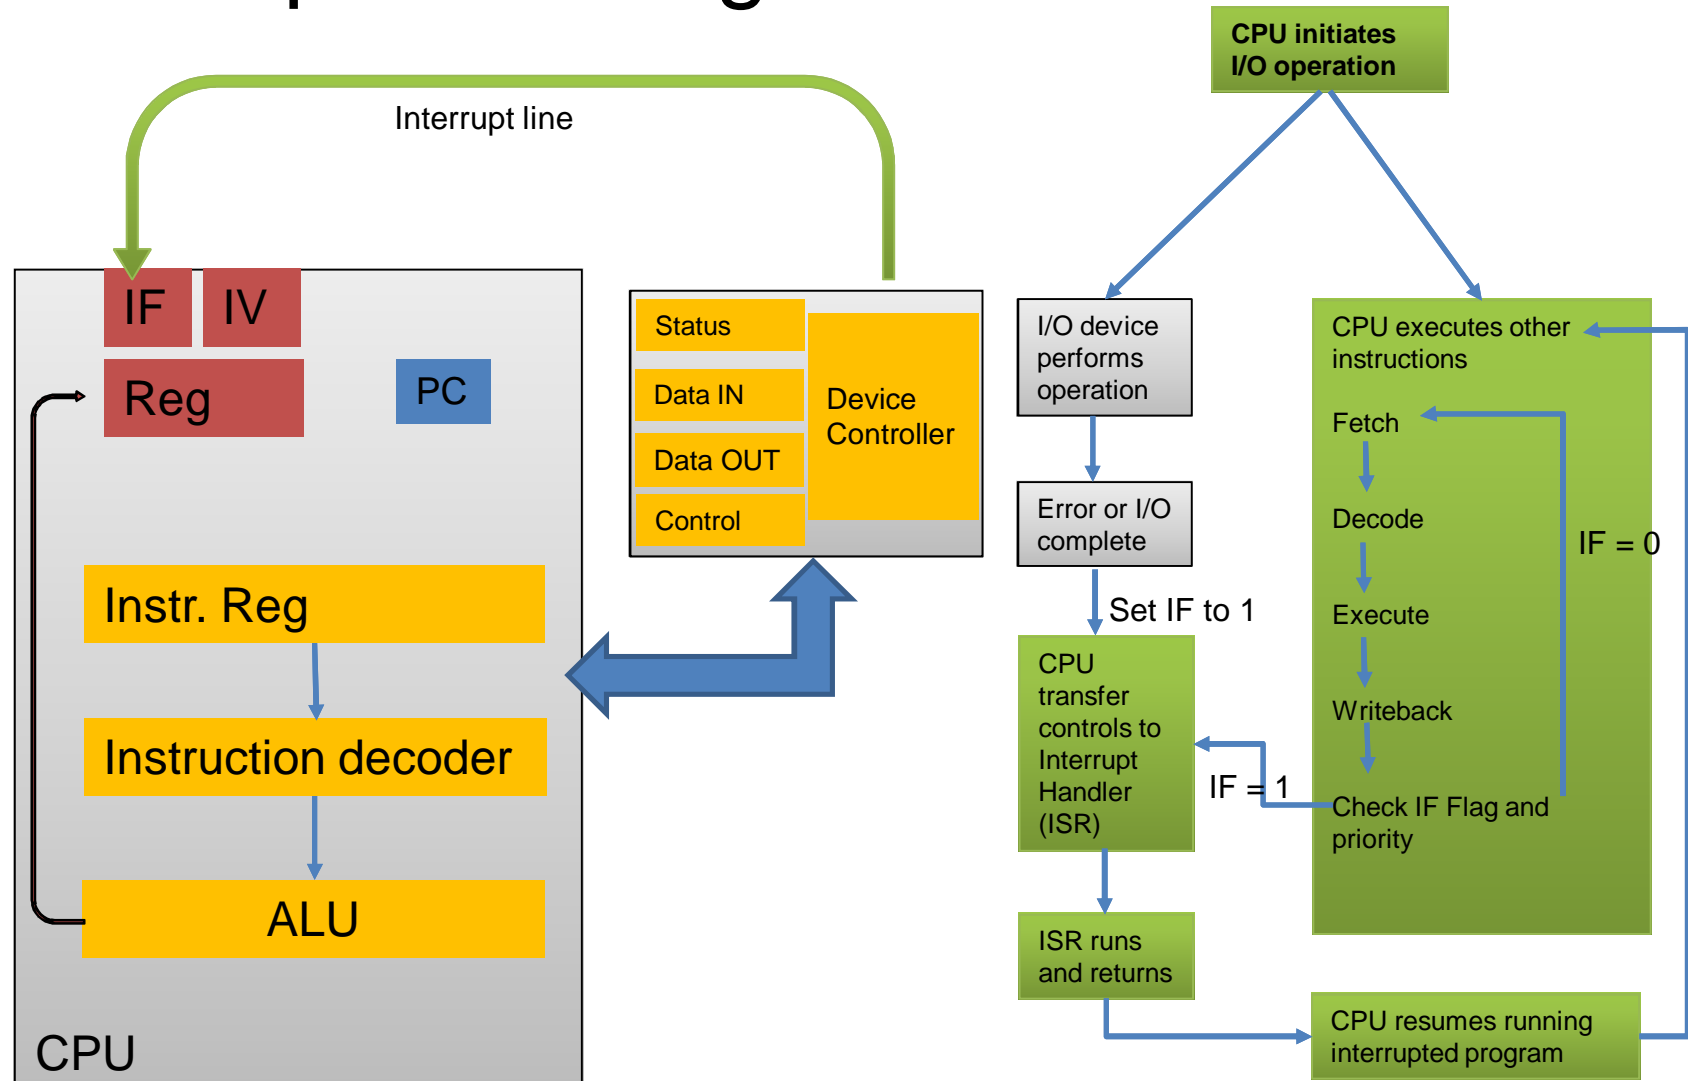
- Using software to wait for I/O `//wait for I/O to be done`  
`//done bit is on status register`
- Can check status register at fixed times `while (done_bit==0) ;`  
`//process further data for I/O`
- Pros
  - Fast response time (for program waiting)
- Cons
  - Low CPU utilization

# HW Setup for Interrupt Handling

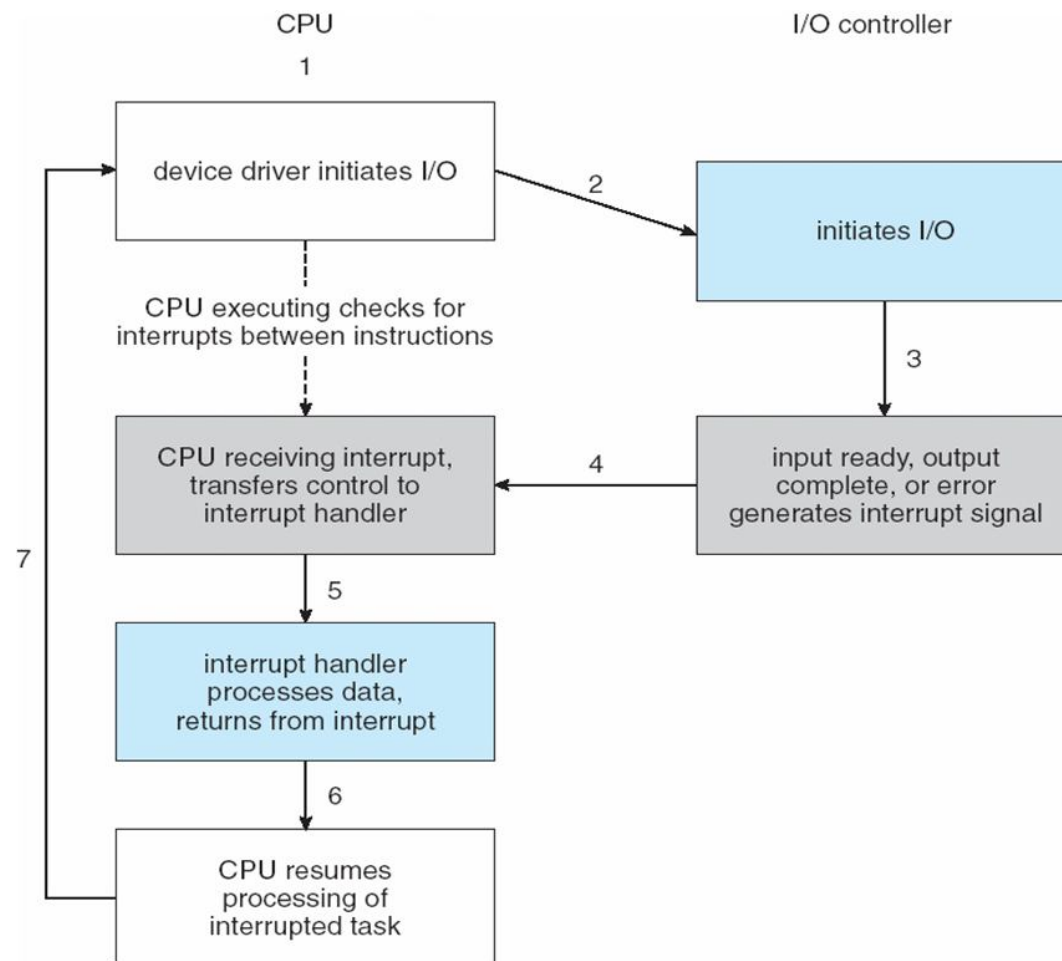
- While waiting for I/O, do other jobs first
- Wait for I/O device to assert interrupt signal
- IF and IV register
  - Interrupt flag
  - Interrupt vector
    - A number indicating I/O controller type



# Interrupt Handling



# Interrupt-Driven I/O Cycle

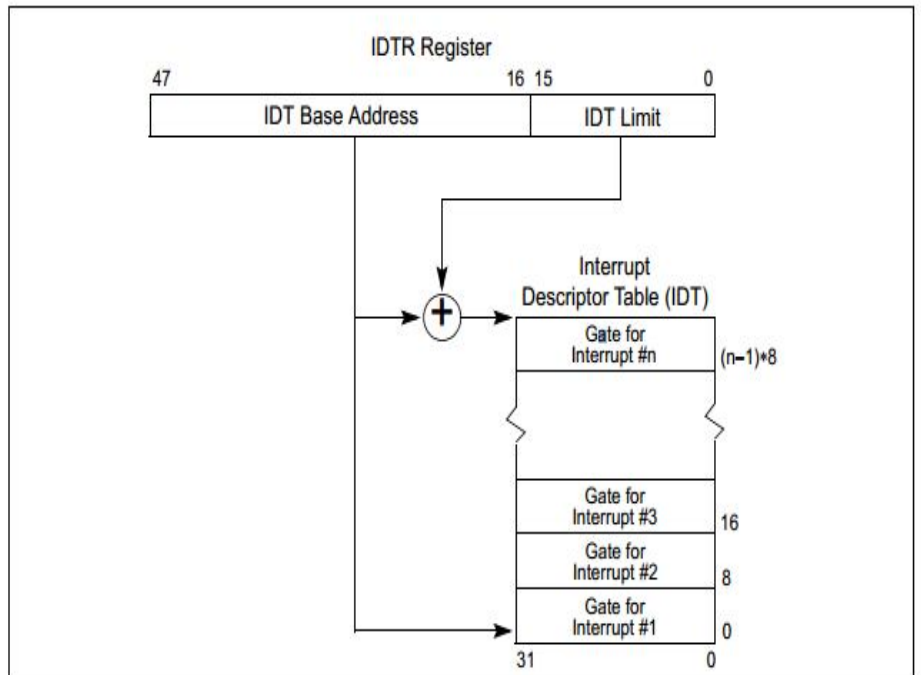


# Interrupts (Cont.)

- 3 kinds of interrupt
  - Exceptions/Faults: Errors: e.g., divide by zero
  - Hardware Interrupt: I/O!
  - Software-generated interrupt: System calls
- Maskable and non-maskable
  - Maskable means turned-off
  - Usually exceptions/faults are not maskable.
- Interrupt priority
  - Interrupt handling can be interrupted by higher priority interrupts.

# Interrupt Vector Table

- Interrupt Vector
  - Unique for each interrupt
- ISR
  - SW handling of interrupts
- Interrupt Vector Table
  - Intel – IDT
    - Real addressing mode: IVT
    - Protected mode: IDT
  - Array of function pointers





# Intel Pentium Processor Event-Vector Table


vector number	description
0	divide error
1	debug exception
2	null interrupt
3	breakpoint
4	INTO-detected overflow
5	bound range exception
6	invalid opcode
7	device not available
8	double fault
9	coprocessor segment overrun (reserved)
10	invalid task state segment
11	segment not present
12	stack fault
13	general protection
14	page fault
15	(Intel reserved, do not use)
16	floating-point error
17	alignment check
18	machine check
19–31	(Intel reserved, do not use)
32–255	maskable interrupts

## Interrupts (Cont.)

- Interrupt mechanism also used for **exceptions**
  - Terminate process, crash system due to hardware error
- Page fault executes when memory access error
- System call executes via **trap** to trigger kernel to execute request
- Multi-CPU systems can process interrupts concurrently
  - If operating system designed to handle it
- Used for time-sensitive processing, frequent, must be fast

# How does CPU transfer control to ISR?

- Must know the address of a function to call a function.
- Where is the interrupt vector table located?
  - Memory
- We want to know address containing the starting address of ISR for vector N
  - $\text{Base Address} + 8 * N$  for 64 bit
  - $\text{Base Address} + 4 * N$  for 32 bit
- Question: how does the HW know where is the base?
- Hardcoded (0x0000 to 0x03ff) during real mode
- IDTR (Interrupt Descriptor Table Register) (Protected mode)
  - Who sets up the IDTR?
  - Who sets up the IDT?

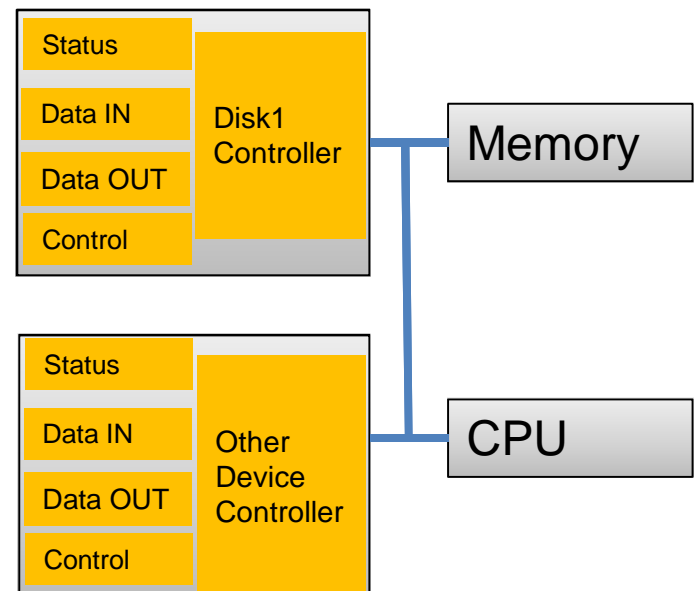


# DMA: Motivation Scenario

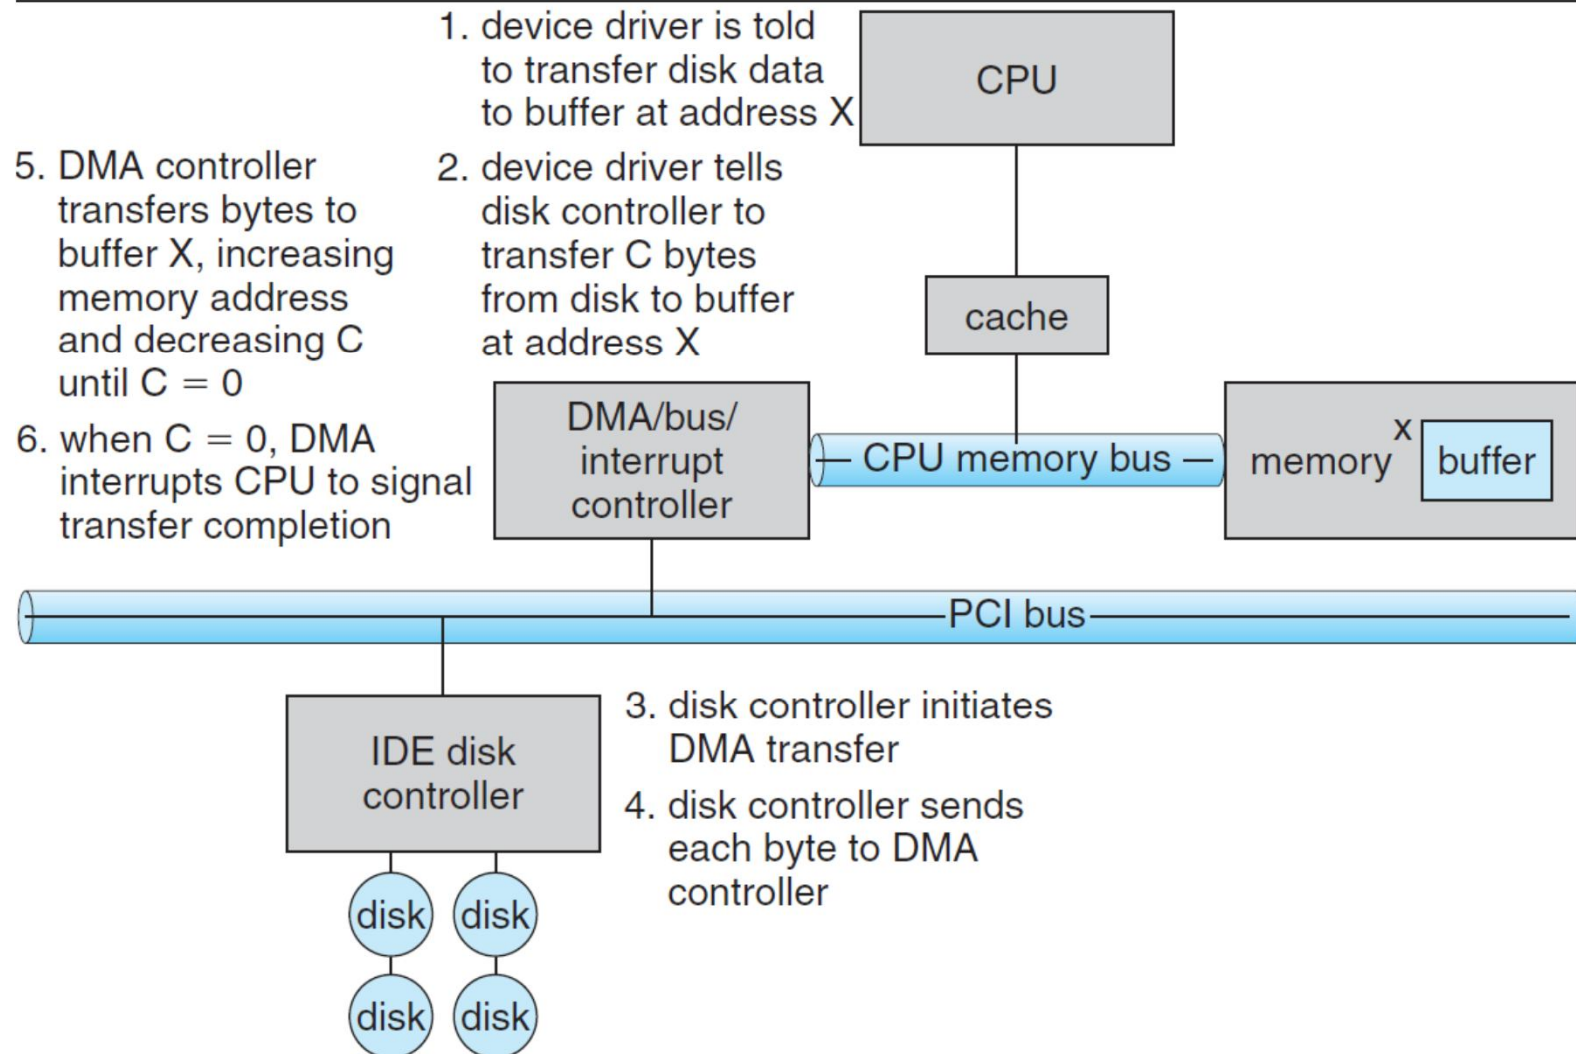
Copy 1KB from Disk 1 to  
Memory (Disk Read)

A simple way:

1. CPU req Disk1 controller for 1<sup>st</sup> byte
2. Disk1 prepares the byte. Put it in Data-IN register.
3. Interrupt CPU
4. ISR reads the byte in Data-IN



# DMA Scenario (Fig 13.5 Textbook)



# Direct Memory Access - I

- Allow devices to read/write memory directly without CPU intervention
- Does not mean accessing the memory in the same cycle
  - Whether multiple devices (e.g. CPU) can access the memory at the same time depends on the memory design
- Allows CPU processing to be in parallel with device to memory communication
  - Reduce number of interrupts

# Direct Memory Access - II

- Used to avoid **programmed I/O** (one byte at a time) for large data movement
- Requires **DMA** controller
- Bypasses CPU to transfer data directly between I/O device and memory
- OS writes DMA command block into memory
  - Source and destination addresses
  - Read or write mode
  - Count of bytes
  - Writes location of command block to DMA controller
  - Bus mastering of DMA controller – grabs bus from CPU
    - **Cycle stealing** from CPU but still much more efficient
  - When done, interrupts to signal completion
- Version that is aware of virtual addresses can be even more efficient - **DVMA**

# Access HW device

- User program access device directly?
  - User program talks with device driver via system calls `open()`, `read()`, `write()` etc
  - Protect data from access-control violations
  - Erroneous use leads to crash



# How does this code execute?

```
int main()
```

```
{
```

```
    int add;
```

```
    int *add_p;
```

```
    scanf("%d", &add);
```

```
    add_p = (int*) add;
```

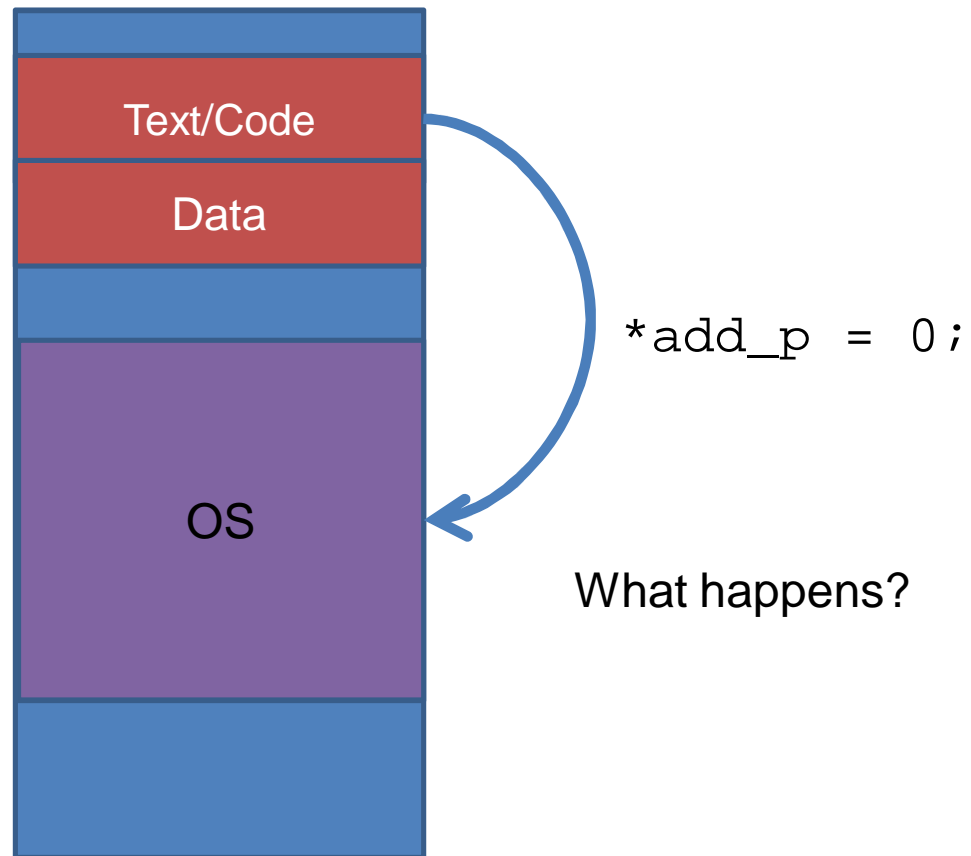
```
    *add_p = 0;
```

```
}
```



libc functions : `read()` and `write()`  
Context switch from user level  
to kernel mode

# MS-DOS and Intel 8088



# CPU privilege modes

- Kernel (Supervisor mode) and user mode
- Privileged instructions
  - e.g. change the idtr value
- Other instructions
  - e.g. add, sub etc

