Lecture 12 - Database Modelling & Database Design (contd…)

CS211 - Introduction to Database

Participation Constraints

- **Total participation** (**indicated by double line**): every entity in the entity set participates in at least one relationship in the relationship set. It represents **NOT NULL** constraint.

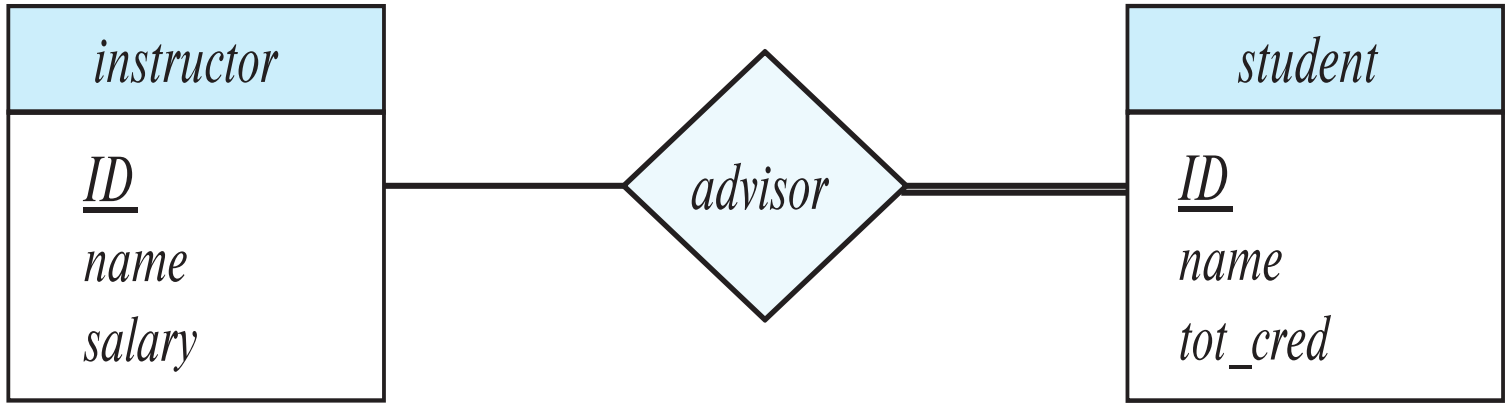

  - participation of *student* in *advisor* relation is total

    Example: every *student* must have an associated instructor

- **Partial participation**: some entities may not participate in any relationship in the relationship set. It represents the attribute values can be **NULL**.

  Example: participation of *instructor* in *advisor* is partial

Notation for Expressing More Complex Constraints

- A line may have an associated **minimum and maximum cardinality**, shown in the form ***l..h***, where ***l*** is the **minimum** and ***h*** the **maximum** cardinality
  - A minimum value of 0 indicates partial participation.
  - A minimum value of 1 indicates total participation.
  - A maximum value of 1 indicates that the entity participates in at most one relationship.
  - A maximum value of * indicates no limit.

- Example

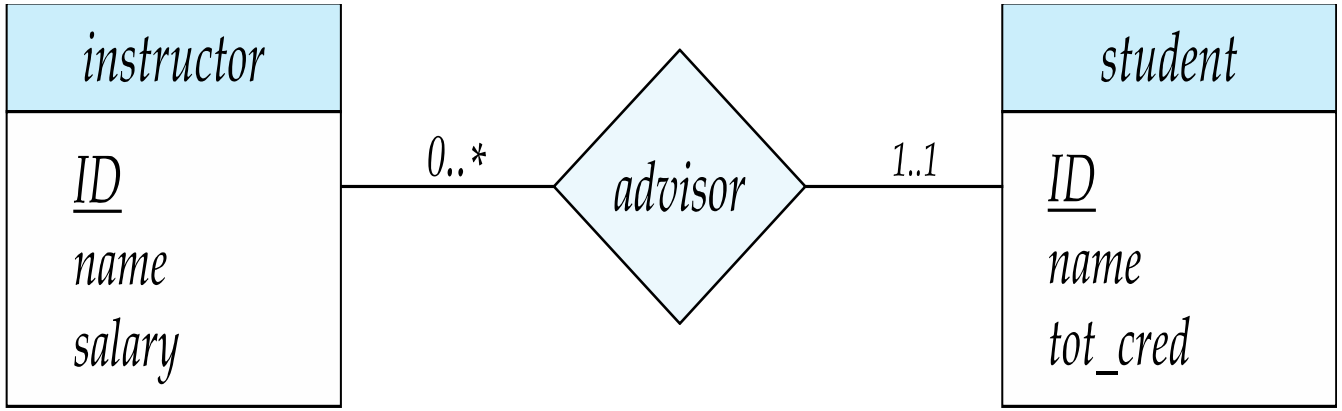

  - Instructor can advise ZERO or more students. A student must have 1 advisor; cannot have multiple advisors

Primary Key Constraints

- Primary keys provide a way to specify how entities and relations are distinguished. We will consider:
  - Entity sets
  - Relationship sets.
  - Weak entity sets

Primary Key for Entity Sets

- By definition, individual entities are distinct.
- From database perspective, the differences among entities must be expressed in terms of their attributes.
- A **key** for an entity is a set of attributes that suffice to distinguish entities from each other.

Primary Key for Relationship Sets

- To distinguish among the various relationships of a relationship set we use the individual **primary keys** of the entities in the relationship set.
  - Let ***R*** be a relationship set involving entity sets $E_1, E_2, \ldots, E_n$.
  - Let ***primary-key***(E_i) denote the set of attributes that forms the primary key for entity set E_i .

Case-1: If the relationship set ***R*** **has no attributes** associated with it, then the set of attributes ***primary-key***$(E_1) \cup$ ***primary-key***$(E_2) \cup \cdots \cup$ ***primary-key***(E_n) makes the primary key for R.

Case-2: If the relationship set ***R*** **has attributes** $a_1, a_2, \ldots, a_m$ associated with it, then the set of attributes ***primary-key***$(E_1) \cup$ ***primary-key***$(E_2) \cup \cdots \cup$ ***primary-key***$(E_n) \cup \{a_1, a_2, \ldots, a_m\}$ makes the primary key for R.

Weak Entity Sets

- An entity set that has a primary key is termed a **strong entity set**.
- An entity set that does not have sufficient attributes to form a primary key is termed a **weak entity set**.
- For a weak entity set to be meaningful, it must be associated with another entity set, called the **identifying** or **owner entity set**.
- The weak entity set is said to be **existence dependent** on the identifying entity set.
- The relationship associating the weak entity set with the identifying entity set is called the **identifying relationship**.

Weak Entity Sets (Discriminator)

- The **discriminator** of a weak entity set is a set of its attributes which combined with the **primary key** of the identifying entity set makes the primary key of weak entity set.

Primary-key of Weak entity set = Discriminator + Primary-key of owner

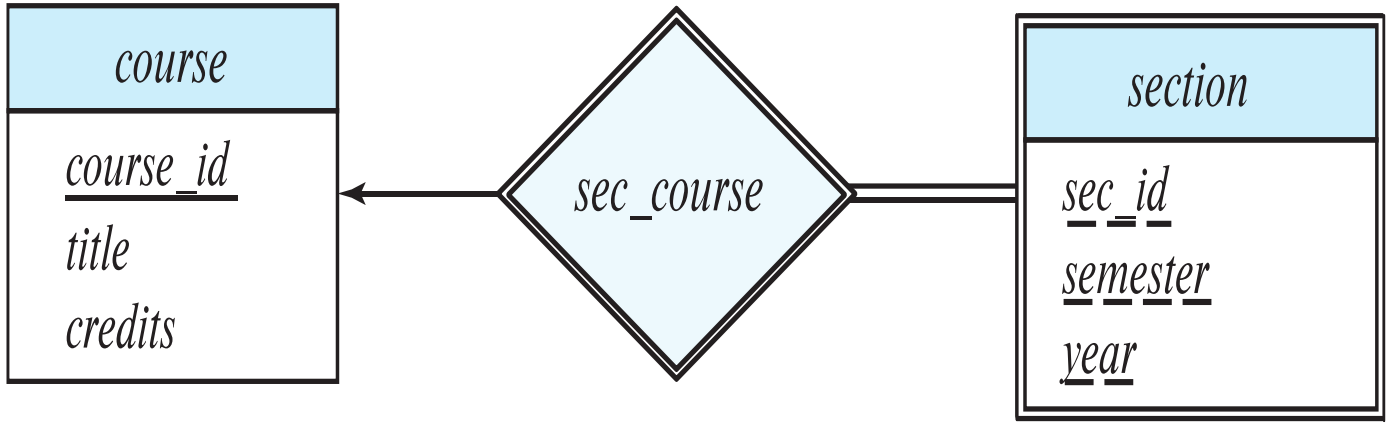

Representing Weak Entity Sets in E-R Diagram

- In E-R diagrams, a **weak entity set** is depicted via a **double rectangle**.
- We underline the **discriminator** of a weak entity set with a **dashed line**.
- The **identifying relationship set** connecting the weak entity set to the identifying strong entity set is depicted by a **double diamond**.
- **Primary key** for the weak entity set "*section*" is: (***course_id, sec_id, semester, year***)

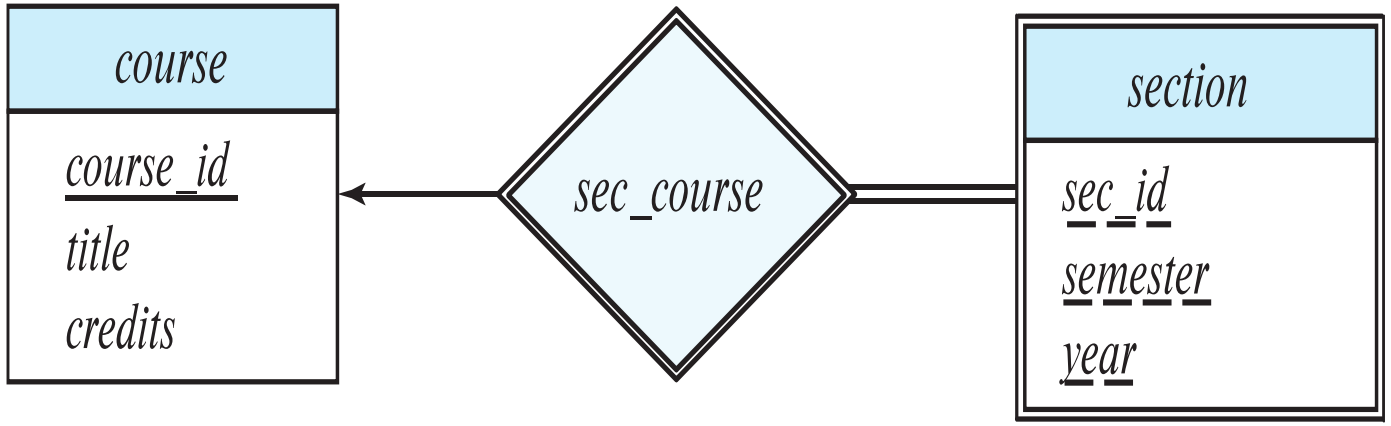

Specialization

- **Specialization** is a **top-down** approach in which a higher-level entity is **divided** into multiple ***specialized*** lower-level entities.

- In addition to **sharing the attributes** of the higher-level entity, these lower-level entities have ***specific*** **attributes** of their own.

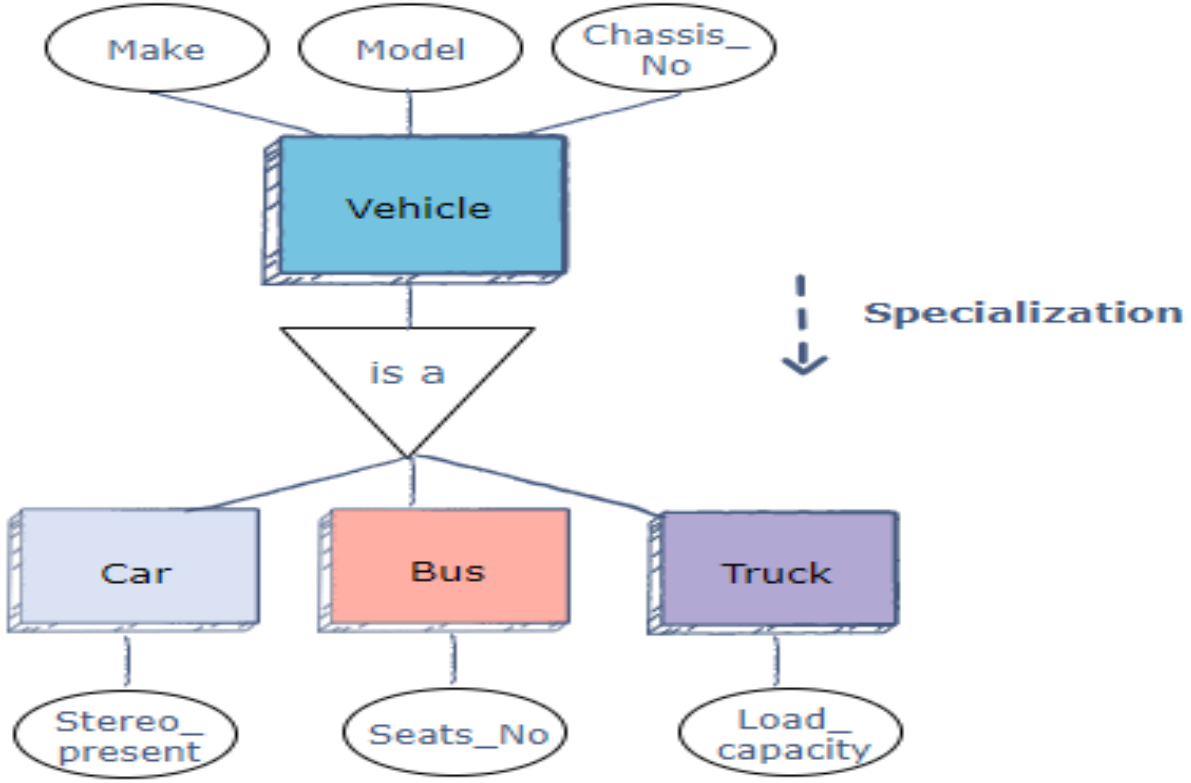

Generalization

- **Generalization** is a **bottom-up** approach in which multiple specialized lower-level entities are combined to form a single **higher-level entity**.

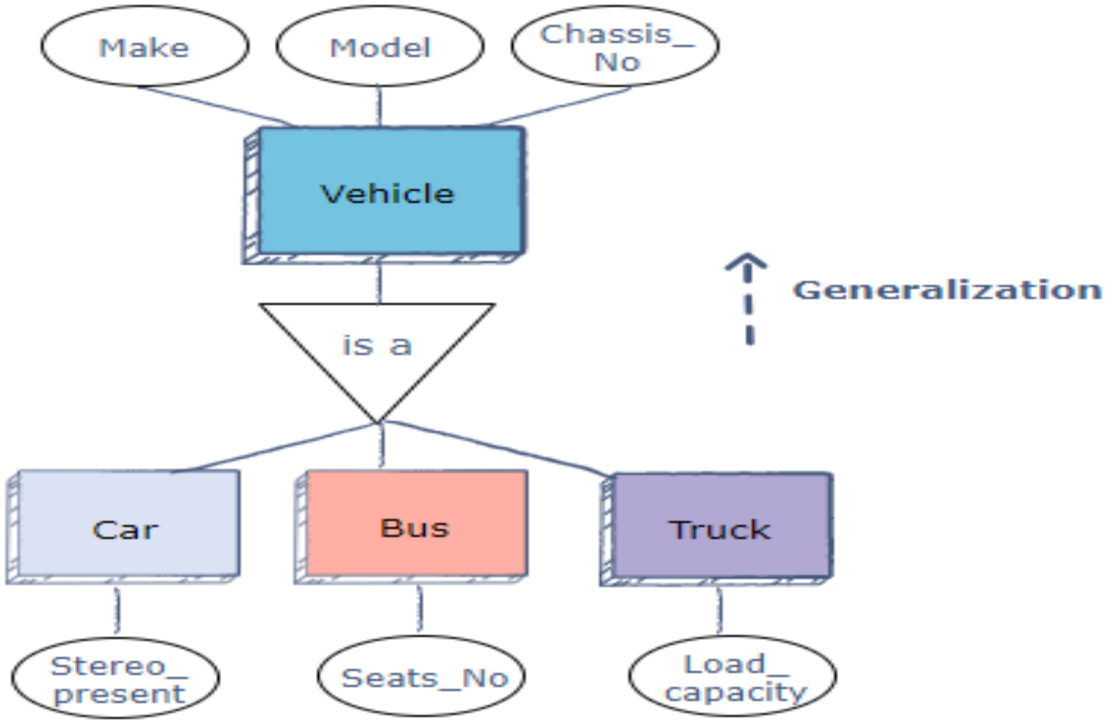

- Generalization is usually used to **find common attributes** among entities to form a generalized entity.

E-R Diagram for a University Enterprise

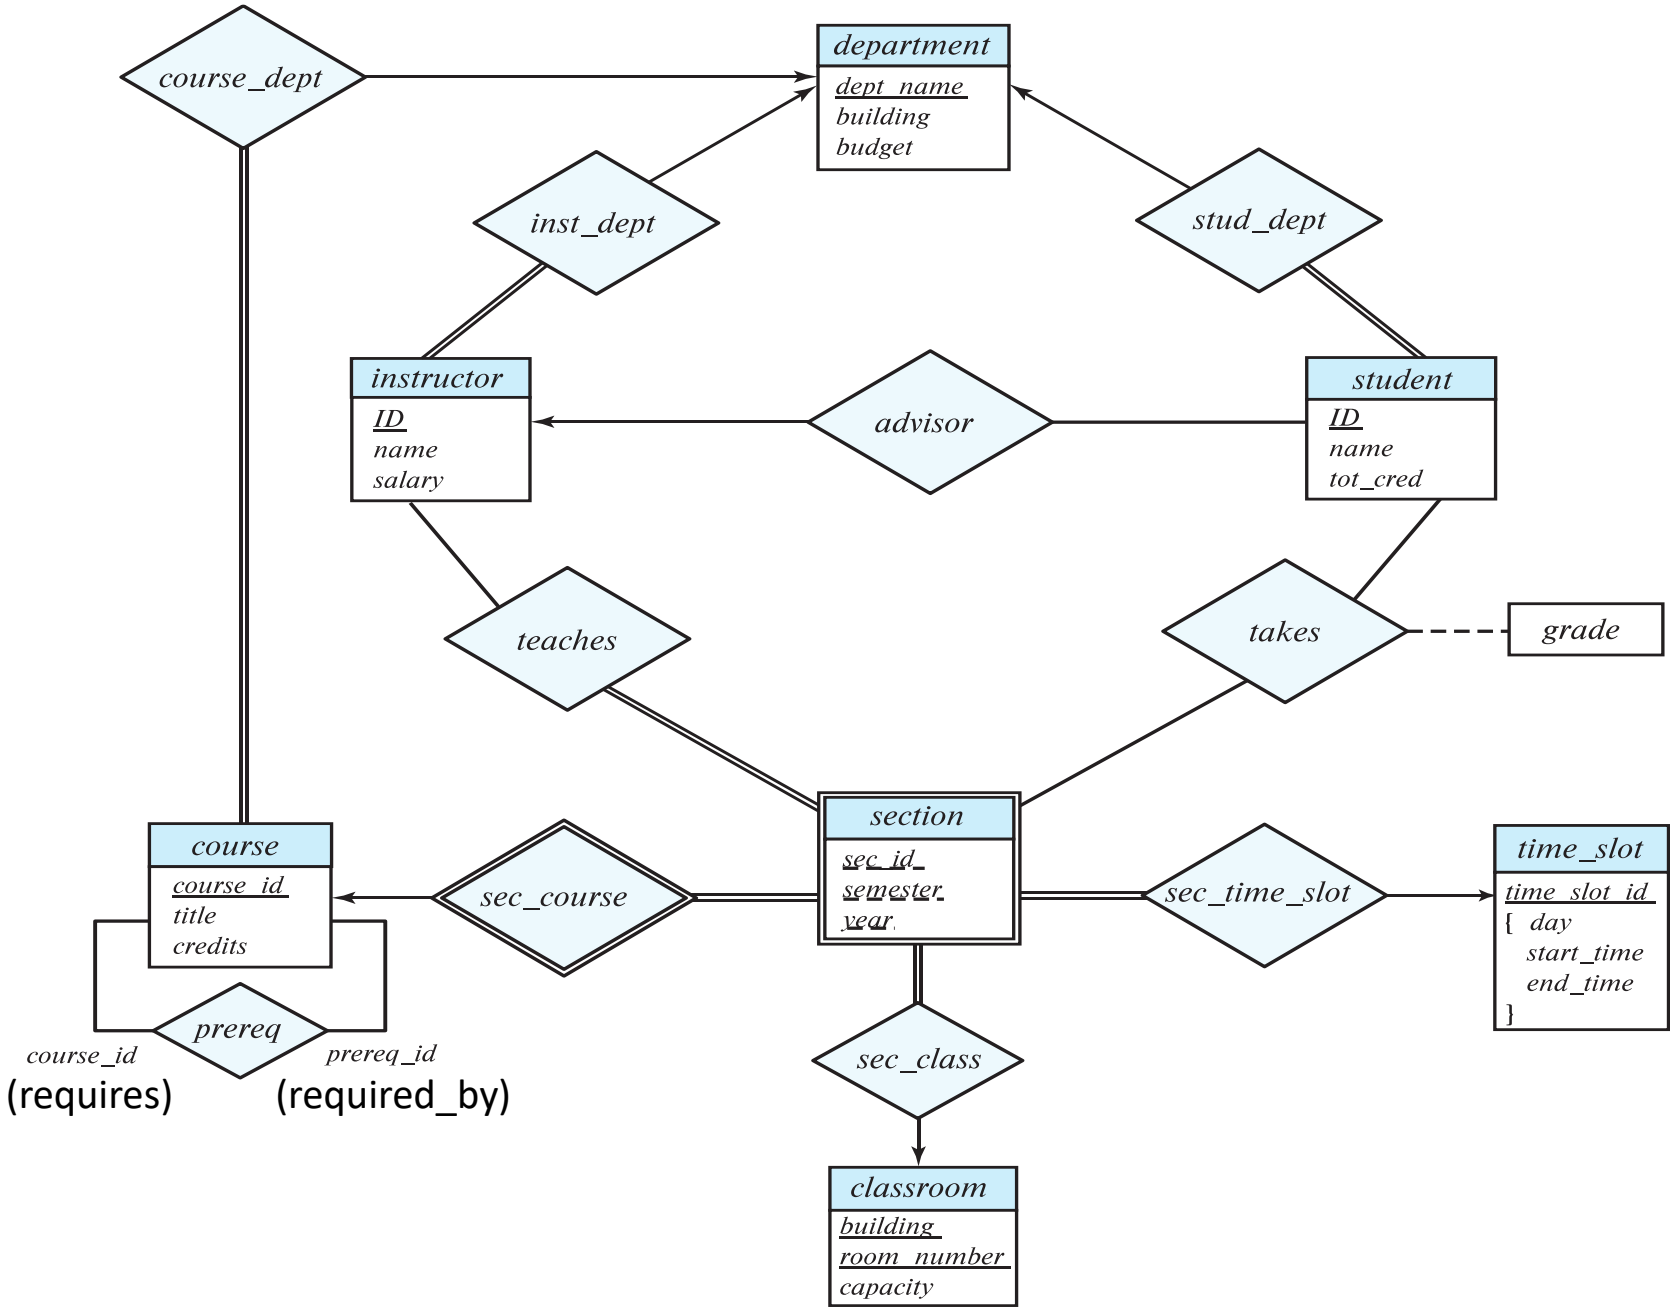

- **Entity set**
- **Attribute type**
- **Entity set primary-key**
- **Relationship set**
- **Relationship set primary-key**
- **Relationship attribute**
- **Entity role**
- **Degree of relationship**
- **Cardinality constraint**
- **Participation constrain**
- **Weak entity**
- **Identifying relationship**
- **Owner entity set**
- **Discriminator**

Database Design

Reduction to Relation Schemas

- **Entity sets** and **relationship sets** are expressed as ***relation schemas*** that represent the contents of the database.

- **For each entity set and relationship set there is a unique schema that is assigned the name of the corresponding entity set or relationship set.**

- **Each schema has a number of columns (generally corresponding to attributes), which have unique names.**

Reduction of Entity Sets

- A **strong entity set** reduces to a schema with the same attributes.

  - the **primary key** of the entity set serves as the primary key of the resulting schema

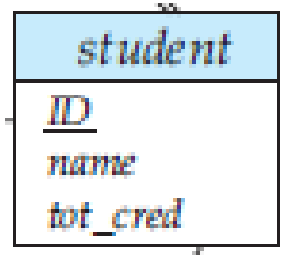

student(ID, name, tot_cred)

- For schemas derived from a **weak entity set**, the combination of the **primary key** of the strong entity set and the **discriminator** of the weak entity set serves as the primary key of the schema.

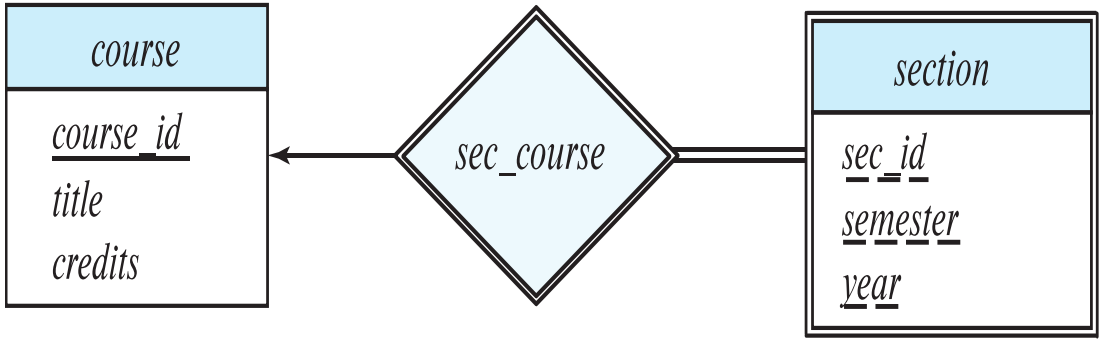

section (course_id, sec_id, sem, year)

Reduction of Entity Sets

- **Composite attributes** are flattened out by creating a separate attribute for each component attribute.

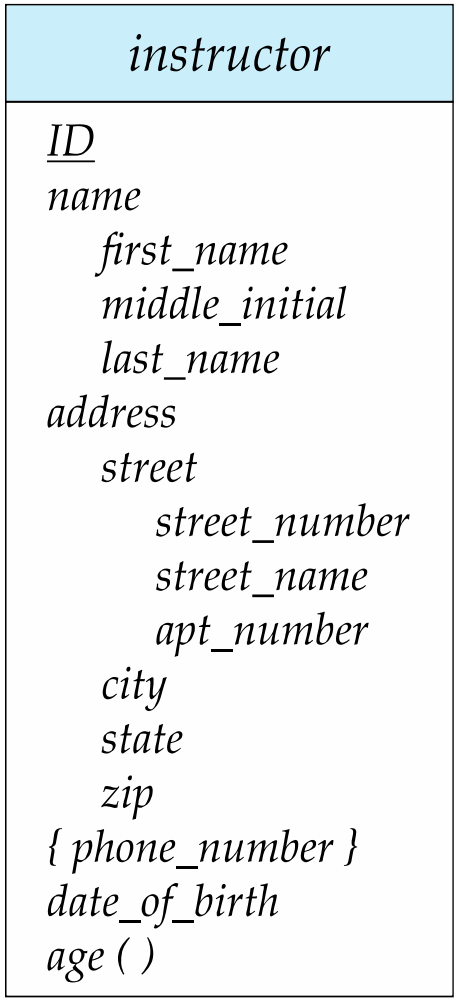

instructor(ID,
first_name, middle_initial, last_name,
street_number, street_name, apt_number, city, state,
zip,
date_of_birth)

Reduction of Entity Sets

- A **multivalued attribute** ***M*** of an entity ***E*** is represented by a **separate schema *EM*.**

- **Schema *EM* has attributes corresponding to the primary key of *E* and an attribute corresponding to multivalued attribute *M*.**

- Example: Multivalued attribute ***phone_number*** of *instructor* is represented by a schema:

  inst_phone= (ID, phone_number)

- Each value of the multivalued attribute maps to a separate tuple of the relation on schema *EM*.

  - For example, an *instructor* entity with primary key 22222 and phone numbers 456-7890 and 123-4567 maps to two tuples:

    (22222, 456-7890) and **(22222, 123-4567)**

Reduction of Relationship Sets

- For a **binary many-to-many relationship**, the union of the primary-key attributes of **all the participating entity sets** becomes the primary key.

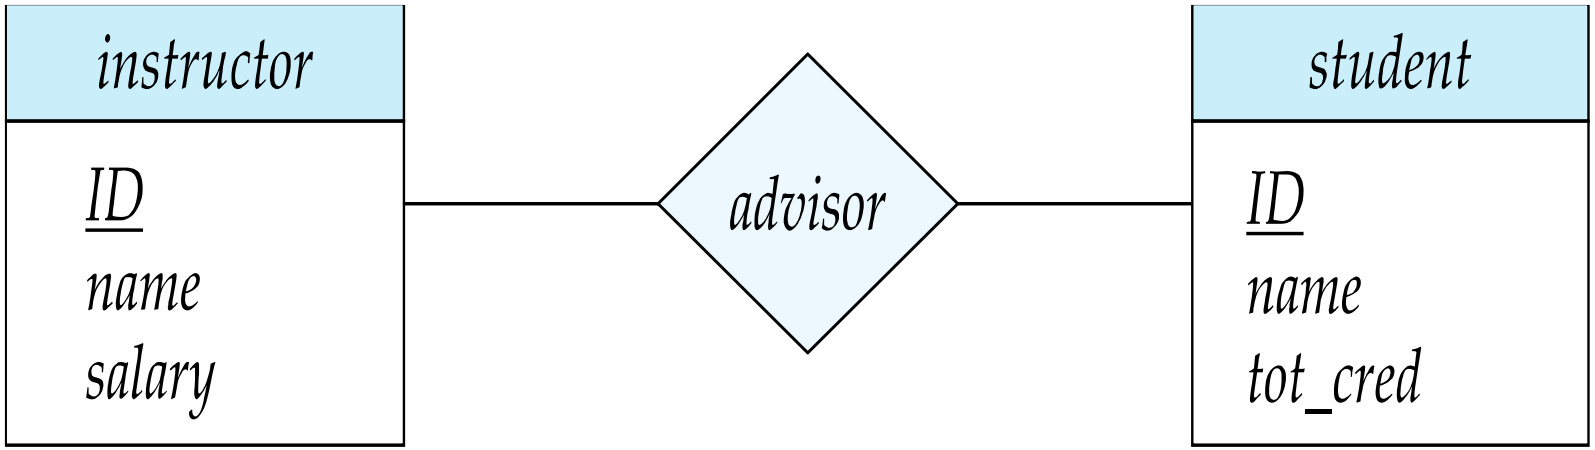

advisor = (i_ID, s_ID)

Reduction of Relationship Sets

- For a **binary one-to-one relationship set**, the primary key of **either entity set** can be chosen as the primary key. The choice can be made arbitrarily.

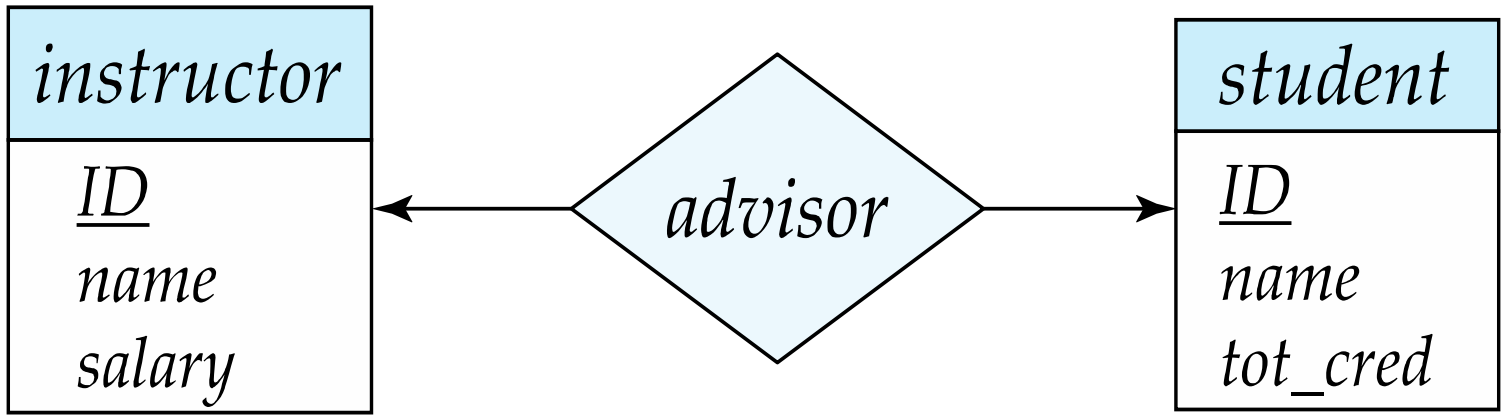

advisor = (s_ID)

Reduction of Relationship Sets

- For **a binary many-to-one or one-to-many relationship set**, the primary key of the **entity set on the "many" side** of the relationship set serves as the primary key.

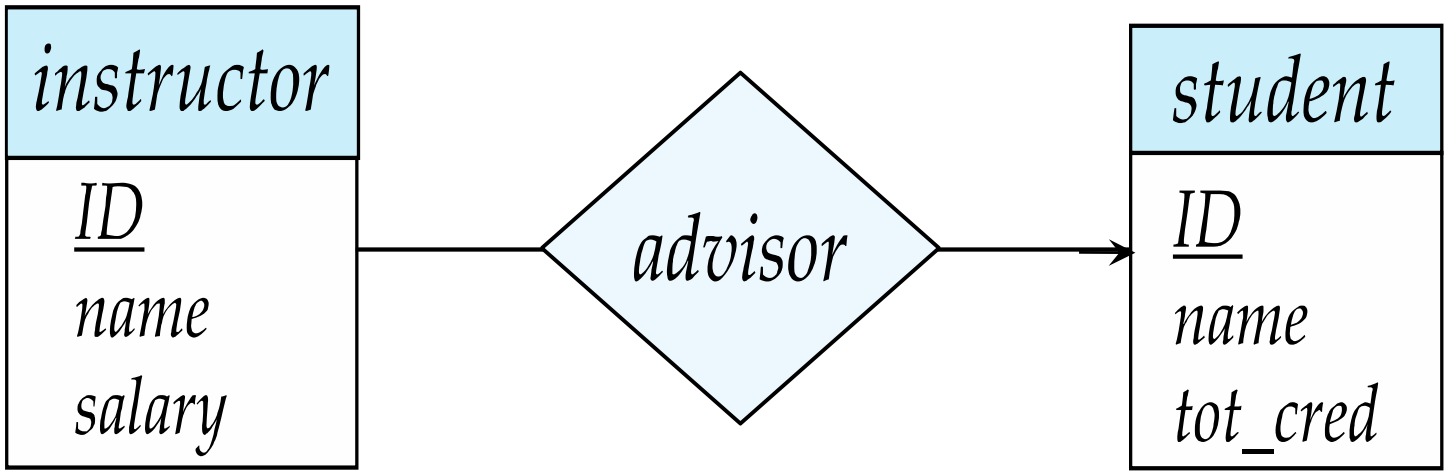

advisor = (i_ID)

Reduction of Relationship Sets

- For **an *n*-ary relationship set without any arrows on its edges**, the union of the primary key-attributes of **all the participating entity sets** becomes the primary key.

- For an ***n*-ary relationship set with an arrow on one of its edges**, the primary keys of the **entity sets not on the "arrow" side** of the relationship set serve as the primary key for the schema.

Reduction of Relationship Sets

Multi-degree relationship → schema (Case-1)

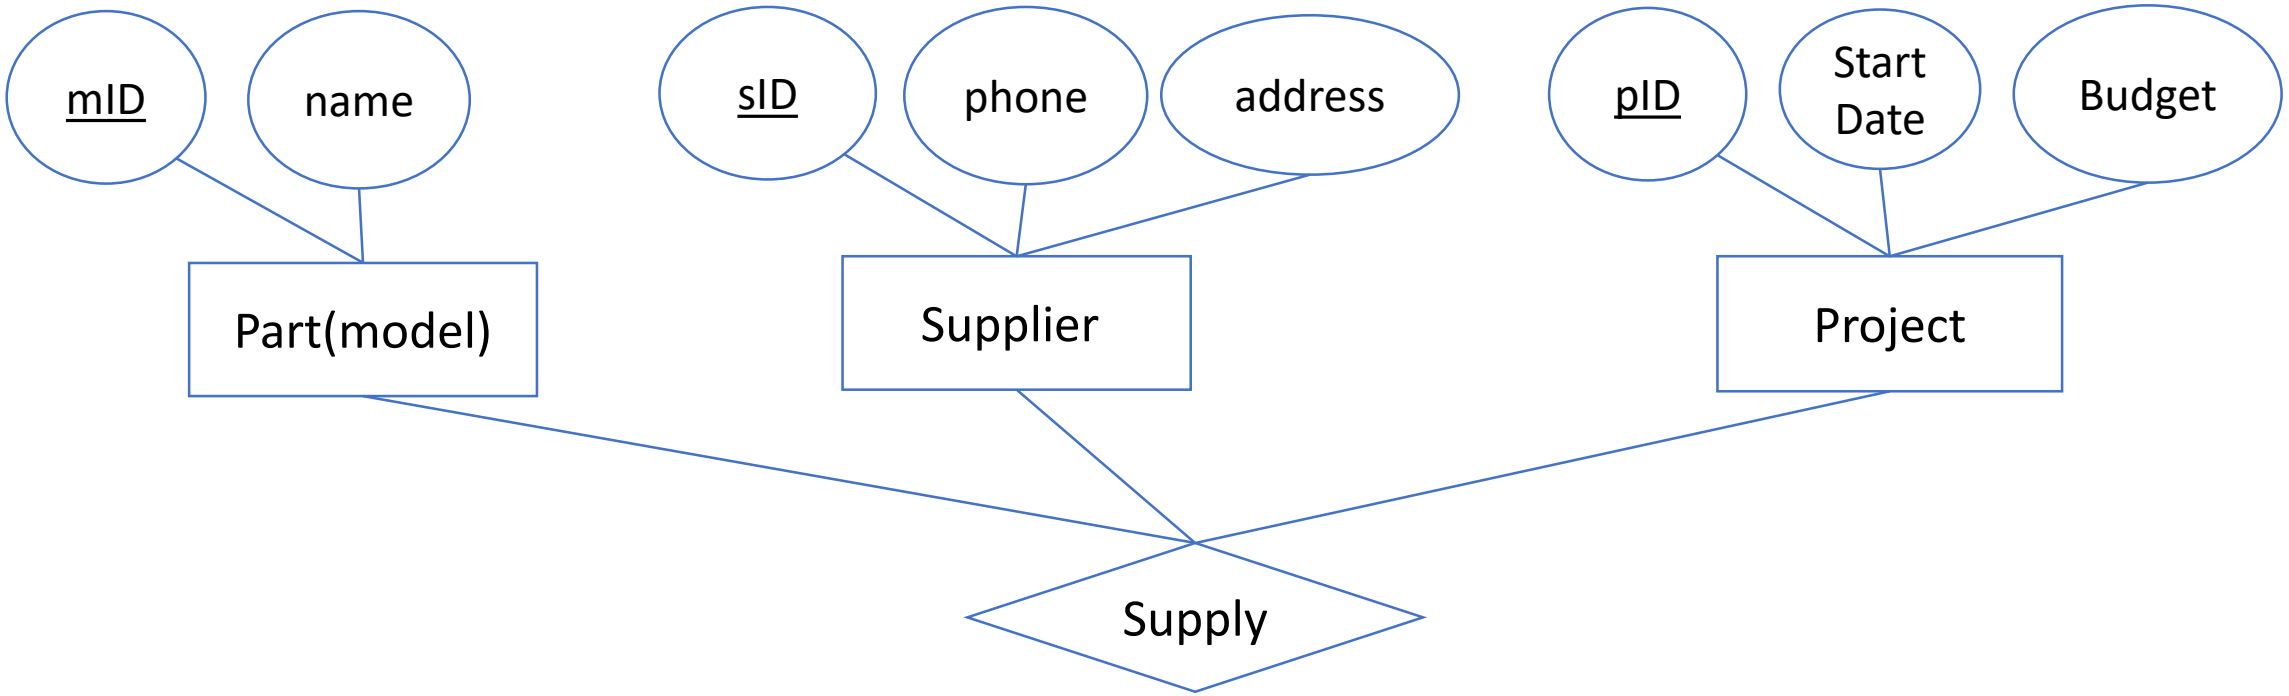

supply(mID, sID, pID)

Reduction of Relationship Sets

Multi-degree relationship → schema (Case-2)

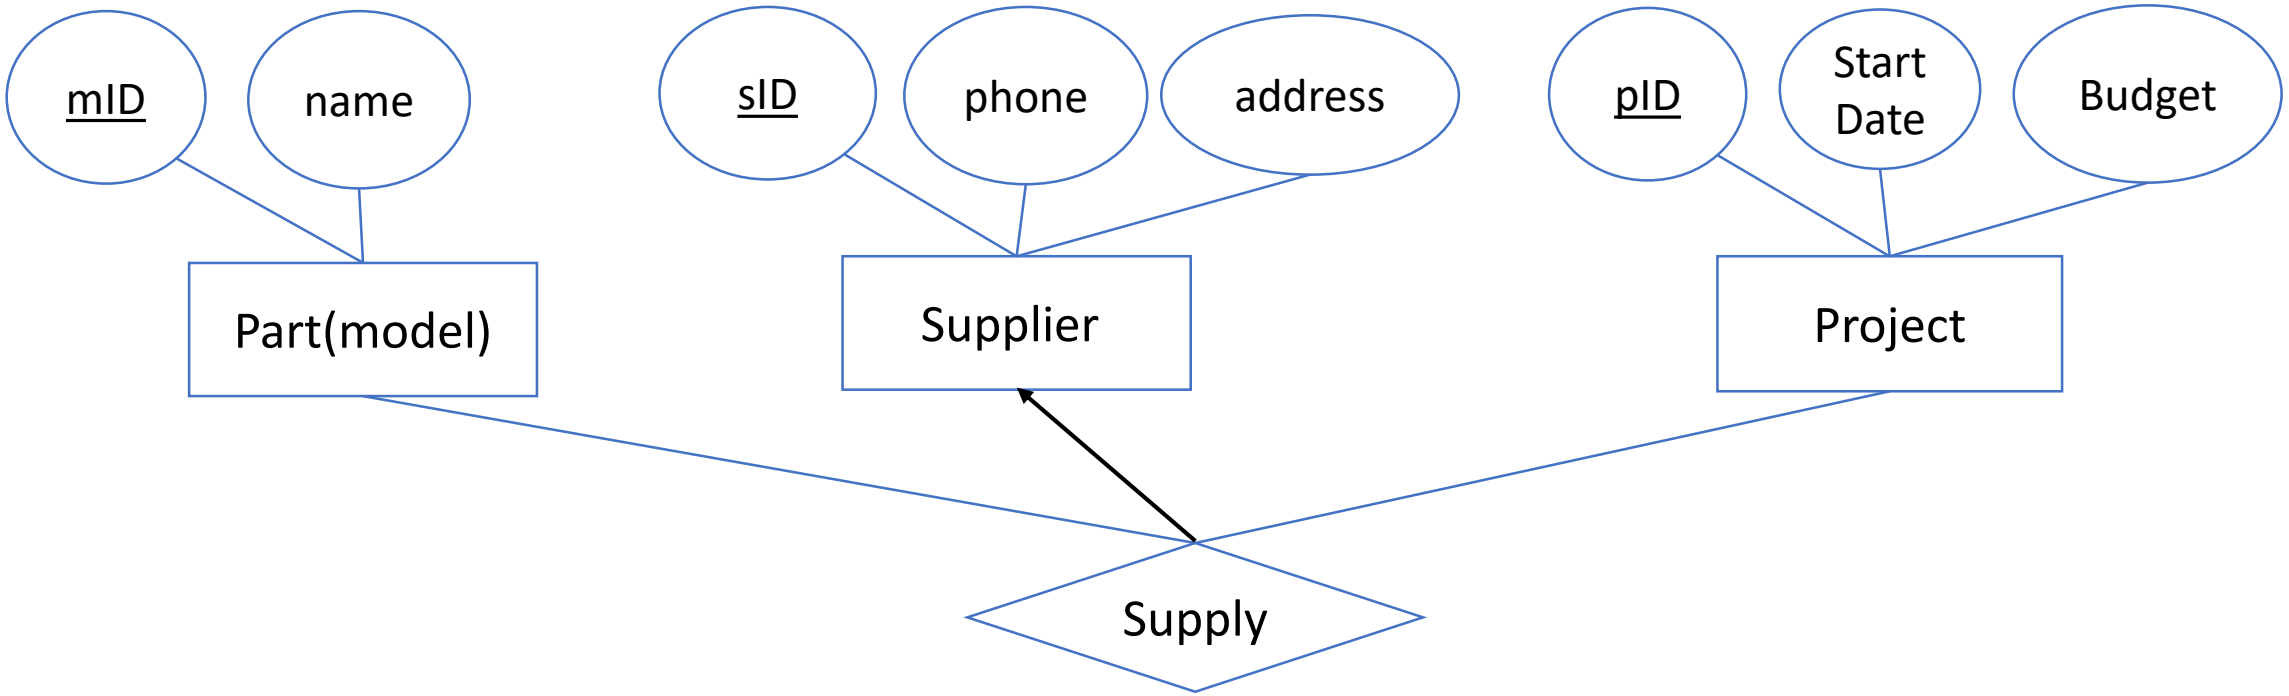

supply(mID, pID)

Reduction of Relationship Sets

Multi-degree relationship → schema (Case-3)

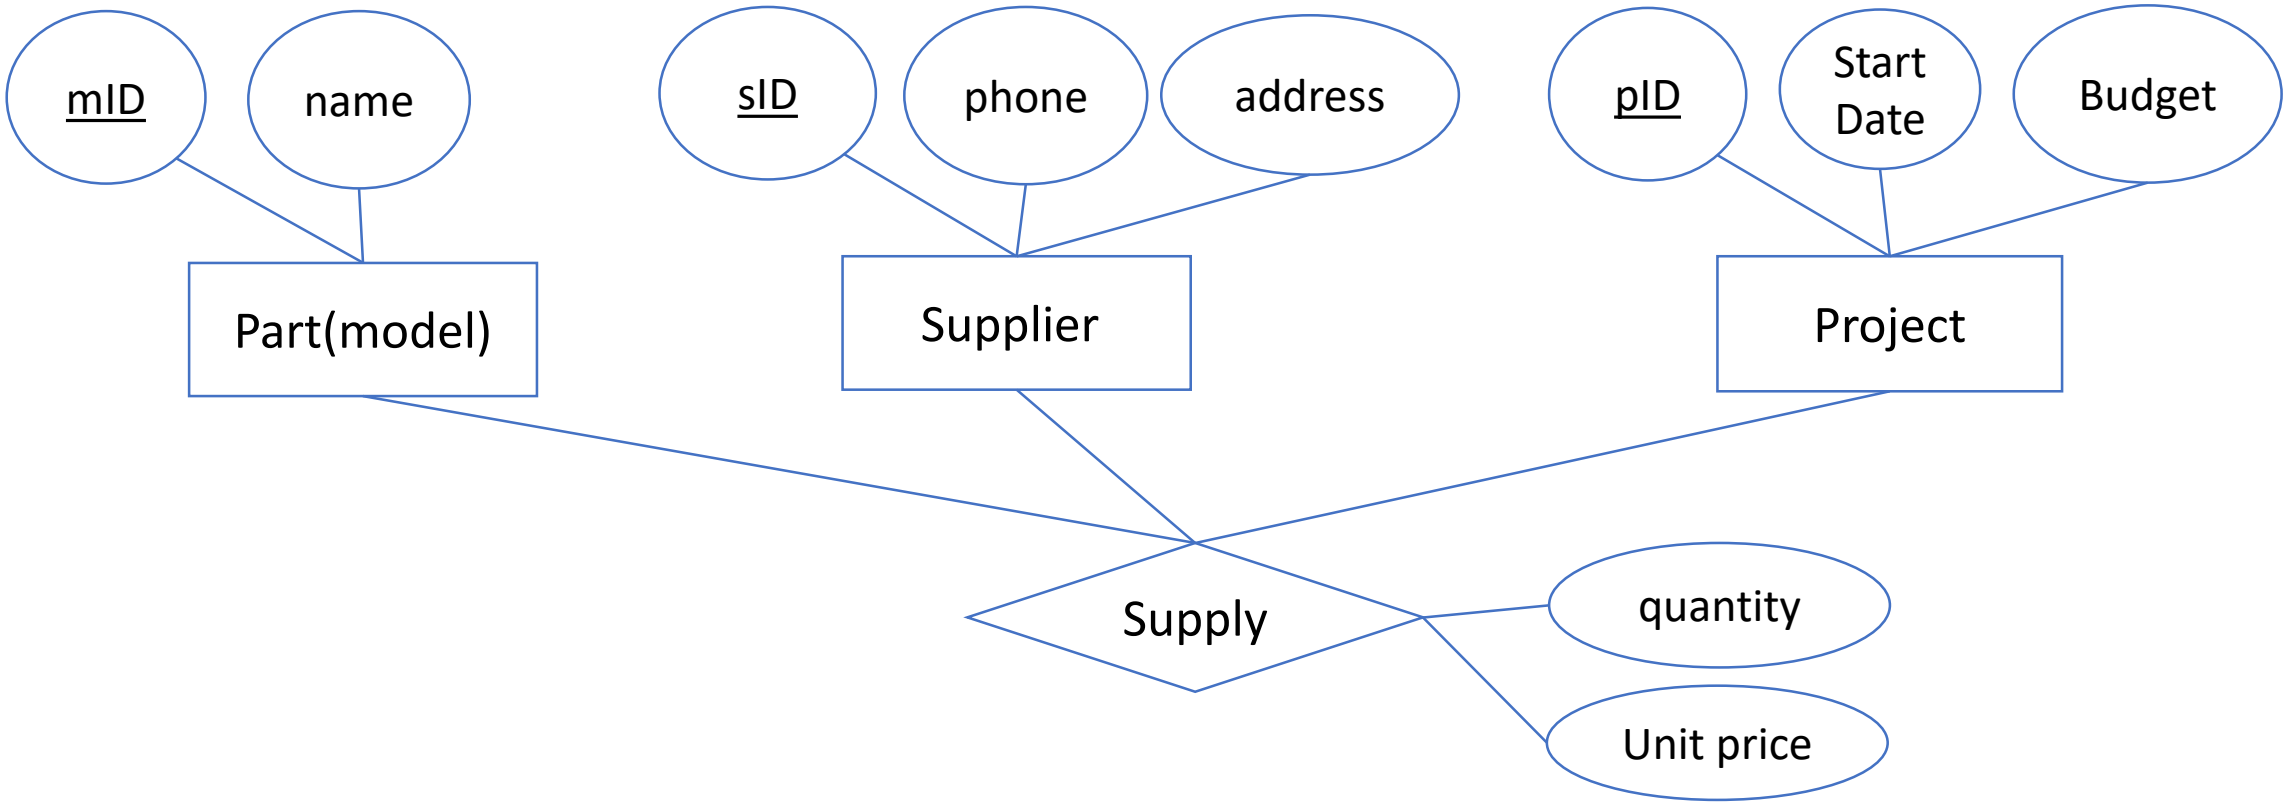

supply(mID, sID, pID, quantity, unitprice)

Reduction of Relationship Sets

Multi-degree relationship → schema (Case-4)

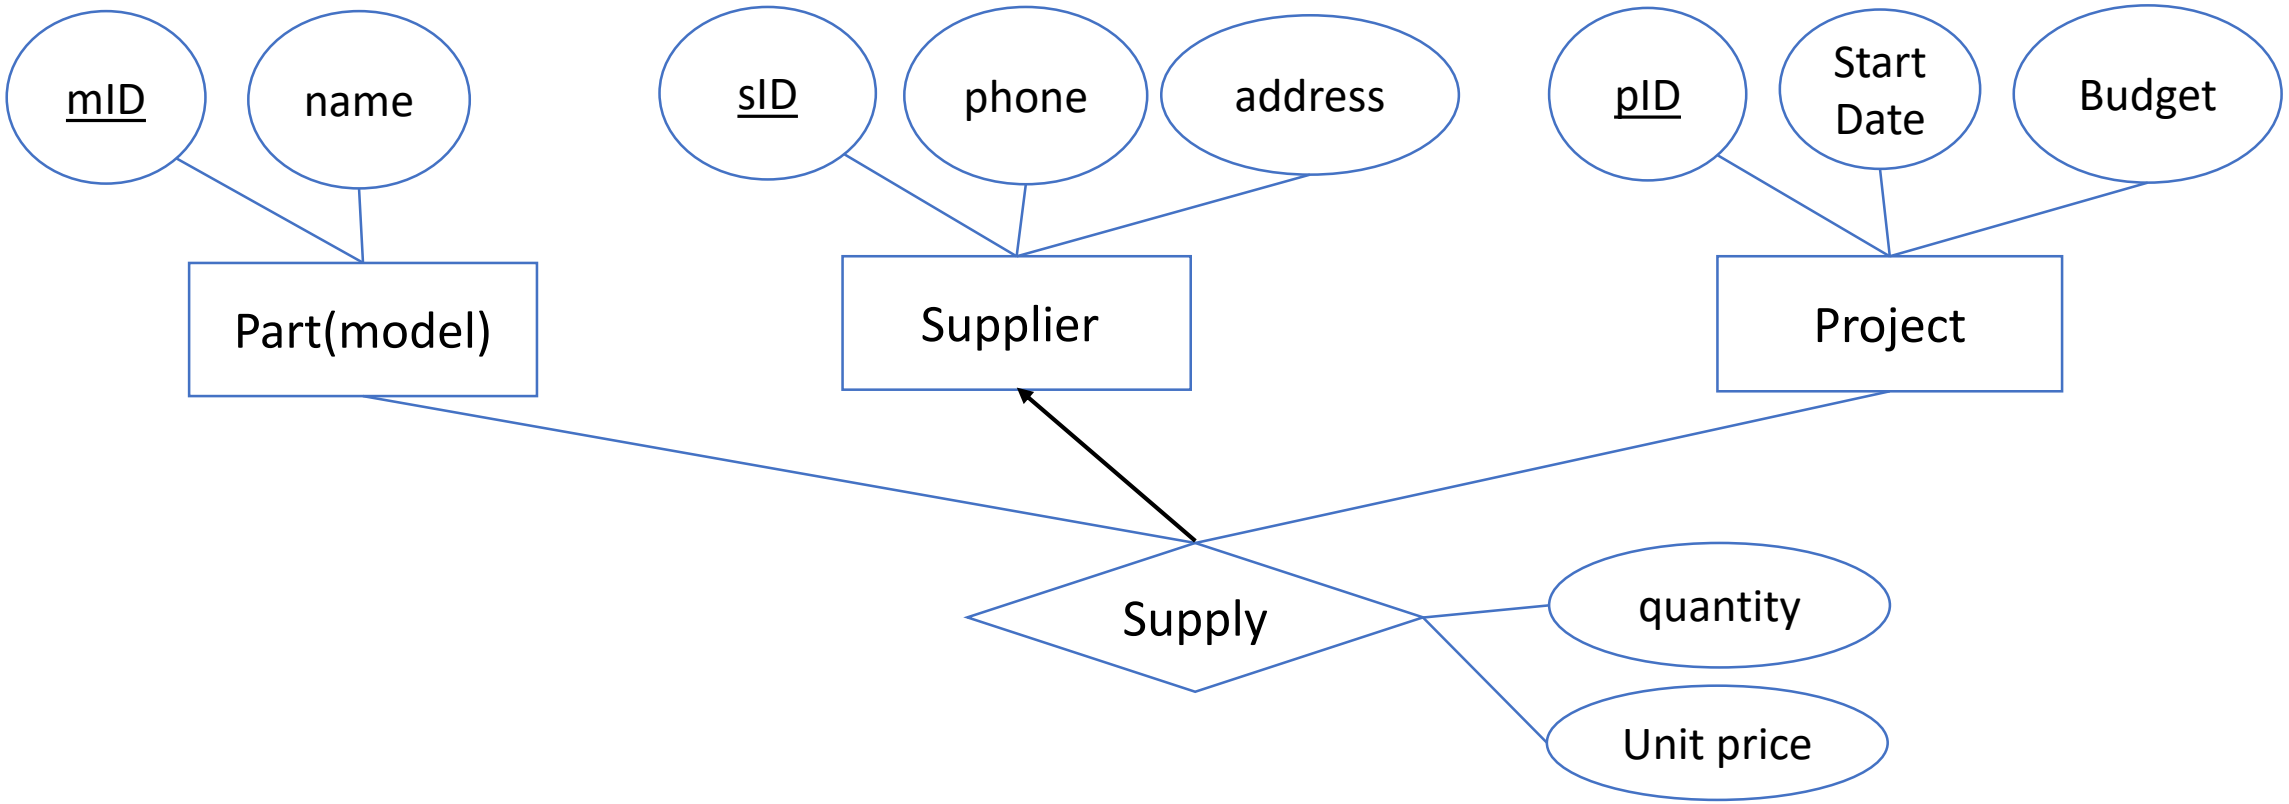

supply(mID, pID, quantity, unitprice)

Summary of Symbols Used in E-R Notation

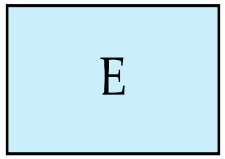

entity set

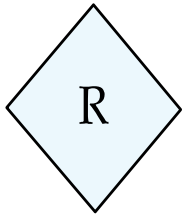

relationship set

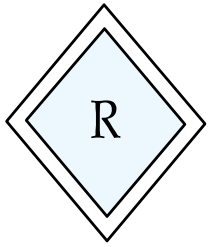

identifying relationship set for weak entity set

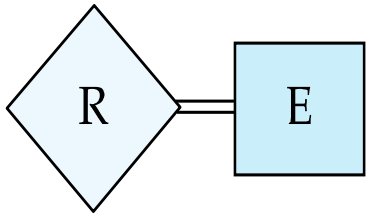

total participation of entity set in relationship

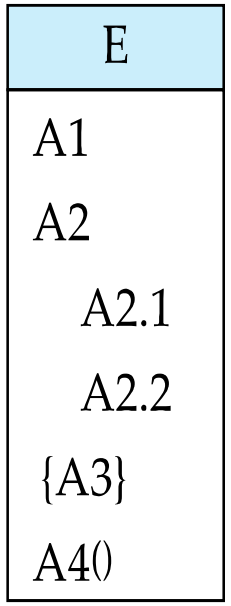

attributes:
simple (A1),
composite (A2) and
multivalued (A3)
derived (A4)

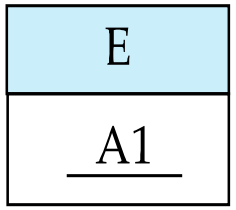

primary key

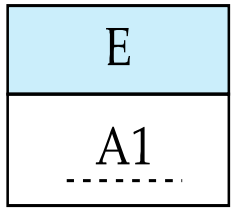

discriminating attribute of weak entity set

Summary of Symbols Used in E-R Notation

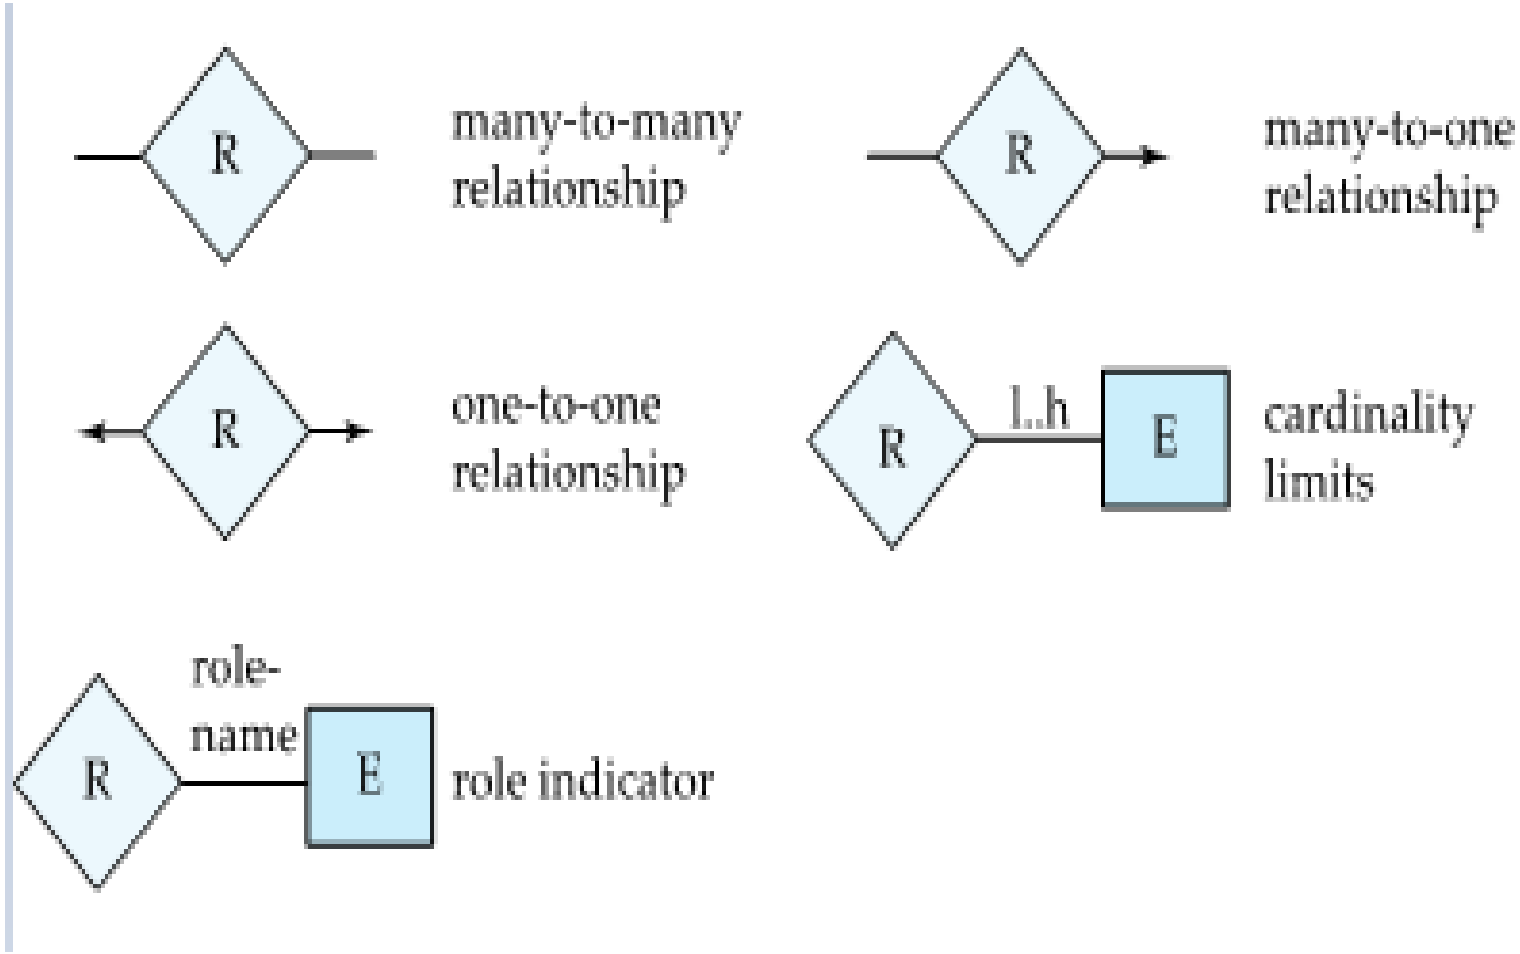

Alternative ER Notations

entity set E with
simple attribute A1,
composite attribute A2,
multivalued attribute A3,
derived attribute A4,
and primary key A1

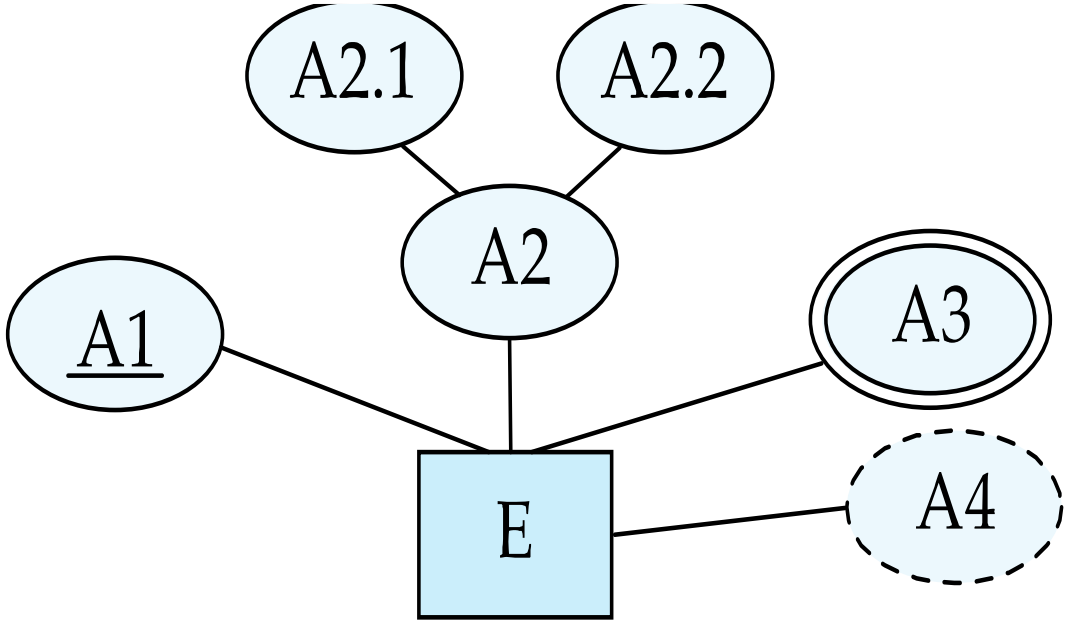

Alternative ER Notations

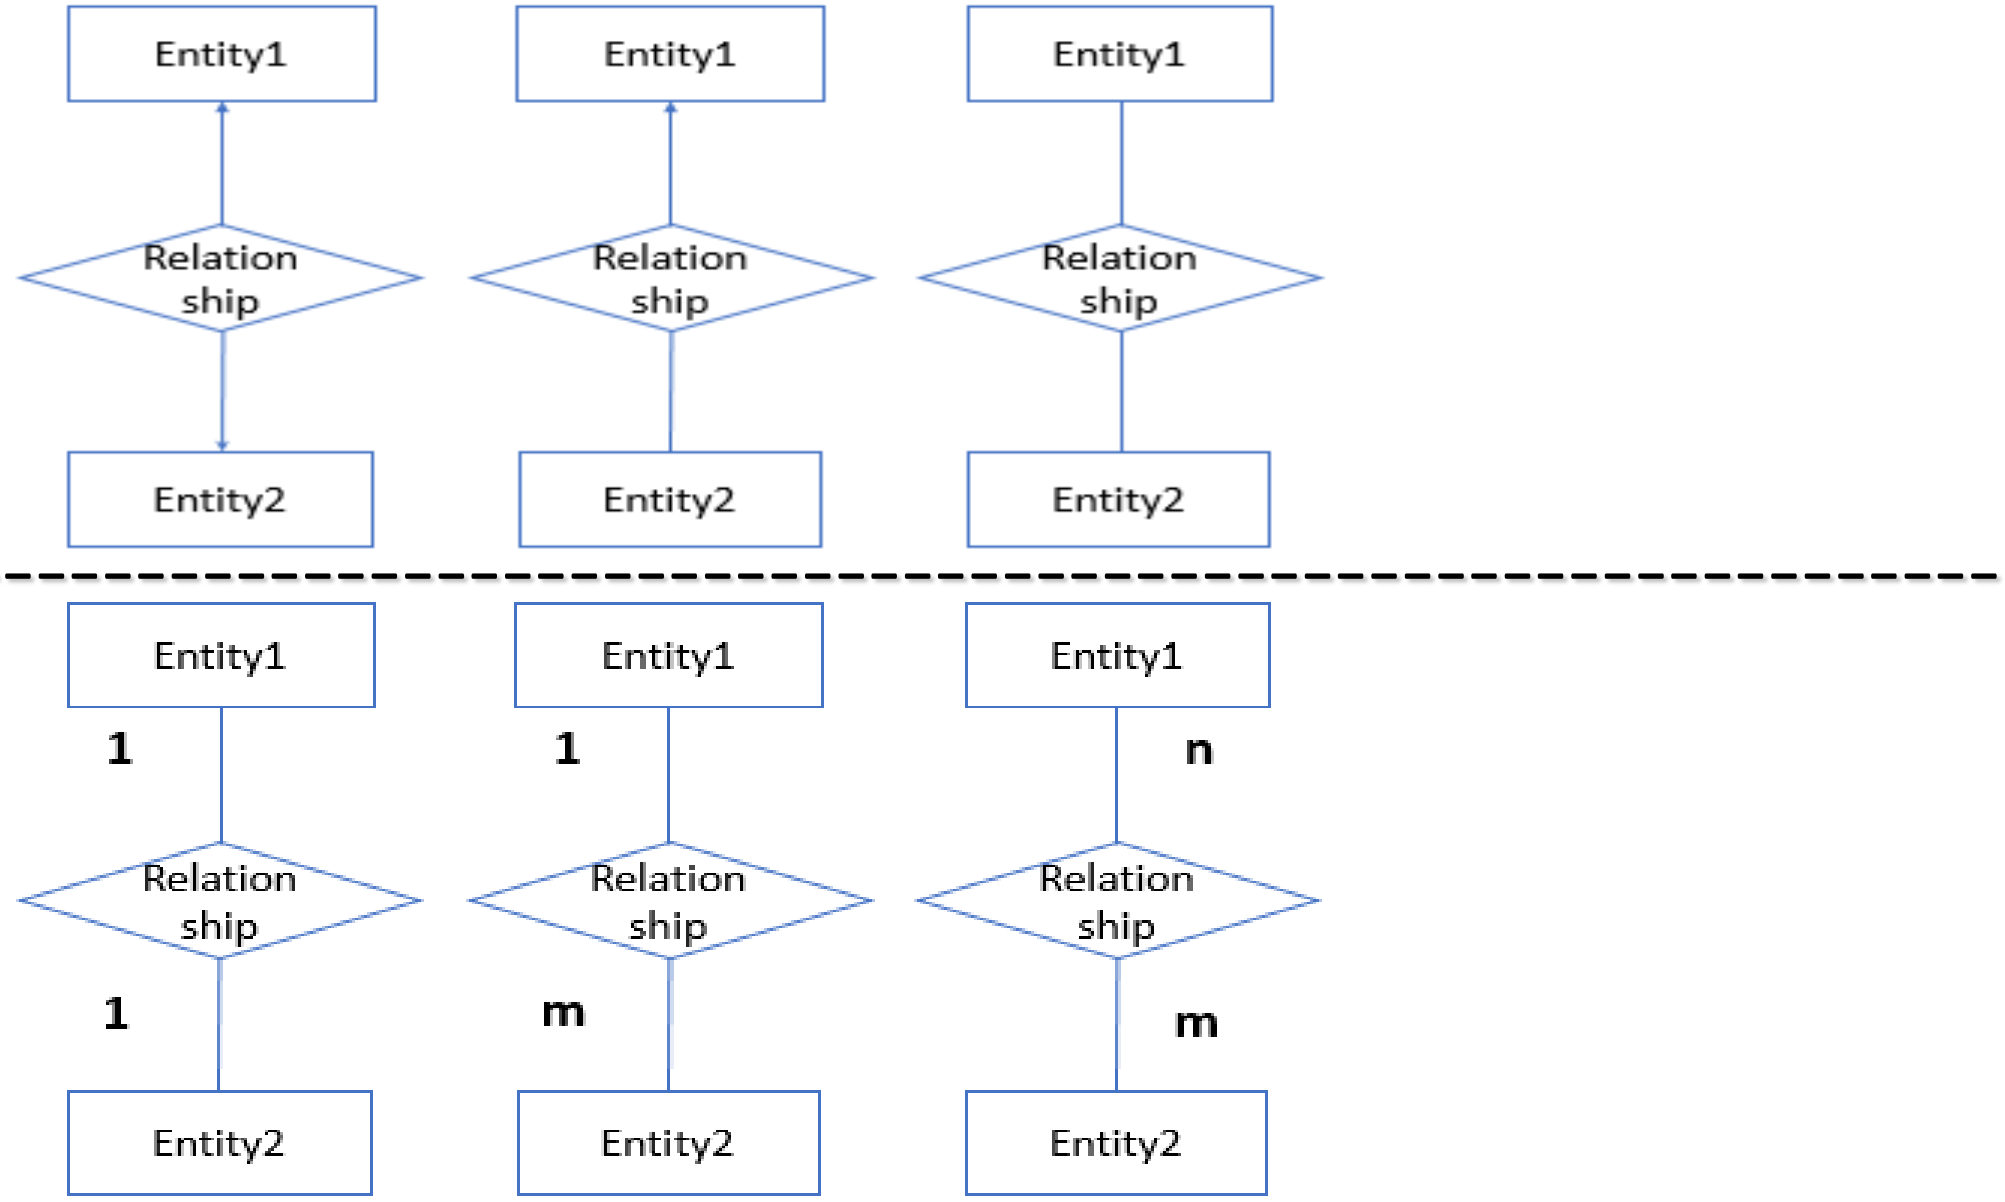

Example – cardinality

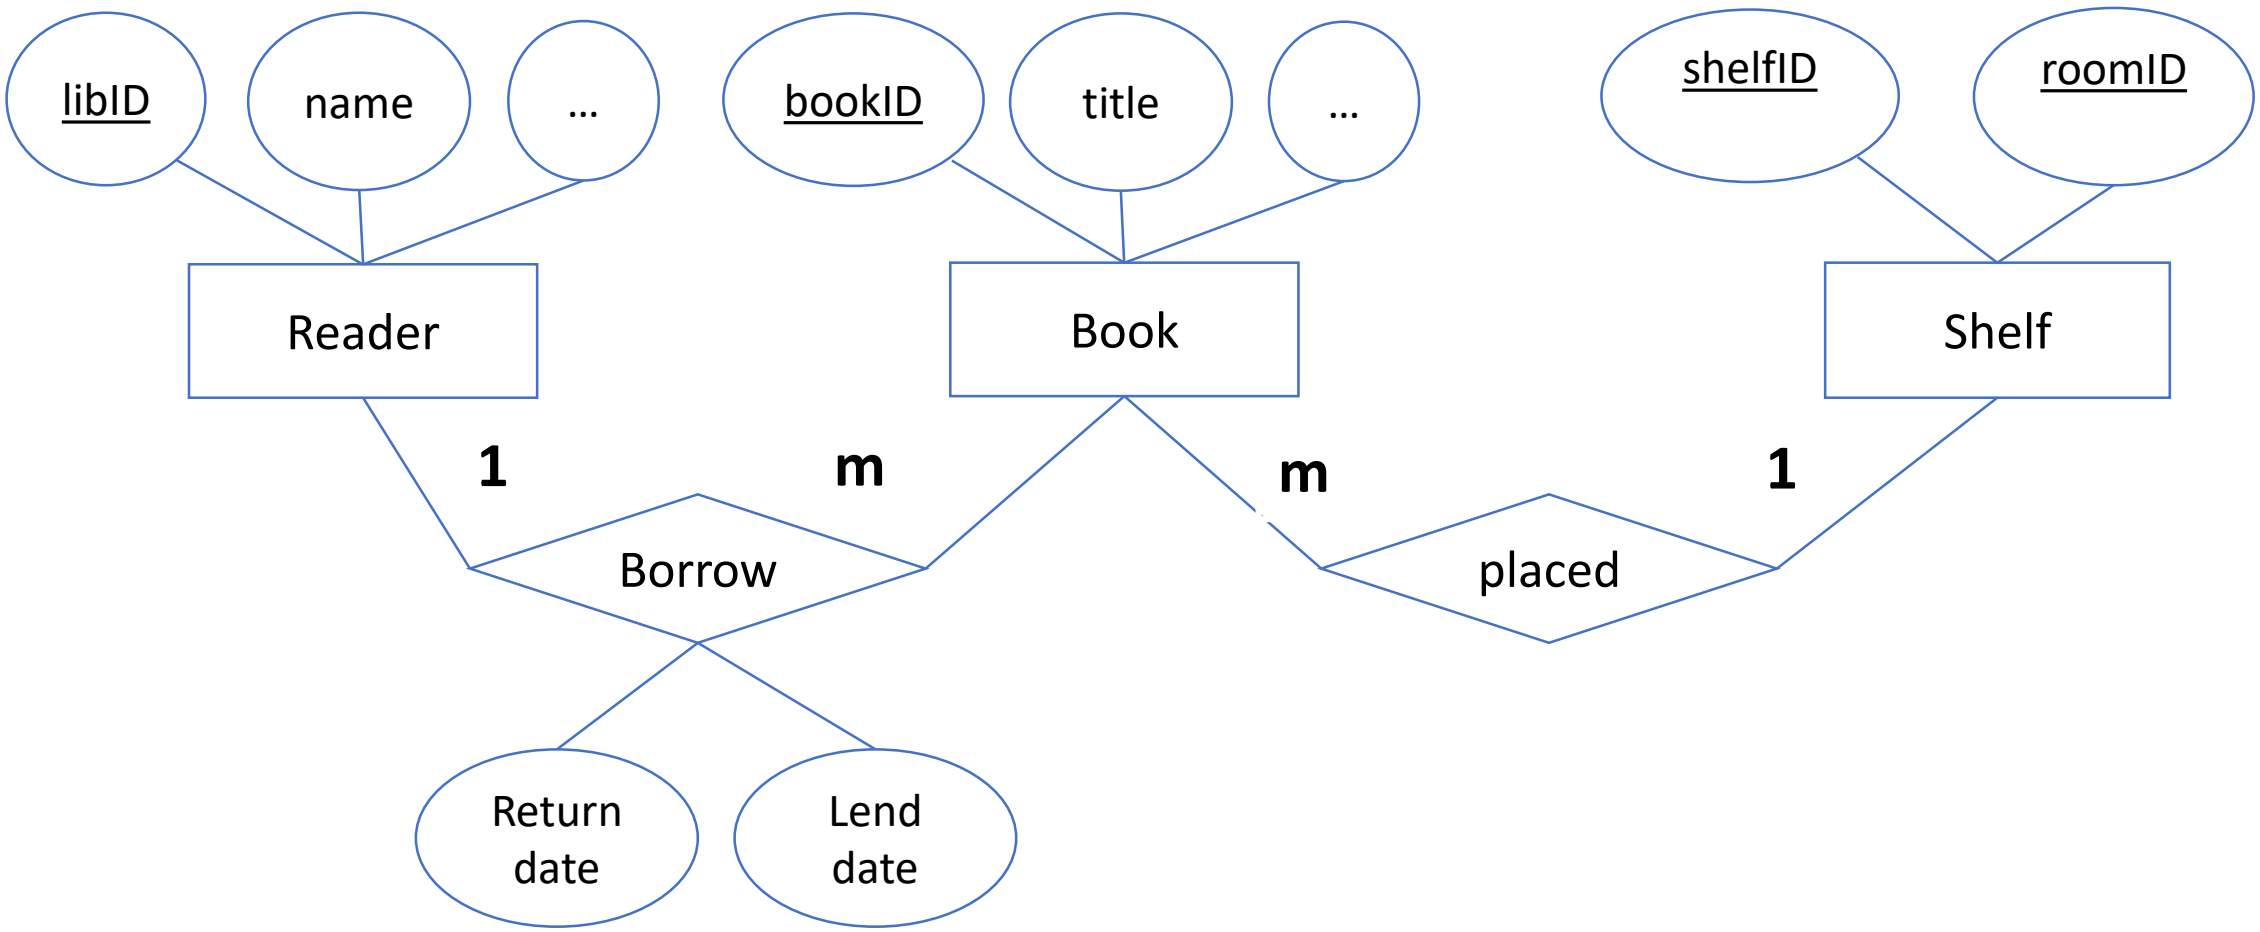

Example – **role**

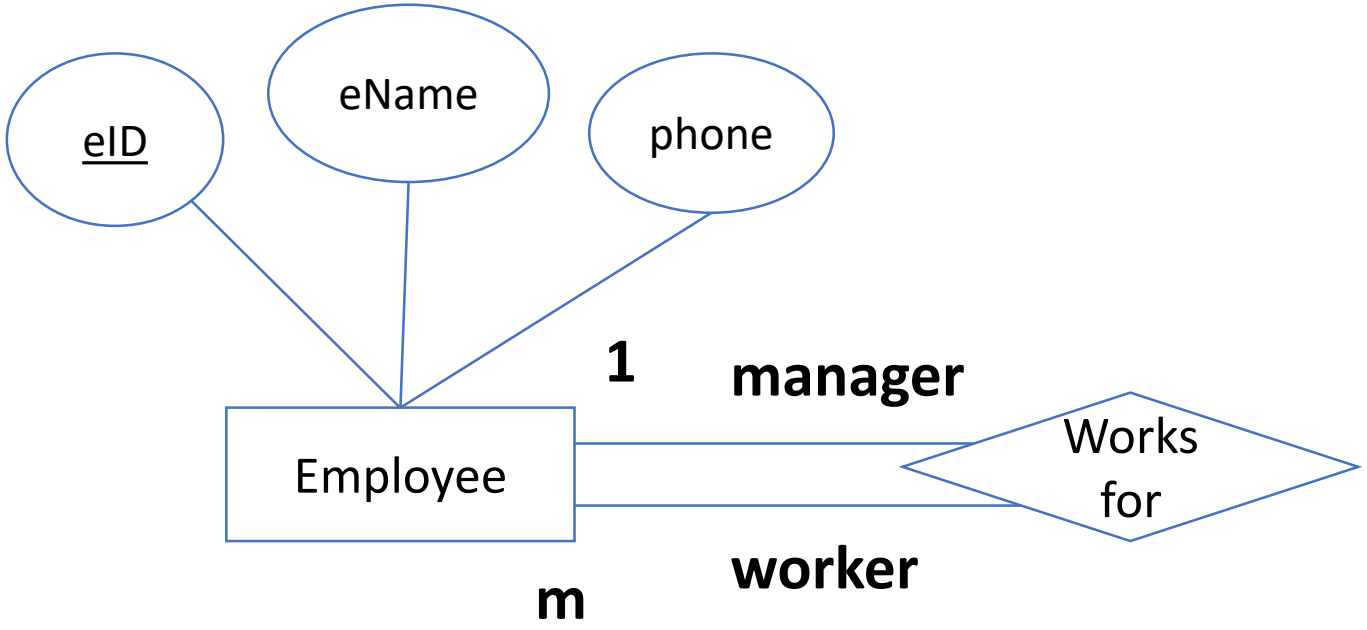

Alternative ER Notations

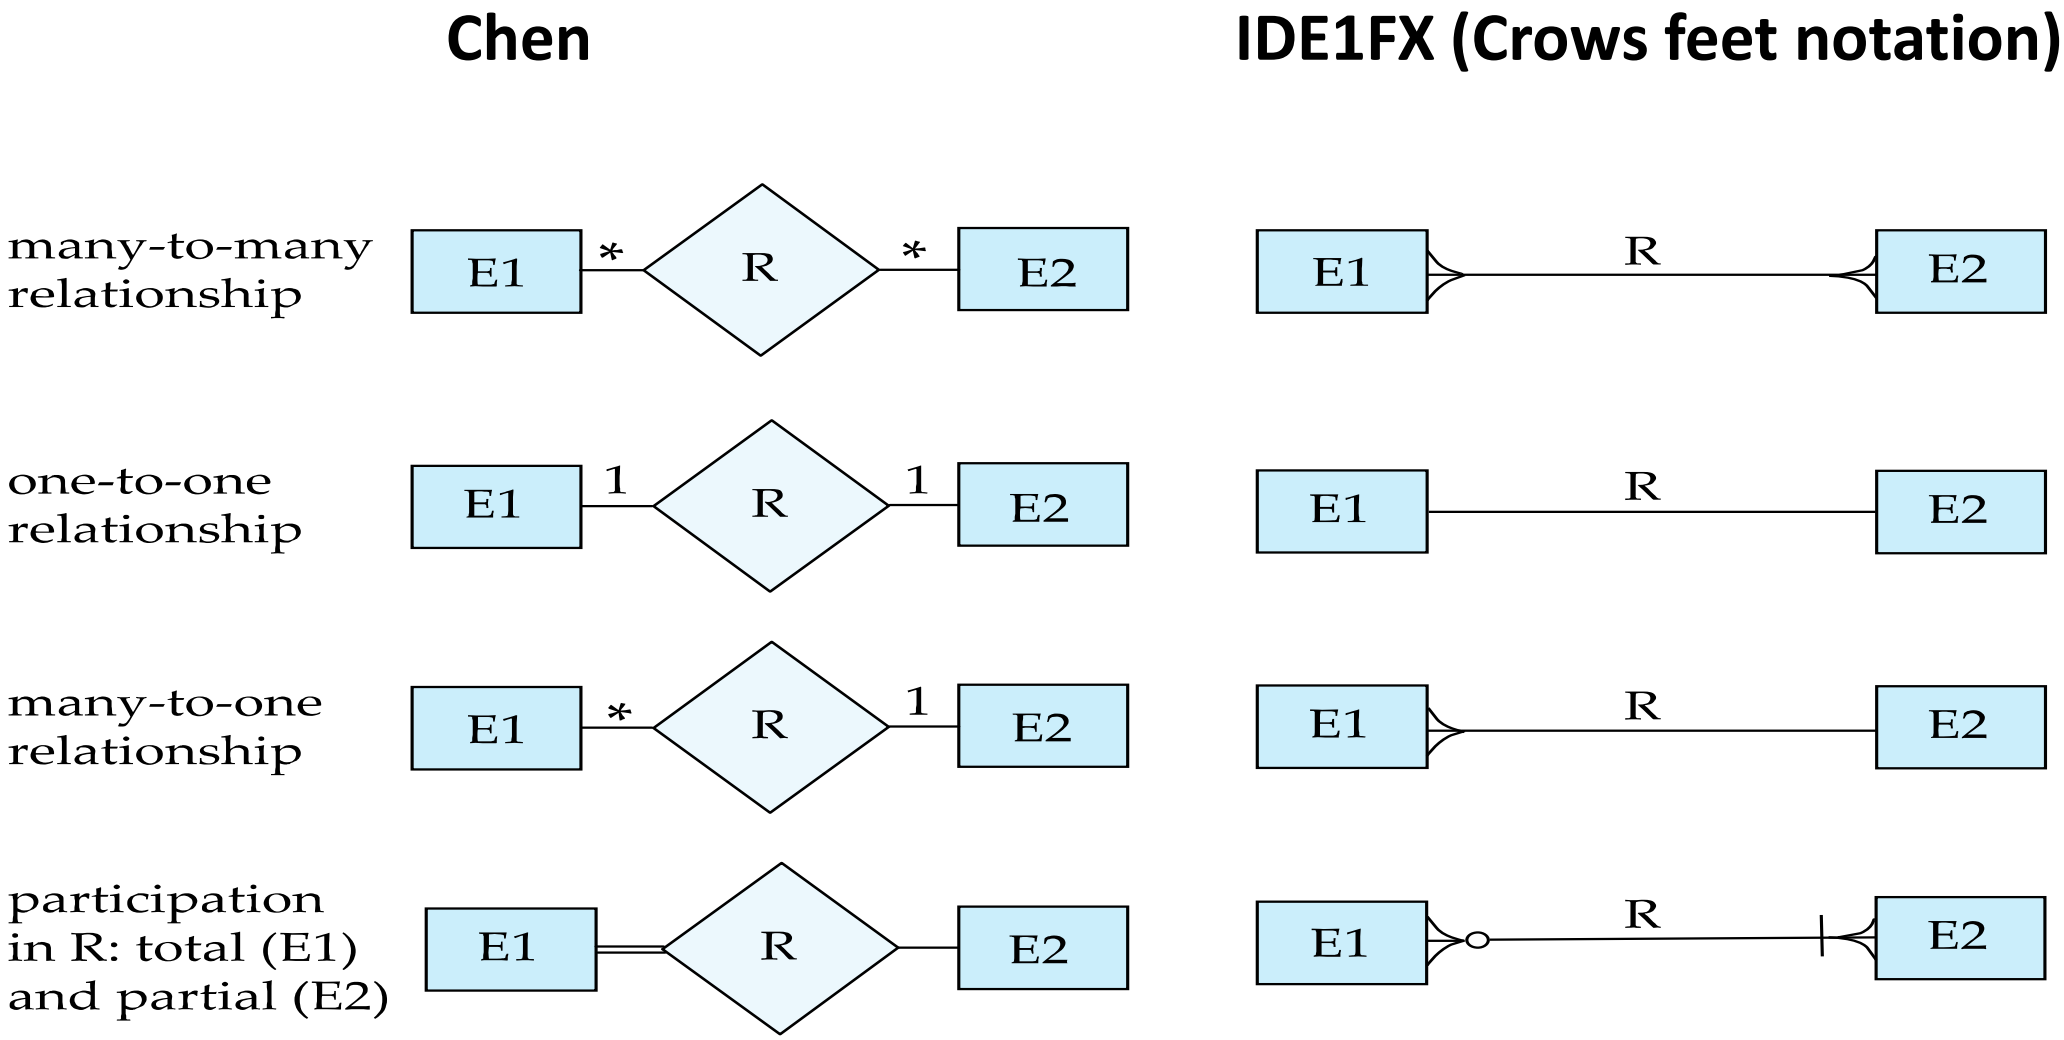

Case study – Inventory management

- Parts
- Supplier
  - Supply parts
- Project
  - Use parts
- Warehouse
  - Store parts
- Warehouse keeper
  - Manage warehouse
  - One keeper manages multiple warehouses

- Parts
- Supplier
- Project
- Warehouse
- Warehouse keeper

Step 1 – determine Entity sets

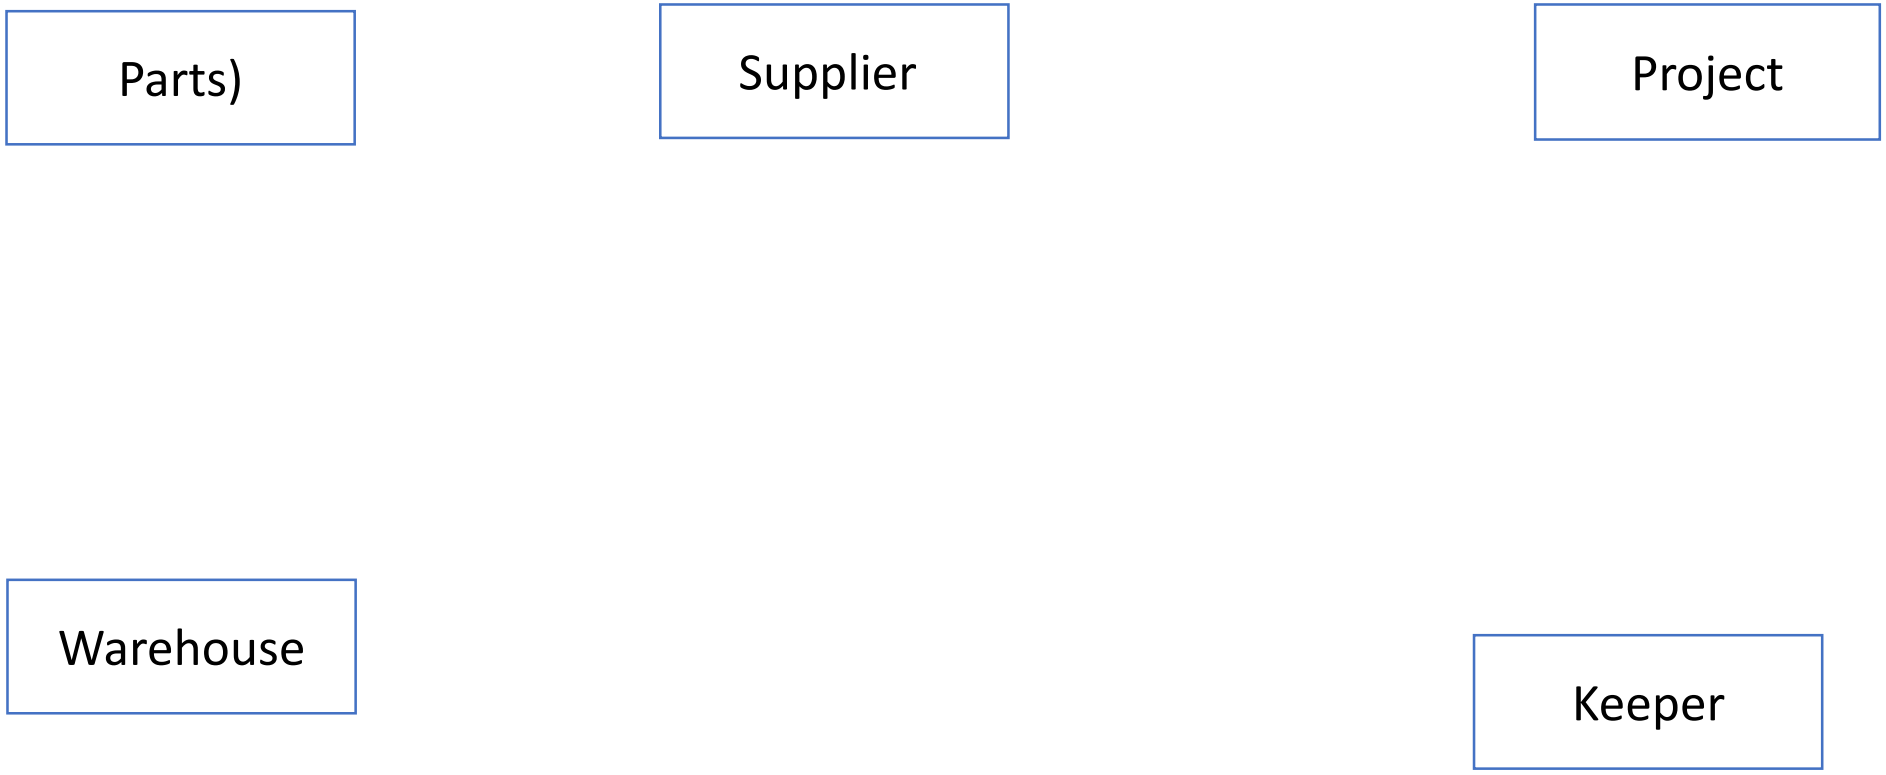

- Parts
- Supplier
- Project
- Warehouse
- Warehouse keeper

Step 2 – determine Attributes

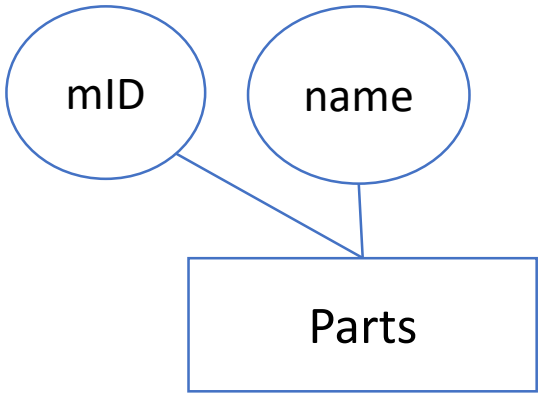

- Parts
- Supplier
- Project
- Warehouse
- Warehouse keeper

Step 2 – determine Attributes

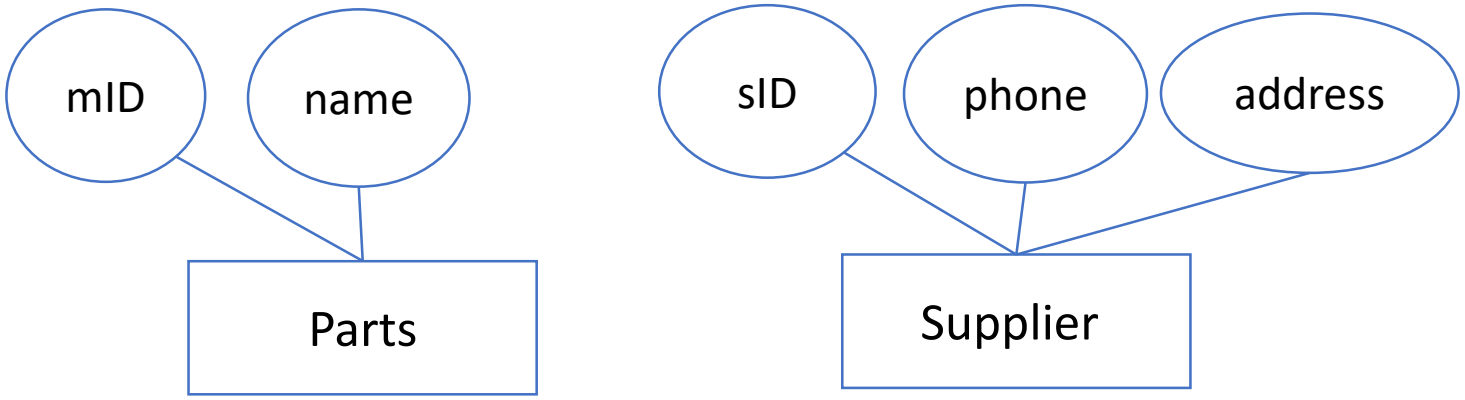

- Parts
- Supplier
- Project
- Warehouse
- Warehouse keeper

Step 2 – determine Attributes

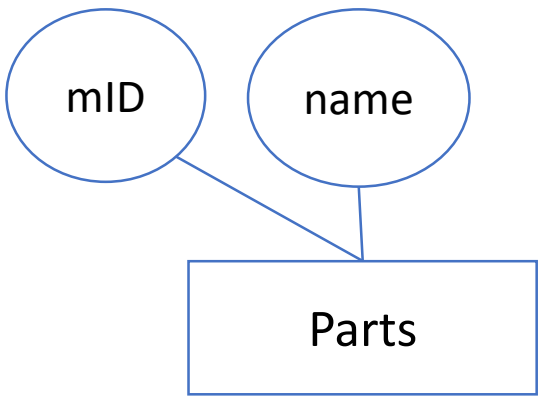

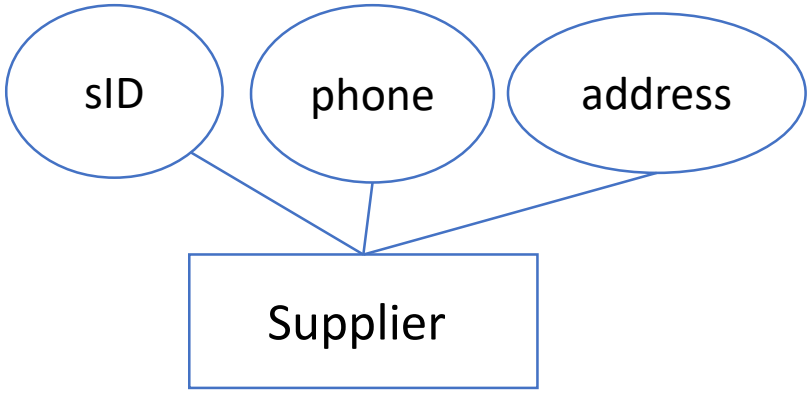

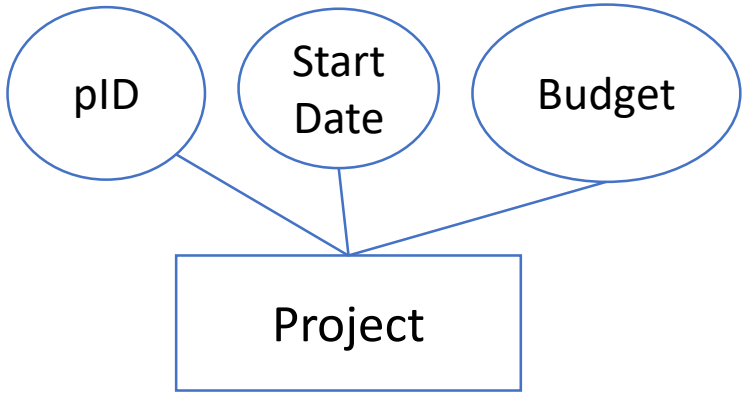

- Parts
- Supplier
- Project
- Warehouse
- Warehouse keeper

Step 2 – determine Attributes

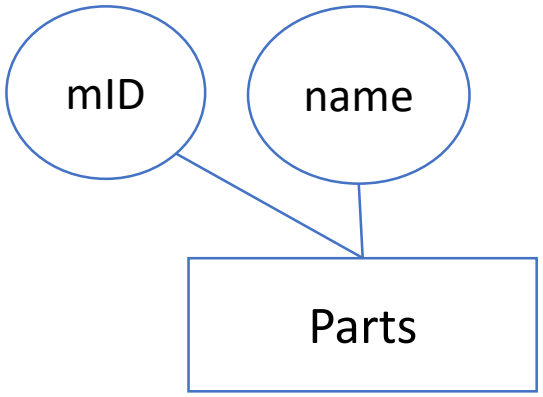

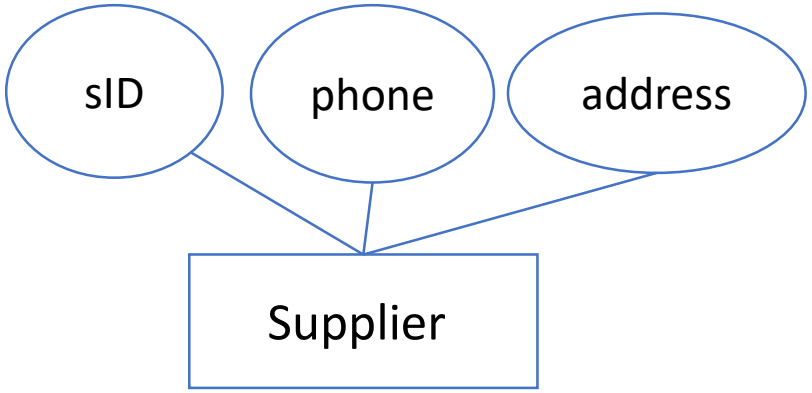

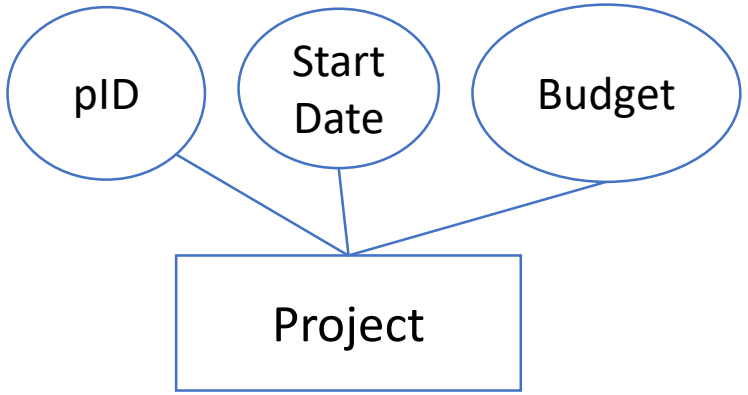

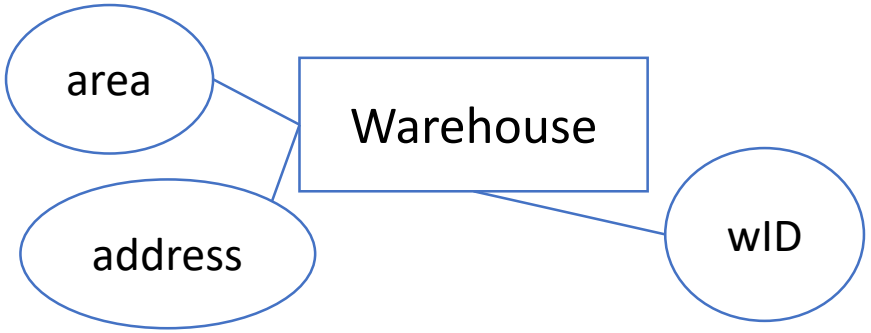

- Parts
- Supplier
- Project
- Warehouse
- Warehouse keeper

Step 2 – determine Attributes

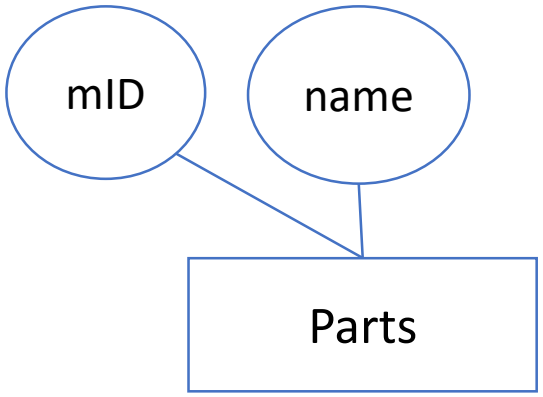

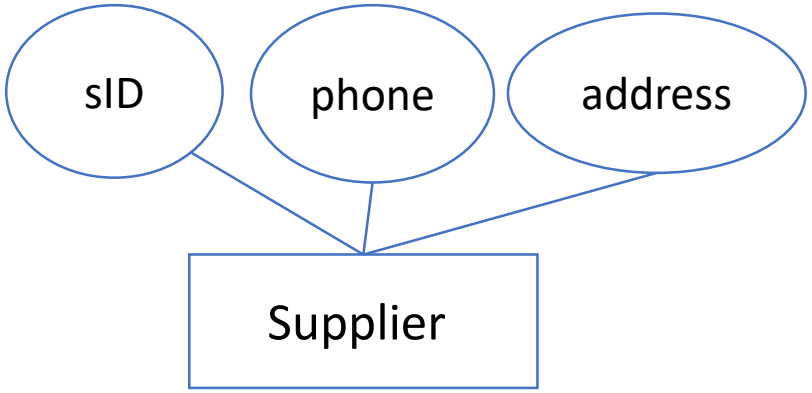

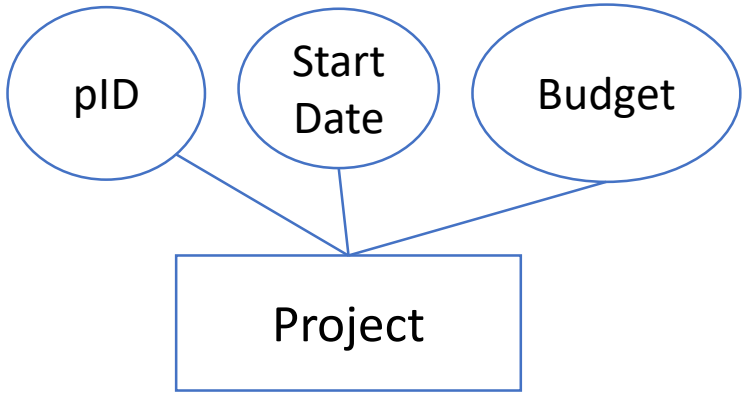

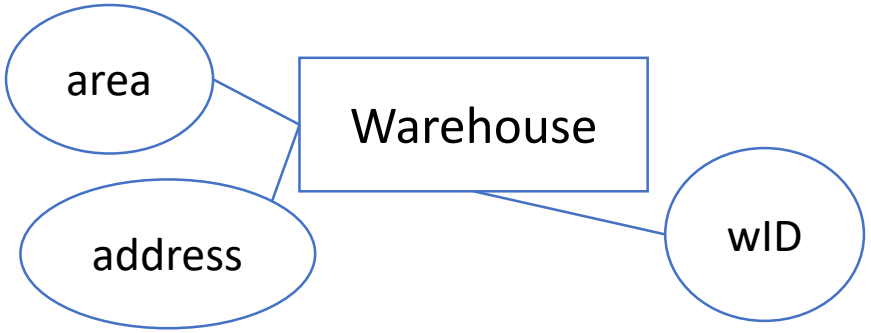

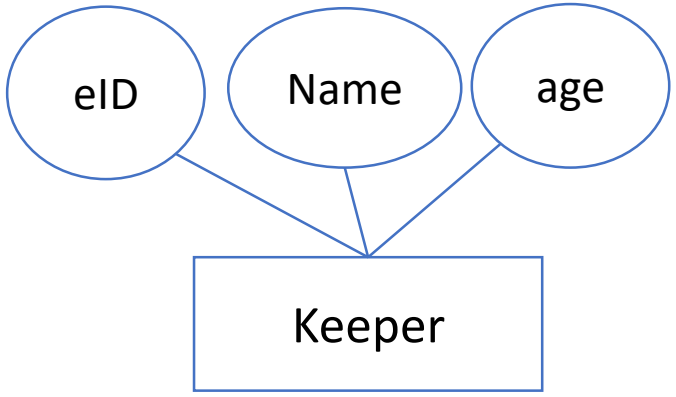

- Parts
- Supplier
- Project
- Warehouse
- Warehouse keeper

Step 3 – determine keys

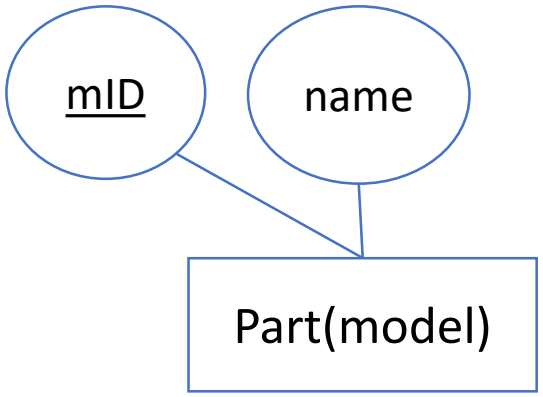

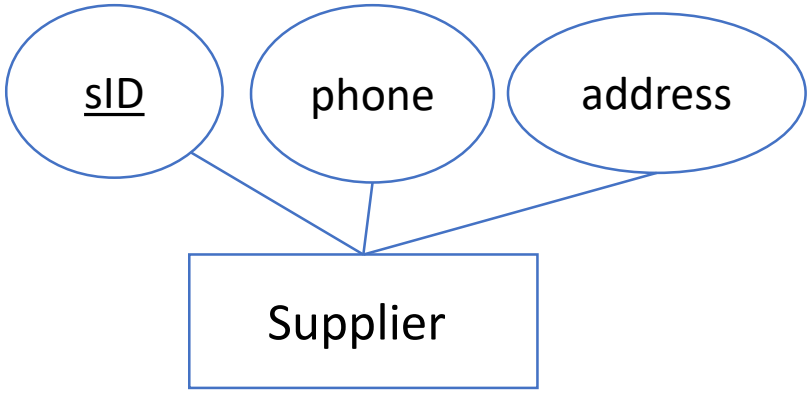

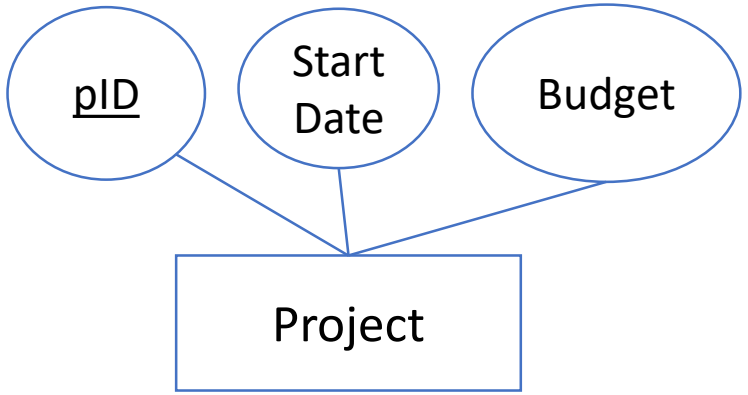

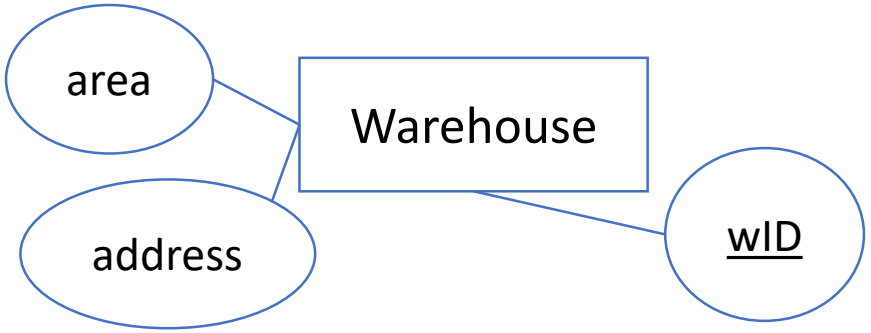

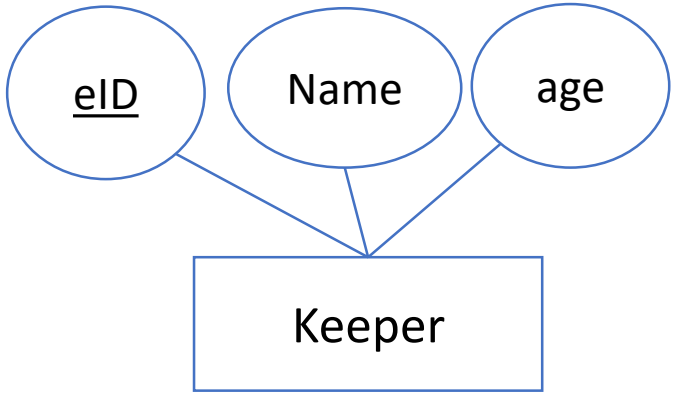

- Parts
- Supplier
- Project
- Warehouse
- Warehouse keeper

Step 4 – establish relationships

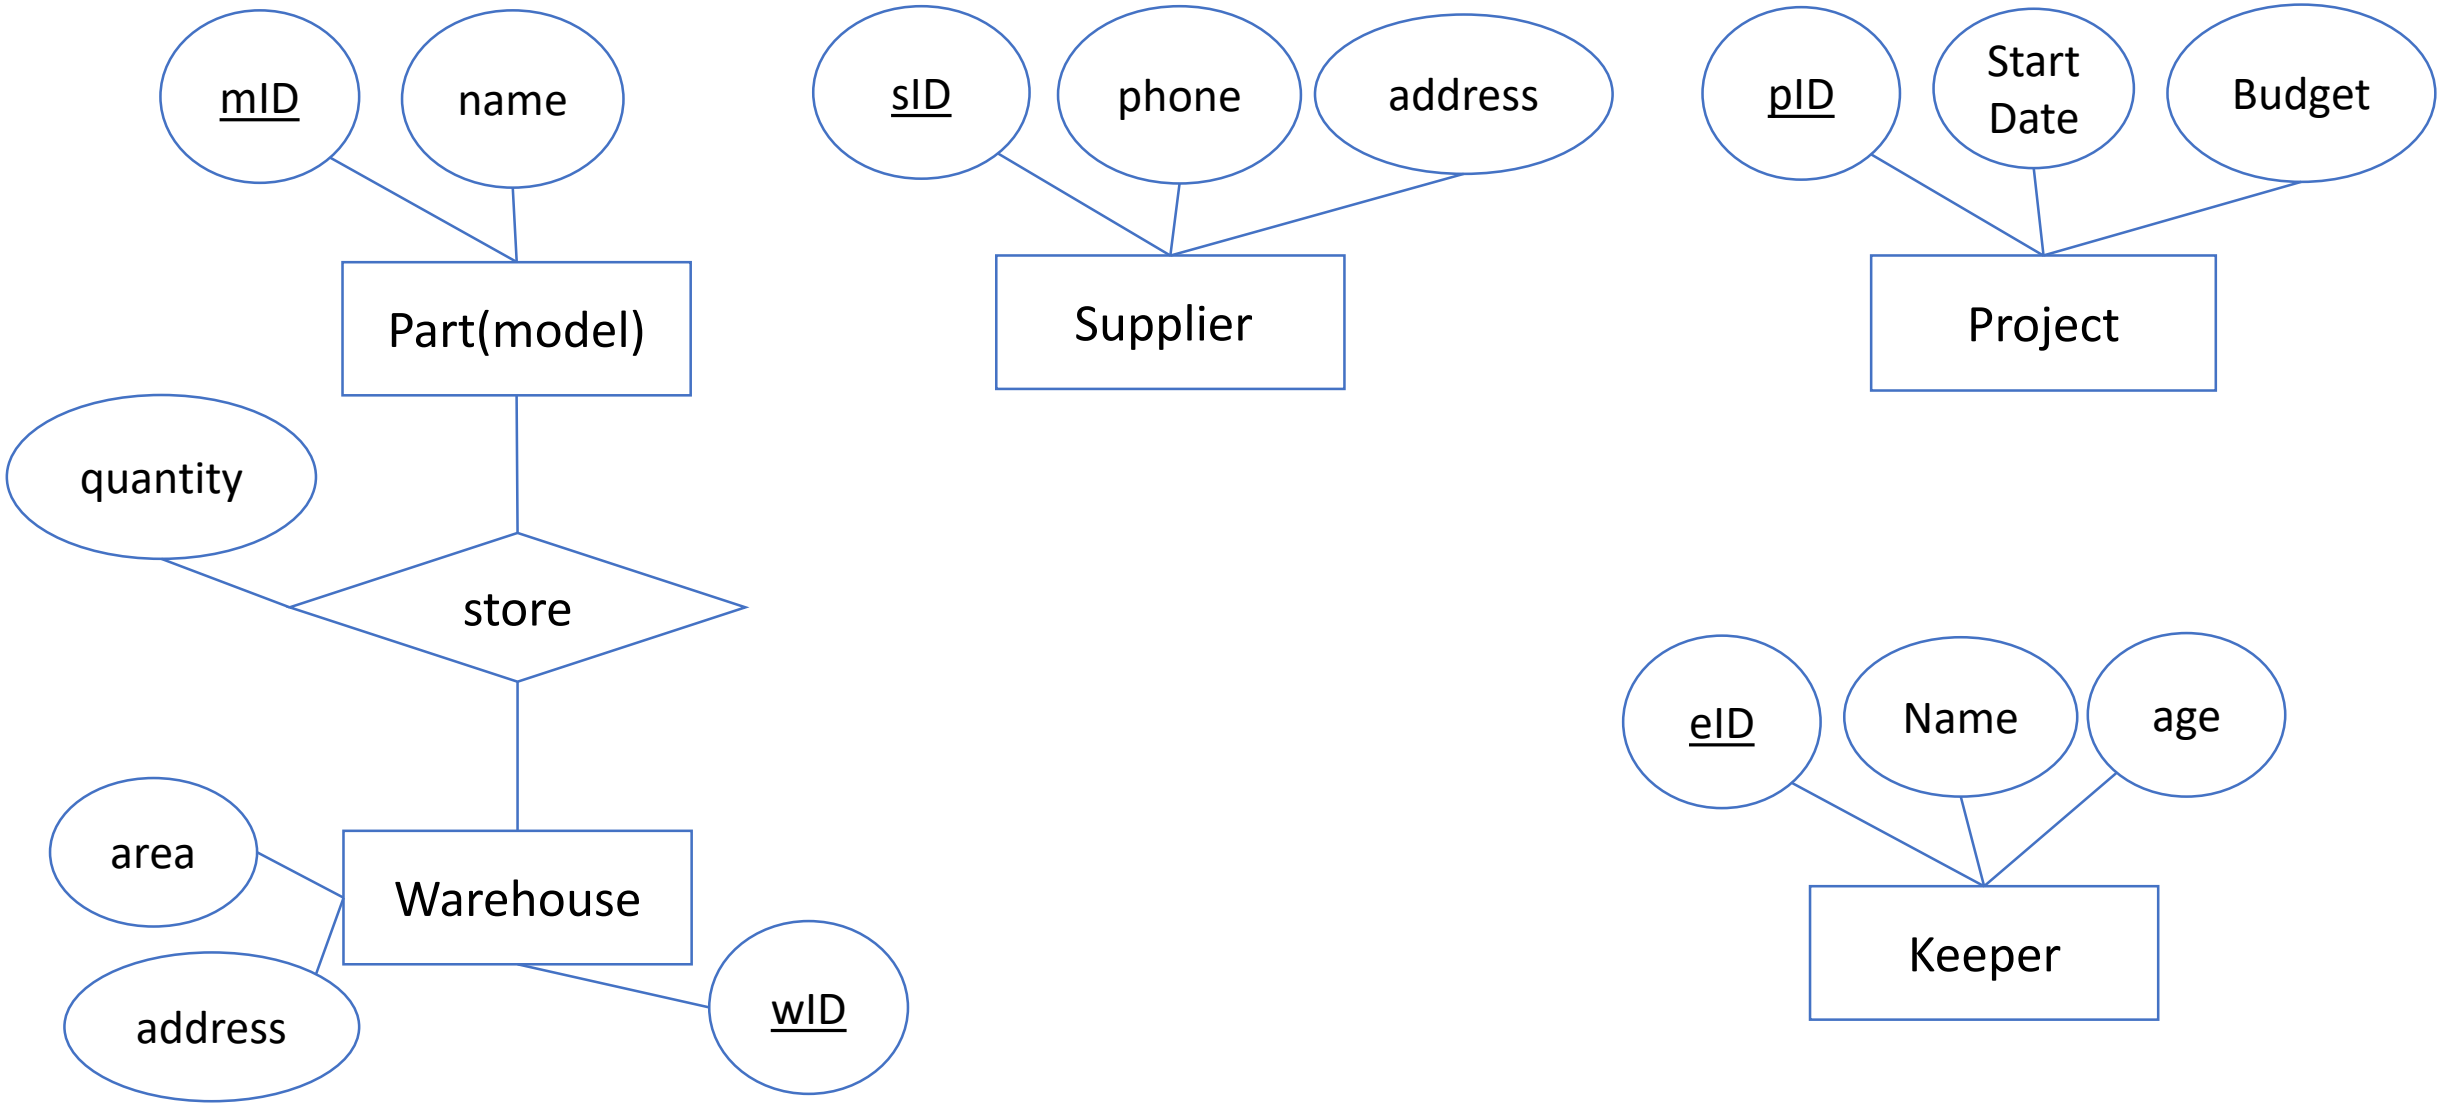

- Parts
- Supplier
- Project
- Warehouse
- Warehouse keeper

Step 4 – establish relationships

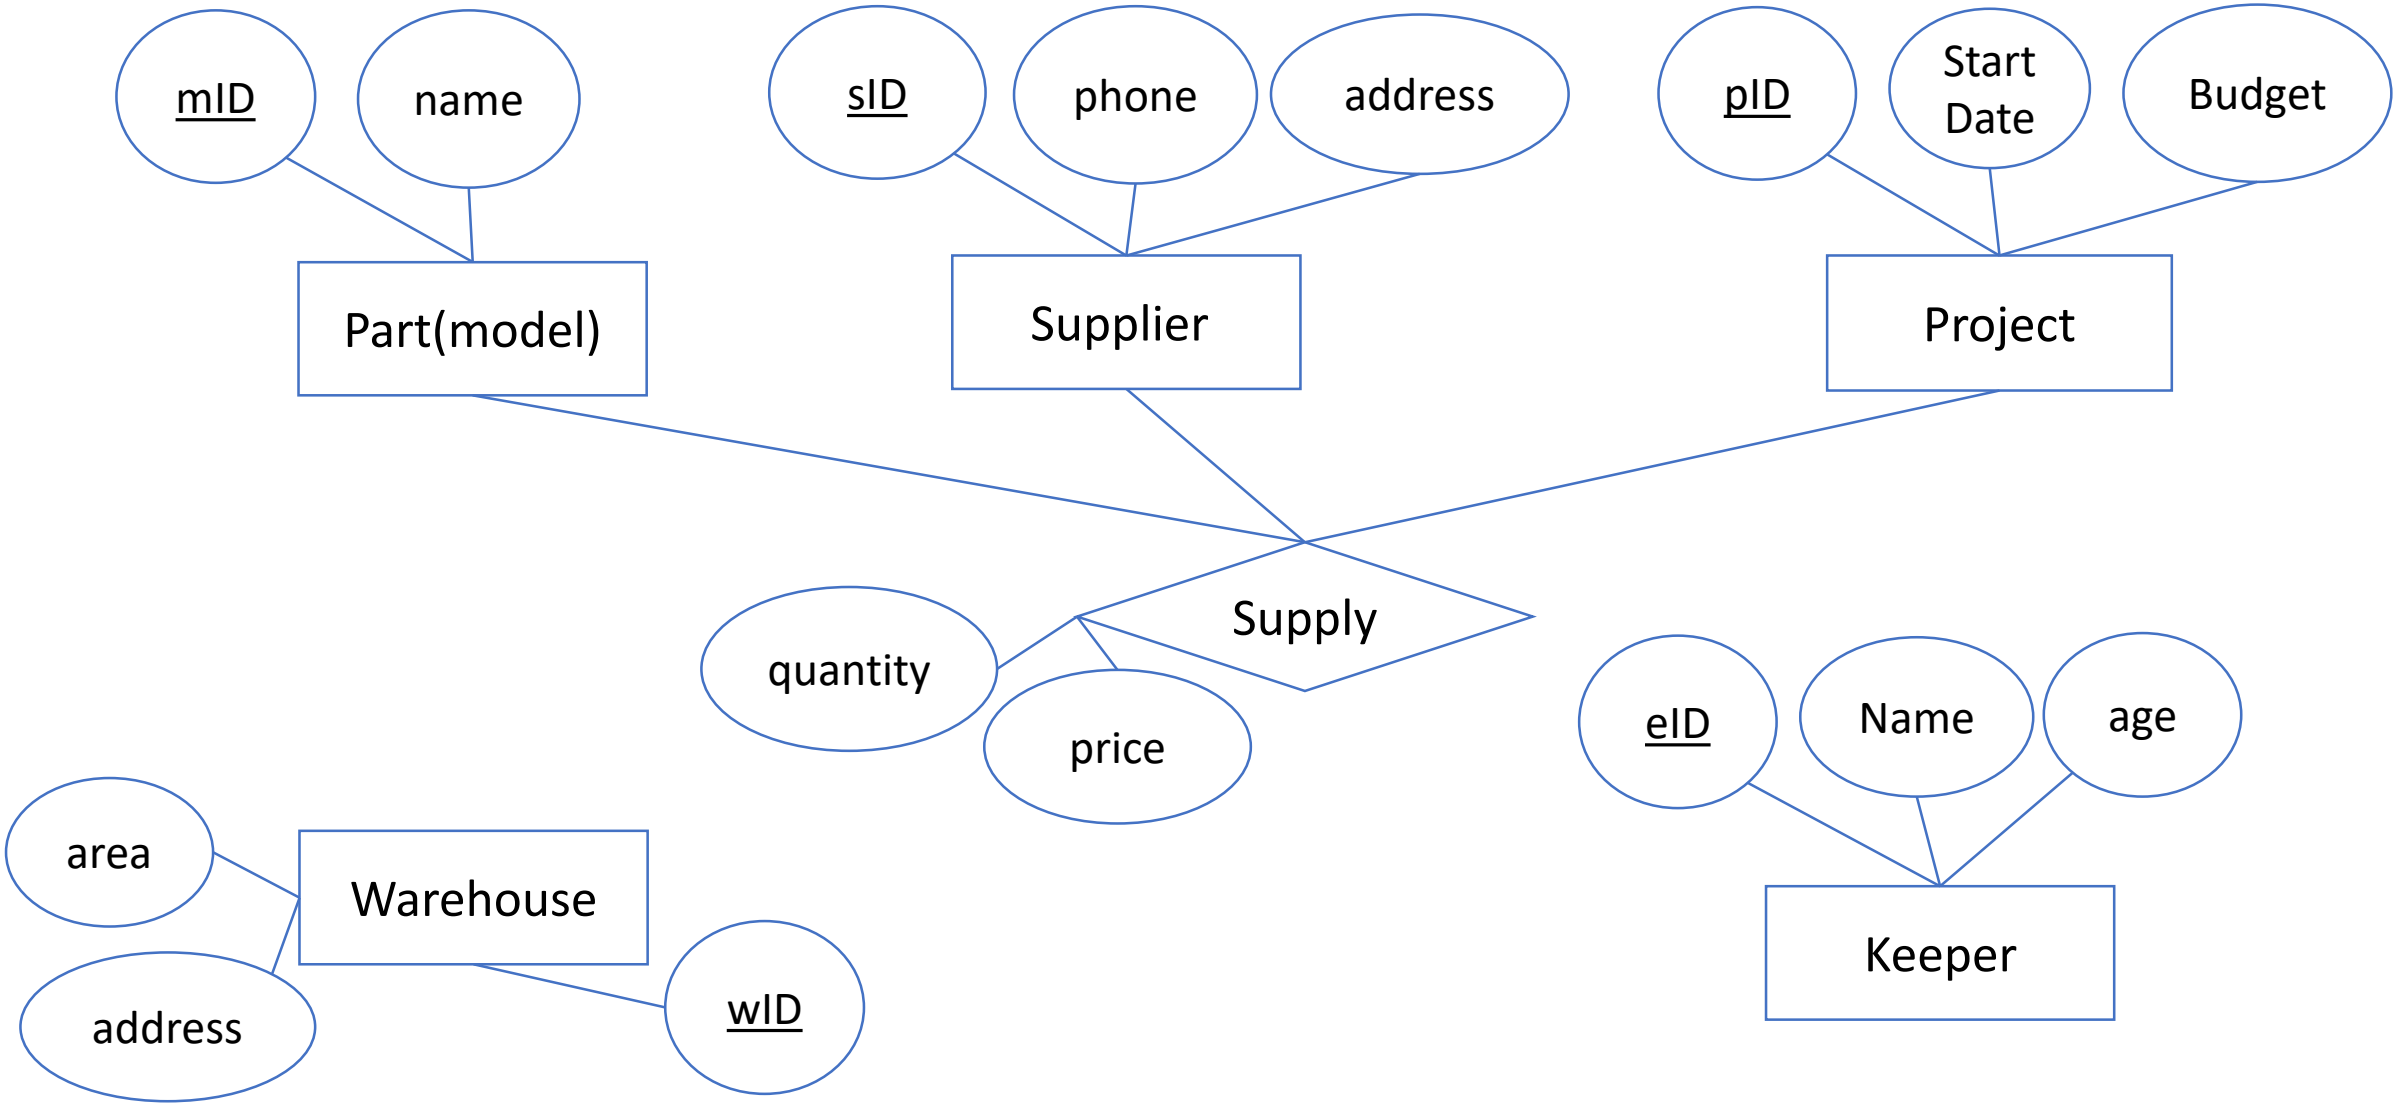

- Parts
- Supplier
- Project
- Warehouse
- Warehouse keeper

Step 4 – establish relationships

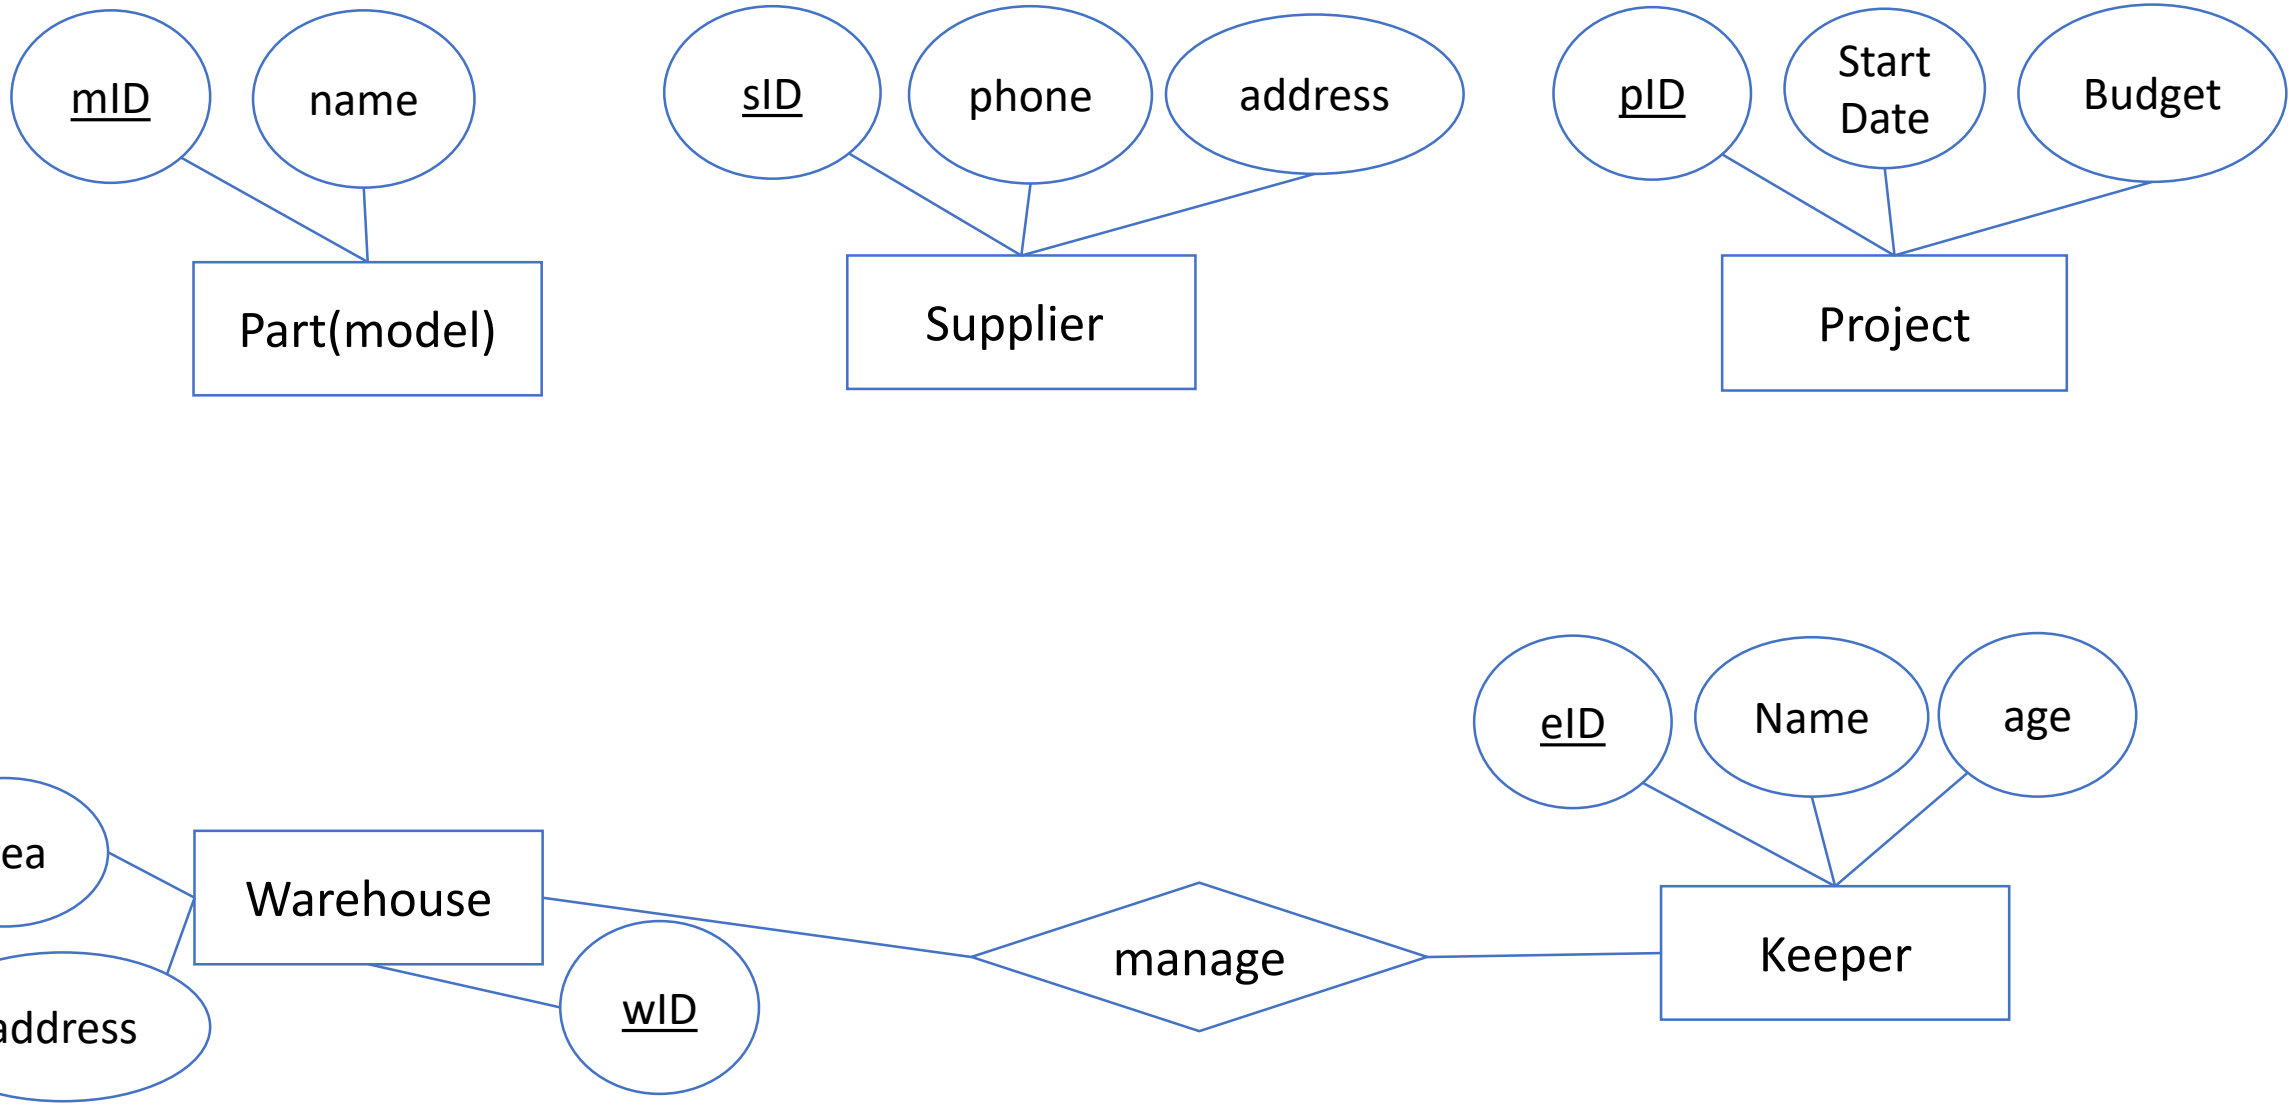

- Parts
- Supplier
- Project
- Warehouse
- Warehouse keeper

Step 4 – establish relationships

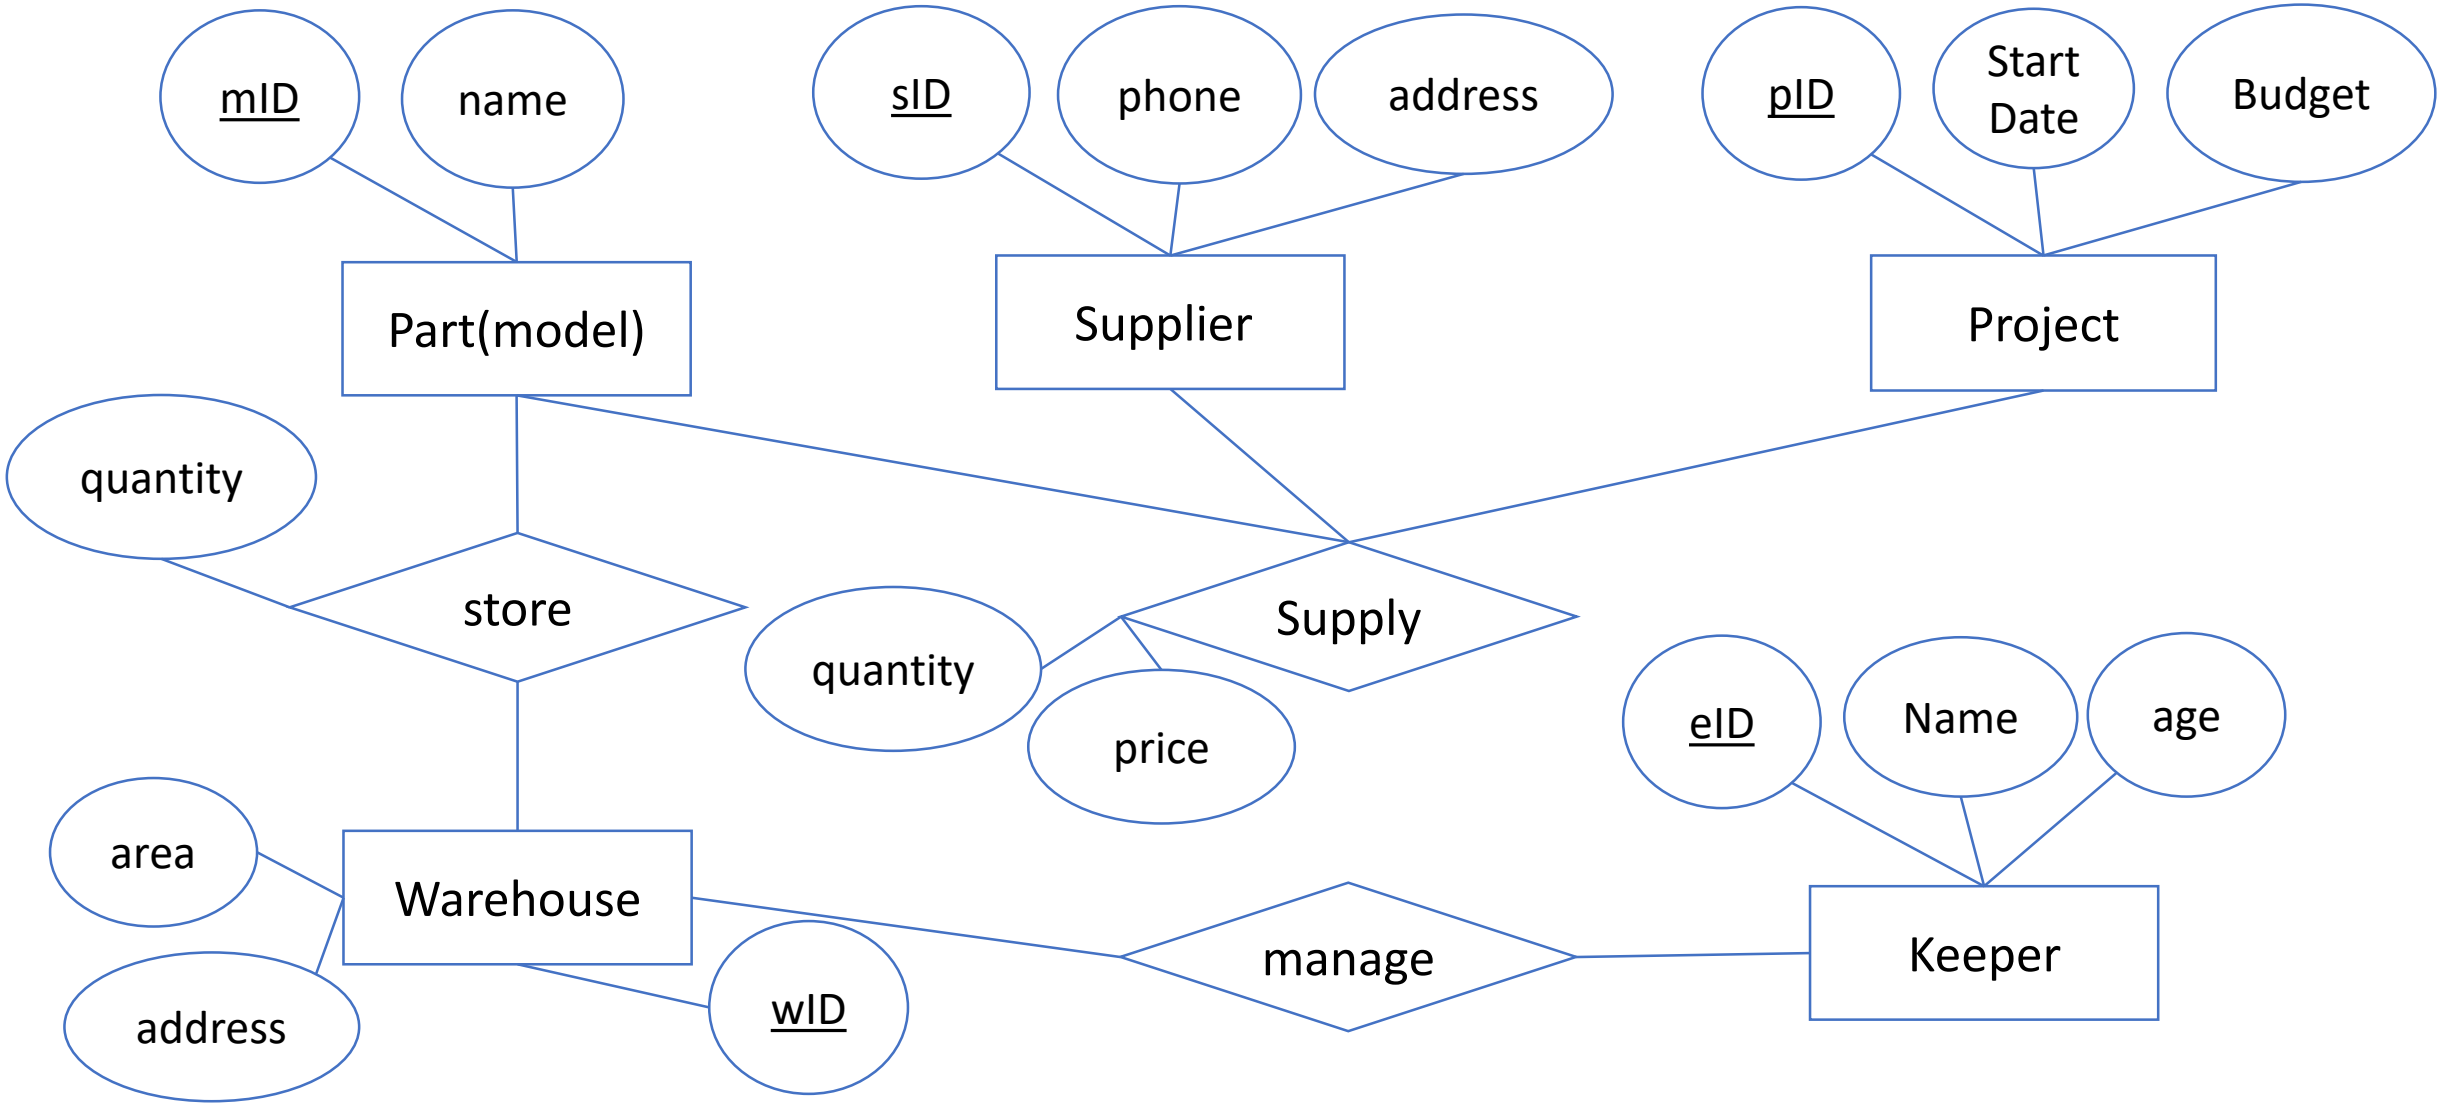

- Parts
- Supplier
- Project
- Warehouse
- Warehouse keeper

Step 5 – determine cardinality

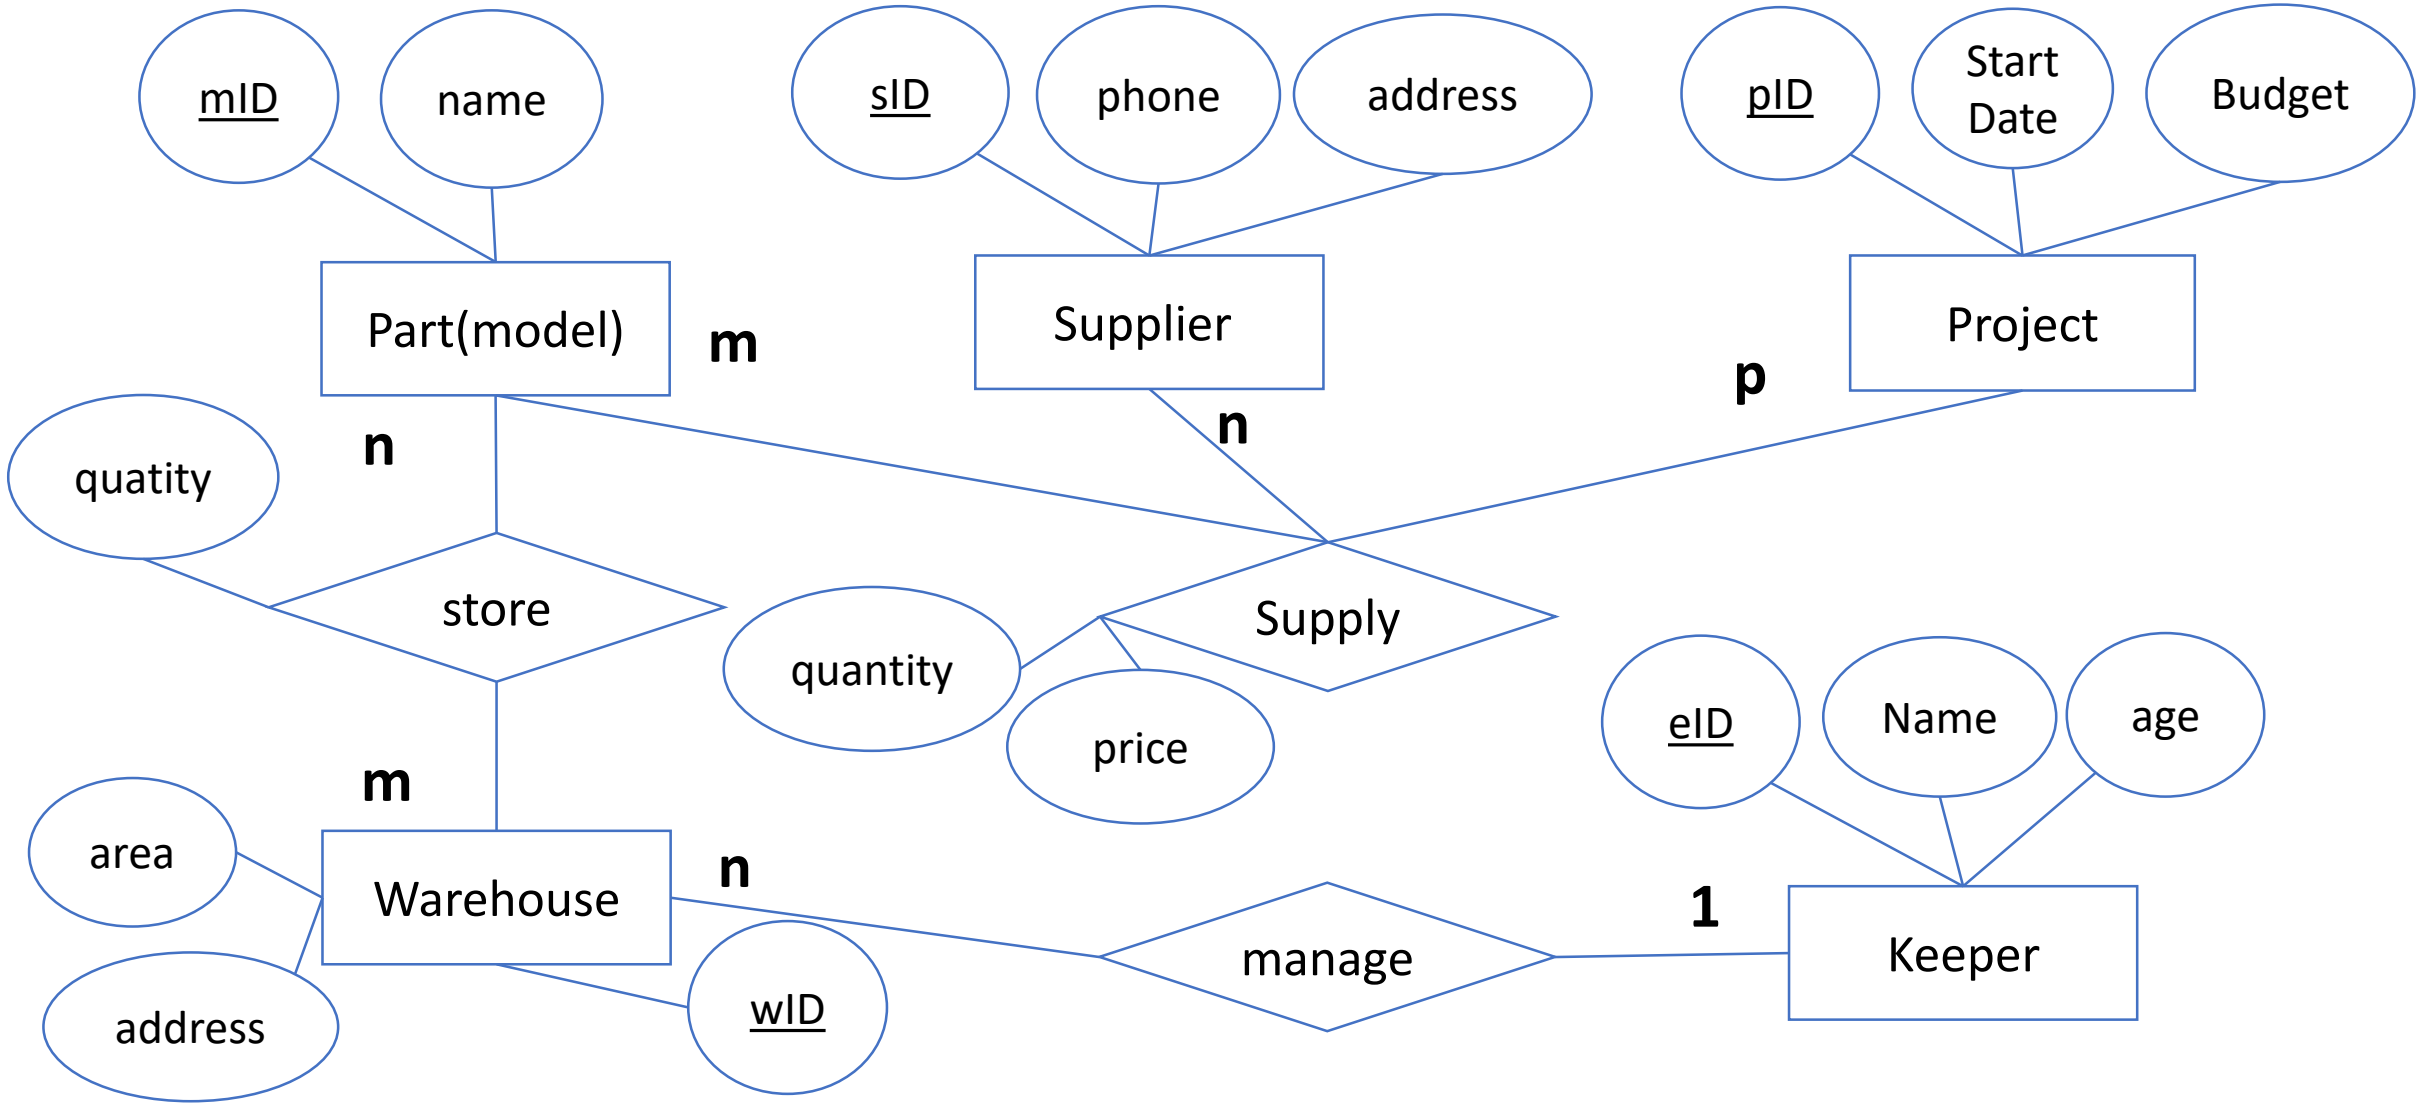

Database for TV Companies (SELF STUDY)

- Store data about the TV series that the company produces
- Data include information about actors who play in the series, and directors who direct the episodes of the series
  - ❑Actors and directors are employed by the company
  - ❑A TV series are divided into episodes
  - ❑Each episode may be transmitted at several occasions
  - ❑An actor is hired to participate in a series, but may participate in many series
  - ❑Each episode of a series is directed by one of the directors, but different episodes may be directed by different directors.

Examples of database queries

- Which actors play in the series The Night Of?
- In which series does the actor Carey Mulligan participate?
- Which actors participate in more than one series?
- How many times has the first episode of the series Wild Lies been transmitted? At what times?
- How many directors are employed by the company?
- Which directors has directed the greatest number of episodes?

Entity sets

- TV company
- TV series
- Episode
- Employee
- Actor
- Director

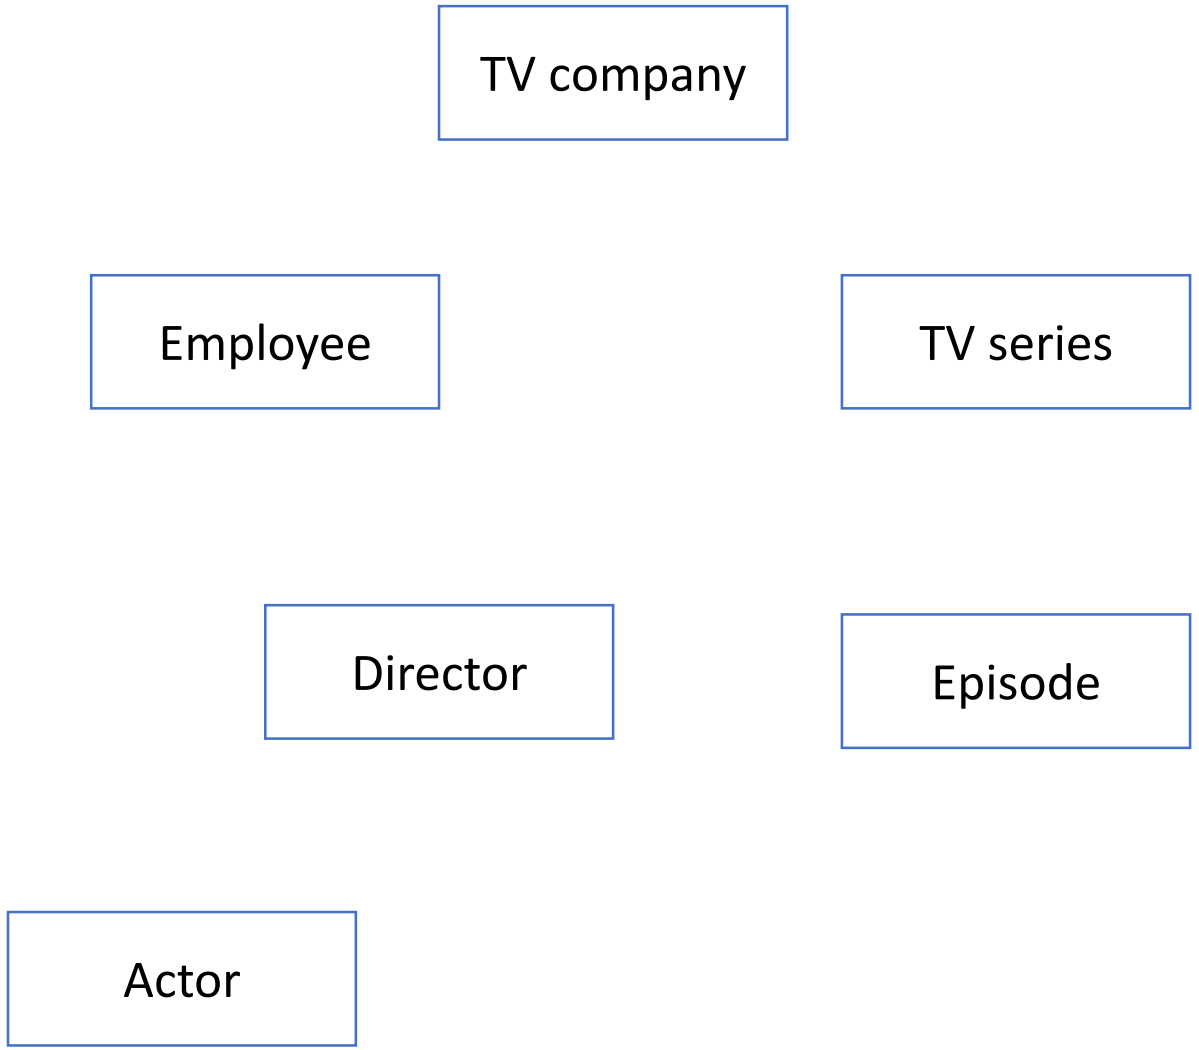

Attributes & Key

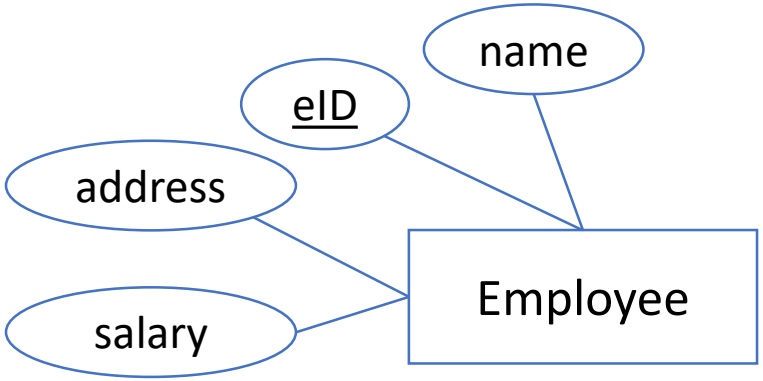

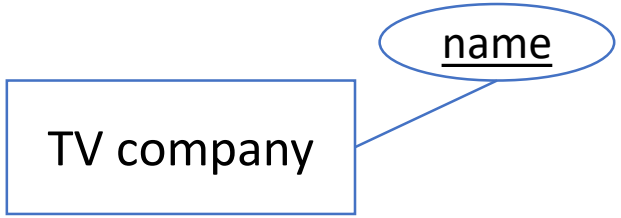

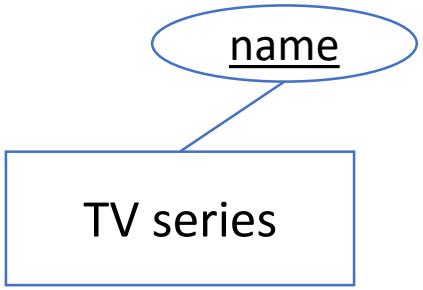

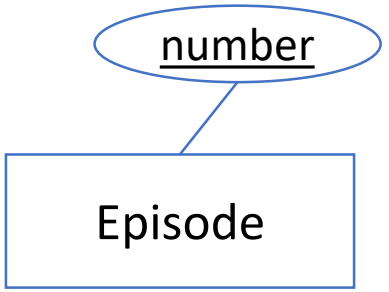

Relationship sets

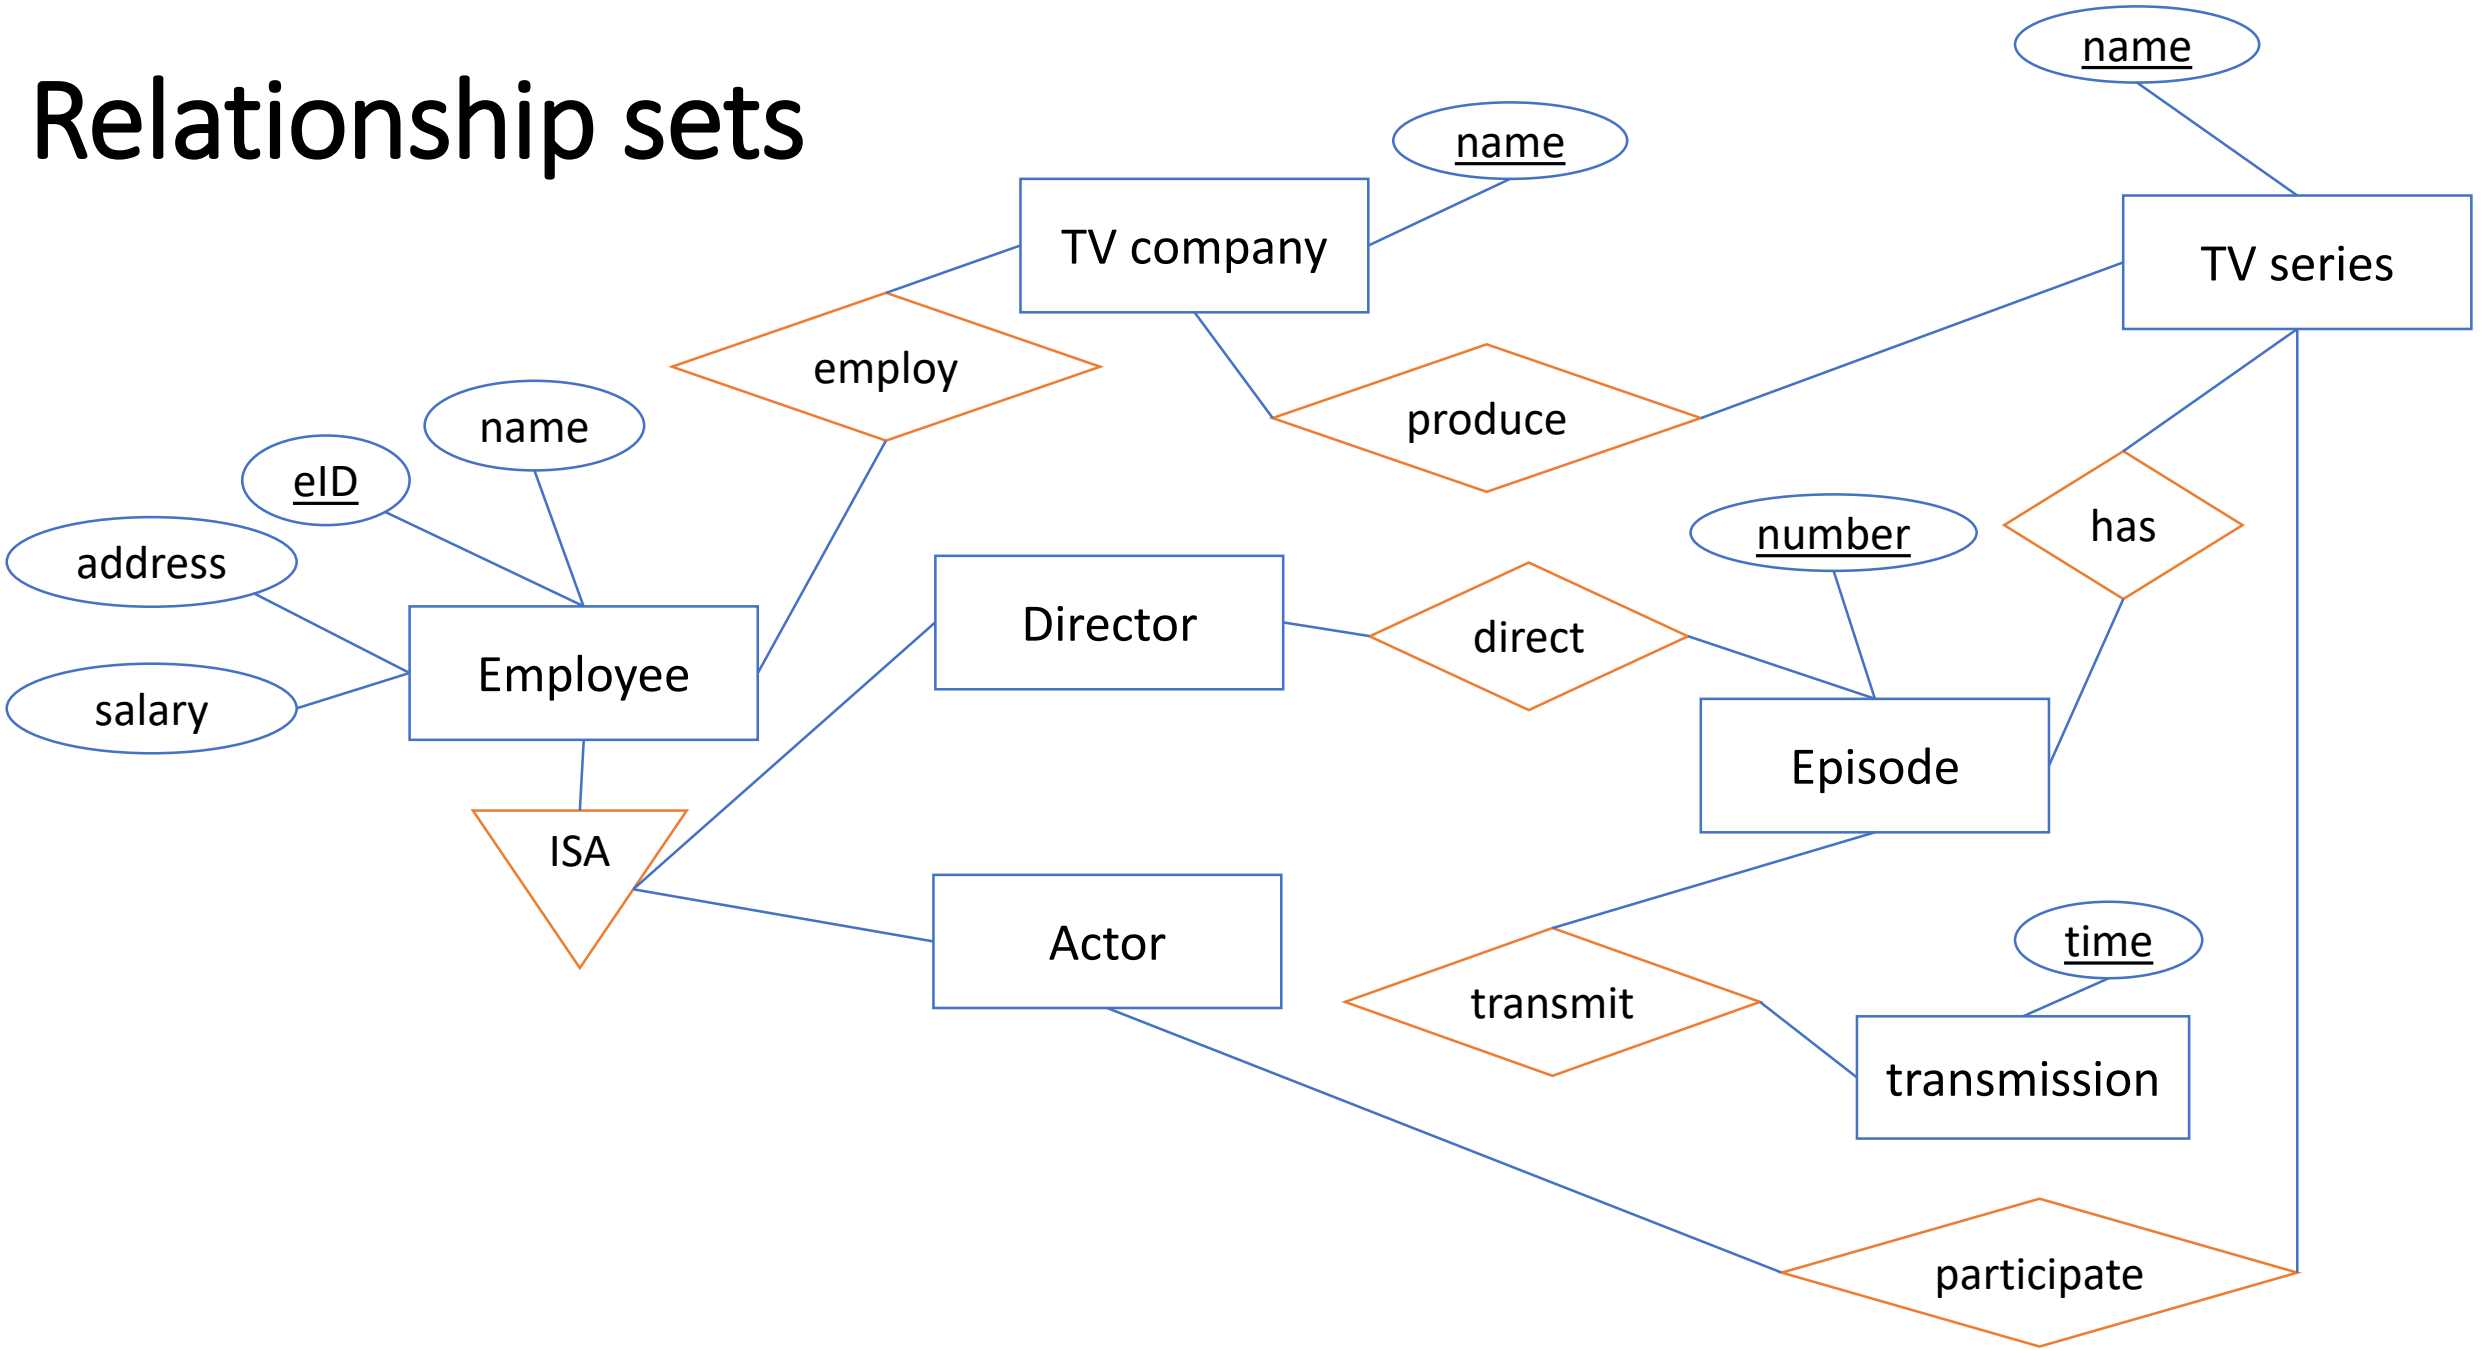

Cardinality ratio