

Fundamentals of Spatial Filtering and Smoothing Filters

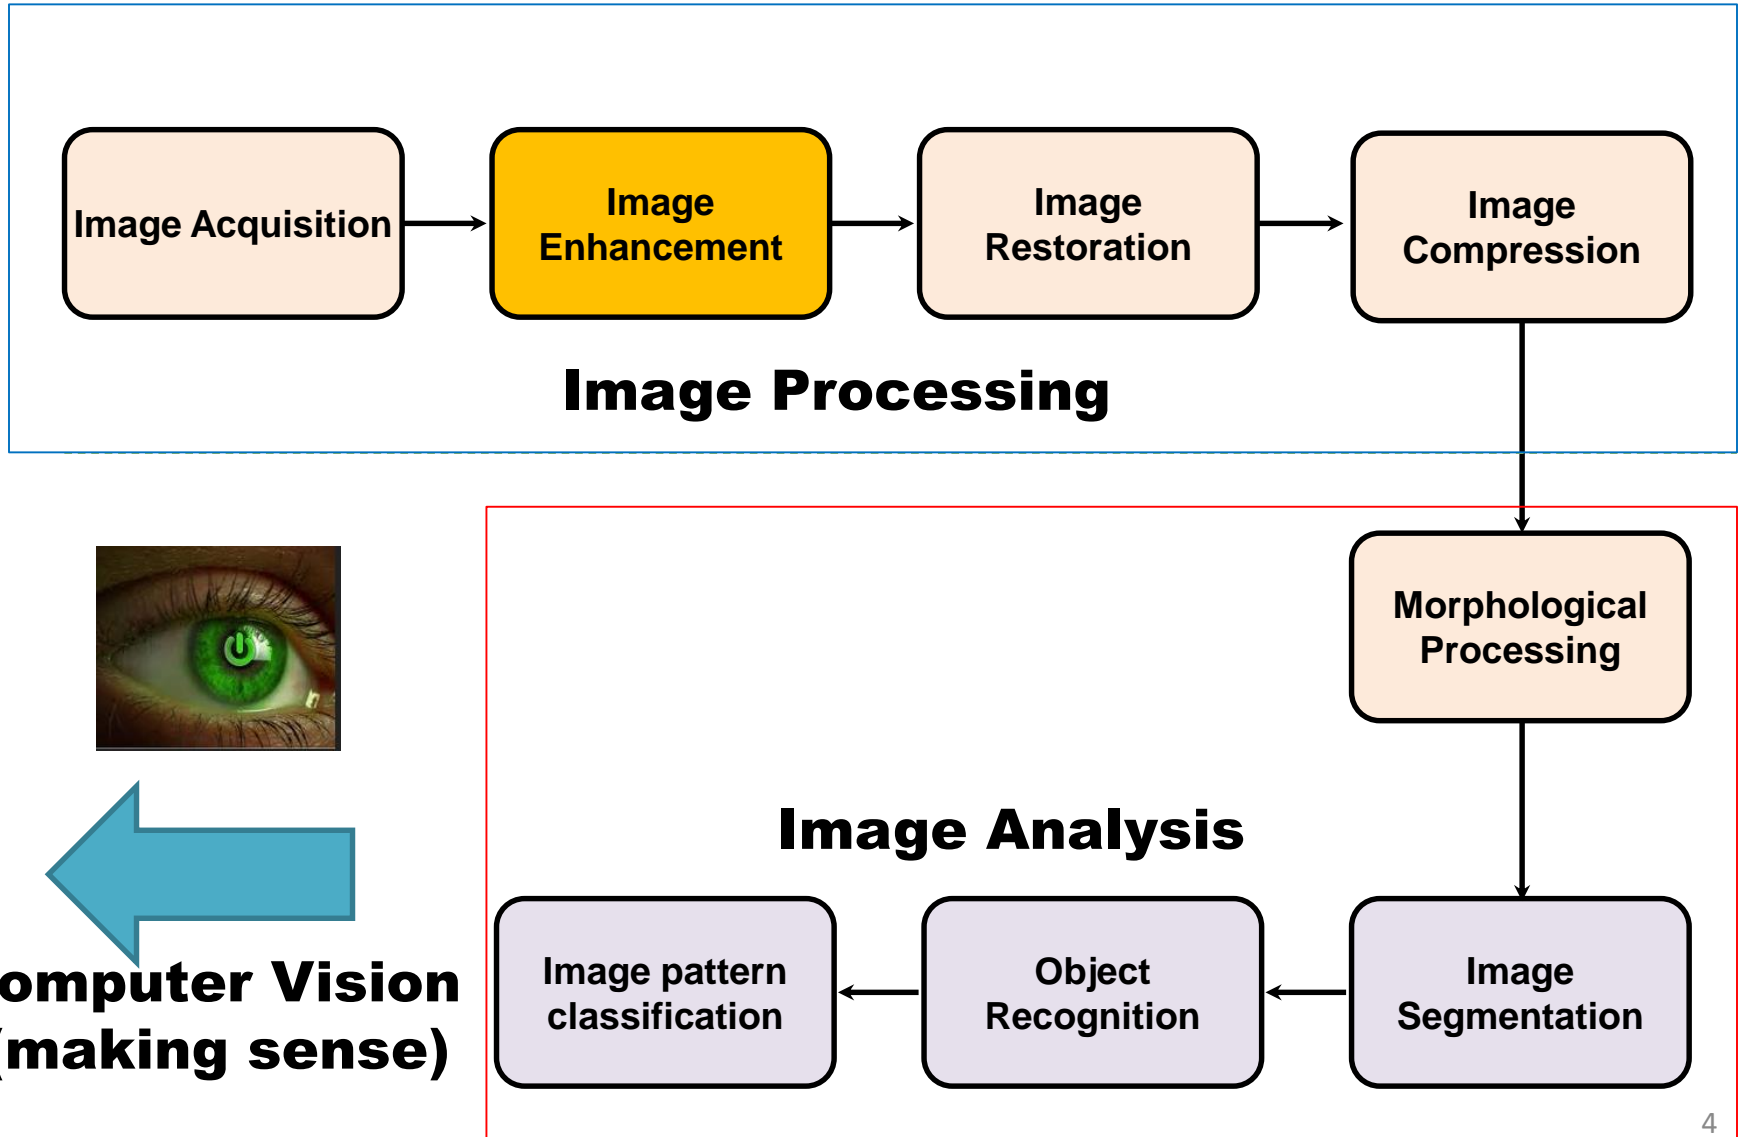
Recap

- Intensity transformation
 - Basic intensity transformation functions
 - Piecewise intensity transformation functions
 - Histogram processing
 - Histogram stretching
 - Histogram equalization
 - Histogram specification
 - Local histogram processing

Lecture Objectives

- Fundamentals of Spatial Filtering
- Correlation and Convolution
- How to construct Spatial filter masks?
- Smoothing (Lowpass) spatial filters
 - Box filter kernels
 - Gaussian filter kernels
 - Smoothing Non-linear Filters

Key Stages in DIP



Fundamentals of Spatial Filtering


- **Filtering** mostly focus on **image enhancement** techniques.
- What is a **spatial filter**?
 - Borrowed from **frequency domain** processing (next topic)
 - **“Filtering”** refers to *passing, modifying, or rejecting specified frequency components of an image*.
 - ❑ For example - filters that pass *low-frequencies/high-frequencies* are called a ***lowpass/highpass filters***
 - Filter is also called “mask”, “kernel”, “template” or “window”

Fundamentals of Spatial Filtering


- Spatial filters are more versatile than their frequency domain counterparts.
 - Spatial domain permits **non-linear filtering**.
- Spatial filtering modifies an image by **replacing the value of each pixel by a function of the values of the pixel and its neighbors**.
- If the operation performed on the image pixels is linear, then the filter is called a *linear spatial filter*. Otherwise, the filter is a *nonlinear spatial filter*.

Fundamentals of Spatial Filtering

- A **linear spatial filter** performs a **sum-of-products operation** between an *image f* and a *filter kernel w* .
- The kernel is an array whose **size** defines the neighborhood of operation, and whose **coefficients** determine the nature of the filter.



Fundamentals of Spatial Filtering



$w(-1,-1)$ $f(x-1,y-1)$	$w(-1,0)$ $f(x-1,y)$	$w(-1,1)$ $f(x-1,y+1)$
$w(0,-1)$ $f(x,y-1)$	$w(0,0)$ $f(x,y)$	$w(0,1)$ $f(x,y+1)$
$w(1,-1)$ $f(x+1,y-1)$	$w(1,0)$ $f(x+1,y)$	$w(1,1)$ $f(x+1,y+1)$

$$g(x,y) = w(-1,-1)f(x-1,y-1) + w(-1,0)f(x-1,y) + \dots + w(0,0)f(x,y) + \dots + w(1,1)f(x+1,y+1)$$

For a kernel of **odd size** $m \times n$, we assume that $m = 2a + 1$ and $n = 2b + 1$, where a and b are **nonnegative integers**

$$g(x,y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s,t)f(x+s,y+t)$$

As coordinates x and y are **varied**, the center of the kernel **moves** from pixel to pixel, generating the filtered image, g , in the process.

Fundamentals of Spatial Filtering

<https://ezyang.github.io/convolution-visualizer/>

Correlation and Convolution (important)

Correlation Vs. Convolution

- ***Spatial correlation*** is the process of moving a filter mask over the image and computing the **sum of products** at each location.
- ***Spatial convolution*** are the same as *Special correlation*, except that the correlation **kernel is rotated by 180°**.
- Thus, when the values of a kernel are **symmetric about its center**, correlation and convolution yield the same result.
- We begin by understanding the Spatial correlation operation with a **1-D illustration**.

Understanding Correlation in 1D

- Let $f(x)$ be the **1D signal** of length L , and $w(x)$ be the **1D kernel** of window size: $2a+1$, where $a \geq 1$.
- The **general linear spatial filtering expression** given by:

$$g(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t) f(x + s, y + t)$$

becomes the following expression for the **1D correlation**:

$$g(x) = \sum_{s=-a}^a w(s) f(x + s)$$

Correlation in 1D - example

- Let $f(x)$ =

Value	5	4	2	3	7	4	6	5
index	0	1	2	3	4	5	6	7

- Kernel: Averaging with two neighbors

Let $w(i)$ =

Value	1	1	1
index	-1	0	1

$$g(x) = \frac{1}{3} \sum_{i=-1}^1 w(i) f(x+i)$$

Correlation in 1D - example

Kernel: Averaging with two neighbors

$$g(x) = \frac{1}{3} \sum_{i=-1}^1 w(i) f(x+i)$$

f(x)=	Value	5	4	2	3	7	4	6	5
	index	0	1	2	3	4	5	6	7
w(i)=		1	1	1					
		-1	0	1					

$$\begin{aligned} g(1) &= \frac{1}{3} [w(-1) \times f(0) + w(0) \times f(1) + w(1) \times f(2)] \\ &= \frac{1}{3} [1 \times 5 + 1 \times 4 + 1 \times 2] \\ &= \text{ceil}\left(\frac{11}{3}\right) \\ &= 4 \end{aligned}$$

Correlation in 1D - example

Kernel: Averaging with two neighbors

$$g(x) = \frac{1}{3} \sum_{i=-1}^1 w(i) f(x+i)$$

$f(x)=$	Value	5	4	2	3	7	4	6	5
	index	0	1	2	3	4	5	6	7
$w(i)=$		$\frac{1}{3}$	$\frac{1}{3}$	$\frac{1}{3}$					
		-1	0	1					

$$\begin{aligned} g(1) &= [w(-1) \times f(0) + w(0) \times f(1) + w(1) \times f(2)] \\ &= \left[\frac{1}{3} \times 5 + \frac{1}{3} \times 4 + \frac{1}{3} \times 2 \right] \\ &= \text{ceil}\left(\frac{11}{3}\right) \\ &= 4 \end{aligned}$$

Boundary Conditions

- What about $g(0)$ and $g(7)$? How to handle **undefined area** during convolution?

$f(x)=$

Value	5	4	2	3	7	4	6	5
index	0	1	2	3	4	5	6	7

$w(i)=$

1	1	1
-1	0	1

- Zero padding** - the 1D signal is **padded with zeros** in both directions.

$f(x)=$

Value	0	5	4	2	3	7	4	6	5	0
index	-1	0	1	2	3	4	5	6	7	8

$w(i)=$

1	1	1
-1	0	1

Boundary Conditions

- Zero padding - the 1D signal is **padded with zeros** in both directions.

$f(x)=$

Value	0	5	4	2	3	7	4	6	5	0
index	-1	0	1	2	3	4	5	6	7	8

$w(i)=$

1	1	1
-1	0	1

Boundary Conditions

- **Replicate padding** - the boundary values are **repeated** as necessary to “complete” the signal.

$f(x) =$

Value	5	5	4	2	3	7	4	6	5	5
index	-1	0	1	2	3	4	5	6	7	8

$w(i) =$

1	1	1
-1	0	1

$f(x) =$

Value	5	5	4	2	3	7	4	6	5	5
index	-1	0	1	2	3	4	5	6	7	8

$w(i) =$

1	1	1
-1	0	1

Boundary Conditions

- **Mirror padding** - the values outside the boundary are obtained by **mirror-reflecting** the image across its border to “complete” the signal.

$f(x) =$

Value	4	5	4	2	3	7	4	6	5	6
index	-1	0	1	2	3	4	5	6	7	8

$w(i) =$

1	1	1
-1	0	1

$f(x) =$

Value	4	5	4	2	3	7	4	6	5	6
index	-1	0	1	2	3	4	5	6	7	8

$w(i) =$

1	1	1
-1	0	1

Boundary Conditions

- **Ignore** - The signal values at the boundary locations is **ignored** during correlation.
 - **Output has less values than the input signal !!!**

$f(x) =$

Value	5	4	2	3	7	4	6	5
index	0	1	2	3	4	5	6	7

$w(i) =$

1	1	1
-1	0	1


$f(x) =$

Value	5	4	2	3	7	4	6	5
index	0	1	2	3	4	5	6	7


$w(i) =$

1	1	1
-1	0	1


Correlation in 1D: Step-1



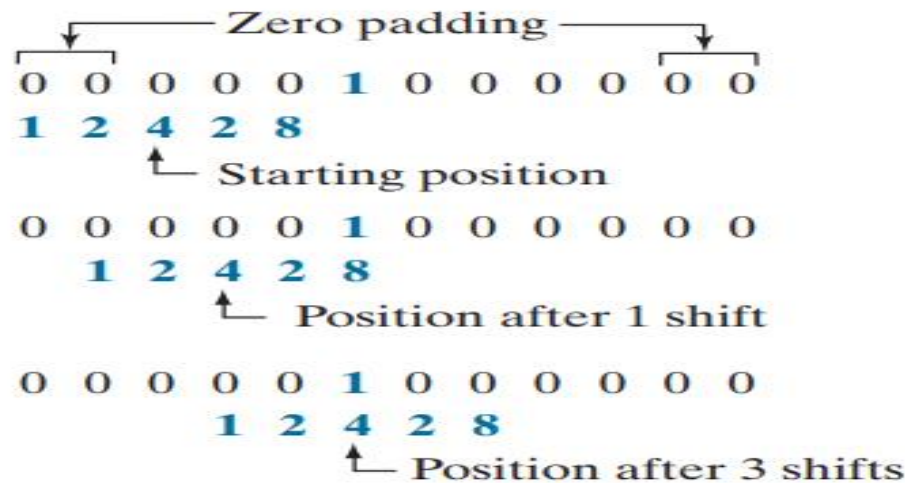
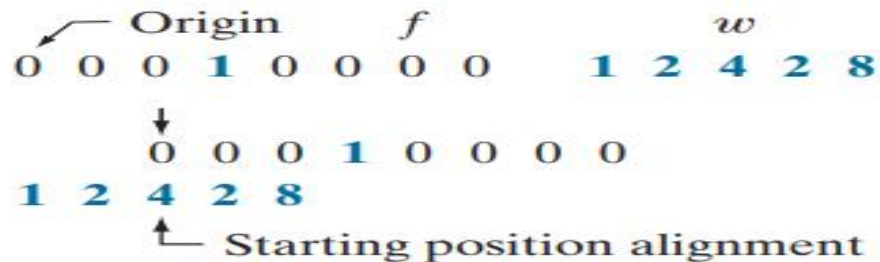
Correlation in 1D: Step-2



Correlation in 1D: Step-8



Correlation in 1D - Example



Correlation result

0 8 2 4 2 1 0 0

Extended (full) correlation result

0 0 0 8 2 4 2 1 0 0 0 0

Correlation in 1D - facts

- Its basically a **sliding window** operation
 - The filter is placed on every input value in turn and the corresponding output value is computed.
- For convenience, assume **odd length** filter window.
- The **center location** on the **filter window** is its **origin** (location 0).
- Assume **$2a+1$** elements in a filter, the indices go from **$[-a, a]$** .
- The **number of padding** required on each side = **a** .

Correlation in 1D - Practice

- $f(x) = [0, 0, 0, 1, 0, 0, 0]$
- $w(x) = [1, 2, 3, 4, 5]$
- $g(x) = ?$ **Use zero padding**

Correlation in 1D - Practice

- $f(x) = [0, 0, 0, 1, 0, 0, 0]$
- $w(x) = [1, 2, 3, 4, 5]$
- $g(x) = [0, 5, 4, 3, 2, 1, 0]$

Correlation in 1D - Practice

- $f(x) = [1, 2, 3, 4, 5, 6, 7]$
- $w(x) = [0, 0, 1, 0, 0]$
- $g(x) = ?$ **Use zero padding**

Correlation in 1D - Practice

- $f(x) = [1, 2, 3, 4, 5, 6, 7]$
- $w(x) = [0, 0, 1, 0, 0]$
- **$g(x) = [1, 2, 3, 4, 5, 6, 7]$**

Correlation in 1D - facts

- $f(x) = [0, 0, 0, 1, 0, 0, 0]$

- $w(x) = [1, 2, 3, 4, 5]$

- $g(x) = [0, 5, 4, 3, 2, 1, 0]$

- $f(x) = [1, 2, 3, 4, 5, 6, 7]$

- $w(x) = [0, 0, 1, 0, 0]$

- $g(x) = [1, 2, 3, 4, 5, 6, 7]$

Discrete unit
impulse

— A function that contains a single 1 with the rest being 0's is called a **discrete unit impulse**.

— Correlating a kernel with a discrete unit impulse yields a **rotated version of the kernel** at the location of the impulse.

— Correlating a function with a discrete unit impulse yields the **same function** at the location of the impulse.

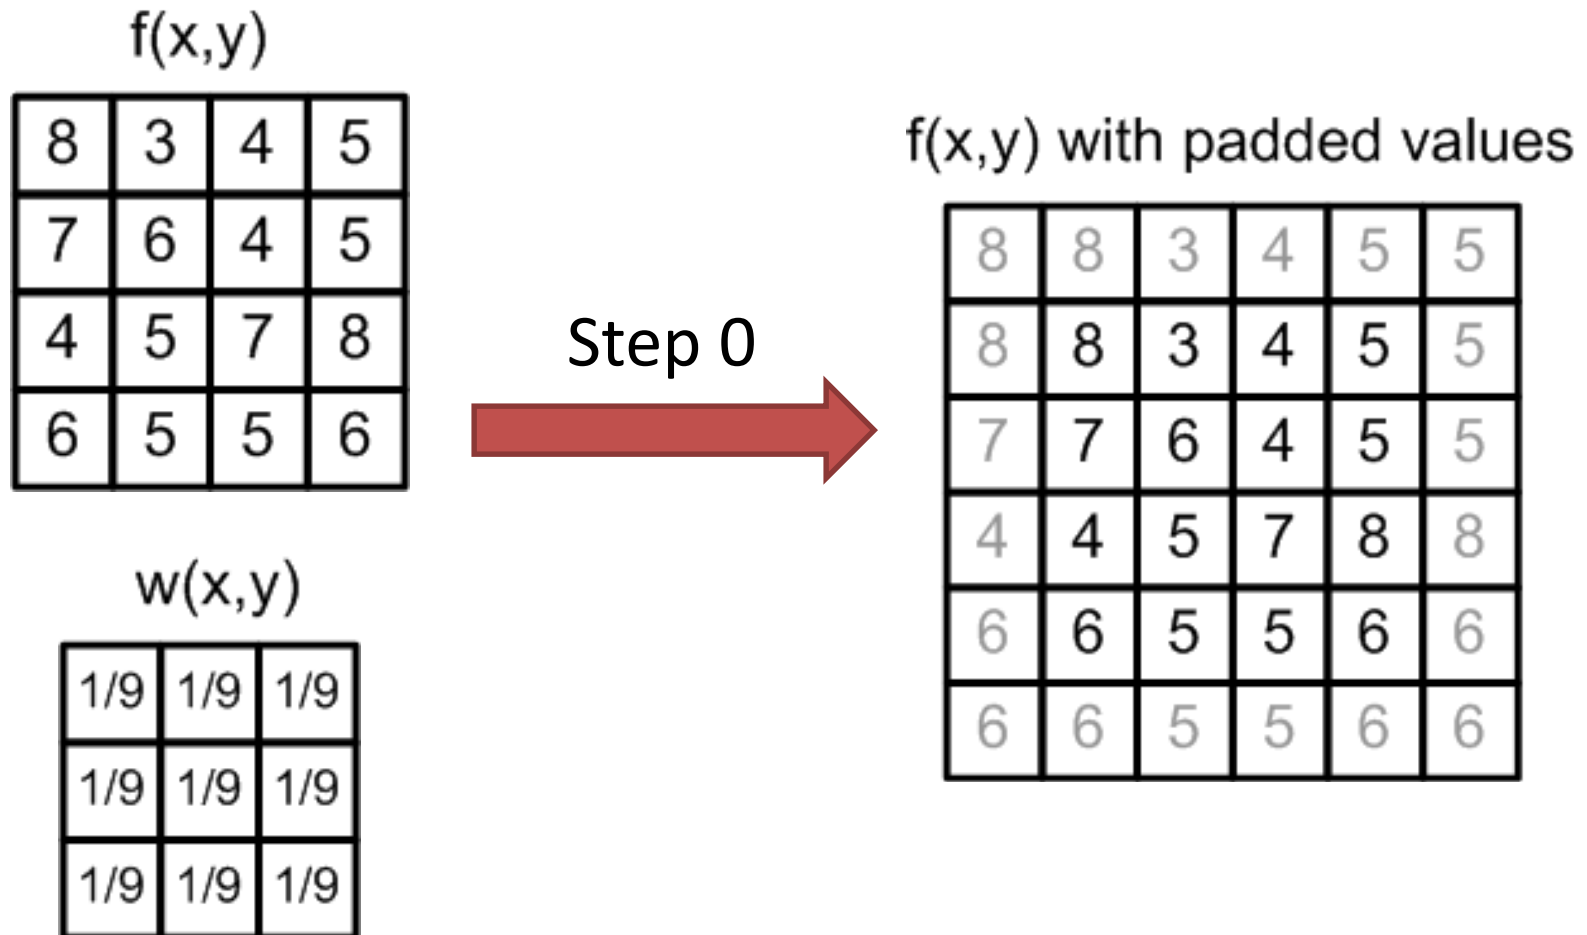
Correlation in 2D

- Let $f(x,y)$ be a **2D signal/image** of size $L \times L$ and $w(x,y)$ be a **2D kernel** of window size $(2a+1, 2b+1)$, where $a \geq 1$ and $b \geq 1$.
- The **2D correlation** is represented by:

$$(w \star f)(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t) f(x + s, y + t)$$

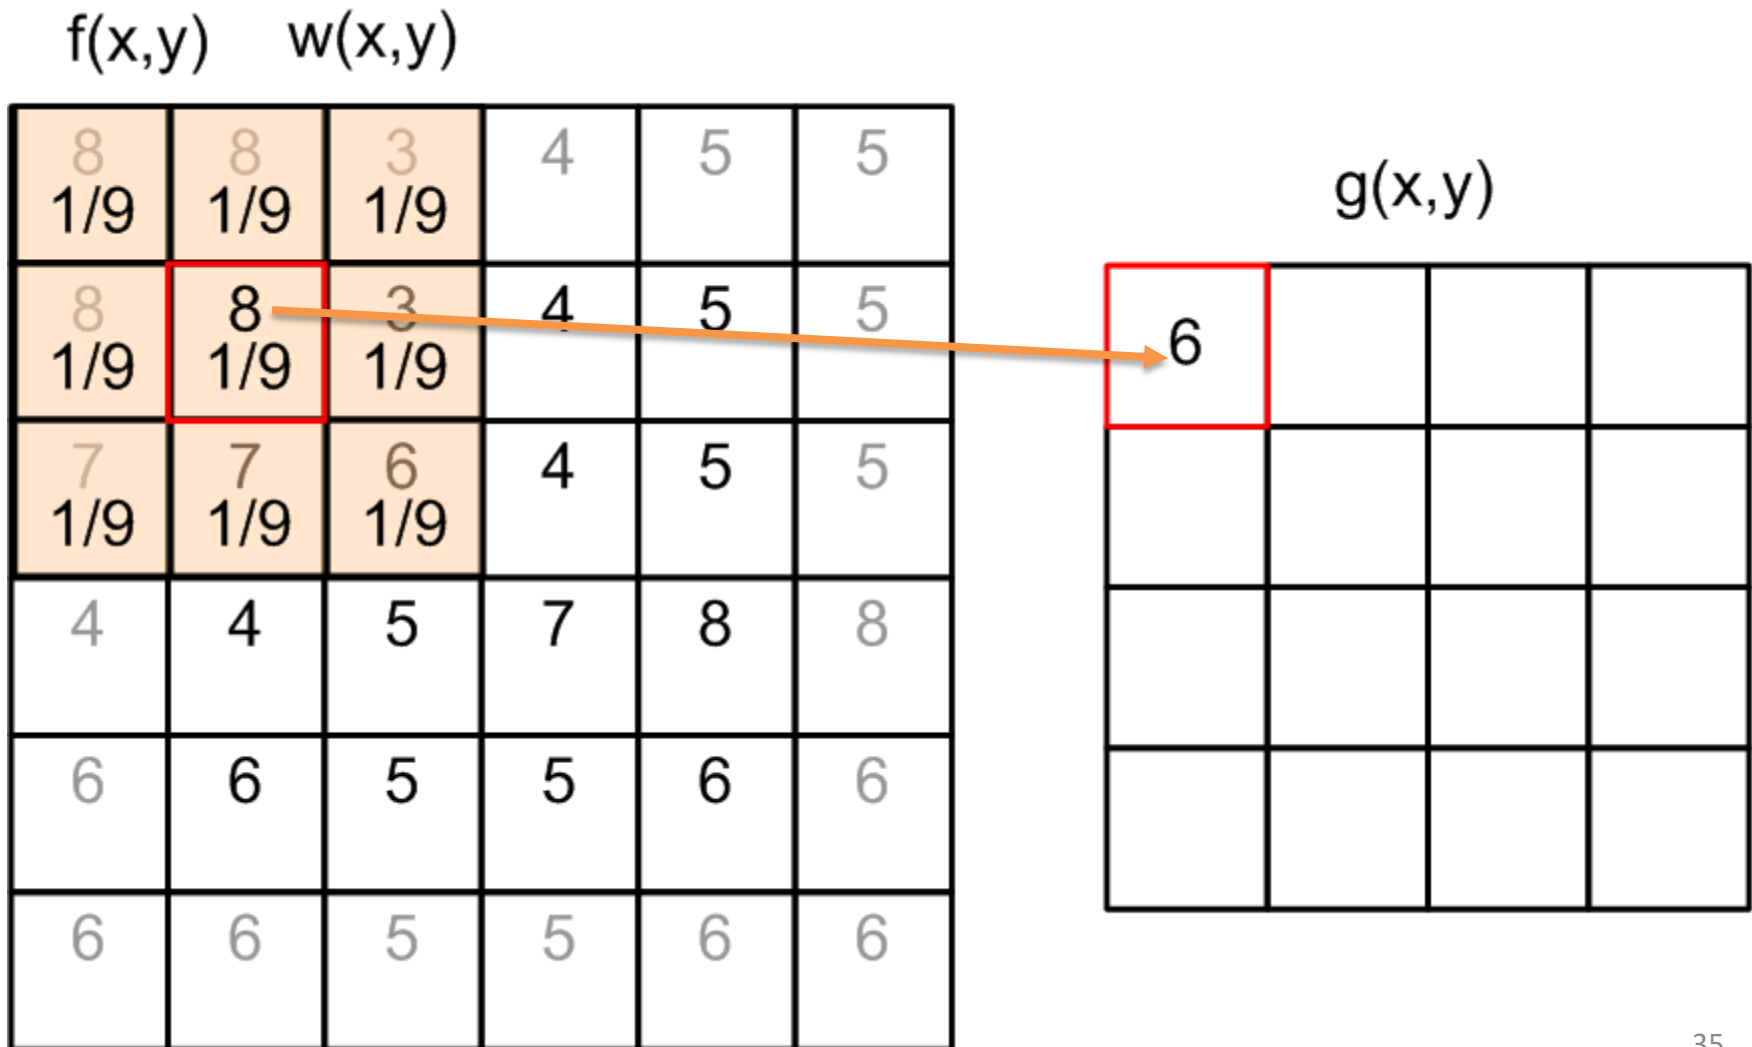
- For a kernel of size $m \times n$, using **zero padding**, we pad the image with a minimum of $(m - 1)/2$ rows of 0's at the **top** and **bottom** and $(n - 1)/2$ columns of 0's on the **left** and **right**.

2D Correlation - Replicate padding example



2D Correlation - **Replicate padding** example

step-1



2D Correlation - **Replicate padding** example

step-2

$f(x,y)$ $w(x,y)$									
8	8 1/9	3 1/9	4 1/9	5	5				
8	8 1/9	3 1/9	4 1/9	5	5				
7	7 1/9	6 1/9	4 1/9	5	5				
4	4	5	7	8	8				
6	6	5	5	6	6				
6	6	5	5	6	6				

$g(x,y)$			
6	5		

2D Correlation - **Replicate padding example**

step-16

$f(x,y)$ $w(x,y)$

8	8	3	4	5	5
8	8	3	4	5	5
7	7	6	4	5	5
4	4	5	7 1/9	8 1/9	8 1/9
6	6	5	5 1/9	6 1/9	6 1/9
6	6	5	5 1/9	6 1/9	6 1/9

$g(x,y)$

6	5	4	4
5	5	5	5
5	5	5	6
5	5	5	6

2D Correlation - **Replicate padding example** practice

$f(x,y)$

0	0	0	0
0	1	1	0
0	1	1	0
0	0	0	0

$w(x,y)$

1	2	3
4	5	6
7	8	9

$g(x,y)$

2D Correlation - **Replicate padding example** practice

$f(x,y)$

0	0	0	0
0	1	1	0
0	1	1	0
0	0	0	0

$w(x,y)$

1	2	3
4	5	6
7	8	9

$g(x,y)$

9	17	15	7
15	28	24	11
9	16	12	5
3	5	3	1

Convolution

- Convolution is similar to correlation, except that the **filter is first flipped by 180°**

- **1D** - reverse the values in the filter

- $w_f[i] = w[N-i-1]$

- Example of $W=[1\ 2\ 3\ 4\ 5]$ becomes $[5\ 4\ 3\ 2\ 1]$

- **2D** - flip the filter about the centre

- $w_f[i][i] = w[N-i-1][N-i-1]$

- Example of $W = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{bmatrix}$ becomes $\begin{bmatrix} 9 & 8 & 7 \\ 6 & 5 & 4 \\ 3 & 2 & 1 \end{bmatrix}$

Convolution in 1D

- Let $f(x)$ be the **1D** signal of length L , and $w(x)$ be the **1D** kernel of window size: $2a+1$, where $a \geq 1$.
- The **1D convolution** is represented by :

$$g(x) = \sum_{s=-a}^a w(s)f(x-s)$$

where the **minus signs** align the coordinates of f and w when one of the functions is rotated by **180°**

Correlation and **Convolution** **yield the same result** if the kernel values are **symmetric** about the center.

Convolution in 1D - example

Kernel: Averaging with two neighbors

$$g(x) = \frac{1}{3} \sum_{i=-1}^1 w(i) f(x-i)$$

$f(x)=$	Value	5	4	2	3	7	4	6	5
	index	0	1	2	3	4	5	6	7
$w(i)=$		$\frac{1}{3}$	$\frac{1}{3}$	$\frac{1}{3}$					
		-1	0	1					

$$\begin{aligned} g(1) &= [w(-1) \times f(2) + w(0) \times f(1) + w(1) \times f(0)] \\ &= \left[\frac{1}{3} \times 2 + \frac{1}{3} \times 4 + \frac{1}{3} \times 5 \right] \\ &= \text{ceil}\left(\frac{11}{3}\right) \\ &= 4 \end{aligned}$$

Convolution in 1D - Practice

- $f(x) = [0, 0, 0, 1, 0, 0, 0]$
- $w(x) = [1, 2, 3, 4, 5]$
- $g(x) = ?$ **Use zero padding**

Convolution in 1D - Practice

- $f(x) = [0, 0, 0, 1, 0, 0, 0]$
- $w(x) = [1, 2, 3, 4, 5]$

flip the kernel by **180°**

$$w(x) = [5, 4, 3, 2, 1]$$

- **$g(x) = [0, 1, 2, 3, 4, 5, 0]$**

Vs. Correlation in **1D**

- $f(x) = [0, 0, 0, 1, 0, 0, 0]$
- $w(x) = [1, 2, 3, 4, 5]$
- **$g(x) = [0, 5, 4, 3, 2, 1, 0]$**

Convolution in 1D - Practice

- $f(x) = [1, 2, 3, 4, 5, 6, 7]$
- $w(x) = [0, 0, 1, 0, 0]$
- $g(x) = ?$ **Use zero padding**

Convolution in 1D - Practice

- $f(x) = [1, 2, 3, 4, 5, 6, 7]$
- $w(x) = [0, 0, 1, 0, 0]$

flipping the kernel by **180°** would result in the same kernel

$$w(x) = [0, 0, 1, 0, 0]$$

Vs. Correlation in **1D**

- **$g(x) = [1, 2, 3, 4, 5, 6, 7]$**

- $f(x) = [1, 2, 3, 4, 5, 6, 7]$






- $w(x) = [0, 0, 1, 0, 0]$

- **$g(x) = [1, 2, 3, 4, 5, 6, 7]$**

Convolution in 1D - facts

- $f(x) = [0, 0, 0, 1, 0, 0, 0]$
 - $w(x) = [1, 2, 3, 4, 5]$ **flipped** $= [5, 4, 3, 2, 1]$
 - $g(x) = [0, 1, 2, 3, 4, 5, 0]$
- Discrete unit impulse
- $f(x) = [1, 2, 3, 4, 5, 6, 7]$
 - $w(x) = [0, 0, 1, 0, 0]$ **flipped** $= [0, 0, 1, 0, 0]$
 - $g(x) = [1, 2, 3, 4, 5, 6, 7]$
- A function that contains a single 1 with the rest being 0's is called a **discrete unit impulse**.
 - Convolving a kernel with a discrete unit impulse yields the **same kernel** at the **location of the impulse**.
 - Convolving a function with a discrete unit impulse yields the **same function** at the **location of the impulse**.

Convolution in 1D - Example



Convolution result

0 1 2 4 2 8 0 0

Extended (full) convolution result

0 0 0 1 2 4 2 8 0 0 0 0

Convolution in 2D

- Let $f(x,y)$ be a **2D signal/image** of size $L \times L$ and $w(x,y)$ be a **2D kernel** of window size $(2a+1, 2b+1)$, where $a \geq 1$ and $b \geq 1$.
- The **2D convolution** is represented by:

$$(w \star f)(x, y) = \sum_{s=-a}^a \sum_{t=-b}^b w(s, t) f(x - s, y - t)$$

- For a kernel of size $m \times n$, using **zero padding**, we pad the image with a minimum of $(m - 1)/2$ rows of 0's at the **top** and **bottom** and $(n - 1)/2$ columns of 0's on the **left** and **right**.

Convolution Vs. Correlation in 2D - Example

										Padded f							
										0	0	0	0	0	0	0	0
\swarrow	Origin				f					0	0	0	0	0	0	0	0
0	0	0	0	0	0					0	0	0	0	0	0	0	0
0	0	0	0	0	0	w				0	0	0	1	0	0	0	0
0	0	1	0	0	1	2	3			0	0	0	0	0	0	0	0
0	0	0	0	0	0	4	5	6		0	0	0	0	0	0	0	0
0	0	0	0	0	0	7	8	9		0	0	0	0	0	0	0	0

\searrow	Initial position for w						Correlation result					Full correlation result						
1	2	3	0	0	0	0						0	0	0	0	0	0	0
4	5	6	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
7	8	9	0	0	0	0	0	9	8	7	0	0	0	9	8	7	0	0
0	0	0	1	0	0	0	0	6	5	4	0	0	0	6	5	4	0	0
0	0	0	0	0	0	0	0	3	2	1	0	0	0	3	2	1	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0						0	0	0	0	0	0	0

\searrow Rotated w	Convolution result					Full convolution result						
9 8 7	0	0	0	0		0	0	0	0	0	0	0
6 5 4	0	0	0	0	0	0	0	0	0	0	0	0
3 2 1	0	0	0	0	0	0	1	2	3	0	0	
0 0 0	1	0	0	0		0	4	5	6	0	0	
0 0 0	0	0	0	0		0	7	8	9	0	0	
0 0 0	0	0	0	0		0	0	0	0	0	0	
0 0 0	0	0	0	0		0	0	0	0	0	0	0

Important Properties

- **Correlation** (\star) of a kernel with a discrete unit impulse gives a **rotated version of the kernel** centered at the impulse location.
- **Convolution** (\star) of a kernel with a discrete unit impulse gives an **exact copy of the kernel** centered at the impulse location.

Property	Convolution	Correlation
Commutative	$f \star g = g \star f$	—
Associative	$f \star (g \star h) = (f \star g) \star h$	—
Distributive	$f \star (g + h) = (f \star g) + (f \star h)$	$f \star (g + h) = (f \star g) + (f \star h)$

Multi-stage filtering with Convolution

- Sometimes an image is filtered (i.e., convolved) sequentially, in stages, using a different kernel in each stage.
- For example, suppose that an image f is first filtered with a kernel w_1 , then the result is filtered with kernel w_2 , that result is filtered with a kernel w_3 , and so on, for Q stages.
 - Because of the **commutative property** of convolution, this multistage filtering can be done in a single filtering operation $w \star f$ where,

$$w = w_1 \star w_2 \star w_3 \star \cdots \star w_Q$$

Separable filter kernels

- A function $\mathbf{G}(\mathbf{x}, \mathbf{y})$ is said to be *separable* if it can be written as the product of two 1-D functions, $\mathbf{G}_1(\mathbf{x})$ and $\mathbf{G}_2(\mathbf{y})$; that is,

$$\mathbf{G}(\mathbf{x}, \mathbf{y}) = \mathbf{G}_1(\mathbf{x}) \mathbf{G}_2(\mathbf{y})$$

- A spatial filter kernel is a matrix, and a separable kernel is a matrix that can be expressed as the outer product of two vectors:

$$\mathbf{w} = \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

$$\mathbf{c} = \begin{bmatrix} 1 \\ 1 \end{bmatrix} \quad \text{and} \quad \mathbf{r} = \begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix}$$

That is,

$$\mathbf{c} \mathbf{r}^T = \begin{bmatrix} 1 \\ 1 \end{bmatrix} \begin{bmatrix} 1 & 1 & 1 \end{bmatrix} = \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} = \mathbf{w}$$

- Separable kernels have computational advantage (and thus execution-time advantage).

Separable filter kernels

- For an *image* of size $M \times N$ and a *kernel* of size $m \times n$, the convolution operation requires the order of **MNmn** **multiplications** and **additions**.
- If the kernel is *separable*, then the convolution operation can be split as:

$$w \star f = (w_1 \star w_2) \star f = (w_2 \star w_1) \star f = w_2 \star (w_1 \star f) = (w_1 \star f) \star w_2$$

Step-1: The *first convolution* $(w_1 \star f)$ requires the order of **MNm** **multiplications** and **additions** because **w1** is of size $m \times 1$. This result is of size $M \times N$.

Step-2: The convolution of **w2** with the result of step-1 requires another **MNn** operations because **w2** is of size $1 \times n$. The **total multiplication and addition operations** in *step-1* and *step-2* are just **MN(m+n)**.

Separable filter kernels

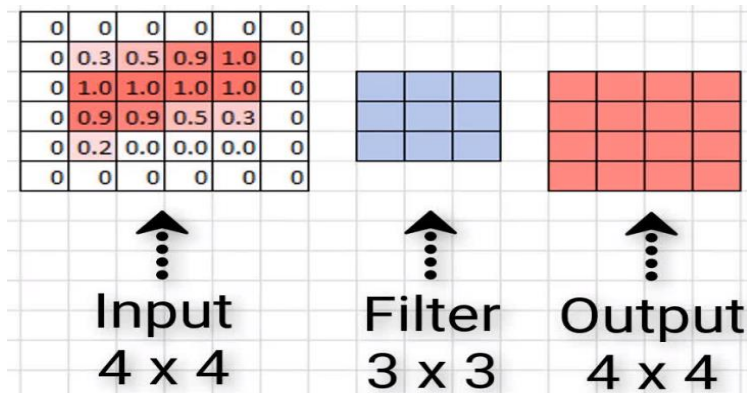
- The *computational advantage* of performing convolution with a separable, as opposed to a non-separable, kernel is defined as:

$$C = \frac{MNmn}{MN(m+n)} = \frac{mn}{m+n}$$


Example: For a kernel of size 11×11 , the computational advantage is **5.2**.

Strides, Padding, Kernel-size, and Convolution

- **Stride:** denotes how many steps we are moving in each steps during convolution.
- **Padding:** maintains the dimension of output as in input.




Padding with stride=1




No-padding with stride=1

Strides, Padding, Kernel-size, and Convolution

- Even-sized kernel: **usually not used** for convolution



No-padding with stride=1



No-padding with stride=2

Vector Representation of Linear Filtering


- Sum of product

$$\begin{aligned}
 R &= w_1 z_1 + w_2 z_2 + \cdots + w_{mn} z_{mn} \\
 &= \sum_{k=1}^{mn} w_k z_k \\
 &= \mathbf{w}^T \mathbf{z},
 \end{aligned}$$

where w s are the coefficients of an $m \times n$ filter and z s are the corresponding image intensities encompassed by the filter

$$\begin{aligned}
 \mathbf{w} &= \left[\frac{1}{9}, \frac{1}{9}, \frac{1}{9}, \frac{1}{9}, \frac{1}{9}, \frac{1}{9}, \frac{1}{9}, \frac{1}{9}, \frac{1}{9} \right] \\
 \mathbf{z} &= [8, 3, 4, 8, 3, 4, 7, 6, 4]
 \end{aligned}$$

f(x,y) w(x,y)					
8	8 1/9	3 1/9	4 1/9	5	5
8	8 1/9	3 1/9	4 1/9	5	5
7	7 1/9	6 1/9	4 1/9	5	5
4	4	5	7	8	8
6	6	5	5	6	6
6	6	5	5	6	6



Shift Invariant Operations

- Let H be a general operator that produces $g(x,y)$ for a given image $f(x,y)$

$$f(x,y): H[f(x,y)] = g(x,y)$$

- H is **shift invariant** if:

$$H[f(x+x_0, y+y_0)] = g(x+x_0, y+y_0)$$

Shift Invariant Operations - Examples

- $H[f(x,y)] = f(x-a, y-b)$
- $H[f(x,y)] = [f(x,y)]^2$
- $H[f(x,y)] = f(M \times x, N \times y)$, where $M, N \in \mathbb{Z}^+$
- $H[f(x,y)] = af(x,y) + b$, where a, b are arbitrary scalars

Summary

- Both Correlation & Convolution are the **simplest image operations** that can be performed in the spatial domain.
- **Linear** - Output pixel value is a linear combination of input pixel values.
- **Shift invariant** - Have same behavior at any location on the image plane

How to construct Spatial filter
masks?

Constructing Spatial Filter Masks

Mainly there are **three** basic approaches for **constructing spatial filters**:

1. Formulating filters based on mathematical properties.
 - Averaging filter
2. Based on sampling a 2-D spatial function whose shape has a desired property.
 - Gaussian filter
3. Based on a specified frequency response.
 - max/median/min filters

Smoothing (Lowpass) Linear Spatial Filters

Smoothing spatial filters

- **Smoothing** (also called *averaging*) spatial filters are used to reduce sharp transitions in intensity.
- **Applications:**
 - Noise reduction in an image.
 - Reduce irrelevant detail in an image where “irrelevant” refers to pixel regions that are small with respect to the size of the filter kernel.
 - Smoothing the false contours that result from using an insufficient number of intensity levels in an image
 - Used in combination with other techniques for image enhancement

Box filter (Arithmetic Mean Filter) kernels

- These are the **simplest**, **separable**, **lowpass averaging** filter kernels whose coefficients have the **same value** (typically **1**).
- An $m \times n$ box filter is an $m \times n$ array of **1's**, with a *normalizing constant in front*, whose value is 1 divided by the sum of the values of the coefficients.

$$\frac{1}{9} \times \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array}$$

$$\frac{1}{4.8976} \times \begin{array}{|c|c|c|} \hline 0.3679 & 0.6065 & 0.3679 \\ \hline 0.6065 & 1.0000 & 0.6065 \\ \hline 0.3679 & 0.6065 & 0.3679 \\ \hline \end{array}$$


Box filter kernels

- The *normalizing constant* has two purposes:
 1. The average value of an area of constant intensity will **remain the same in the filtered image**, as it should be.
 2. It prevents introducing a bias during filtering; that is, **the sum of the pixels in the original and filtered images will be the same.**

Box filter kernels



Test pattern of size
1024 x 1024 pixels



Result-Box kernel of size 3×3

Zero padding is used



Result-Box kernel of size 11×11



Result-Box kernel of size 21×21

Box filter kernels - Limitations

1. Poor approximations to the blurring characteristics of lenses.
 2. Favor blurring along perpendicular directions.
- **The kernels of choice in situations such as these** are *circularly symmetric kernels* (also called *isotropic*, meaning their response is independent of orientation) – **Gaussian Kernel**

Gaussian filter kernels

- Gaussian kernels* of the following form are the only **averaging** circularly symmetric kernels that are also separable:

$$w(s, t) = G(s, t) = Ke^{-\frac{s^2 + t^2}{2\sigma^2}}$$

where s and t are **kernel coordinates**, K is a **constant** and σ is the **standard deviation**.

$w(-1, -1)$	$w(-1, 0)$	$w(-1, 1)$
$w(0, -1)$	$w(0, 0)$	$w(0, 1)$
$w(1, -1)$	$w(1, 0)$	$w(1, 1)$

$$\frac{1}{4.8976} \times$$

0.3679	0.6065	0.3679
0.6065	1.0000	0.6065
0.3679	0.6065	0.3679

$$K = \sigma = 1$$

Gaussian filter kernels

A Gaussian kernel gives less weight to pixels further from the center of the window and vice-versa.

$$w(s,t) = G(s,t) = Ke^{-\frac{s^2 + t^2}{2\sigma^2}}$$

1/16

1	2	1
2	4	2
1	2	1

1/273

1	4	7	4	1
4	16	26	16	4
7	26	41	26	7
4	16	26	16	4
1	4	7	4	1



1/1003

0	0	1	2	1	0	0
0	3	13	22	13	3	0
1	13	59	97	59	13	1
2	22	97	159	97	22	2
1	13	59	97	59	13	1
0	3	13	22	13	3	0
0	0	1	2	1	0	0

Gaussian filter kernels

A Gaussian kernel gives less weight to pixels further from the center of the window and vice-versa.




$$w(s,t) = G(s,t) = Ke^{-\frac{s^2 + t^2}{2\sigma^2}}$$



Gaussian filter kernels

$$w(s, t) = G(s, t) = Ke^{-\frac{s^2 + t^2}{2\sigma^2}}$$

The σ term controls the “tightness” of the function




Low pass filtering with a Gaussian kernel

- Gaussian kernels have to be **larger** than box filters to achieve the same degree of blurring.
- This is because, whereas a **box kernel** assigns the **same weight** to all pixels, the values of **Gaussian kernel** coefficients (and hence their effect) **decreases** as a function of distance from the kernel center.



Test pattern of size
1024 x 1024 pixels



Box kernels of sizes
 21×21



Gaussian kernels of
sizes 21×21
and $\sigma=3.5$



Gaussian kernels of
sizes 43×43
and $\sigma=7$


K=1 and Zero padding is used in all cases

Comparison of Gaussian and Box filter smoothing characteristics


- Box filter:
 - The box filter produces linear smoothing, with the transition from black to white (i.e., at an edge) having the shape of a ramp.
 - The important features here are hard transitions at the onset and end of the ramp.
 - We would use this type of filter when **less smoothing of edges is desired**.
- Gaussian filter:
 - The Gaussian filter yields significantly smoother results around the edge transitions.
 - We would use this type of filter when generally **uniform smoothing is desired**.



Test image



Box kernel of size 71×71




Gaussian kernel of size 151×151 ,
with $K = 1$ and $\sigma = 25$.

Comparison of Zero/Mirror/Replicate padding


- *Replicate padding* is useful when the areas near the border of the image are constant.
- *Mirror padding* is more applicable when the areas near the border contain image details.
- These two types of padding attempt to “**extend**” the characteristics of an image past its borders.
- *Zero padding* introduces **dark borders** after filtering which can be eliminated by other two types of padding.




Test image



Replicate padding



Mirror padding



Zero padding


A Gaussian kernel of size 187×187 , with $K = 1$ and $\sigma = 31$ was used in all three cases

Smoothing performance as a function of kernel and image size


- The amount of relative blurring produced by a smoothing kernel of a given size **depends directly on image size**.
- **Not understanding** the relationship between kernel size and the size of objects in an image can lead to ineffective performance of spatial filtering algorithms.



Test image of size
1024 x 1024 (a)



Gaussian kernel of
size 187 x 187, with
 $K = 1$ and $\sigma = 31$ (b)



Test image of size
4096 x 4096 (c)



Gaussian kernel of
size 187 x 187, with
 $K = 1$ and $\sigma = 31$ (d)

To obtain results that are comparable to Fig. (b) we have to **increase the size and standard deviation** of the Gaussian kernel by **four**, the same factor as the increase in image dimensions (i.e. **745 x 745** (with $K = 1$ and $\sigma = 124$)).

Using lowpass filtering and thresholding for region extraction

- *Lowpass filtering* can be combined with *Intensity thresholding* for eliminating *irrelevant detail* in this image. In the present context, “irrelevant” refers to pixel regions that are small compared to kernel size.




A 2566×2758 image from the Hubble telescope


Using lowpass filtering and thresholding for region extraction



2566 × 2758 Image
from the Hubble
telescope




Gaussian kernel of
size **151 × 151**, $K=1$,
 $\sigma = 25$.



Result of blurring

Using lowpass filtering and thresholding for region extraction




Using lowpass filtering and thresholding for region extraction




Shading correction using lowpass filtering


- When shading pattern is **unknown**, we have to estimate the pattern directly from available samples of shaded images.
- *Lowpass filtering* is a simple method for *estimating shading patterns*.



2048 × 2048 Image shaded by a shading pattern oriented in the **-45°** direction (a)



Estimate of the shading patterns using a **512 × 512** Gaussian kernel (four times the size of the squares), **$K = 1$** , and **$\sigma = 128$** (b)



Result of dividing (a) by (b)

Smoothing Non-linear Filters (Order-static Non-linear Filters)

Order Statistics

- The **k^{th} order statistic** of a sequence of intensity values is equal to its **k^{th} smallest value**.
- **E.g.** [5,7,6,4,2]
Order statistics: $x_{(1)}=2, x_{(2)}=4, x_{(3)}=5, x_{(4)}=6, x_{(5)}=7$
- The **first** order statistic is the **minimum**.
- The **n^{th}** order statistic is the **maximum**.

Order-Statistic Filters

- *Order-statistic filters* are **nonlinear spatial filters** whose response is based on **ordering** (ranking) the pixels contained in the region encompassed by the filter.
- Median filter:
 - Replaces the value of the center pixel by the **median** of the intensity values in the neighborhood of that pixel (the value of the center pixel is included in computing the median).
 - Particularly effective in the presence of **impulse noise** (salt and pepper noise).
 - We first **sort** the values of the pixels in the neighborhood, determine their median, and assign that value to the pixel in the filtered image corresponding to the center of the neighborhood.
 - When several values in a neighborhood are the same, all equal values are grouped.

Order-Statistic Filters

- Median filter:

- For example, suppose that a 3×3 neighborhood has values (10, 20, 20, 20, 15, 20, 20, 25, 100). These values are sorted as (10, 15, 20, 20, 20, 20, 20, 25, 100), which results in a **median of 20**.
- The main function of median filter is to *force points to be more like their neighbors*.


- Max filter:

- Useful for finding the **brightest points** in an image or for **eroding** dark areas adjacent to bright regions.

- Min filter:


- Useful for finding the **darkest points** in an image or for **eroding** bright areas adjacent to dark regions.

Median Filter Example




X-ray image of circuit board corrupted by impulse noise (salt-pepper noise)


Median Filter Example



X-ray image of circuit board corrupted by impulse noise



A 19×19 Gaussian filter with $K = 1$, and $\sigma = 3$



A 7×7 Median filter

Next Lecture

- Sharpening spatial filters
 - Foundation of image sharpening
 - Using second-order derivative for image sharpening (Laplacian)
 - Unsharp masking and high boost filtering
 - Using first-order derivative for image sharpening (Gradient)
 - Highpass, Bandreject, and Bandpass filters from lowpass filters
 - Combining special enhancement methods