

# Image Compression

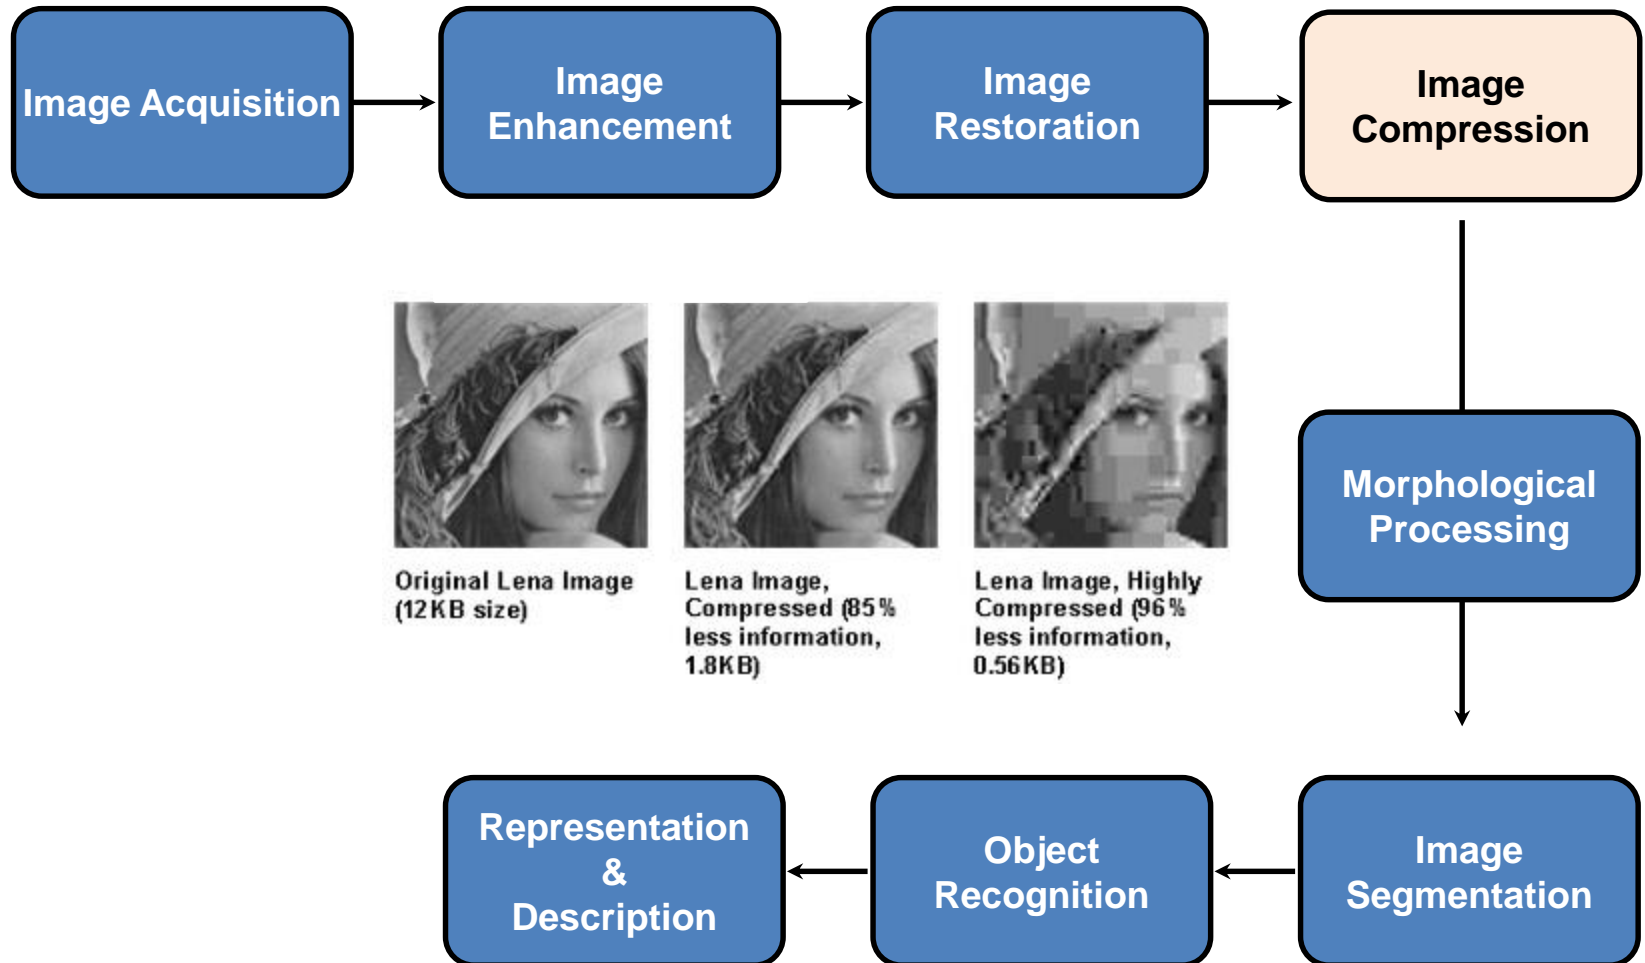
# Recap

- The 2-D DFT - Some Observations
- Separability of Fourier Transform
- IDFT in terms of DFT
- Fast Fourier Transform (FFT)
  - FFT Process in 1-D
  - Special Properties of  $W_M$
  - FFT even-odd approach
  - FFT "Butterfly" Method
  - FFT – time complexity
  - Can we speed it up??
  - FFT Algorithm

# Lecture Objectives

- Introduction to Image Compression
- Types of Data Redundancy
  - Coding redundancy
  - Spatial and Temporal Redundancy
  - Irrelevant Information
- Measuring Image Information
- Fidelity Criteria
- General Image Compression Model
  - Encoding/Compression Process
  - Decoding/Decompression Process
- Lossless Compression
  - Huffman Coding

# Key Stages in DIP



# Introduction to Image Compression

# Illustrating Example



8.9 MB



68.34 KB

# Illustrating Example

- Let a 2-hour, standard definition (SD) movie using  $720 \times 480 \times 24$  bit pixel arrays to represent each frame and the Frame rate = 30 fps.
- What is the size of the movie file?
- It requires  $\approx 223$  GB memory to store this movie file !!!.
- We need 25 DVDs each of 8.5 GB (single-sided, double-layer) to save this movie.
- Hence, *we need to compress each frame.*
- Compression must be even higher for high definition(HD), where image resolution is  $1920 \times 1080 \times 24$  bits/frame.


# Illustrating Examples

- **Image:** 6.0 megapixels camera, 3000×2000 px/inch
  - 13804 MB per image → 13.48 1GB/image
  - One **megapixel** refers to one million pixels/inch, which are small squares of information that combine to make up an image. So, if a camera has a resolution of six megapixels, it would be able to capture images with about six million pixels of information per inch.
- **Video:** DVD Disc 4.7 GB
  - video 720×480×24/frame, RGB, 30 frames/sec → 31.1MB/sec
  - audio 16bits × 44.1KHz stereo → 176.4KB/sec
    - 1.5 min per DVD disc
- **Send video from cellphone**
  - 352×240×24, RGB, 15 frames / sec
    - 3.8 MB/sec



# Objective of Image Compression

- The goal of image compression is to reduce the amount of data required to represent a digital image.



# Approaches

- Lossless
  - Information preserving
  - Low compression ratios
- Lossy
  - Not information preserving
  - High compression ratios
- Trade-off: image quality V.S. compression ratio

# Data $\neq$ Information

- Data and information **are not** synonymous terms!
- **Data** is the means by which **information** is conveyed.
- Data compression aims to **reduce** the *amount of data required to represent a given quantity of information* while **preserving** as much information as possible.


# Data V.S. Information

- The **same amount** of **information** can be represented by various amount of **data**.

## Examples:

- Your girlfriend, Helen, will meet you at *Common Man Coffee Roasters*, 22 Martin Rd, #01-00, Singapore 239058 at 5 minutes past 4:00 pm tomorrow afternoon.
- Your girlfriend will meet you at *Common Man Coffee Roasters* at 5 minutes past 4:00 pm tomorrow afternoon.
- Helen will meet you tomorrow at the regular café at the regular time.
- Helen will meet you tomorrow.

# Definitions: Compression Ratio



Compression ratio:

$$C_R = \frac{n_1}{n_2}$$

# Definitions: Data Redundancy

- Relative data redundancy:

Given the compression ratio  $C_R$  ,  $R_D = 1 - \frac{1}{C_R}$

- **Example:** let a dataset-1 *before the compression* has 10 bits of data, and *after the compression* it has only 1 bit of data.

$$\text{If } C_R = \frac{10}{1}, \text{ then } R_D = 1 - \frac{1}{10} = 0.9$$

(90% of the data in dataset 1 is redundant)

if  $n_2 = n_1$ , then  $C_R=1$ ,  $R_D=0$

if  $n_2 \ll n_1$ , then  $C_R \rightarrow \infty$ ,  $R_D \rightarrow 1$

$$C_R = \frac{n_1}{n_2}$$

# Types of Data Redundancy

# Types of Data Redundancy

- Compression attempts to reduce one or more of these redundancy types:
  - **Coding Redundancy**
  - **Spatial/Temporal Redundancy**
  - **Irrelevant Information**



# Coding Redundancy

- **Code**: a list of symbols (*letters, numbers, bits* etc).
- **Code word**: a sequence of symbols used to represent a piece of information or an event (e.g., *gray levels/intensity levels*).
- **Code word length**: number/length of symbols in each code word.

Example: (binary code, symbols: 0,1, length: 3)

0: 000	4: 100
1: 001	5: 101
2: 010	6: 110
3: 011	7: 111

# Coding Redundancy

- Recall: *Intensity as random variables*.

$$p_r(r_k) = \frac{n_k}{MN} \quad k = 0, 1, 2, \dots, L-1$$

Where,

- $L$  is the number of intensity values.
  - $n_k$  is the number of times that the  $k^{\text{th}}$  intensity appears in the image.
  - $r_k$  represent the intensities of an  $M \times N$  image.
  - $P_r(r_k)$  is the probability of occurrence of pixels with intensity  $r_k$  in the image.
- The *average number of bits* required to represent *each pixel* is:

$$L_{\text{avg}} = \sum_{k=0}^{L-1} l(r_k) p_r(r_k)$$


• The total number of bits required to represent  $M \times N$  image is:  
 $MNL_{\text{avg}}$

Where,

- $l(r_k)$  is the number of bits used to represent  $r_k$ .

# Coding Redundancy - Example

- Case 1:  $l(r_k) = \text{fixed}$  length coding



256×256

$r_k$	$p_r(r_k)$	Code 1	$l_I(r_k)$
$r_{87} = 87$	0.25	01010111	8
$r_{128} = 128$	0.47	10000000	8
$r_{186} = 186$	0.25	11000100	8
$r_{255} = 255$	0.03	11111111	8
$r_k$ for $k \neq 87, 128, 186, 255$	0	—	8


$$L_{\text{avg}} = \sum_{k=0}^{L-1} l(r_k) p_r(r_k)$$

$$L_{\text{avg}} = 0.25 \times 8 + 0.47 \times 8 + 0.25 \times 8 + 0.03 \times 8 = 8 \text{ bits}$$

$$\text{Total Number of bits: } 256 \times 256 \times 8 = 64 \text{ KB}$$

# Coding Redundancy - Example

- Case 2:  $l(r_k)$  = variable length coding



256×256×8

$r_k$	$p_r(r_k)$	Code 1	$l_1(r_k)$	Code 2	$l_2(r_k)$
$r_{87} = 87$	0.25	01010111	8	01	2
$r_{128} = 128$	0.47	10000000	8	1	1
$r_{186} = 186$	0.25	11000100	8	000	3
$r_{255} = 255$	0.03	11111111	8	001	3
$r_k$ for $k \neq 87, 128, 186, 255$	0	—	8	—	0

$$L_{\text{avg}} = \sum_{k=0}^{L-1} l(r_k) p_r(r_k)$$

011001000  
2130

$$L_{\text{avg}} = 0.25 \times 2 + 0.47 \times 1 + 0.25 \times 3 + 0.03 \times 3 = 1.81 \text{ bits}$$


$$\text{Total Number of bits: } 256 \times 256 \times 1.81 = 14.4 \text{ KB}$$

$$C_R = 8 / 1.81 = 4.42$$


$$R_D = 1 - 1 / 4.42 = 0.774$$

# Spatial and Temporal Redundancy (mappings)

- In **majority of cases**, any pixel value can be reasonably predicted by its neighbors (i.e., **if correlated**).



256×256×8



The intensity histogram

cannot be compressed by variable-length coding

11111  
3,7 OR 7,3 ??

- All 256 intensities are **equally probable**.
- Pixels are **independent** of one another in the *vertical direction*.
- Pixels are **completely dependent** (correlated) on one another in the *horizontal direction*.

# Spatial and Temporal Redundancy

- Run-length pairs:

- Specifies the *start of a new intensity* and the *number of consecutive pixels that have that intensity*.

- = (8-bit intensity value per pixel  $\times$  256 pixels) + number of pixels per line + number of lines

- =  $(8 \times 256) + 8 + 8$

- = 2064 bits

- $C_R = 256 \times 256 \times 8 / 2064$   
= 254

- $R_D = 1 - 1/128 = 0.996$




# Irrelevant Information (quantization)

- **The human eye does not respond with equal sensitivity to all visual information.**
- It is more sensitive to the **lower frequencies** than to the **higher frequencies** in the visual spectrum.
- **Idea:** discard data that is **perceptually insignificant!**

# Irrelevant Information



256×256×8




Intensities **125** through **131** are present, but the human visual system averages these intensities, perceives only the average value.



# Irrelevant Information



Original image



Histogram equalized version

**The histogram equalized version** of the original image makes the intensity changes visible *and* reveals two previously undetected regions of constant intensity—one oriented vertically, and the other horizontally.

# Irrelevant Information

- This image can be represented by its *average intensity* alone
  - a single **8-bit** value.
  - The original **256 \* 256 \* 8** bit intensity array is represented by:  
(#Rows, #Cols, Average intensity)
- $C_R = (256 \times 256 \times 8) / (8 + 8 + 8) = 21845$
- $R_D = 1 - 1/21845 = 0.99995422$

# Measuring Image Information

# Measuring Image Information

- What is the **minimum amount of data** that is actually needed to describe an image completely, without losing information?
- We assume that information generation is a **probabilistic process**.
- **Idea: associate information with probability !**
- A random event  $E$  with probability  $P(E)$  contains  $I(E)$  units of information.

$$I(E) = \log \frac{1}{P(E)} = -\log P(E)$$

# How much information does a **pixel** contain?

- Suppose that the *gray level* values are generated by a random variable  $r_k$ .
- The amount of information in pixel is given by :

$$I(r_k) = -\log(P_r(r_k))$$

# How much information does an **image** contain?

- Average information content of an image is given by:

$$L_{\text{avg}} = \sum_{k=0}^{L-1} l(r_k) p_r(r_k)$$


- Since  $I(r_k) = -\log(P_r(r_k))$ , we have:

$$H = - \sum_{k=0}^{L-1} P_r(r_k) \log(P_r(r_k)) \text{ units/pixel}$$

- **H** is called the *Entropy* of the image.

It is not possible to code the *intensity values* of an image with fewer than **H** bits/pixel.

# Entropy - Example



$r_k$	$p_r(r_k)$	Code 1	$l_1(r_k)$	Code 2	$l_2(r_k)$
$r_{87} = 87$	0.25	01010111	8	01	2
$r_{128} = 128$	0.47	10000000	8	1	1
$r_{186} = 186$	0.25	11000100	8	000	3
$r_{255} = 255$	0.03	11111111	8	001	3
$r_k$ for $k \neq 87, 128, 186, 255$	0	—	8	—	0

$$H = - \sum_{k=0}^{L-1} P_r(r_k) \log(P_r(r_k))$$

$$\begin{aligned}
 H &= -[0.25 \log_2 0.25 + 0.47 \log_2 0.47 + 0.25 \log_2 0.25 + 0.03 \log_2 0.03] \\
 &= -[0.25(-2) + 0.47(-1.09) + 0.25(-2) + 0.03(-5.06)] \\
 &\approx 1.6614 \text{ bits/pixel}
 \end{aligned}$$

It is not possible to code the *intensity values* of this image with fewer than **1.6614** bits/pixel.

# Redundancy Revisited

$$R = L_{avg} - H$$

Where,

$$L_{avg} = \sum_{k=0}^{L-1} l(r_k) p_r(r_k) \quad \text{and} \quad H = - \sum_{k=0}^{L-1} P_r(r_k) \log(P_r(r_k))$$


- **Note:** if  $L_{avg} = H$ , then  $R=0$  (no redundancy)



# Fidelity Criteria

# Fidelity Criteria

- Fidelity criteria provides a means of *quantifying the nature/amount of the loss*.



- The **error** (ex:- **root-mean-squared error**) between *original image*  $f(x,y)$  and the compressed-decompressed *approximation* of it  $\hat{f}(x,y)$  is given as:

$$e(x,y) = \hat{f}(x,y) - f(x,y)$$

- How close is  $f(x,y)$  to  $\hat{f}(x,y)$  ?
- Criteria
  - **Subjective**: based on human observers
  - **Objective**: based on mathematical model

# Subjective Fidelity Criteria

Value	Rating	Description
1	Excellent	An image of extremely high quality, as good as you could desire.
2	Fine	An image of high quality, providing enjoyable viewing. Interference is not objectionable.
3	Passable	An image of acceptable quality. Interference is not objectionable.
4	Marginal	An image of poor quality; you wish you could improve it. Interference is somewhat objectionable.
5	Inferior	A very poor image, but you could watch it. Objectionable interference is definitely present.
6	Unusable	An image so bad that you could not watch it.

# Objective Fidelity Criteria

- The **total error** between the two images  $f(x, y)$  and  $\hat{f}(x, y)$  is:

$$\sum_{x=0}^{M-1} \sum_{y=0}^{N-1} [\hat{f}(x, y) - f(x, y)]$$


- The **root-mean-squared-error** ( $e_{\text{rms}}$ ) between  $f(x, y)$  and  $\hat{f}(x, y)$  is:

$$e_{\text{rms}} = \left[ \frac{1}{MN} \sum_{x=0}^{M-1} \sum_{y=0}^{N-1} [\hat{f}(x, y) - f(x, y)]^2 \right]^{1/2}$$

- The **mean-squared signal-to-noise ratio** ( $\text{SNR}_{\text{rms}}$ ) of the output image is:

$$\text{SNR}_{\text{ms}} = \frac{\sum_{x=0}^{M-1} \sum_{y=0}^{N-1} \hat{f}(x, y)^2}{\sum_{x=0}^{M-1} \sum_{y=0}^{N-1} [\hat{f}(x, y) - f(x, y)]^2}$$


# Subjective Vs. Objective Fidelity Criteria



$e_{\text{rms}} = 5.17$



$e_{\text{rms}} = 15.67$

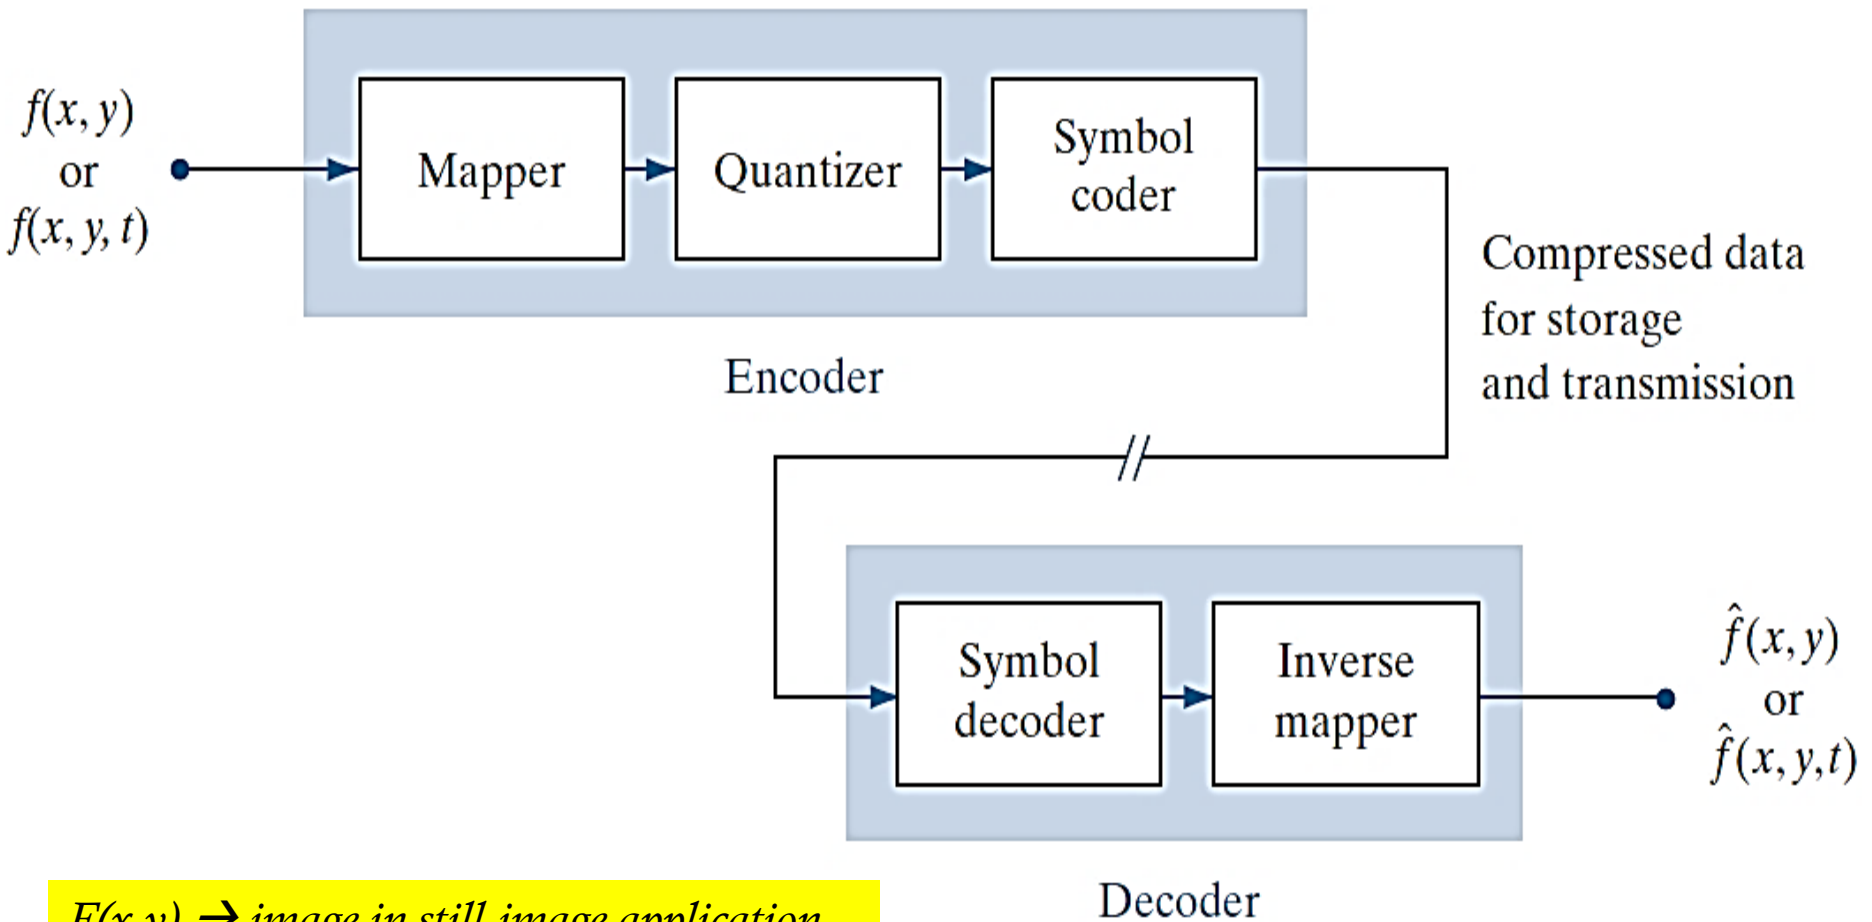


$e_{\text{rms}} = 14.17$

$e_{\text{rms}}$  should be **less** and  $\text{SNR}_{\text{rms}}$  should be **high**

# General Image Compression Model

# General Image Compression Model



$F(x, y) \rightarrow$  image in still-image application

$F(x, y, t) \rightarrow$  image in video application where 't' specifies time

# General Image Compression Model

- **Mapper**

- Transforms input into a (usually nonvisual) format designed to **reduce spatial and temporal redundancy**.
- This operation generally is **reversible**, and may or may not directly reduce the amount of data required to represent the image.

- **Quantizer**

- **Reduces the accuracy of the mapper's output** in accordance with a pre-established fidelity criterion.
- The goal is to **keep irrelevant information out of the compressed representation**.
- This operation is **irreversible**.
- It must be **omitted** when error-free compression is desired.

- **Symbol coder**

- Generates a **fixed-length or variable-length code** to represent the quantizer output, and maps the output in accordance with the code.
- This operation is **reversible**.



# Lossless Compression

# Lossless Compression



$$e(x, y) = \hat{f}(x, y) - f(x, y) = 0$$

# Binary Codes

**Binary code:** Maps each character of an alphabet  $\Sigma$  to a binary string.

Example:  $\Sigma = \text{a-z}$  and various punctuation (size **32** overall, say).

- Use the **32, 5-bit binary strings** to encode this  $\Sigma$ . (a fixed-length coding)  
i.e. 00000 to 11111, so we need  $32 \times 5 = 160$  bits
- **Can we do better?** Yes, if some characters of  $\Sigma$  are **much more frequent** than others, use a **variable-length code**.

# Ambiguity in Variable-length Encoding

Suppose  $\Sigma = \{A, B, C, D\}$ .

- Suppose, the variable-length encoding used is  $\{0, 01, 10, 1\}$ .
- What does **001** represent?
  - A) AB
  - B) CD
  - C) AAD
  - D) None

# Ambiguity in Variable-length Encoding

Suppose  $\Sigma = \{A, B, C, D\}$ .

- Suppose, the variable-length encoding used is  $\{0, 01, 10, 1\}$ .
- What does **001** represent?
  - A) AB  $\rightarrow$  Leads to 001
  - B) CD
  - C) AAD  $\rightarrow$  Also leads to 001
  - D) None

# Prefix-Free Codes

- **Problem:** With variable-length codes, it is not clear where one character ends & the next one begins.

**Example:** **001** string in the previous slide

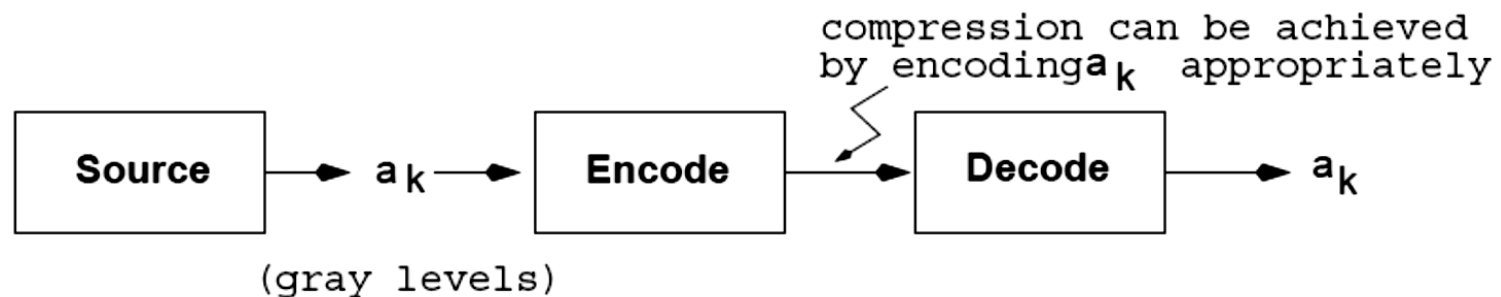
- **Solution:** *Prefix-free codes* - make sure that for every pair  $(i, j) \in \Sigma$ , neither of the encoding  $f(i)$ ,  $f(j)$  is a prefix of the other.

Example: **{0,10,110,111}**

# Huffman Coding

## (Reducing Coding Redundancy)

- A **variable-length coding** technique.
- Symbols are encoded **one at a time!**
  - There is a one-to-one correspondence between source symbols and code words
- **Optimal code** (i.e., it yields the smallest possible number of code symbols per source symbol).



# Huffman Coding

- Create series of source reductions

**Step-1:** Order the probabilities of the symbols under consideration.

**Step-2:** Combine the lowest two probabilities symbols.

**Step-3:** Repeat *Step-1* and *Step-2* until only two probabilities remain.

Original source		Source reduction			
Symbol	Probability	1	2	3	4
$a_2$	0.4	0.4	0.4	0.4	0.6
$a_6$	0.3	0.3	0.3	0.3	0.4
$a_1$	0.1	0.1	0.2	0.3	
$a_4$	0.1	0.1	0.1		
$a_3$	0.06	0.1			
$a_5$	0.04				



# Huffman Coding

- Code each reduced source

**Step-1:** Starting with the *highest reduced source* and working back to the *original source*, assign the unique codes to the symbols in the increasing order of code word length.

**Step-2:** Repeat *Step-1* for each reduced source until the original source is reached.

Original source			Source reduction							
Symbol	Probability	Code	1		2		3		4	
$a_2$	0.4	1	0.4	1	0.4	1	0.4	1	0.6	0
$a_6$	0.3	00	0.3	00	0.3	00	0.3	00	0.4	1
$a_1$	0.1	011	0.1	011	0.2	010	0.3	01		
$a_4$	0.1	0100	0.1	0100	0.1	011				
$a_3$	0.06	01010	0.1	0101						
$a_5$	0.04	01011								

51

# Huffman Coding

Symbols (like intensity levels)	Probabilities (sorted)	Source Reduction (do till two values are left) (Maintain in sorted order here as well)			
		1	2	3	4
a2	0.4    1	0.4	0.4	0.4	0.6    0
a6	0.3    00	0.3	0.3	0.3    00	0.4    1
a1	0.1    011	0.1	0.2    010	0.3    01	
a4	0.1    0100	0.1    0100	0.1    011		
a3	0.06    01010	0.1    0101			
a5	0.04    01011				

# Huffman Coding

- The **average length** of Huffman code for the previous example is:

$$\begin{aligned} L_{\text{avg}} &= (0.4)(1) + (0.3)(2) + (0.1)(3) + (0.1)(4) + (0.06)(5) + (0.04)(5) \\ &= 2.2 \text{ bits/pixel} \end{aligned}$$

- The Huffman coding table *must be known* to the decoder to retrieve (decompress) the original data.
- Any string of Huffman encoded symbols can be decoded by examining the individual symbols of the string in a *left-to-right manner*.

**Example:** encoded string is 010100111100

decoded string is  $a_3 a_1 a_2 a_2 a_6$

Symbol	Code
$a_2$	1
$a_6$	00
$a_1$	011
$a_4$	0100
$a_3$	01010
$a_5$	01011

# Huffman Coding

Message	Stage - I	Stage - II	Stage - III	Stage - IV	Stage - V
$x_1$	0.3	0.3	0.3	0.45	0.55
$x_2$	0.25	0.25	0.25	0.3	0.45
$x_3$	0.2	0.2	0.25	0.25	
$x_4$	0.12	0.13	0.2		
$x_5$	0.08	0.12			
$x_6$	0.05				

Write the codes for each symbols

# Huffman Coding/Decoding

- **Coding/decoding** can be implemented using a predefined look-up table.
- Decoding can be done unambiguously.

# Next Lecture

- The image degradation/restoration model
- Noise models
  - Important noise probability density functions
  - Periodic noise
  - Estimating noise parameters
- Restoration using spatial filters
  - Mean filters
  - Order-static filters
  - Adaptive filters