

CS380
Artificial Intelligence for Games

Adversarial Search




Game Playing State-of-the-Arts

- **Chess:** 1997: Deep Blue defeats human champion **Gary Kasparov** in a six-game match. Deep Blue examined 200 million positions per second, used very sophisticated evaluation and undisclosed methods for extending some lines of search up to 40 ply



Game Playing State-of-the-Arts

- **Go:** Human champions are now starting to be challenged by machines, though the best humans still beat the best machines. In Go, number of terminal states $> 300!$ Classic programs use pattern knowledge bases, but big recent advances use Monte Carlo (randomized) expansion methods



Applications

- Two player
 - Player or Human, Computer or AI
- Players take turns when playing
 - Somebody goes first. Who? Flip a coin
- Zero-sum
 - if one player loses, the other player wins
- Perfect information
 - Each player is completely informed of previous moves
- Deterministic
 - No randomness, follows a strict pattern
- Have small number of possible actions
- Precise, formal rules

Applications

	Perfect Information	Imperfect Information
Deterministic	Tic-Tac-Toe, Go, Chess, Checkers, Othello, Kalah, NIM	Monopoly, Aeroplane Chess, Stratego
Non-deterministic	Backgammon	Poker, Scrabble

Challenge


- Challenging due to huge search space
 - Chess:
 - branching factor = 35,
 - number of moves per player = 50
- Search optimization
 - Alpha-Beta pruning
- If unlikely to find goal, time limit is introduced
 - so a terminal state is approximated

Games as a Search Problem

- Game has
 - **initial state**
 - configuration + player to move
 - **successor function**
 - defines the set of possible moves/actions/states from a state
 - **terminal test**
 - game over or not
 - **utility function**
 - maps terminal states to numeric values
 - win : 1, 10, +INF, loss : -1, -10, -INF, draw : 0
- Players have to find **contingent strategies**
 - Devised for a specific situation where things could go wrong
 - An optimal strategy assumes an infallible opponent

Game Tree

- Game tree = initial state + legal moves or actions
- Given tree is one move deep, consisting of two halfmoves, each of which is called a ply

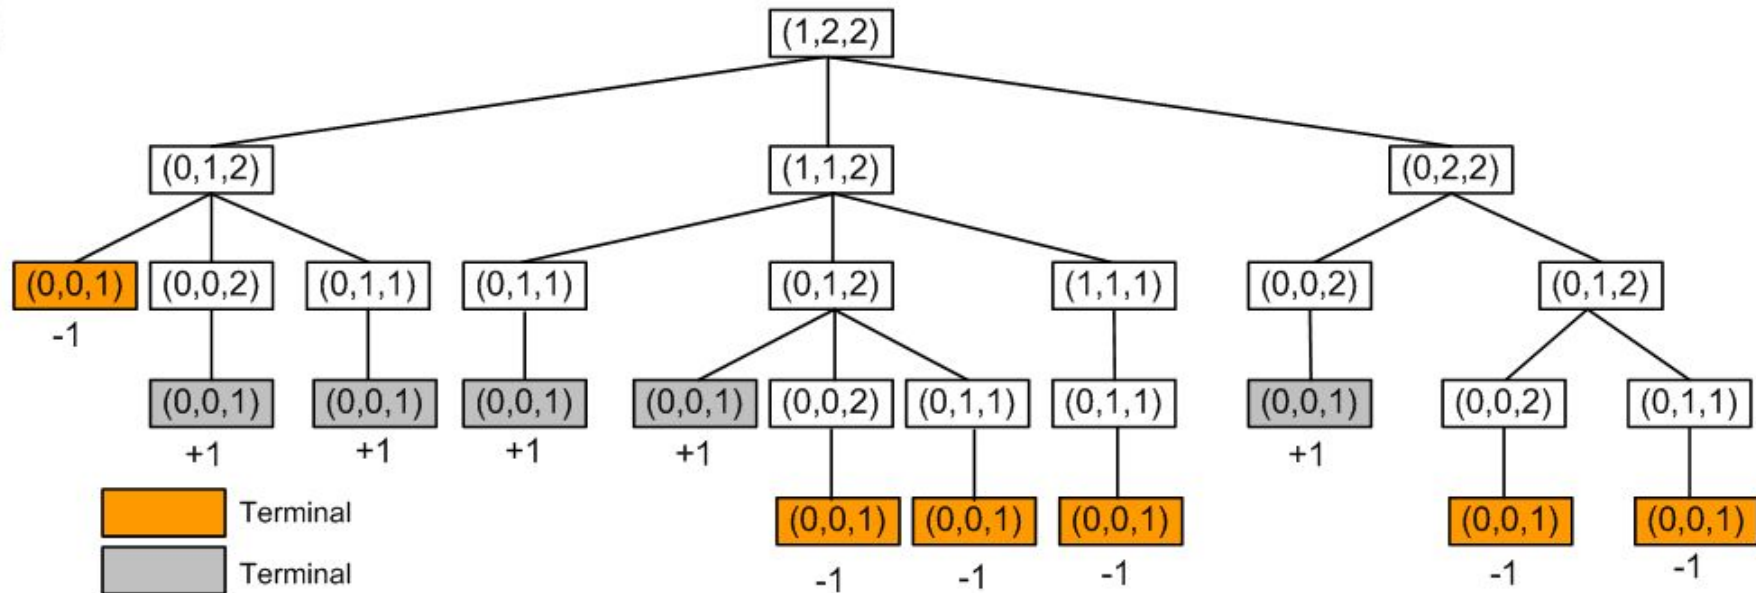


Game of NIM

- Several piles of sticks are given. We represent the configuration of the piles by a **monotone sequence** of integers, such as (1,3,5) or (0,0,1)
- A player may remove, in one turn, any number of sticks from one pile.
 - Thus, (1,3,5) would become (1,1,3) if the player were to remove 4 sticks from the last pile. The player who takes the last stick loses.
- Let's represent the game with initial state (1, 2, 2) as a game tree



Example: Game of NIM




Tic-Tac-Toe Game

- Tic-Tac-Toe is well-known pencil-and-paper game for two players, who take turns marking the spaces in a 3×3 grid.
- The player who succeeds in placing three of their marks in a diagonal, horizontal, or vertical row is the winner.

x	o	o
	o	x
o	x	x

Example: Tic-Tac-Toe



How to Choose the Optimal Decision?

Minimax Search

How to Choose the Optimal Decision?

- Player has always a role in complicating the trajectory of Computer to reach a win
- Computer needs to consider every possible response by Player
- Computer must take in consideration that Player is always playing his best moves when he has turn


Minimax Search

- Minimax search guarantees the perfect play for deterministic, perfect-information games
- Idea: choose move to position with highest **minimax value** = best achievable payoff against best play

$$\text{MINIMAX}(s) = \begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s,a)) & \text{if } \text{PLAYER-TURN}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s,a)) & \text{if } \text{PLAYER-TURN}(s) = \text{MIN} \end{cases}$$

Minimax Search

$$\text{MINIMAX}(s) = \begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s,a)) & \text{if } \text{PLAYER-TURN}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s,a)) & \text{if } \text{PLAYER-TURN}(s) = \text{MIN} \end{cases}$$



Minimax Search Algorithm

function MINIMAX-DECISION(*state*) *returns an action*

$v \leftarrow \text{MAX-VALUE}(\textit{state})$

return the *action* in SUCCESSORS(*state*) with value *v*

function MAX-VALUE(*state*) *returns a utility value*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

for *a, s* in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s))$

return *v*

function MIN-VALUE(*state*) *returns a utility value*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

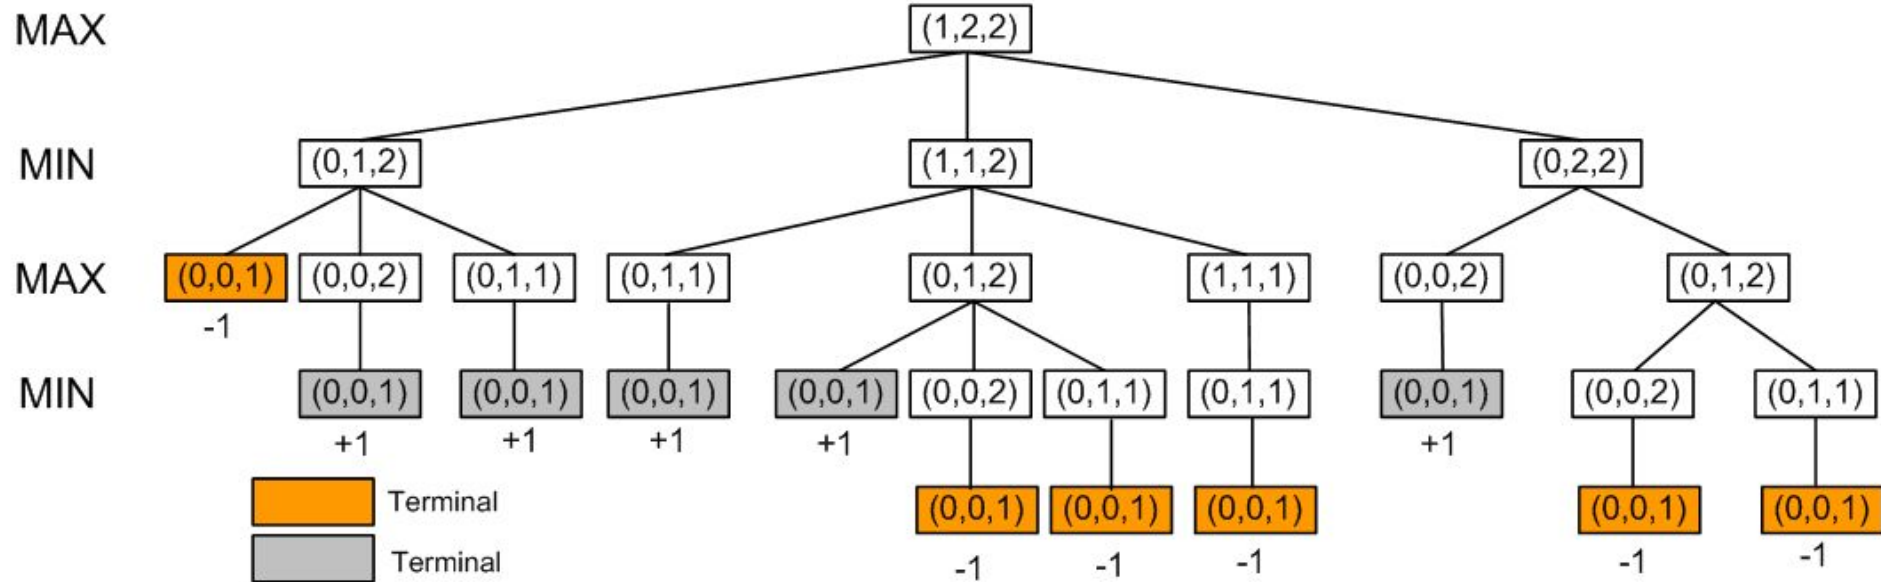
$v \leftarrow \infty$

for *a, s* in SUCCESSORS(*state*) **do**

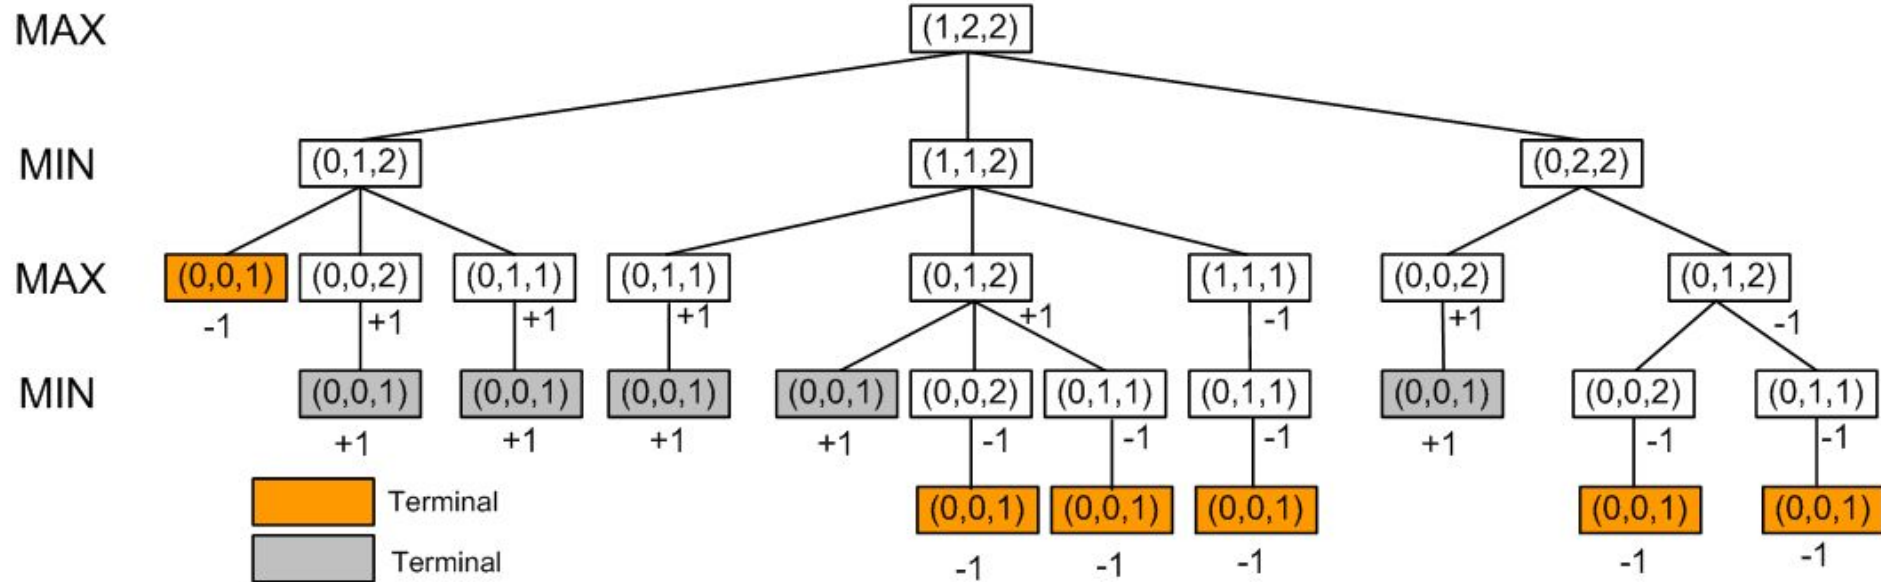
$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s))$

return *v*

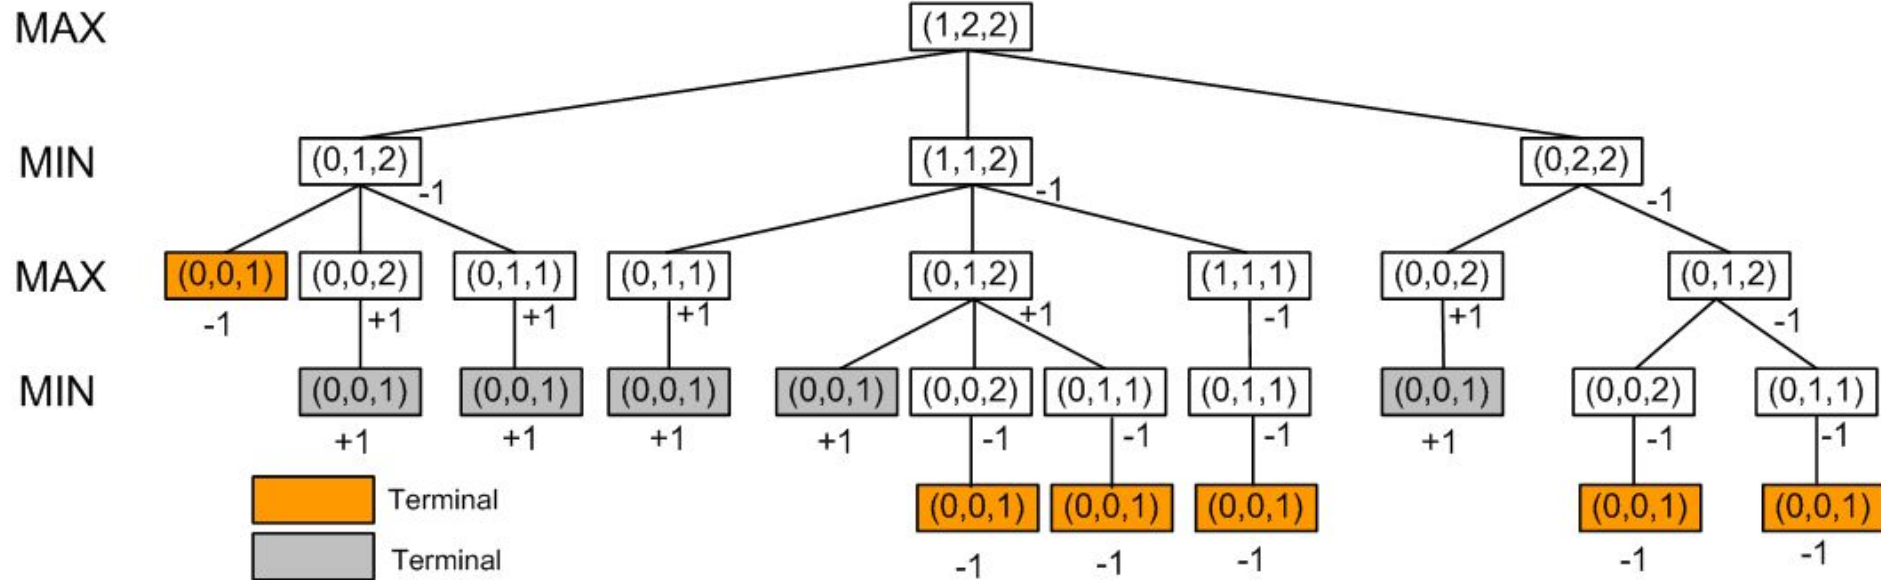
Example: Game of NIM



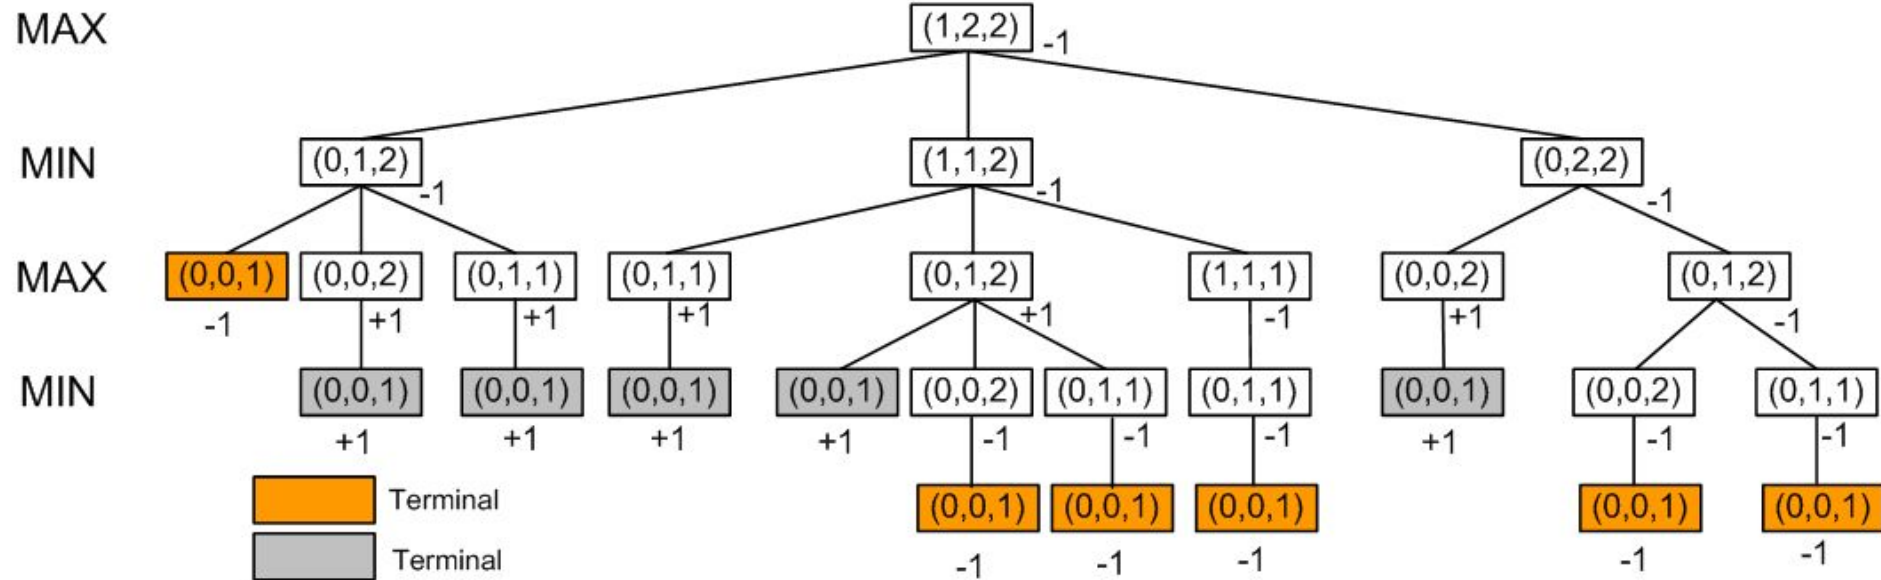
Example: Game of NIM




Example: Game of NIM



Example: Game of NIM



Example: A Partial Game Tree for Tic-Tac-Toe



Properties of Minimax

- Complete?
 - Yes, if the tree is finite
- Optimal?
 - Yes, if against an optimal opponent.
 - Otherwise?
- Assume the branching factor is k and the depth of the tree is d
 - Time complexity?
 - $O(k^d)$
 - Space complexity?
 - $O(kd)$ (depth-first exploration)

Properties of Minimax

- For tic-tac-toe
 - $k = ?, d = ?$
- For chess
 - $k \approx 35, d \approx 100$ for “reasonable” games
 - Exact solution completely infeasible
- For GO
 - $k \approx 200, d \approx 100$ for “reasonable” games
- Can we do better?


α - β Pruning

Search optimisation techniques for choosing a good move (may not be optimal) when time is limited

α - β Pruning

- The number of game states with minimax search is exponential in the # of moves
- Is it possible to compute the correct minimax decision without looking at every node in the game tree?
- Need to **prune** away branches that cannot possibly influence the final decision

Motivating Example




```

Line 1  Minimax(MAX)=
Line 2      MAX_Value(
Line 3          MIN(7,50,25),    //Child1
Line 4          MIN(5,10,15),    //Child2
Line 5          MIN(75,12,5)     //Child3
Line 6      )
    
```

$\text{MAX}(7, 5 \text{ or less}, \dots) = \text{MAX}(7, \dots)$ //after expanding 5

Basic Idea of α - β Pruning

- Consider a node **n** such that Player has a choice of moving to
- If Player has a **better choice m** either at the parent of n or at any choice point further up, then **n will never be reached in actual play**
- α - β pruning gets its name from the two parameters that describe bounds on the backed-up values



α - β Pruning

- α = the value of the best (**highest-value**) choice we have found so far at any choice point along the path for **MAX**
- β = the value of the best (**lowest-value**) choice we have found so far at any choice point along the path for **MIN**
- α - β search updates the values of α and β as it goes along and **prunes the remaining branches** at a node as soon as the value of the current node is **worse than the current α or β for MAX or MIN** respectively

function ALPHA-BETA-SEARCH(*state*) *returns an action*

inputs: *state*, current state in game

$v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$

return the *action* in SUCCESSORS(*state*) with value *v*

function MAX-VALUE(*state*, α , β) *returns a utility value*

inputs: *state*, current state in game

α , the value of the best alternative for MAX along the path to *state*

β , the value of the best alternative for MIN along the path to *state*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow -\infty$

for *a, s* in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(s, \alpha, \beta))$

if $v \geq \beta$ **then return** *v*

$\alpha \leftarrow \text{MAX}(\alpha, v)$

return *v*

function MIN-VALUE(*state*, α , β) *returns a utility value*

inputs: *state*, current state in game

α , the value of the best alternative for MAX along the path to *state*

β , the value of the best alternative for MIN along the path to *state*

if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)

$v \leftarrow +\infty$

for *a, s* in SUCCESSORS(*state*) **do**

$v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(s, \alpha, \beta))$

if $v \leq \alpha$ **then return** *v*

$\beta \leftarrow \text{MIN}(\beta, v)$

return *v*

α - β Pruning Example

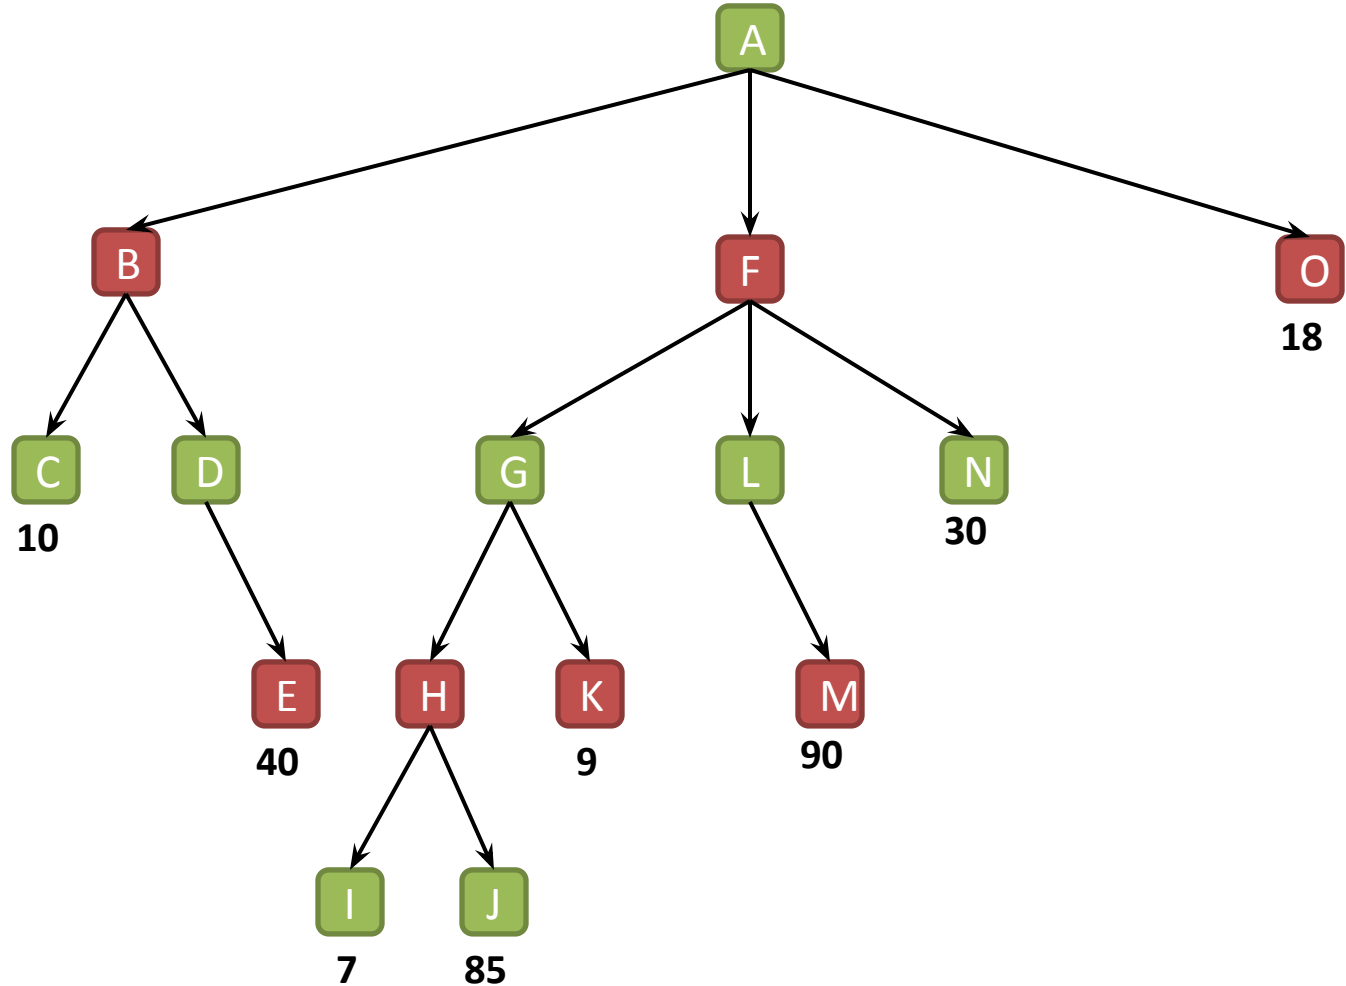
α - β Pruning

MAX

MIN

MAX

MIN



α - β Pruning


$\alpha: -\infty / \beta: +\infty$

MAX

MIN

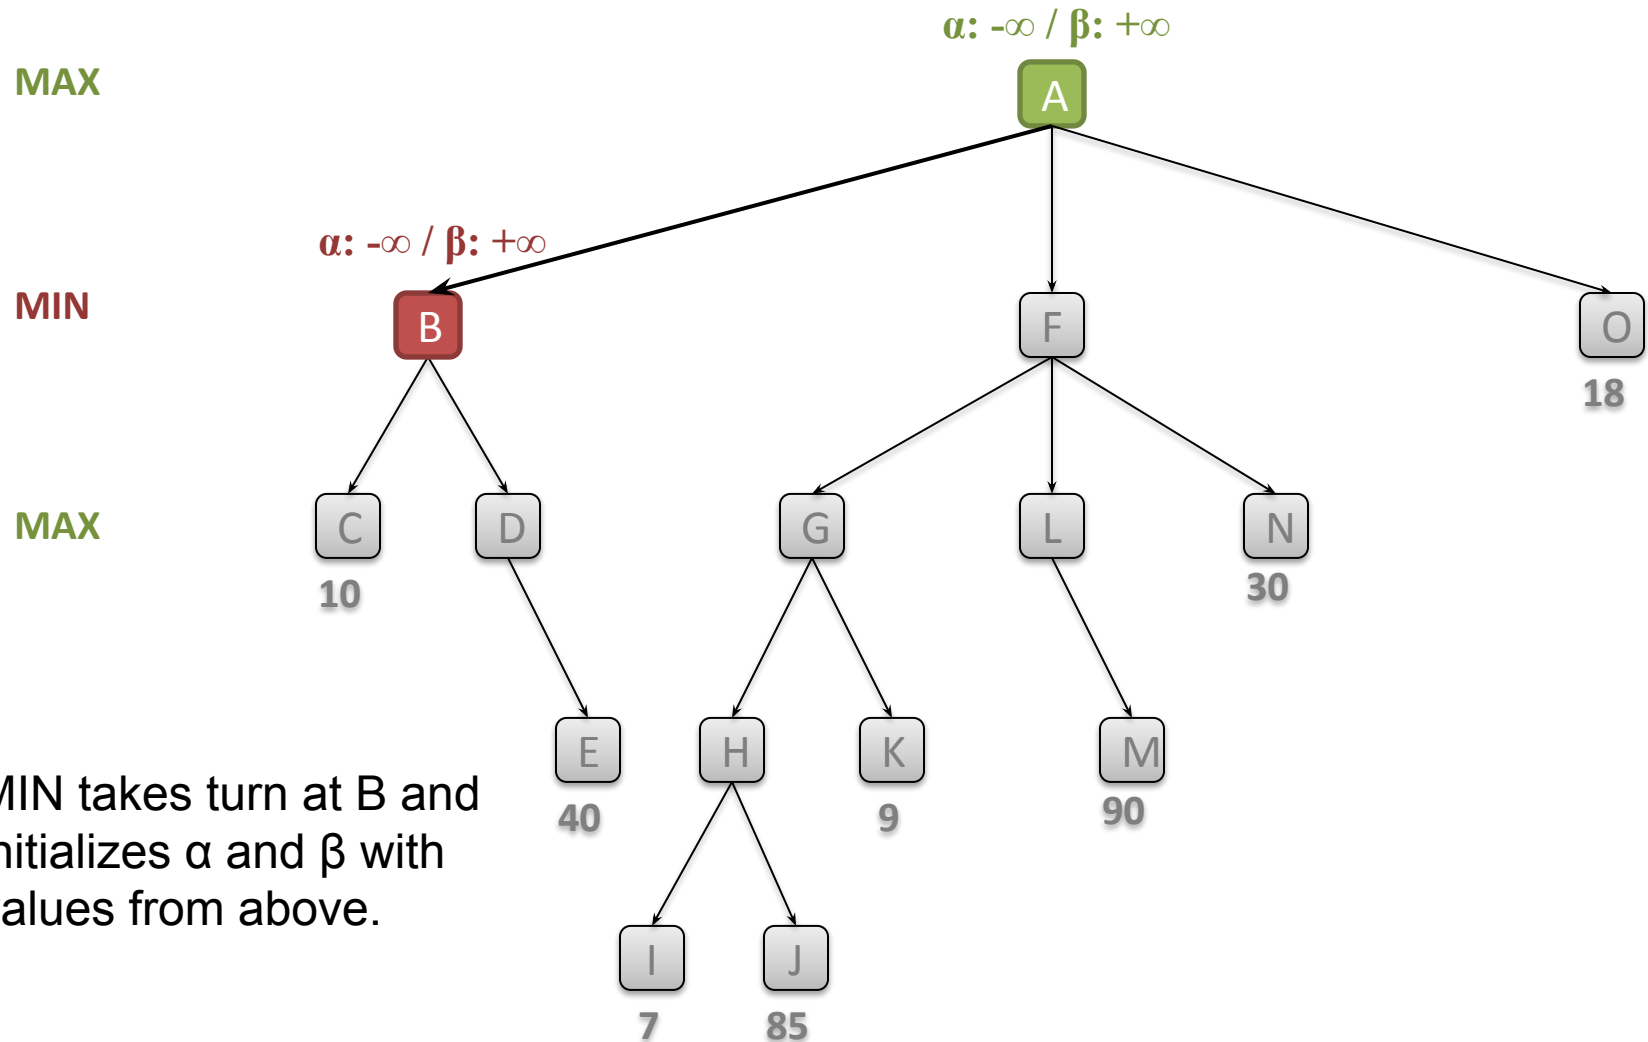
MAX

MIN

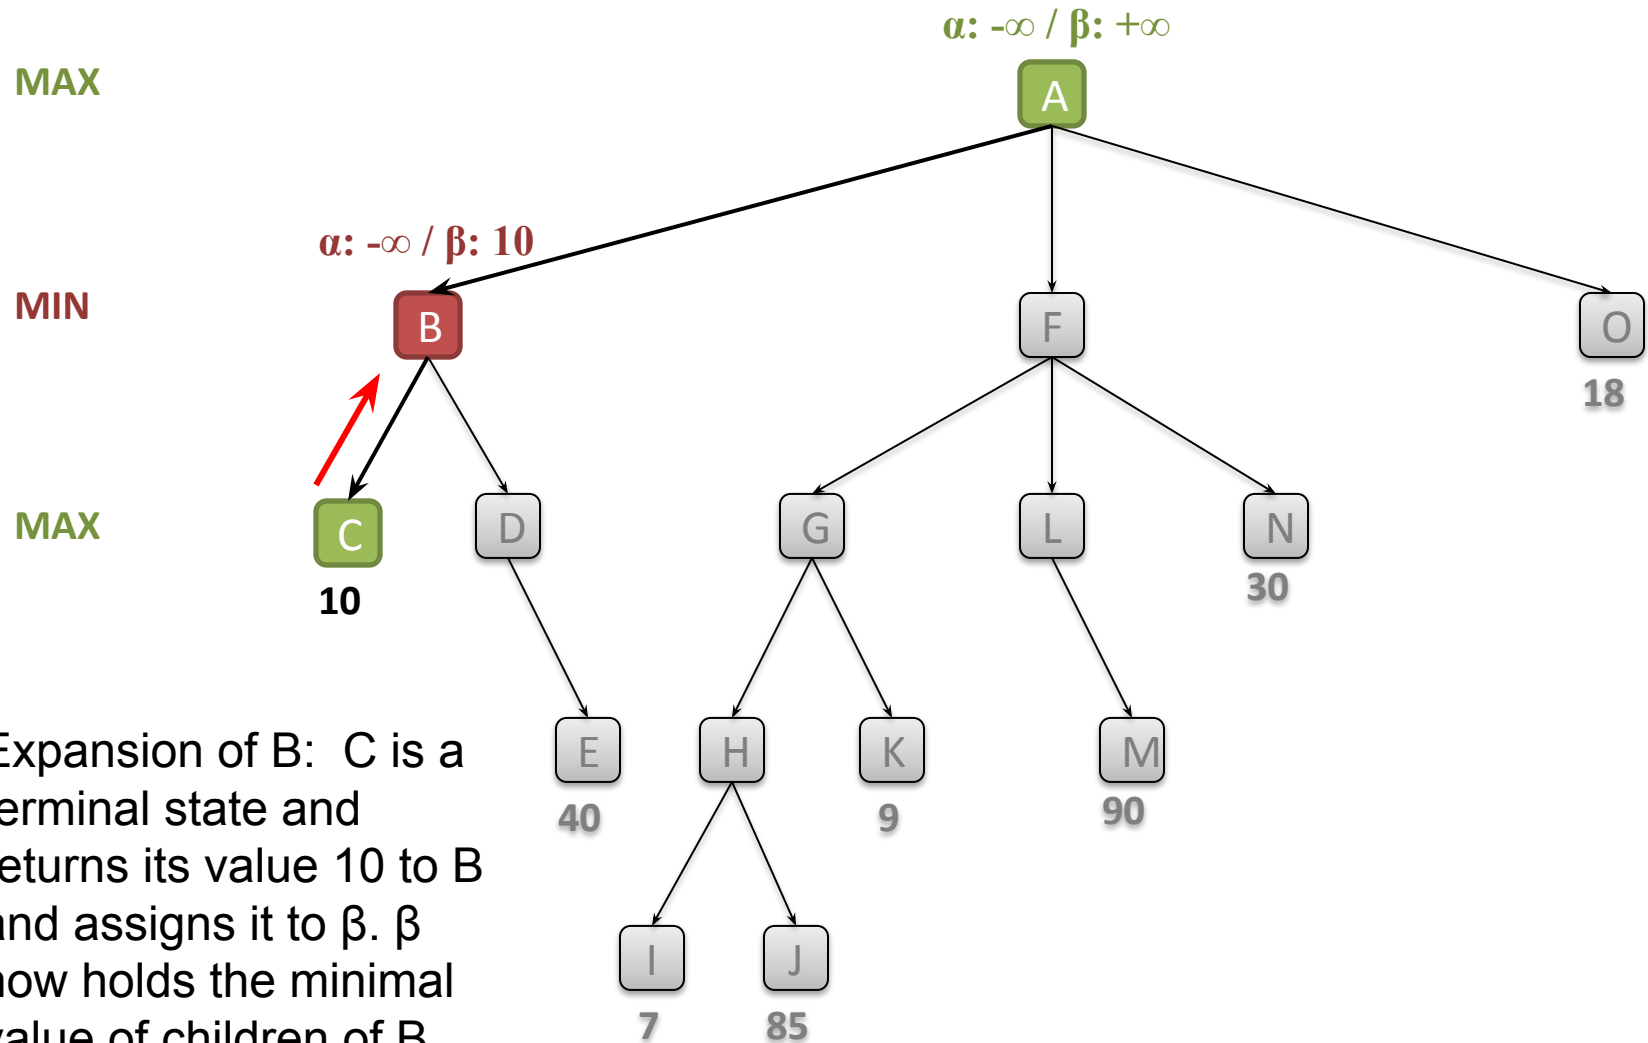


Expansion of A is started with initialization of α and β .

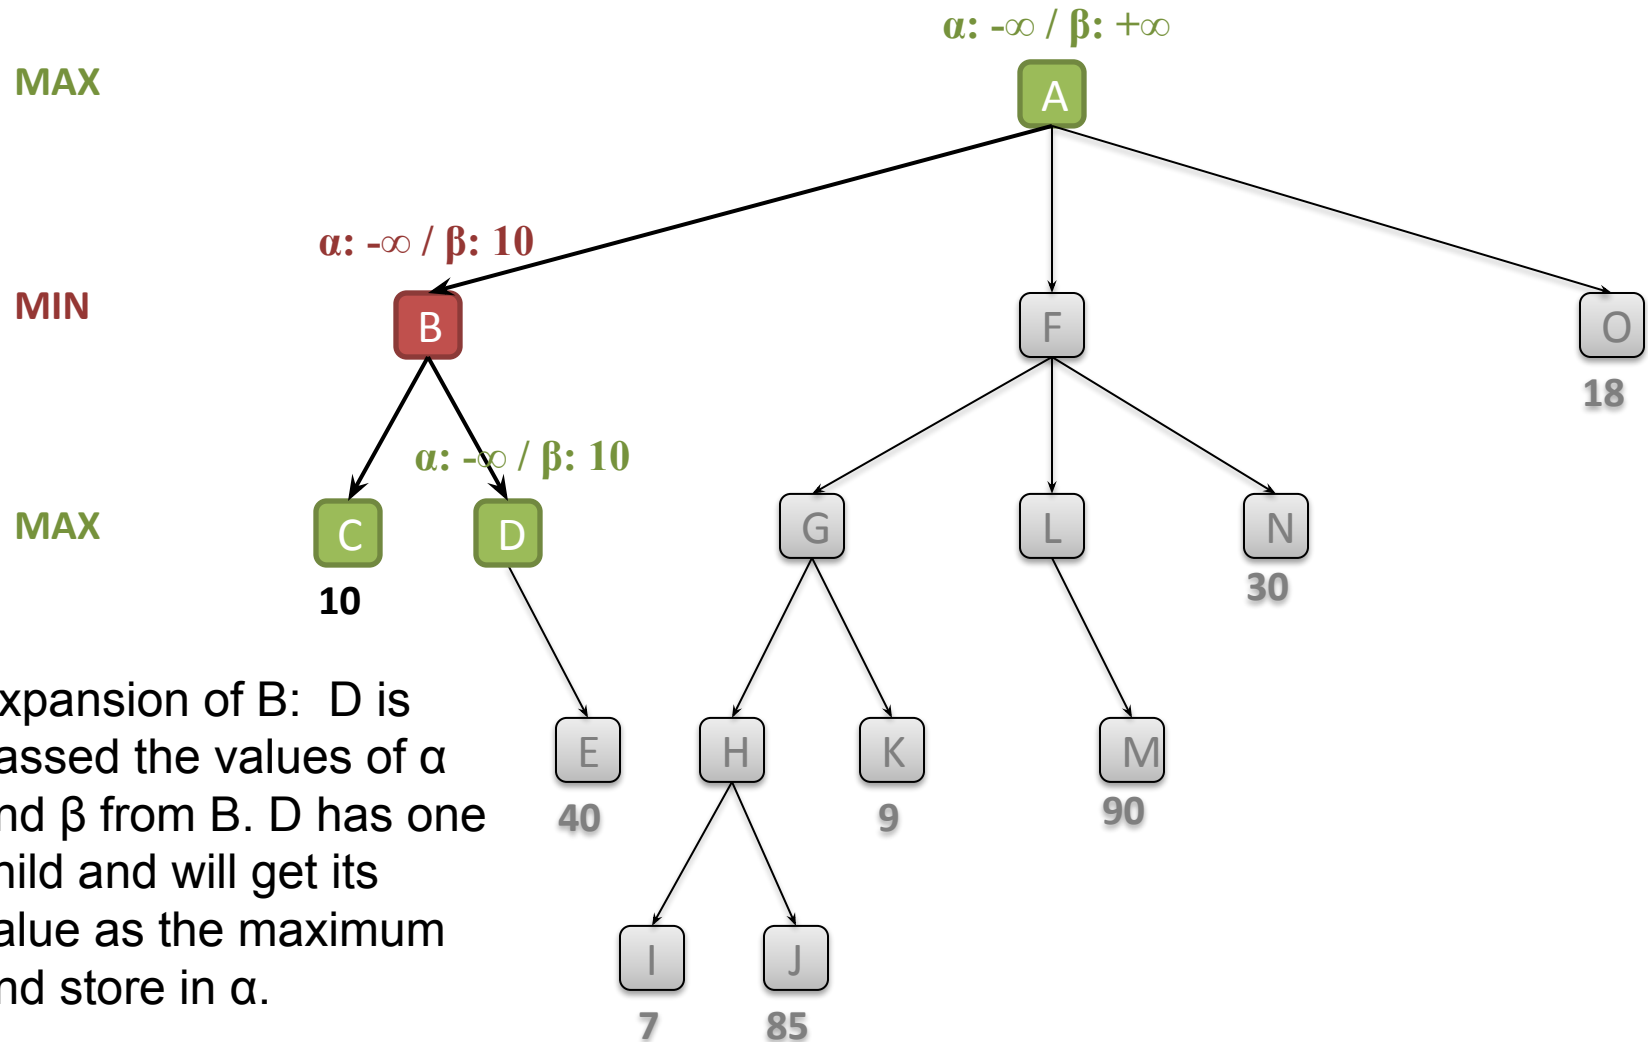
α - β Pruning



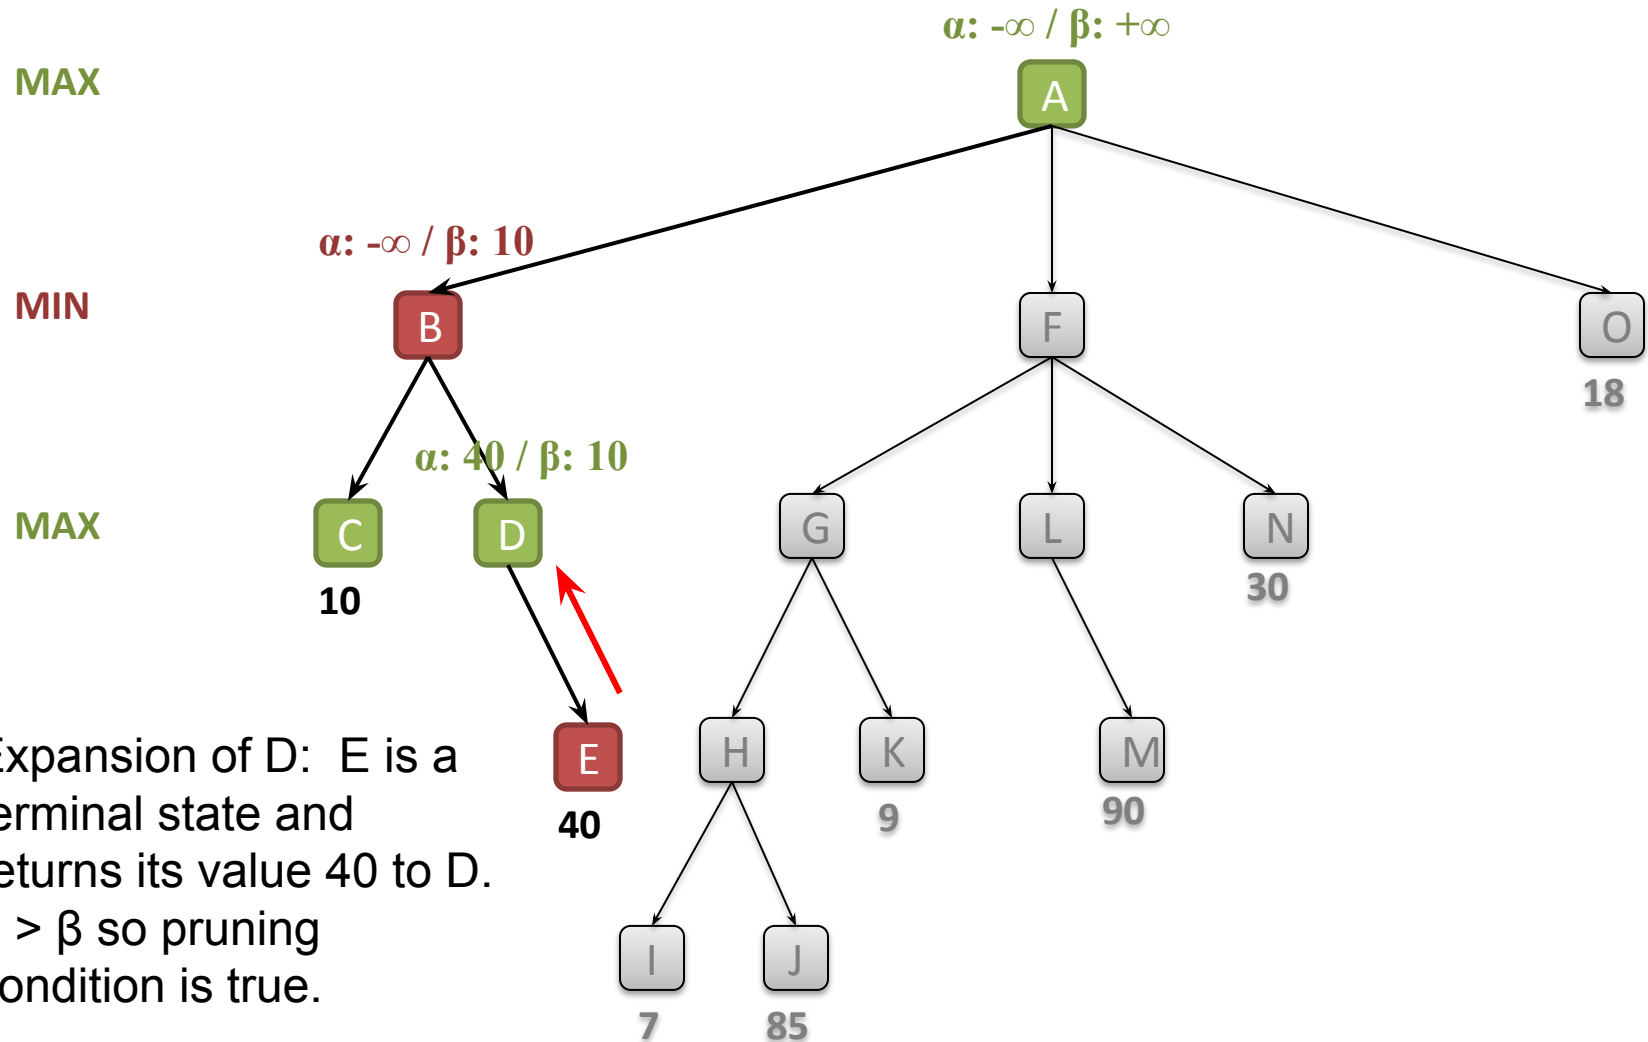
α - β Pruning



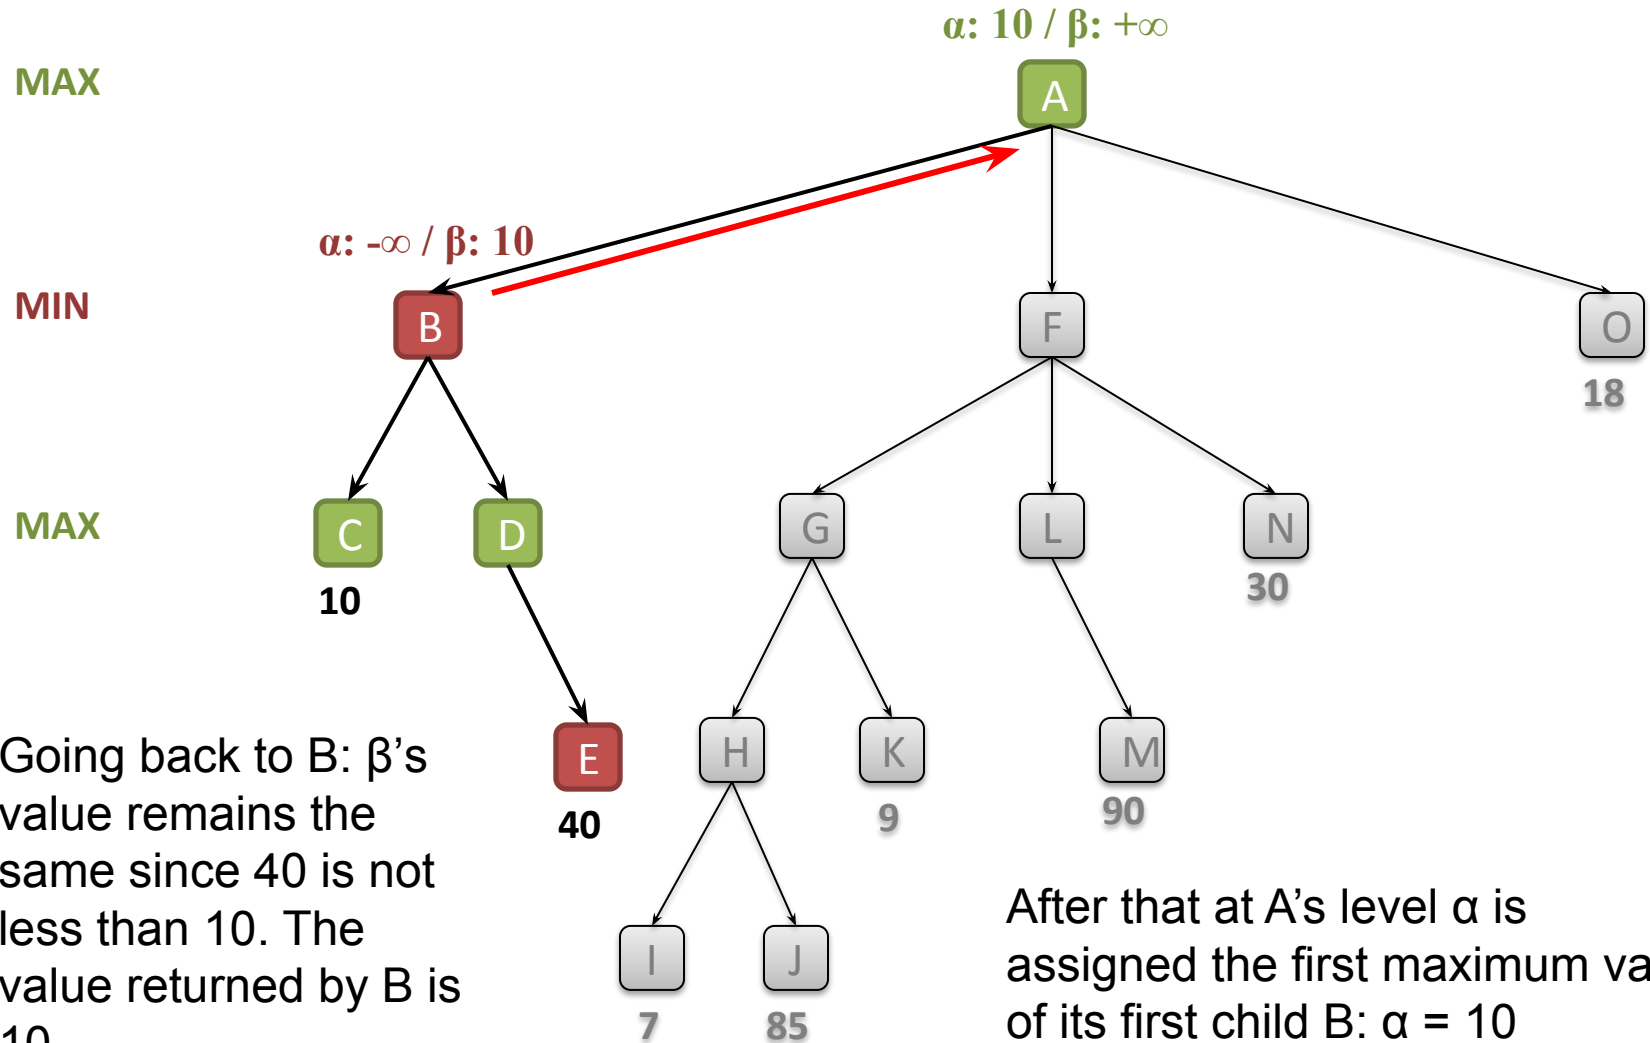
α - β Pruning



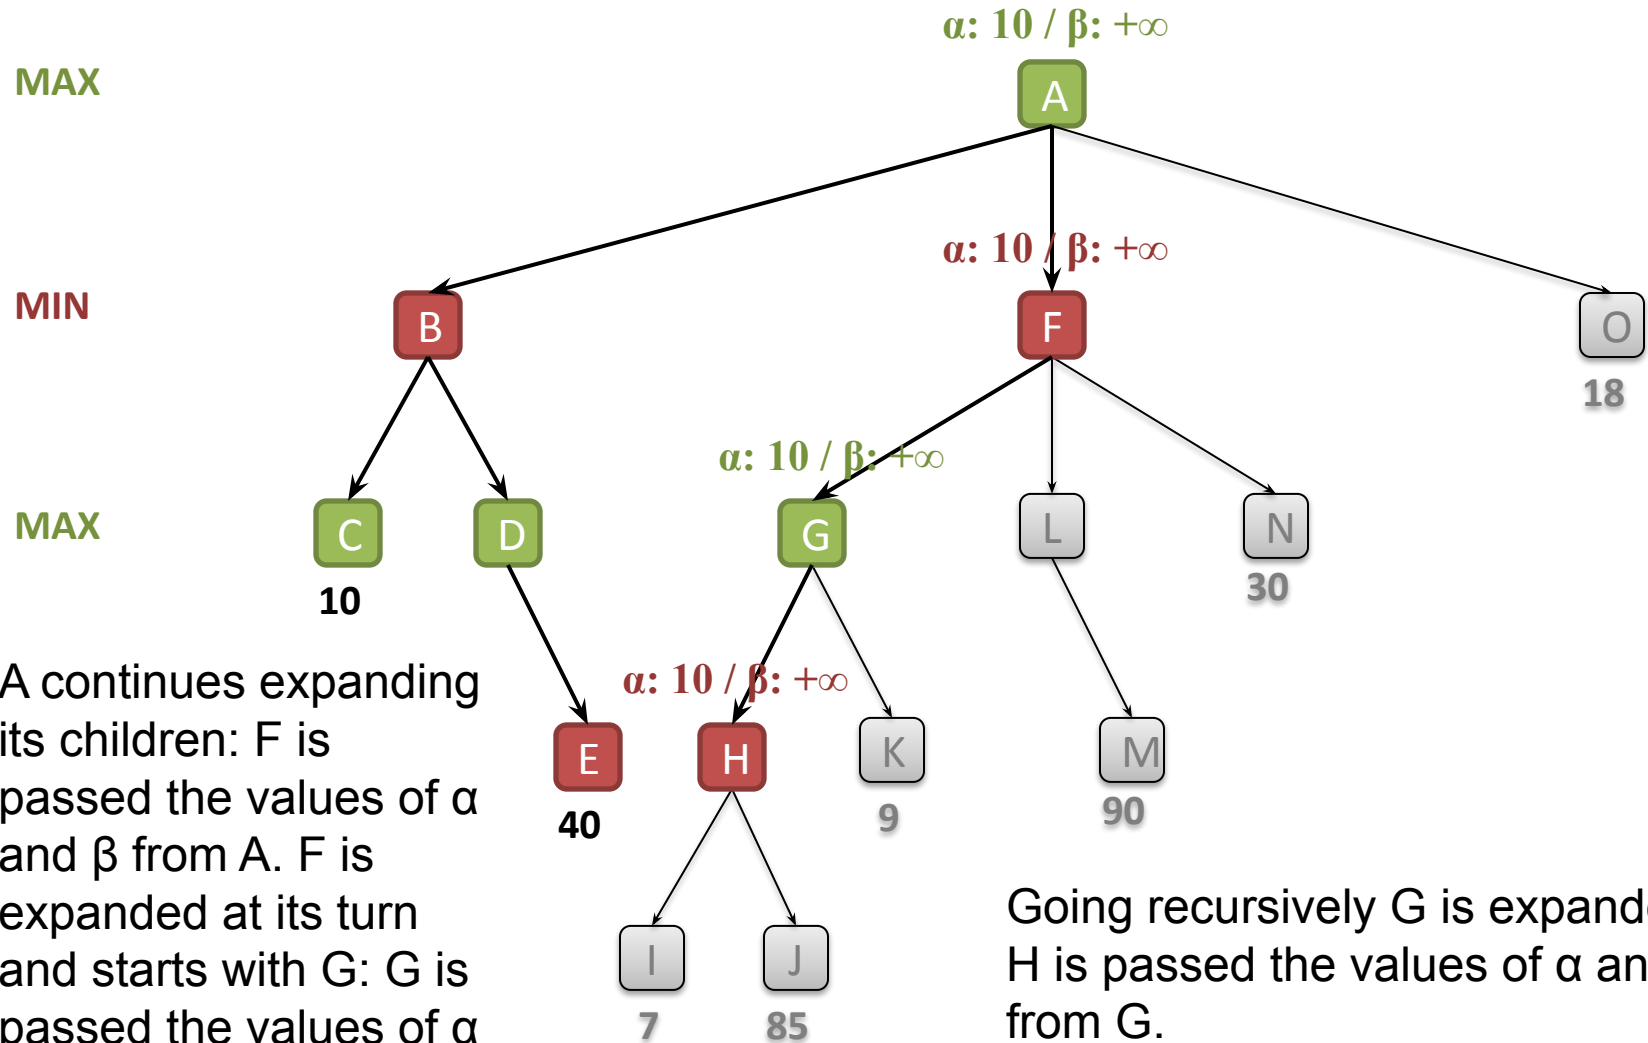
α - β Pruning



α - β Pruning



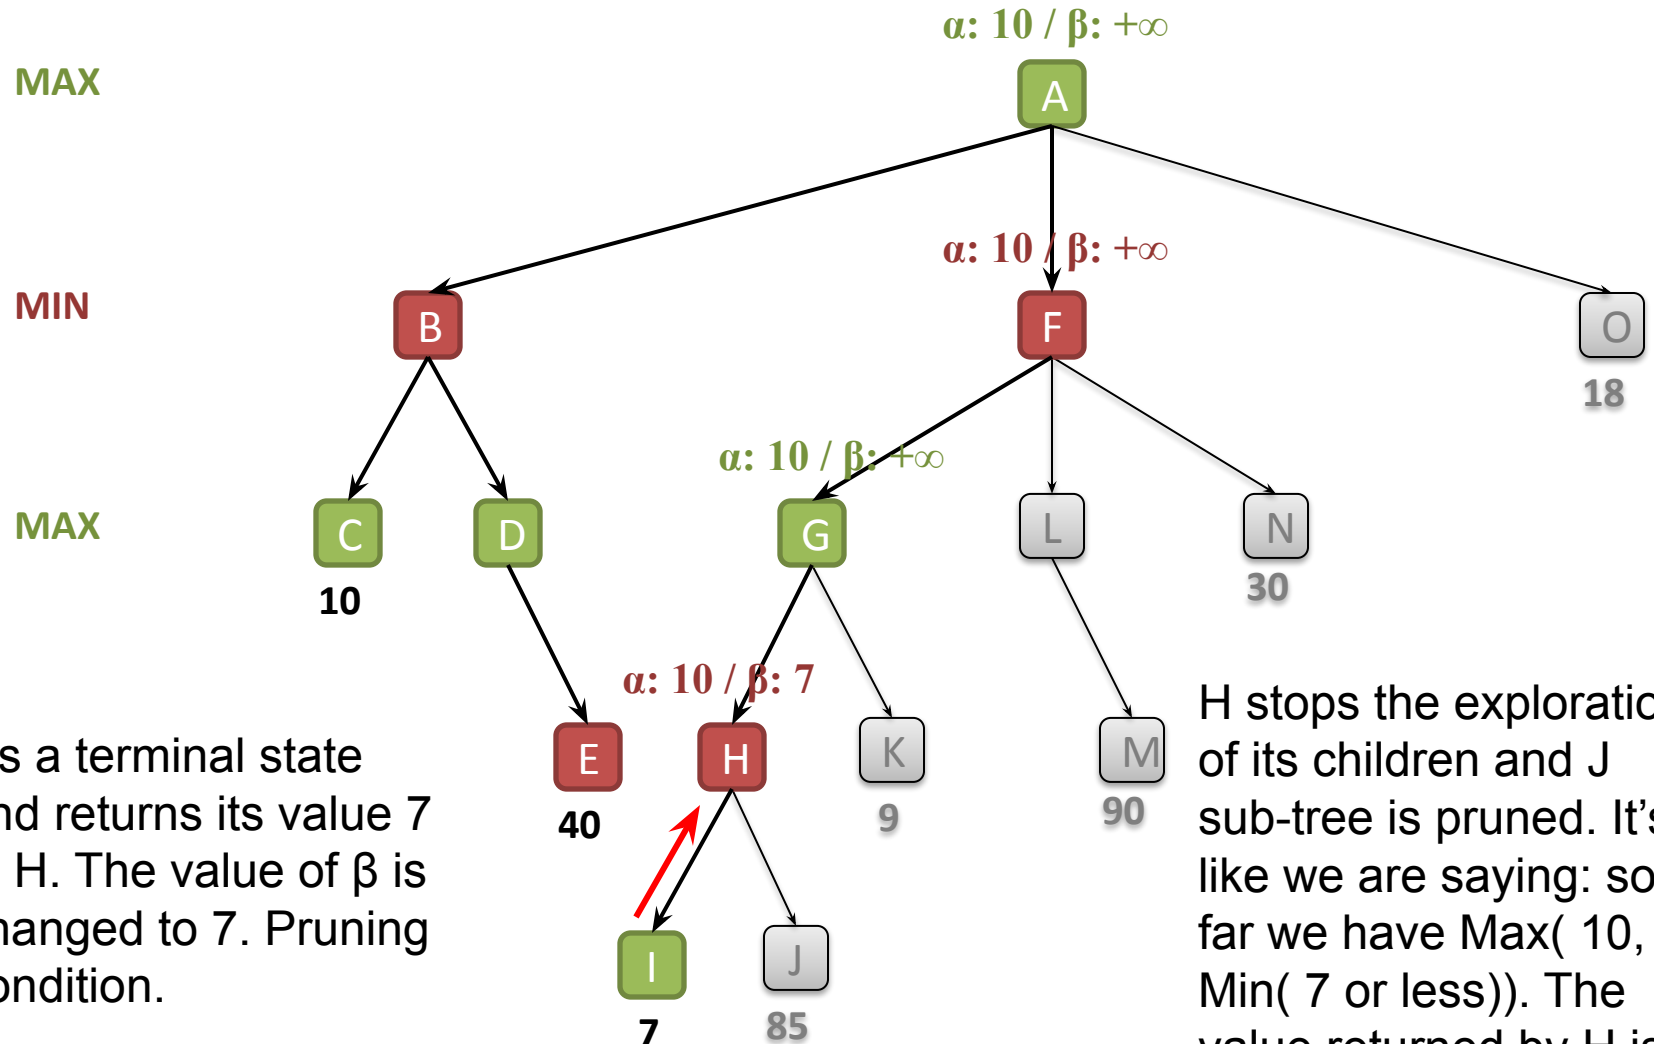
α - β Pruning



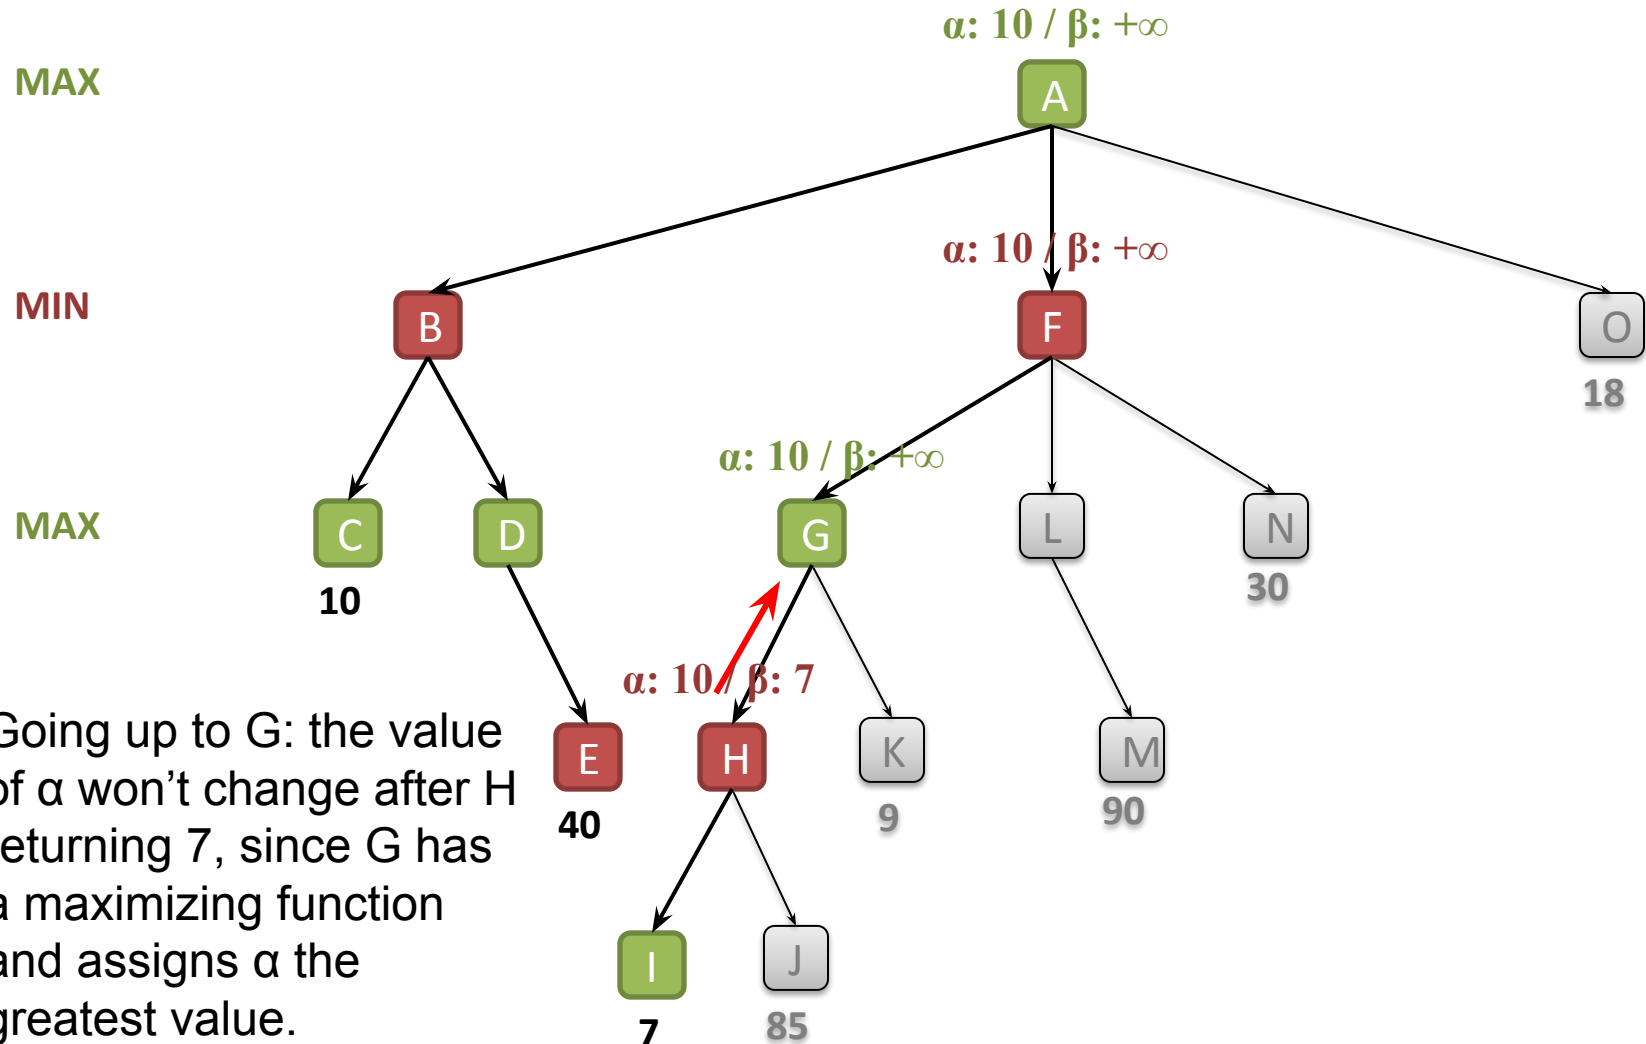
A continues expanding its children: F is passed the values of α and β from A. F is expanded at its turn and starts with G: G is passed the values of α and β from F.

Going recursively G is expanded: H is passed the values of α and β from G.

α - β Pruning

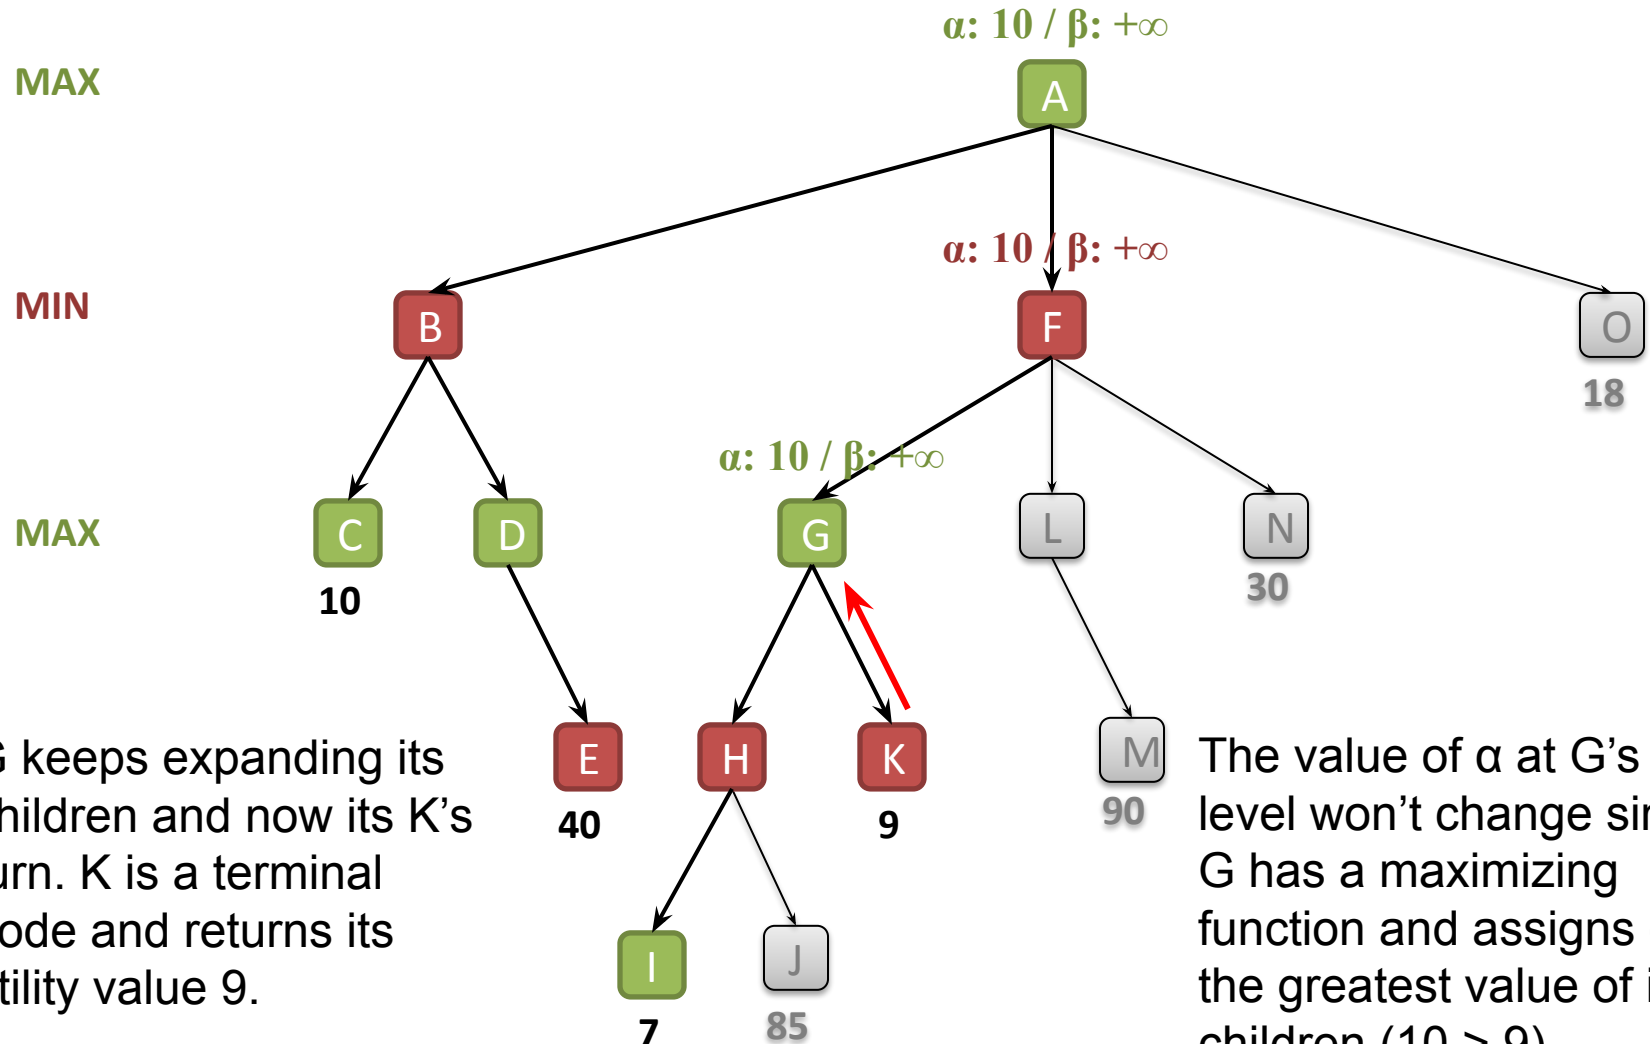


α - β Pruning

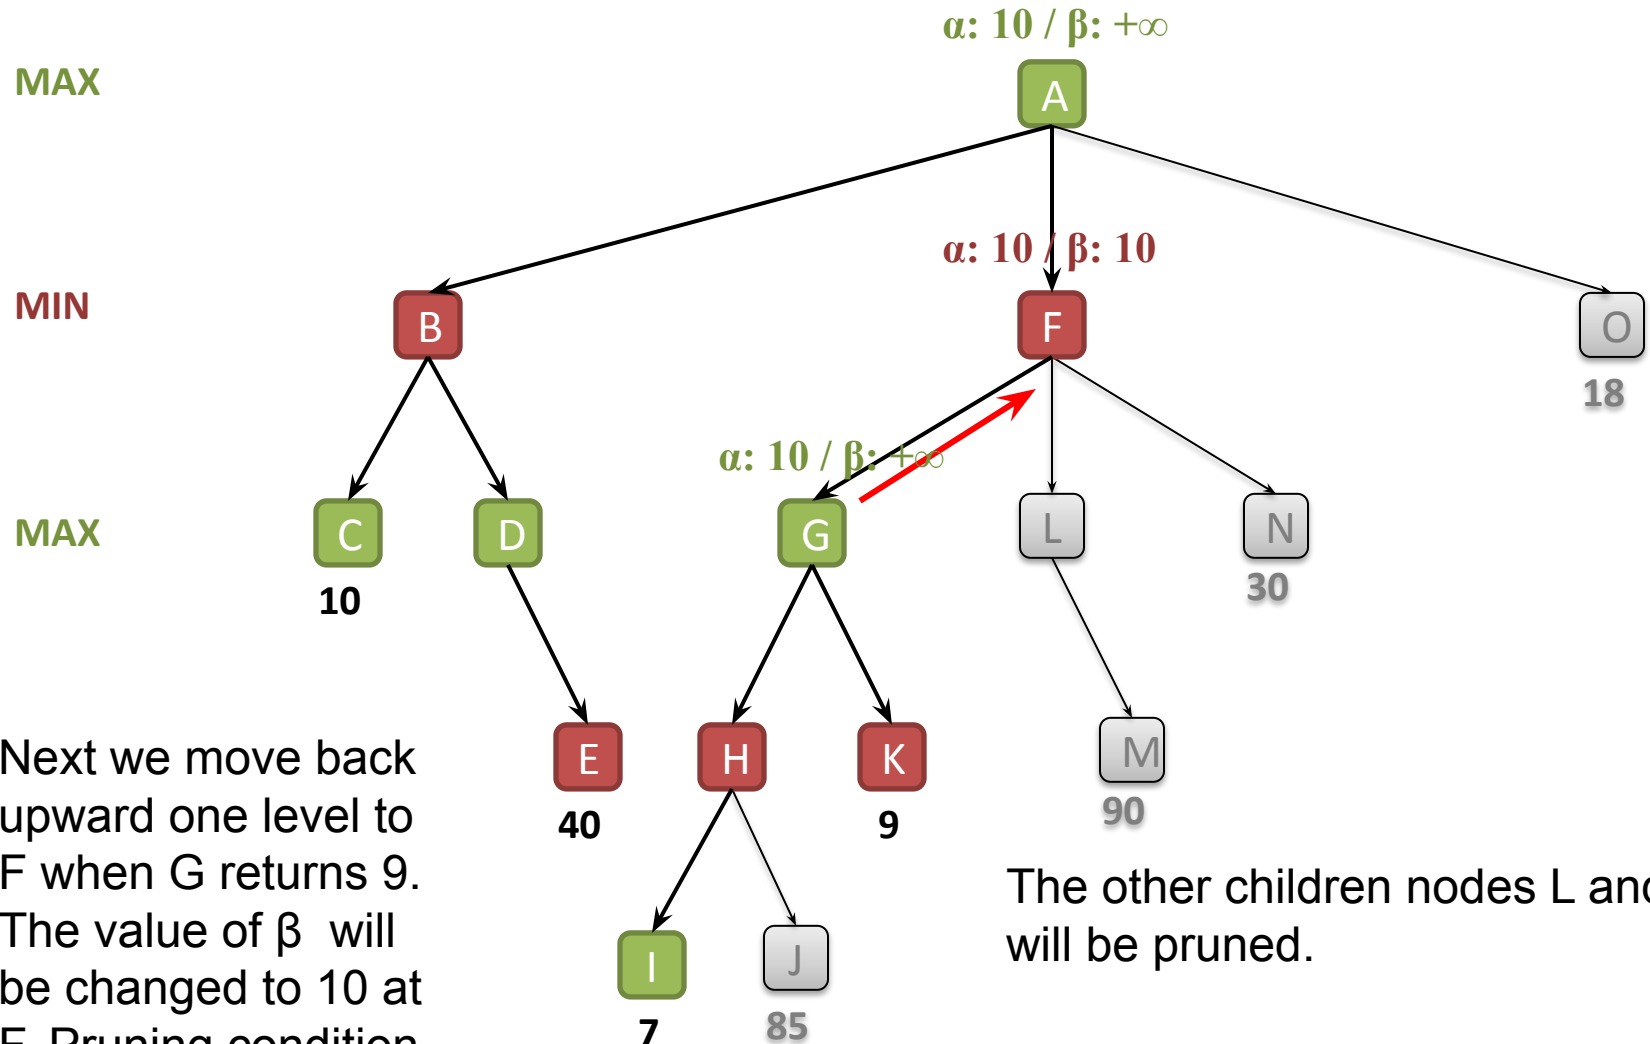


Going up to G: the value of α won't change after H returning 7, since G has a maximizing function and assigns α the greatest value.

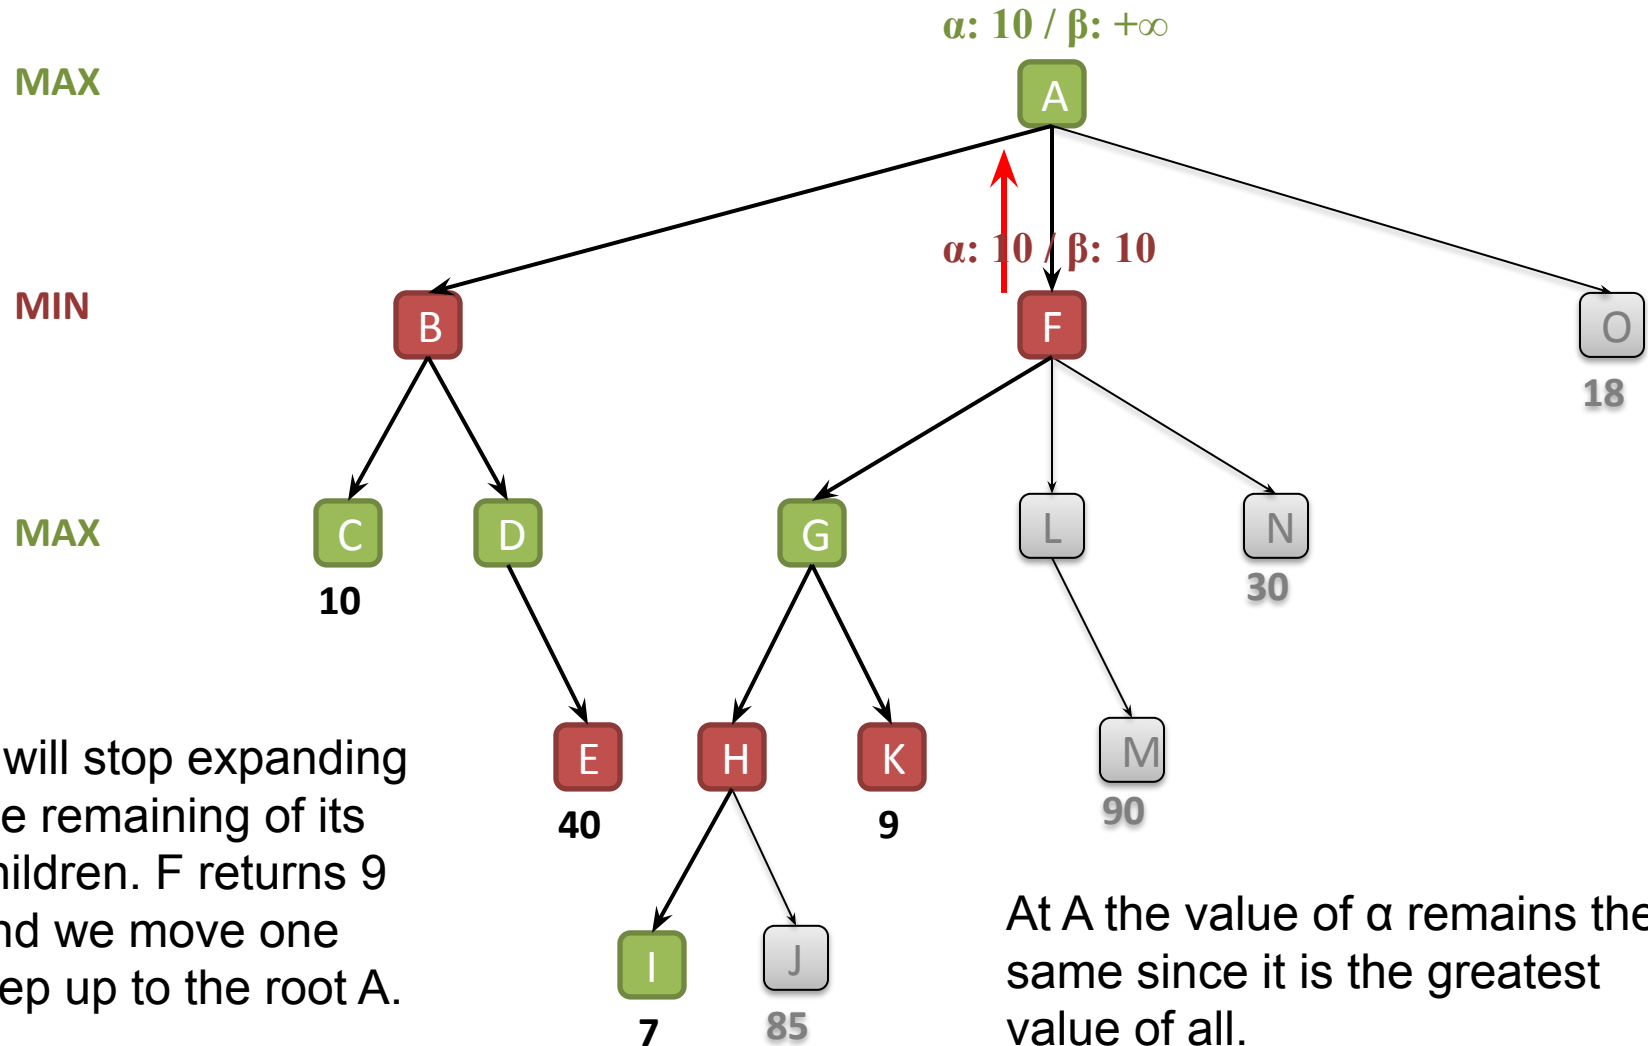
α - β Pruning



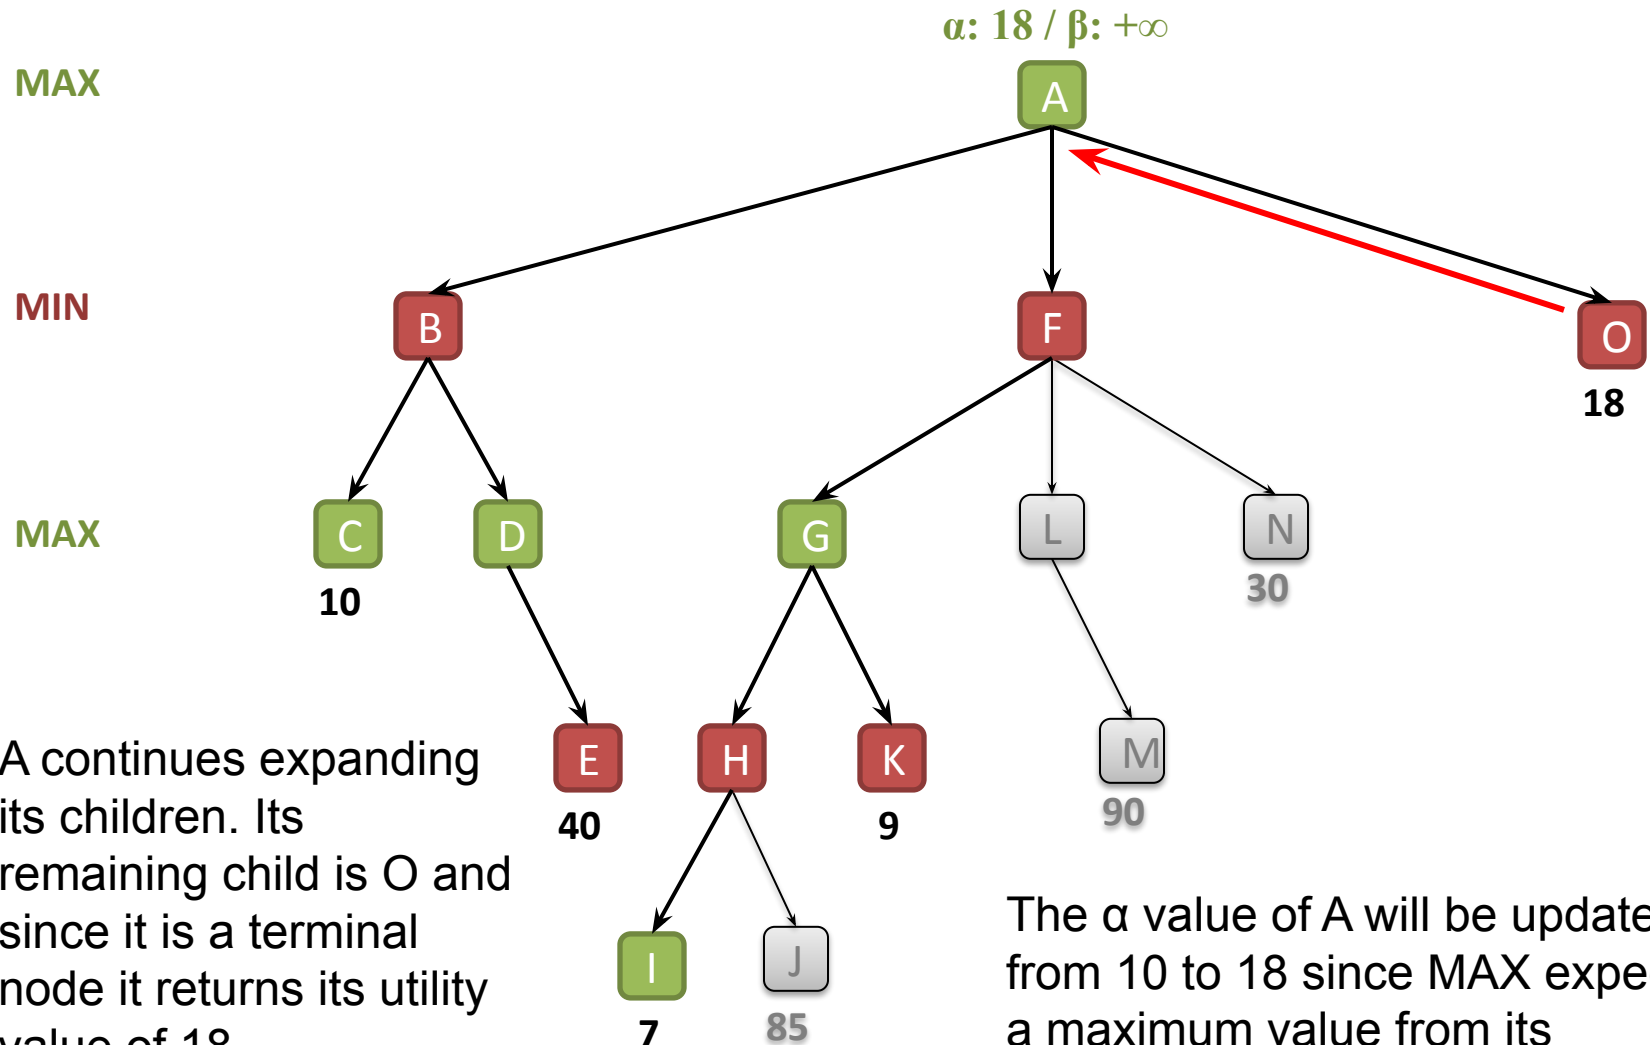
α - β Pruning



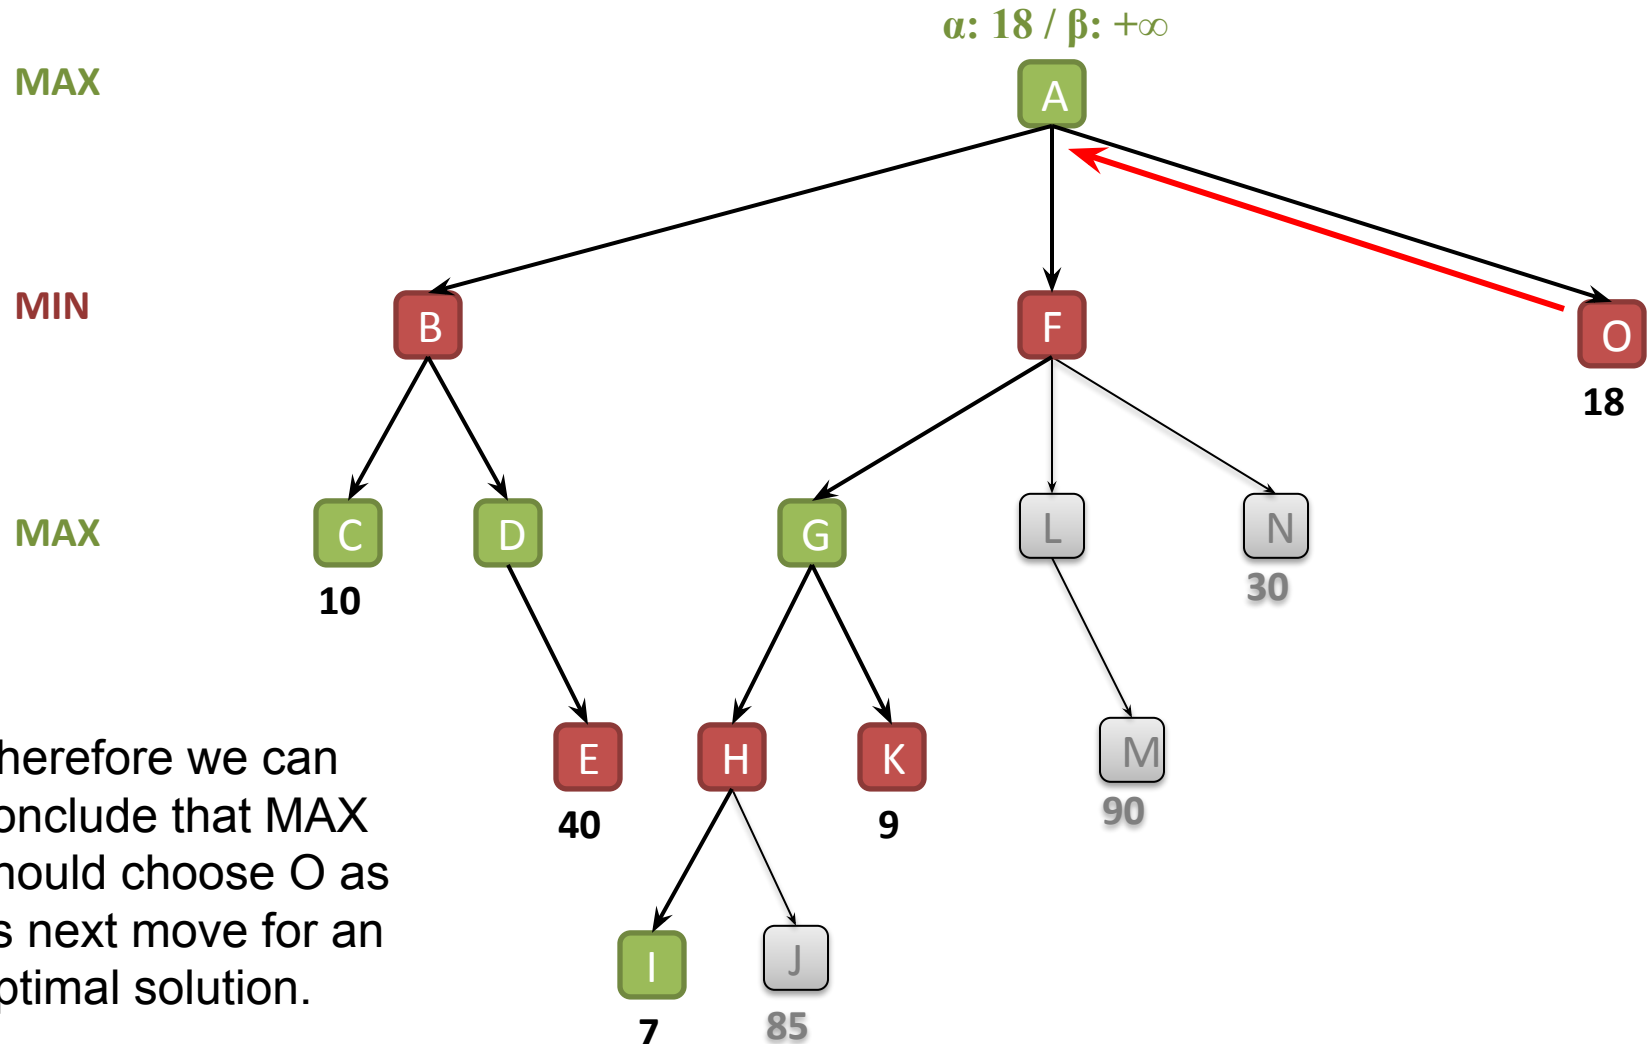
α - β Pruning



α - β Pruning



α - β Pruning



Properties of Alpha-Beta Pruning

- Pruning **does not** affect the final results
- Good move ordering improves effectiveness of pruning
- With “perfect ordering”, time complexity = $O(k^{d/2})$ => doubles depth of search
 - What’s the worse and average case time complexity?
 - Does it make sense then to have good heuristics for which nodes to expand first?

Readings

- Artificial Intelligence a Modern Approach
 - Chapter 5: Adversarial Search